

COMP201 – Software Engineering I

Lecture 9 – Finite State Machines and Petri Nets

Lecturer: T. Carroll

Email: Thomas.Carroll2@Liverpool.ac.uk

Office: G.14

See Vital for All Notes



Today


Overview

- Finite State Machines
- Mealy Machines
- Moore Machines
- Intro to Petri Nets

Finite State Machines

Finite State Machines (Finite State Automata)

- Model behaviour of a system or a complex object
- Limited number of defined states, changing with circumstance
- States change with input
- Machine either accepts a string, or does not accept a string
- *You may recall finite state machines (or automata) from COMP209.*



Finite State Machines - Definition


- *A model of computation* consisting of
 - a set of *states*
 - a *start (initial) state*,
 - *Accepting state*,
 - an *input alphabet*, and
 - a *transition function* that maps input symbols and current states to a *next state*

Finite State Machines - Definition

- Computation begins in the start state with an input string over the input alphabet.
- Automaton transitions to new states depending on the transition function.
 - **states** *define behaviour and may produce actions*
 - **state transitions** *are movement from one state to another*
 - **input events** *are either externally or internally generated, which may possibly trigger rules and lead to state transitions.*

Exercise

- What strings would be accepted by this Automaton?



Variants of FSMs

- There are many variants, for instance,
 - machines having actions (**outputs**) associated with transitions (***Mealy machine***) or states (***Moore machine***),
 - multiple start states,
 - transitions conditioned on no input symbol (a null) or more than one transition for a given symbol and state (***nondeterministic finite state machine***),
 - **one or more** states designated as *accepting states* (***recognizer***), etc.

FSM with Output (Mealy and Moore Machines)

- Finite state automata are like computers:
 - they receive input
 - they process the input by changing states.
- The only output that we have seen finite automata produce so far is a yes/no at the end of processing.
- We will now look at two models of finite automata that produce **more output** than a yes/no.


Moore Machines

Moore Machine

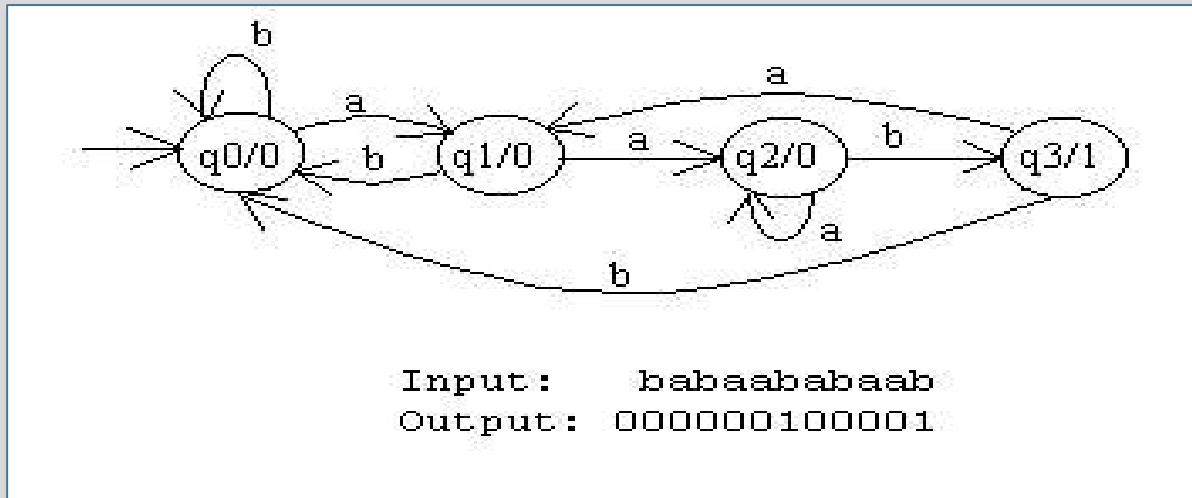
- A **Moore machine** is a FSA with two extra attributes.

1. It has TWO alphabets, an **input** and **output** alphabet.
2. It has an **output letter** associated with each state.

The machine writes the appropriate output letter as it **enters** each state.



Example - Moore Machine




This machine might be considered as a "counting" machine.

The output produced by the machine contains a 1 for each occurrence of the substring **aab** found in the input string.

Mealy Machines

Mealy Machine

- Mealy machines are **computationally equivalent** to Moore machines
- However, Mealy machines move the output function from the state **to the transition**.
- This turns out to be **easier to deal with** in practice, making Mealy machines **more practical**.



A Mealy Machine Produces Output on a Transition

Transitions are labelled i/o where

- i is a character in the input alphabet and
- o is a character in the output alphabet.

The following Mealy machine takes the one's complement of its binary input. In other words, it flips each digit from a 0 to a 1 or from a 1 to a 0.



Input: 010110
Output: 101001

A Mealy Machine Produces Output on a Transition

- Mealy machines are **complete** in the sense that there is a transition for each character in the input alphabet leaving every state.
- There are **no accept states** in a Mealy machine because it is not a language recogniser, it is an *output producer*.
- Its output will be the same length as its input.

Petri Nets


Petri Net Models

- **Petri Nets** were developed originally by *Carl Adam Petri* (when he was 13), and were the subject of his doctoral dissertation in 1962, at Technischen Hochschule Darmstadt.
- Originally created to model chemical reactions
- Extended, developed, and applied in a variety of areas.
- While the mathematical properties of Petri Nets are interesting and useful, the beginner will find that a good approach is to learn to model systems by constructing them graphically.



The Basics

- A **Petri Net** is a collection of directed arcs connecting places and transitions.
- Places may hold tokens.
- The state or marking of a net is its assignment of tokens to places.




Capacity and Connections

- Arcs have **capacity 1** by default; if otherwise, the capacity is marked on the arc, or use multiple arrows
- Places have **infinite capacity** by default.
- Transitions have **no capacity**, and cannot store tokens at all.
- Arcs can only connect **places and transitions**

Enabling Condition

- Transitions are the **active** components and places and tokens are **passive** components.
- A transition is **enabled** if each of the input places contains tokens, and the tokens are at least equal to the arc weight.




Transition t1 is not enabled, transition t2 is enabled.

Firing

An **enabled** transition may **fire**.

Firing corresponds to **consuming** tokens from the input places and **producing** tokens for the output places.




Firing is **atomic** (only one transition fires at a time, even if more than one is enabled)


Transitions with Multiple Inputs and Outputs

Transition t1 has three **input places** (p1, p2 and p3) and two **output places** (p3 and p4).

Place p3 is both an input and an output place of t1.




An Example Petri Net




Creating/Consuming Tokens

A transition without any input can fire at any time and produces tokens in the connected places:



Creating/Consuming Tokens

A transition without any output must be enabled to fire and deletes (or consumes) the incoming token(s):




Non-Determinism in Petri Nets

Two transitions fight for the same token: **conflict**.



Even if there are two tokens, there is still a conflict.

The next transition to fire (t1 or t2) is arbitrary (**non-deterministic**).




When Arcs have Different Weights...

- When fired, the **tokens in the input places are consumed**, according to arc weights and place capacities
- **New tokens are created in output places**, according to arc weights and place capacities.
- This results in a new marking of the net

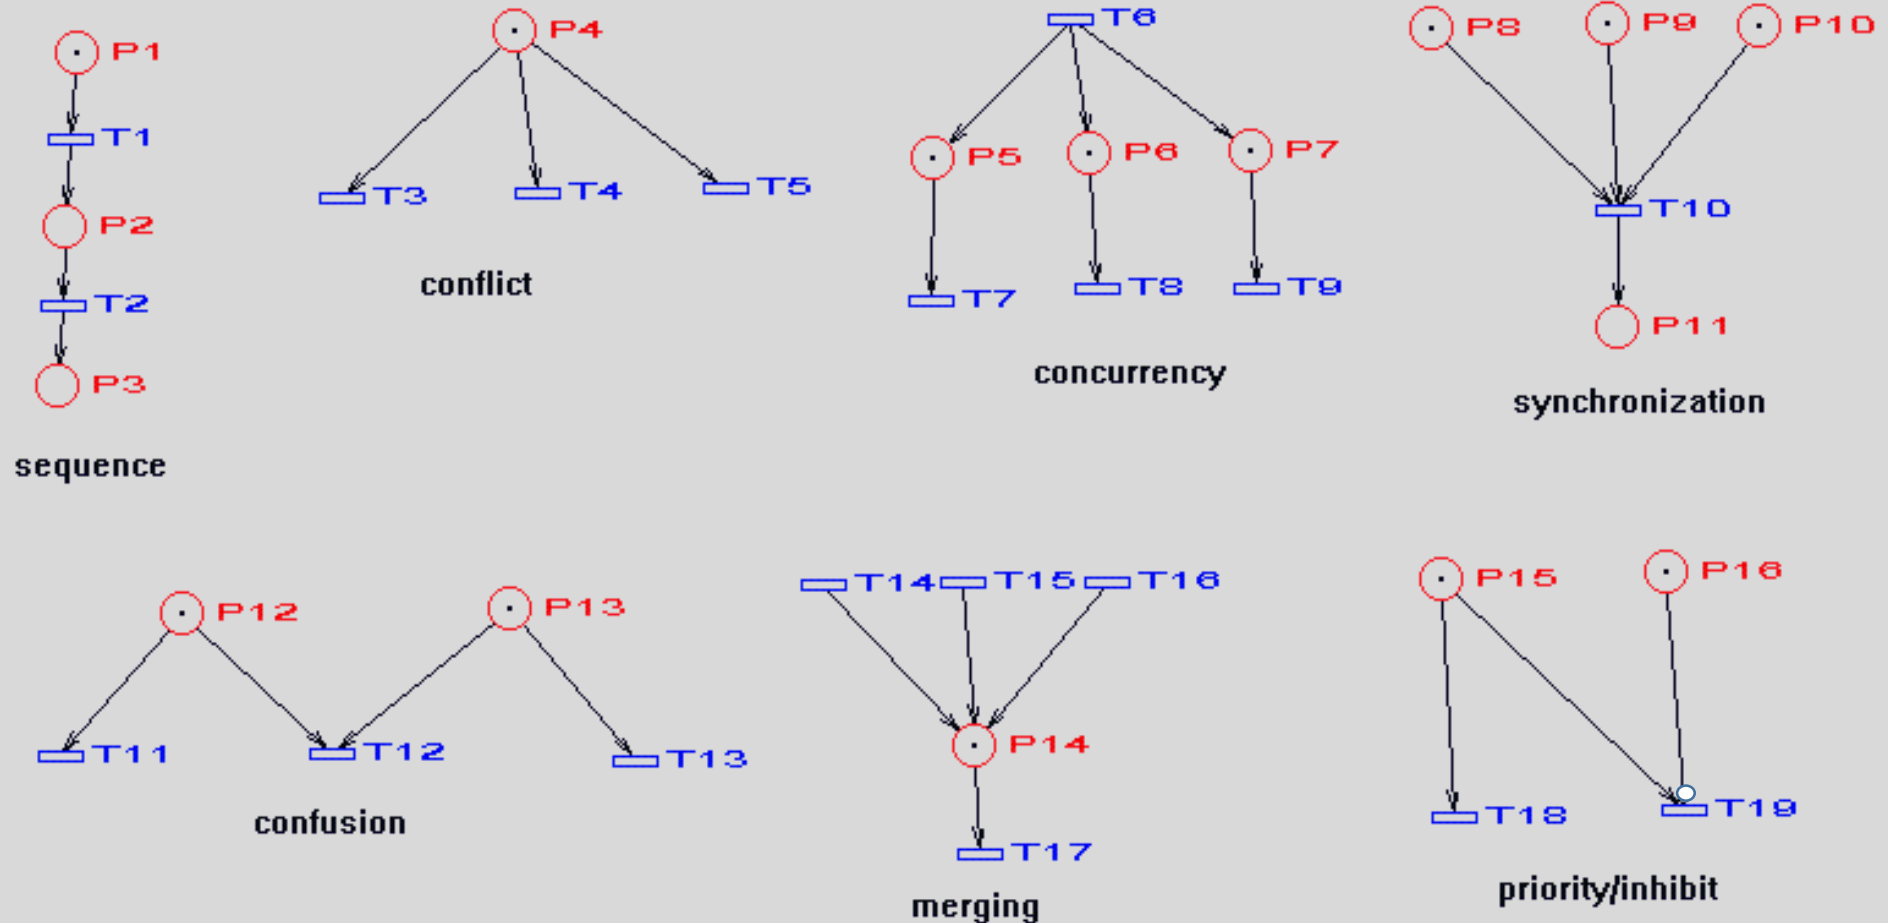


Inhibition


- Inhibition prevents a transition from firing



A Collection of Primitive Structures that Occur in Real Systems



How are they used in Chemistry?



Why use Petri Nets?

- Petri nets are **non-deterministic** and thus may be used to model discrete distributed systems.
 - There may be several arcs which can fire and we do not know in which order this will happen..
- There have been many **variations and extensions** of Petri nets to more effectively model a **wider variety of systems**.
- Since Petri nets have a rigorous mathematical notation, many questions about a system may be verified by studying properties of the Petri net.

Example Extension: High-Level Petri Nets

- The classical Petri net was invented by Carl Adam Petri in 1962.
- **High-level Petri nets** extend the classical Petri nets:
 - colour (for the modelling of attributes)
 - time (for performance analysis)
 - hierarchy (for the structuring of models, DFD's)



Recap

Recap

- Finite State Machines have **states, transition functions, alphabets, and accepting state** – Language acceptors
- Mealy and Moore machines **add output to FSM** – process input, rather than accept languages
- Mealy machines are **more practical**
- Petri nets have **Arcs, Places and Transitions**.
- Petri nets are **non-deterministic** and thus may be used to model discrete distributed systems.
- They have a well defined semantics and many variations and extensions of Petri nets exist.
- The **state or marking** of a net is an assignment of tokens to places.