

COMP207 Database Development

Lecture 3

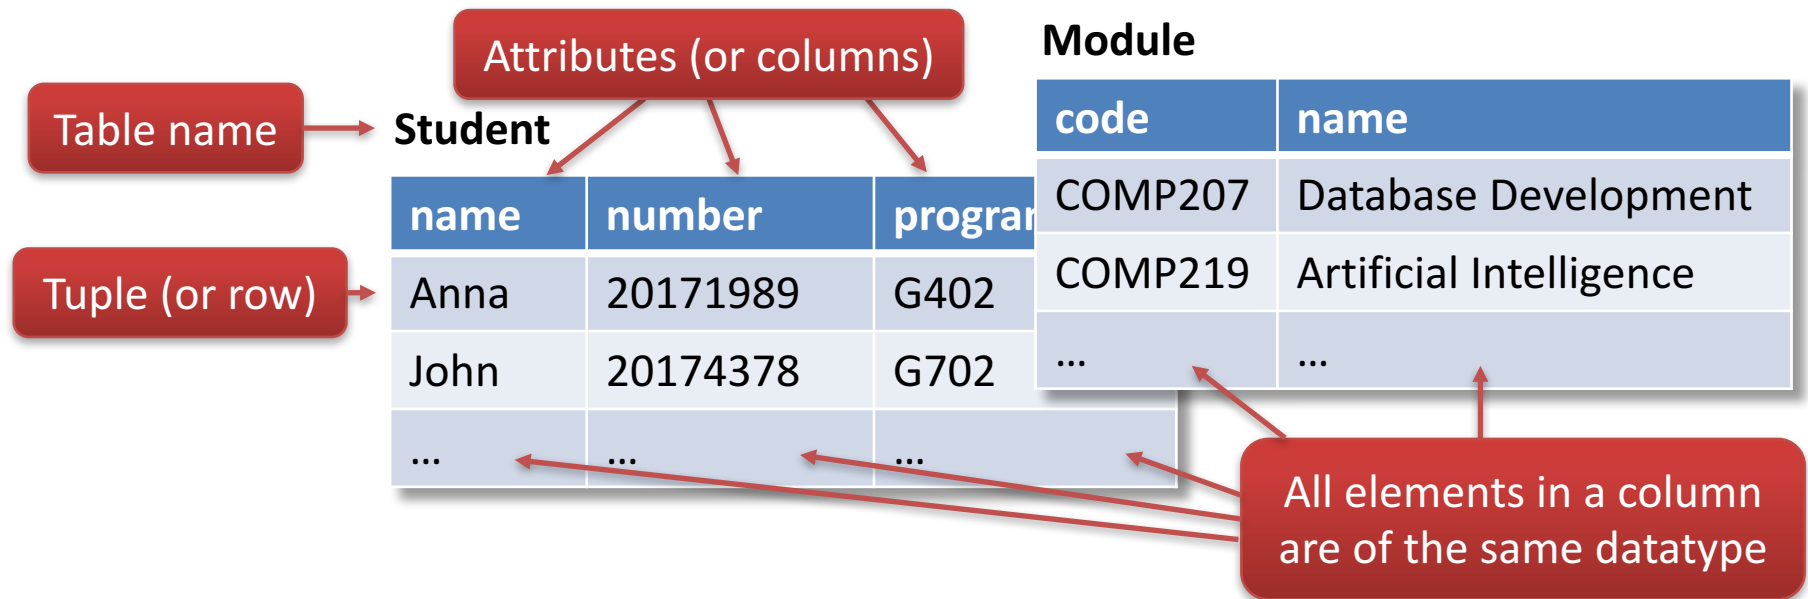
Transaction Management: Introduction

Transaction Management

- Transactions: Sequence of queries
- Ensures that operations on a database are executed without “damaging” the database
- This lecture:
 - Introduces the central concept of a transaction
 - Outlines how transactions help to execute database operations safely
 - Discusses desired properties of transactions
- Lectures 4-11:
 - How DBMS process transactions

Relational Model

- Data is organised in **tables** (also called **relations**)



- Schema:** description of all tables in the database

Student(name, number, programme)

Module(code, name)

Accessing Relational Databases

- Modern DBMS allow it to define & access relational databases using **SQL statements**
 - SQL as a *Data Definition Language (DDL)*

```
CREATE TABLE Student
( name          VARCHAR(15) ,
  number        INT          PRIMARY KEY,
  programme     VARCHAR(15) );
```

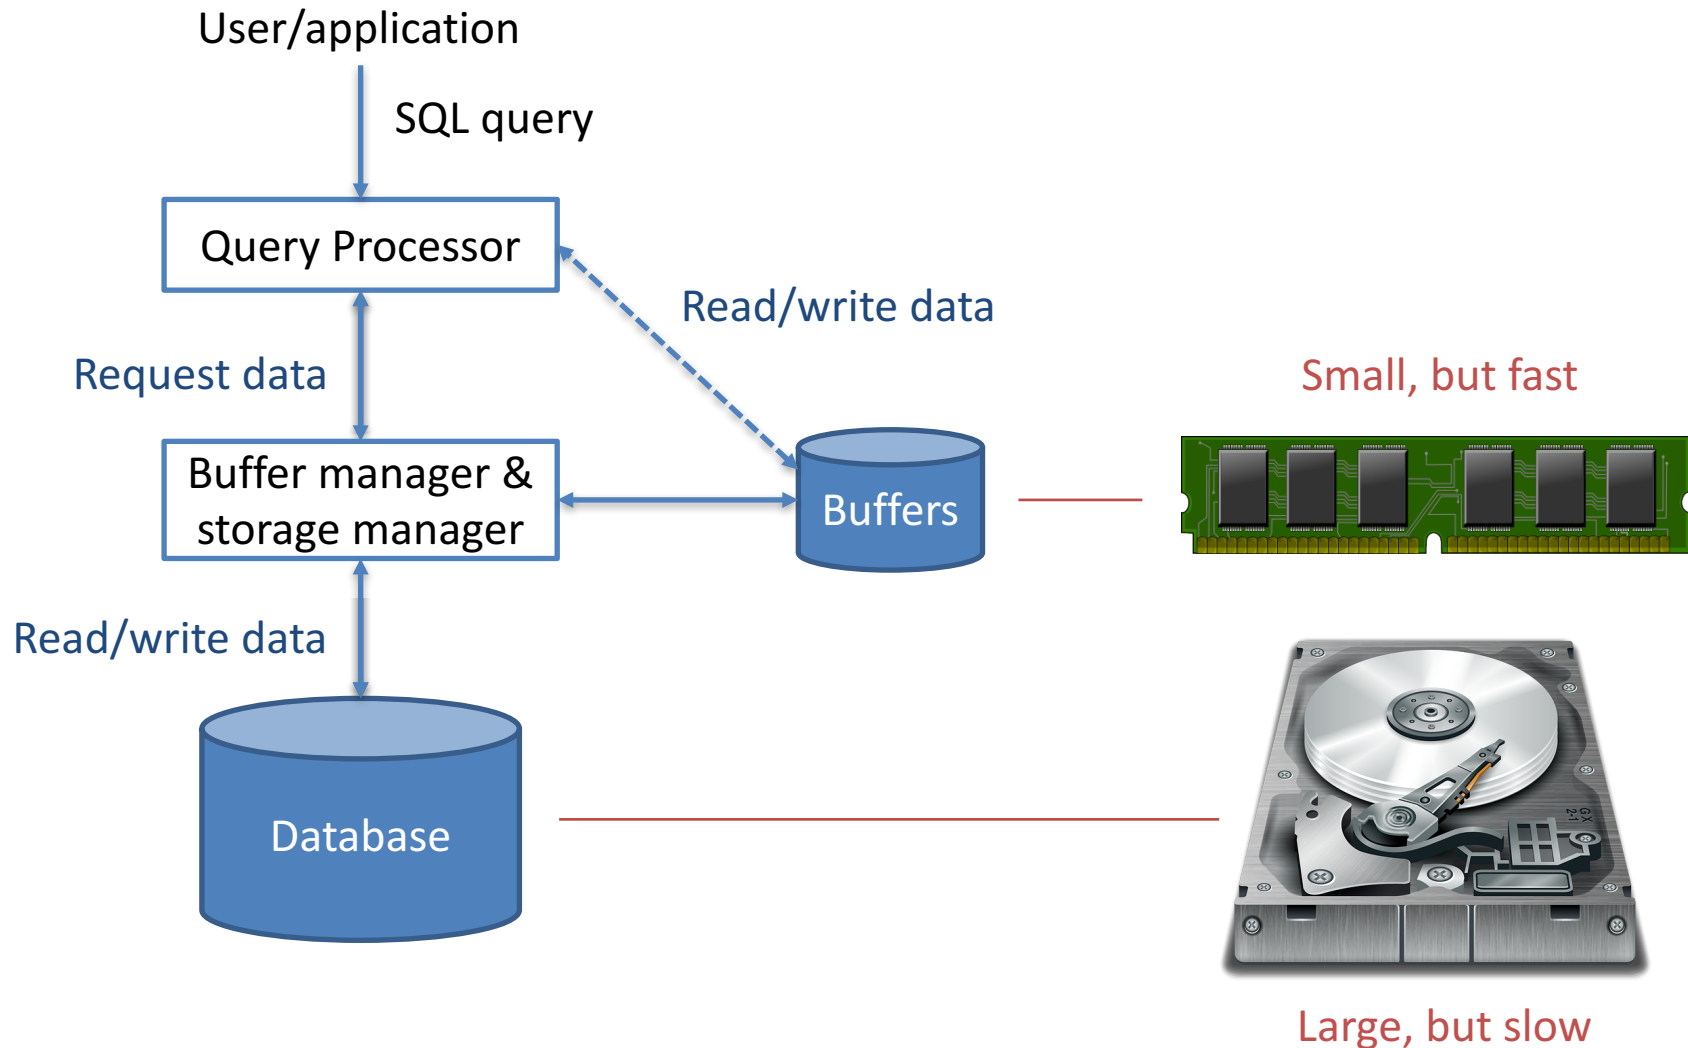
- SQL as a *Data Modification Language (DML)*

```
SELECT number
FROM Student
WHERE programme = 'G402';
```

- **SQL query** = SELECT/INSERT/UPDATE/DELETE statement

Execution of SQL Queries

(Simplified)



Relational Databases on Disk

Relational Model

File Model


Relation	=	“Table”	=	File
Tuple	=	“Row”	=	Record
Attribute	=	“Column”	=	Value

- Further terminology:
 - **Environment**: a relational DBMS on a computer system
 - **Catalog**: a collection of schemas in an environment

Transactions in SQL

Executing SQL Queries in "Real Life"

- So far: SQL queries in isolation




- In practice, problems may arise due to
 - **Concurrency**: SQL statements that overlap in time
 - **Partial execution** of SQL statements (e.g., due to failures)

Problem 1: Concurrency

Flights(flightNo, date, seatNo, seatStatus)

Might lead to an inconsistent database

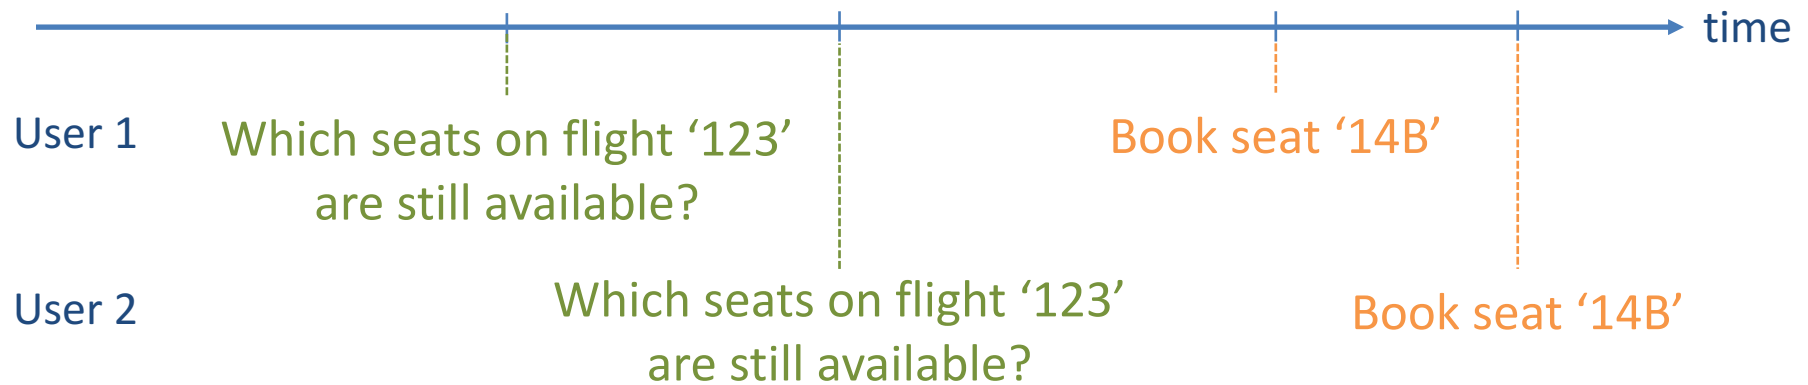


```
SELECT seatNo
FROM Flights
WHERE flightNo = 123
      AND date = '2018-10-2'
      AND seatStatus = 'available';
```

```
UPDATE Flights
SET    seatStatus = 'occupied'
WHERE  flightNo = 123
      AND date = '2018-10-2'
      AND seatNo = '14B';
```

Problem 1: Concurrency

Flights(flightNo, date, seatNo, seatStatus)




```
SELECT seatNo
FROM   Flights
WHERE  flightNo = 123
      AND date = '2018-10-2'
      AND seatStatus = 'available';
```

```
UPDATE Flights
SET    seatStatus = 'occupied'
WHERE  flightNo = 123
      AND date = '2018-10-2'
      AND seatNo = '14B';
```

Problem 1: Concurrency

Flights(flightNo, date, seatNo, seatStatus)

Cannot happen anymore



```
SELECT seatNo
FROM Flights
WHERE flightNo = 123
      AND date = '2018-10-2'
      AND seatStatus = 'available';
```

```
UPDATE Flights
SET    seatStatus = 'occupied'
WHERE  flightNo = 123
      AND date = '2018-10-2'
      AND seatNo = '14B';
```

Transactions to the Rescue

- SQL allows us to state that a group of SQL statements must be executed so that no conflicts arise (we'll see later how this is enforced)

```
START TRANSACTION;
```

Can often be omitted

```
SELECT seatNo  
FROM Flights  
WHERE flightNo = 123 AND date = ' 2018-10-2'  
      AND seatStatus = 'available';
```

```
UPDATE Flights  
SET    seatStatus = 'occupied'  
WHERE  flightNo = 123 AND date = ' 2018-10-2'  
      AND seatNo = '14B';
```

Before this, all changes to the database are tentative.

```
COMMIT;
```

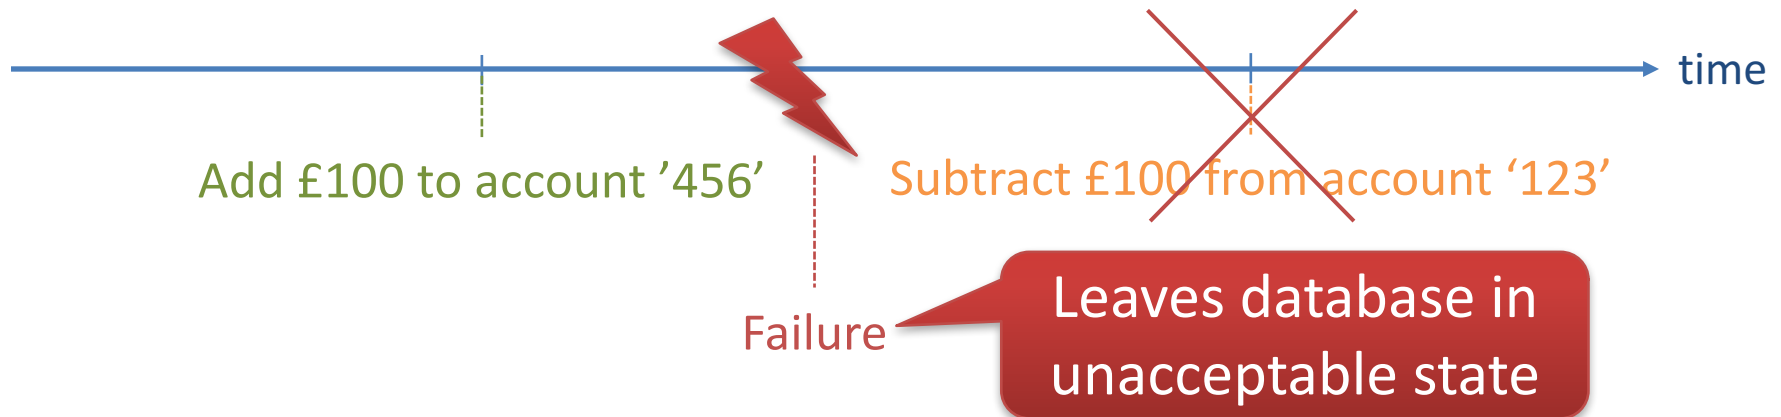
Transactions in SQL

- **Transaction in SQL:** a sequence of SQL statements
 - Special case: each individual SQL statement is a transaction
- By telling a DBMS that a sequence of SQL statements forms a transaction, it ensures *serialisable behaviour*
 - The transaction is executed as if was executed in isolation from all other transactions
 - *Equivalently:* ... as if all transactions were executed one after the other
- It also ensures other properties...

Problem 2: Partial Execution

Accounts(accountNo, accountHolder, balance)

Goal: Transfer £100 from account '123' to account '456'



```
UPDATE Accounts
SET    balance = balance + 100
WHERE  accountNo = 456;
```

```
UPDATE Accounts
SET    balance = balance - 100
WHERE  accountNo = 123;
```

Transactions to the Rescue

- SQL allows us to state that a transaction must be executed atomically (as a whole or not at all)

```
START TRANSACTION;
```

```
UPDATE Accounts  
SET      balance = balance + 100  
WHERE    accountNo = 456;
```

```
UPDATE Accounts  
SET      balance = balance - 100  
WHERE    accountNo = 123;
```

```
COMMIT;
```

Transactions in SQL – Summary

- Transaction in SQL: a sequence of SQL statements
 - Special case: each individual SQL statement is a transaction
 - General form:

`START TRANSACTION;`

Starts a transaction
(can often be omitted)

SQL statements

`COMMIT;`

or

`ROLLBACK;`

Writes all changes to the database

Aborts the transaction

- Most DBMS ensure that transactions are executed
 - as if they were executed in series (“serialisability”)
 - as a whole or not at all (“atomicity”)

Let's see how this works in detail...

Transaction: Overview

- **Transaction:**
 - Is an executing program, often comprising several queries (in **SQL**)
 - Can be submitted interactively or embedded within another programming language
 - **More details:** We consider SQL statements as transaction operations (**read** and **write** – “access operations”)
 - Executed **concurrently** at hundreds per second
 - Operations must leave the database in a valid or consistent state – enforced using **ACID** properties

Translating SQL into Low-Level Operations

STAFF(staffNo, familyName, firstName, jobTitle, gender, DOB, salary, department)

Two SQL Statements  produces Three Transaction Operations

```
SELECT salary
FROM Staff
WHERE staffNo = 1234;
```

```
UPDATE Staff
SET salary = salary*1.1
WHERE staffNo = 1234;
```

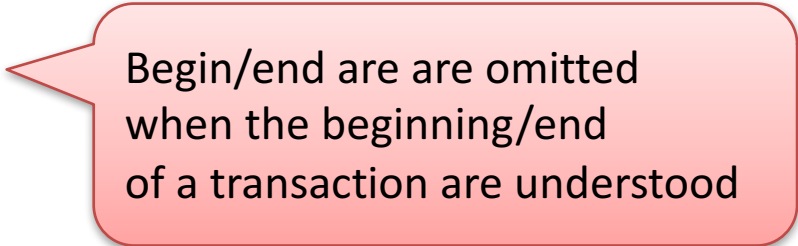
1. read(staffNo=1234, salary);
2. salary=salary*1.1;
3. write(staffNo=1234, salary);

Notes:

- Abstraction (at a high level)
- Read data item 'salary' from tuple with primary key 1234
- **Two database operations**
 - op1 (read) and op3 (write)
- **One non-database operation**
 - op2 (the calculation)

Transaction

- A logical unit of processing using access operations
 - Begin
 - End
 - read (retrieval – SELECT etc.)
 - write (insert, update, or delete)
 - + other non-database operations



Begin/end are omitted when the beginning/end of a transaction are understood

Transactions Preserve Consistency

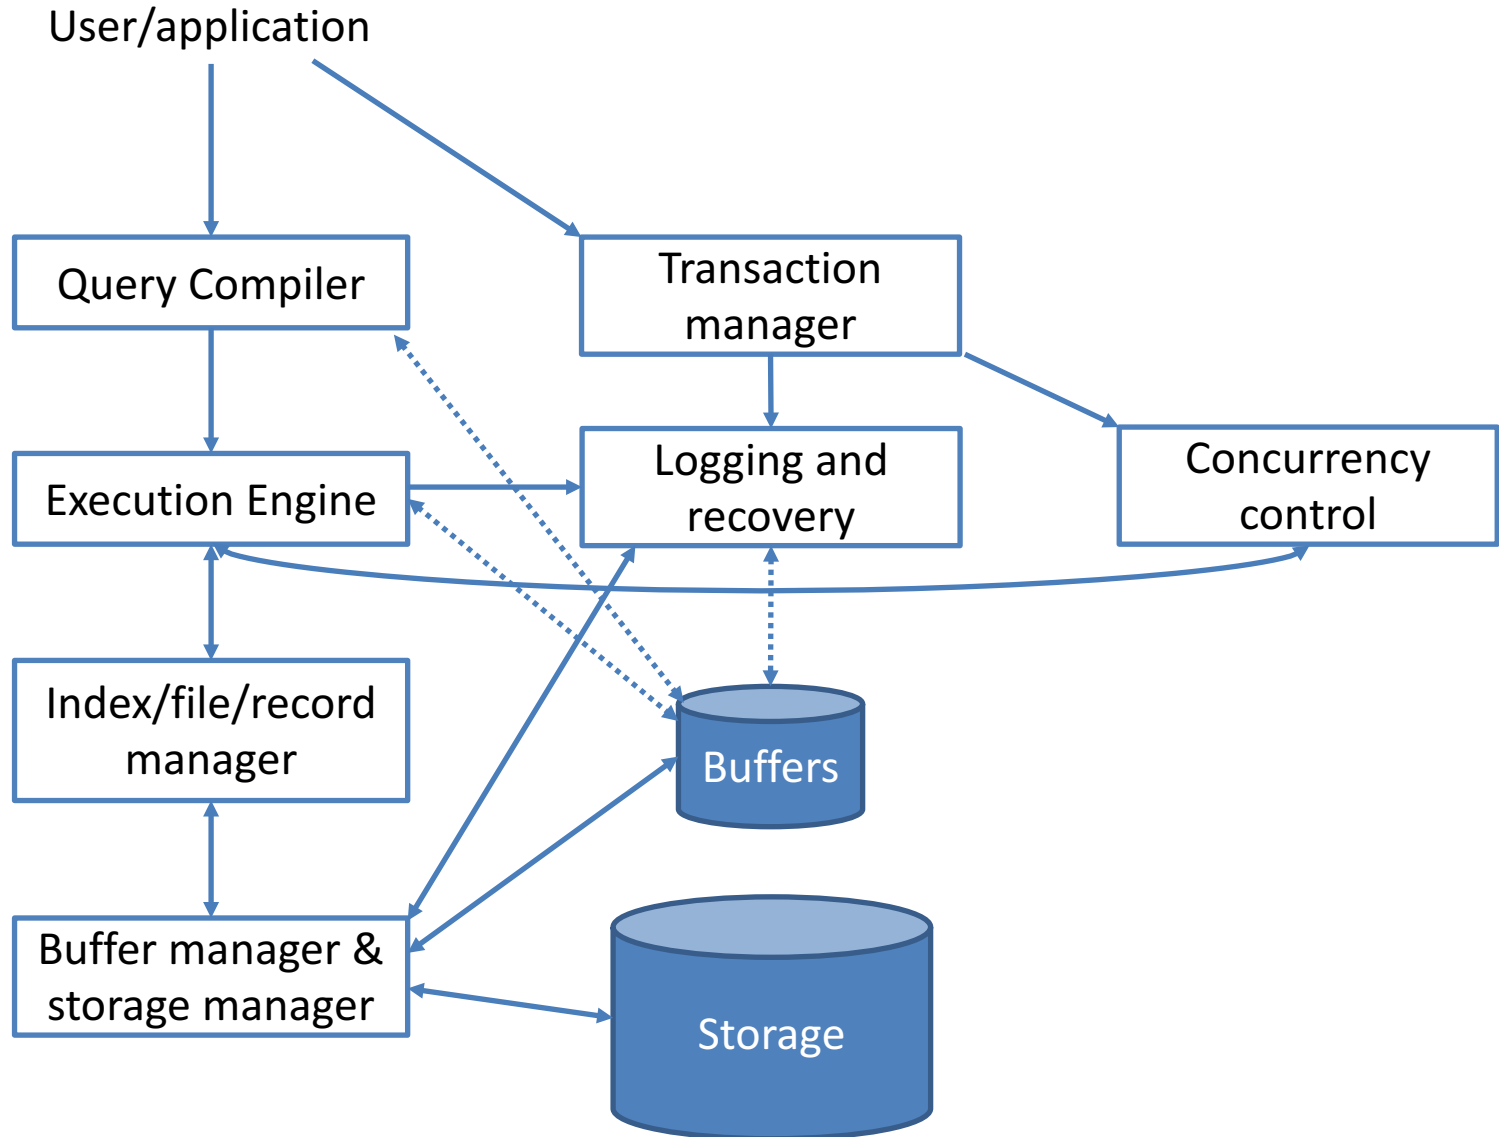
- Fundamental assumption:

Transactions always transform a *consistent* database state into another *consistent* database state.

- They produce one of two outcomes
 - Commit (i.e. Successful)
 - Execution was successful and database is left in a consistent state
 - Abort (i.e. Failed)
 - Execution was not successful and we need to restore the database to the state it was in before execution

Relational DBMS Components

(Simplified, from lecture 1)



Transaction - ACID Properties

- Transactions must maintain the correctness of the database, so we use **ACID** properties to validate transaction execution
- **A: Atomicity**
 - via Recovery Control (Logging and Recovery)
- **C: Consistency**
 - via Scheduler – Concurrency Control
- **I: Isolation**
 - via Scheduler – Concurrency Control
- **D: Durability** (or permanency)
 - via Recovery Control

A - Atomicity

- A transaction is an **atomic unit of processing**
 - An indivisible unit of execution
 - Executed in its entirety or not at all
- Deals with failure (“aborts”)
 - User aborts transaction (e.g., cancel button)
 - System aborts transaction (e.g., deadlock)
 - Transaction aborts itself (e.g., unexpected database state)
 - System crashes, network failure, etc.

A - Atomicity

- Abort - an error prevented full execution
 - We **UNDO** the work done up to the error point
 - System re-creates the database state as it was before the start of the aborted transaction
- Commit - no error, entire transaction executes
 - The system is updated correctly

C - Consistency

- A correct execution of the transaction must take the database from one consistent state to another
 - It should correctly transform the database state to reflect the effect of a real world event
 - Transactions may not violate integrity constraints

I - Isolation

- A transaction only makes its updates visible to other transactions after it has committed
 - The effect of **concurrently** executing a set of transactions is the same as if they had executed **serially** (“*serialisable*”)
 - When enforced strictly, this solves the **temporary update problem** and makes **cascading rollbacks** of transactions unnecessary

D - Durability

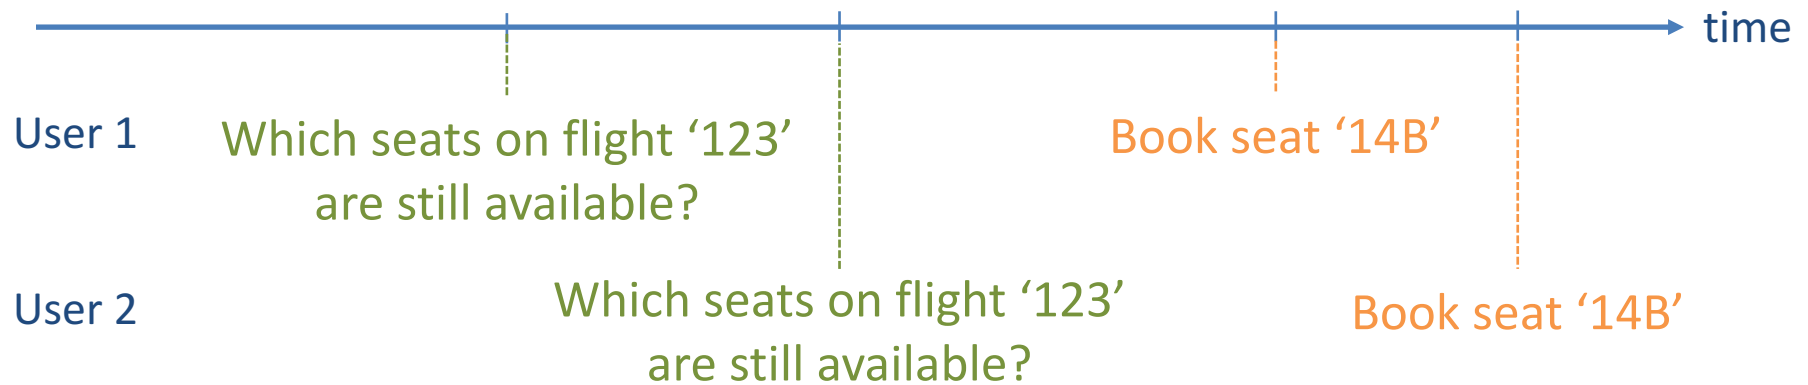
- Once a transaction commits and changes the database, these changes cannot be lost because of subsequent failure
 - The effect of a transaction on the database should not be lost after the commit point
 - We **REDO** the transaction if there are any problems after the update
 - Durability deals with things like media failure

Goal of Transactions

- DBMS's are expected to be reliable and remain in a consistent state
- The components of a DBMS that ensure this are:
 - Concurrency Control: responsible for 'C' and 'I'
 - Recovery Control: responsible for 'A' and 'D'
- These help maintain the 'ACID' Properties

Example 1

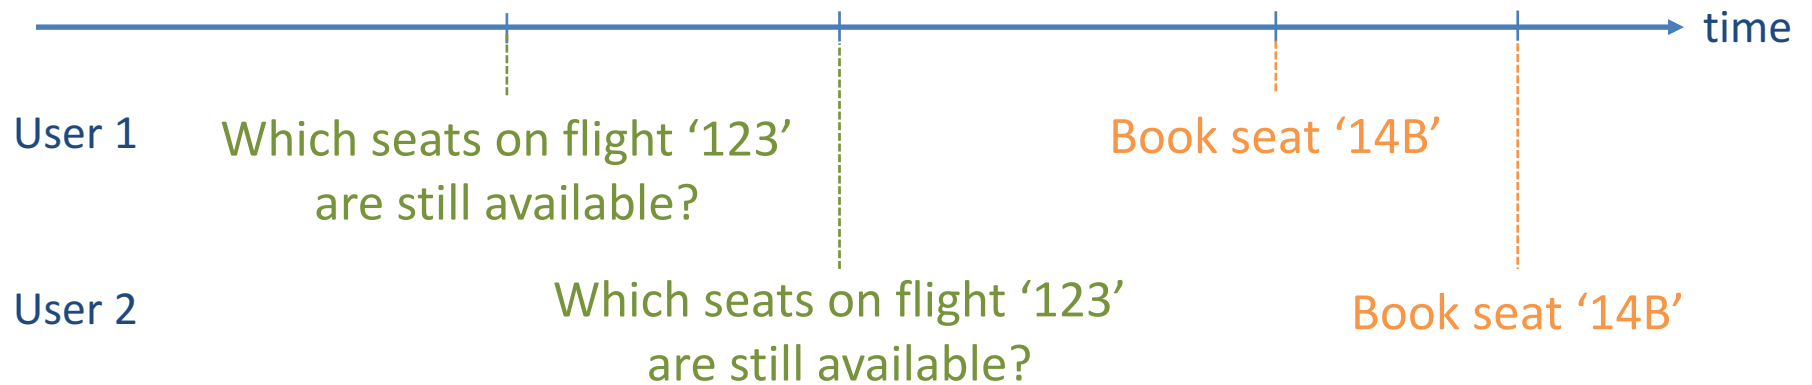
Flights(flightNo, date, seatNo, seatStatus)



```
SELECT seatNo
FROM   Flights
WHERE  flightNo = 123
      AND date = '2017-09-28'
      AND seatStatus = 'available';
```


```
UPDATE Flights
SET    seatStatus = 'occupied'
WHERE  flightNo = 123
      AND date = '2017-09-28'
      AND seatNo = '14B';
```

Example 1



- Which of the ACID properties does this violate?

Example 1




- Which of the ACID properties does this violate?
 - *Consistency*
 - Note: *Isolation* is **not** violated

Example 2

Accounts(accountNo, accountHolder, balance)

Goal: Transfer £100 from account '123' to account '456'




```
UPDATE Accounts
SET    balance = balance + 100
WHERE  accountNo = 456;
```

```
UPDATE Accounts
SET    balance = balance - 100
WHERE  accountNo = 123;
```

Example 2


Goal: Transfer £100 from account '123' to account '456'



- Which of the ACID properties does this violate?


Example 2

Goal: Transfer £100 from account '123' to account '456'




- Which of the ACID properties does this violate?
 - *Atomicity*

Example 3



- Which of the ACID properties does this violate?

Example 3



- Which of the ACID properties does this violate?
 - *Atomicity*
 - *Isolation*

Can you think of any situation where
Durability might be violated?

Summary

- Transactions are sequences of operations on a database (here: read and write operations)
- To avoid “damaging” the database, a DBMS should enforce the ACID properties. In particular:
 - Transactions should be executed as a whole or not at all.
 - Transactions should be executed as if they were executed serially, one after the other.
- The ACID properties are enforced by:
 - Concurrency control: C and I
 - Recovery control: A and D