

COMP207

Database Development

Lecture 5

Transaction Management:
Which Schedules are “Good”?

Review

- Previous lecture:
 - Exemplified ACID properties
 - Started to talk about concurrency control
 - Schedules
 - Serial & concurrent schedules
- What is a **schedule**? When is it **serial/concurrent**?

Schedules

- Schedule:
 - holds operations of one or several transactions for execution
 - operations of individual transactions occur in the order in which they occur in the transaction

- Example:

Transaction T ₁
read_item(X)
X := X + N
write_item(X)
read_item(Y)
Y := Y + N
write_item(Y)
commit

Transaction T ₂
read_item(X)
X := X + M
write_item(X)
commit



Schedule for T ₁ and T ₂	
read_item(X)	
	read_item(X)
	X := X + M
X := X + N	
write_item(X)	
read_item(Y)	
	write_item(X)
	commit
Y := Y + N	
write_item(Y)	
commit	

Is this schedule serial?

No (but concurrent)

Serial Schedules

- Operations of transaction 1 come first, then operations of transaction 2, then operations of transaction 3, ...
- No interleaving of operations

Transaction T ₁
read_item(X)
X := X + N
write_item(X)
read_item(Y)
Y := Y + N
write_item(Y)
commit

Transaction T ₂
read_item(X)
X := X + M
write_item(X)
commit



Serial schedule for T ₁ and T ₂	
read_item(X)	
X := X + N	
write_item(X)	
read_item(Y)	
Y := Y + N	
write_item(Y)	
commit	
	read_item(X)
	X := X + M
	write_item(X)
	commit

- Always guarantee consistency & isolation

Reminder

(see Lecture 3)

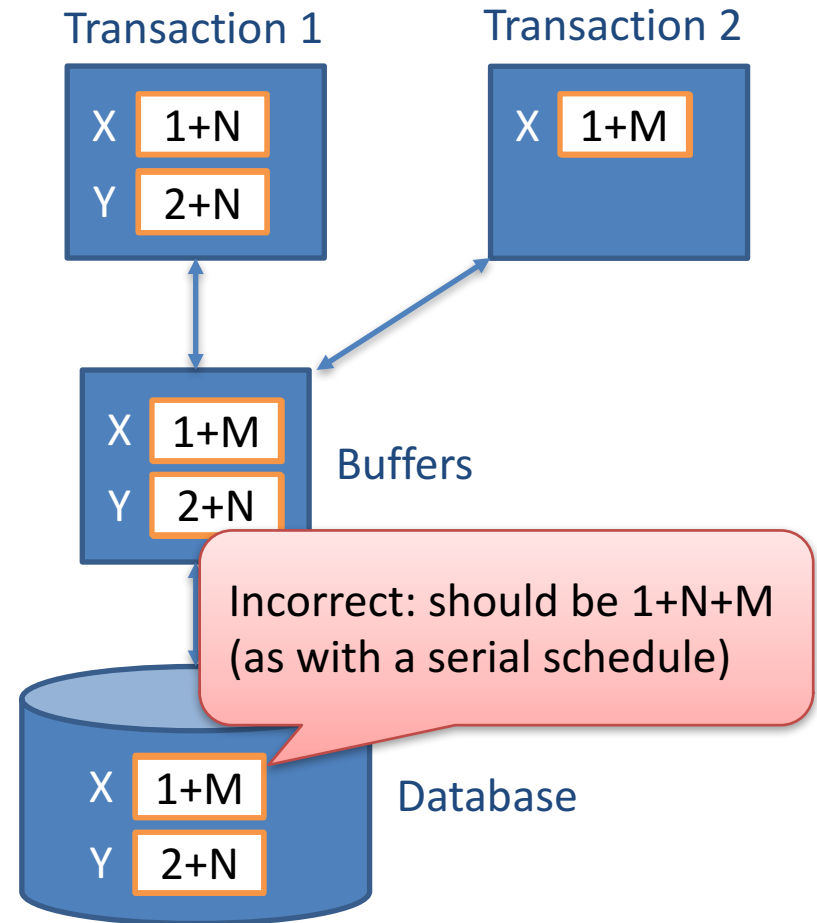
Transactions *always* transform a database from one consistent state to another consistent state.

Concurrent Schedules Problems

- **Lost updates:**
 - The problem demonstrated at the end of the last lecture (see next slide).
 - A successful transaction can be overwritten by another transaction.
- **Dirty reads:** a transaction can see the intermediate results of another transaction before it commits
- **Inconsistent analysis:** a transaction reads several values but another transaction updates some of them in the meantime
- Other problems: see tutorial slides

Concurrent Schedules Do Not Guarantee Consistency

Time	Schedule	
t0		
t1	read_item(X)	
t2		read_item(X)
t3		$X := X + M$
t4	$X := X + N$	
t5	write_item(X)	
t6	read_item(Y)	
t7		write_item(X)
t8		commit
t9	$Y := Y + N$	
t10	write_item(Y)	
t11	commit	



Schedules – Two Extremes

- **Serial Schedules**
 - execute correctly
 - maintain consistency of the database
 - *inefficient* in multi-user environments
- **Concurrent Schedules ('non-serial')**
 - may not execute correctly
 - may not guarantee consistency of the database or isolation
 - *efficient* in multi-user environments

What makes a schedule a “good” schedule?
(one that is *efficient* and guarantees *consistency/isolation*)

Serialisable Schedules

- A **schedule S is serialisable** if there is a **serial schedule S'** that has the same effect as S on every initial database state.

Schedule S	
read_item(X)	
X := X + N	
write_item(X)	
	read_item(X)
	X := X + M
	write_item(X)
	commit
read_item(Y)	
Y := Y + N	
write_item(Y)	
commit	



Equivalent serial schedule S'	
read_item(X)	
X := X + N	
write_item(X)	
read_item(Y)	
Y := Y + N	
write_item(Y)	
commit	
	read_item(X)
	X := X + M
	write_item(X)
	commit

So, S is serialisable

Serialisable Schedules

- A **schedule S is serialisable** if there is a **serial schedule S'** that has the same effect as S on every initial database state.

Schedule S	
read_item(X)	
	read_item(X)
	$X := X + M$
$X := X + N$	
write_item(X)	
read_item(Y)	
	write_item(X)
	commit
$Y := Y + N$	
write_item(Y)	
commit	

In general not serialisable

However, serialisable if $N = 0$
(do you see why?)

Serialisable Schedules

- A **schedule S is serialisable** if there is a **serial schedule S'** that has the same effect as S on every initial database state.

Schedule S	
read_item(X)	
	read_item(X)
	X := X + M
X := X + N	
write_item(X)	
read_item(Y)	
	write_item(X)
	commit
Y := Y + N	
write_item(Y)	
commit	

N = 0
→

Equivalent to S (if N = 0)	
read_item(X)	
X := X + N	
write_item(X)	
read_item(Y)	
Y := Y + N	
write_item(Y)	
commit	
	read_item(X)
	X := X + M
	write_item(X)
	commit

Serialisable Schedules

- Serialisable schedules are essentially those schedules that we are looking for.
- Guarantee:
 - Correctness & consistency (because serial schedules do) ✓
 - No isolation (but could be added later...)
- Problem: serialisability is difficult to test
 - Does not only depend on reads, writes, and commits, but also on the non-database operations
 - Non-database operations can be complex

Assumption from now on: serialisability only depends on read and write operations (still difficult to test)

Conflict-serialisability

Conflict-Serialisability

- Stronger form of serialisability that is used/enforced by most commercial DBMS
- Based on the notion of a **conflict**.

A **conflict** in a schedule is a pair of operations from different transactions *that cannot be swapped* without changing the behaviour of at least one of the transactions.

Do not write commit/abort
(will deal with these later)

- Example: $r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); w_1(Y)$




Observation: Operations can only be in conflict if one of them is a write & they access the same item.

Conflict – Characterisation

- A **conflict** in a schedule is a pair of operations from different transactions such that:
 - the operations access the same item
 - at least one of them is a write operation
- Example (from previous slide):


S: $r_1(X)$; $w_1(X)$; $r_2(X)$; $w_2(X)$; $r_1(Y)$; $w_1(Y)$



conflict in S

- Can you find any other conflicts?

Conflict-Serialisability

- Two schedules S and S' are **conflict-equivalent** if S' can be obtained from S by swapping any number of *consecutive* non-conflicting operations from different transactions. 

- Example:

$S:$ $r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); w_1(Y)$

$r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y)$

$r_1(X); w_1(X); r_1(Y); r_2(X); w_2(X); w_1(Y)$

$r_1(X); w_1(X); r_1(Y); r_2(X); w_1(Y); w_2(X)$

$S':$ $r_1(X); w_1(X); r_1(Y); w_1(Y); r_2(X); w_2(X)$

Conflict-serialisable
schedule

Serial schedule

- A schedule is **conflict-serialisable** if it is conflict-equivalent to a serial schedule.

Example

Is the schedule S below conflict-serialisable?

S: $r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); w_1(Y); r_2(Y); w_2(Y)$

Example

Is the schedule S below conflict-serialisable?

S: $r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); w_1(Y); r_2(Y); w_2(Y)$

Yes: Conflict-serialisable

$r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); r_2(Y); w_2(Y)$

$r_1(X); w_1(X); r_1(Y); r_2(X); w_2(X); w_1(Y); r_2(Y); w_2(Y)$

$r_1(X); w_1(X); r_1(Y); r_2(X); w_1(Y); w_2(X); r_2(Y); w_2(Y)$

S': $r_1(X); w_1(X); r_1(Y); w_1(Y); r_2(X); w_2(X); r_2(Y); w_2(Y)$

Serial

Example 2

What about this one?

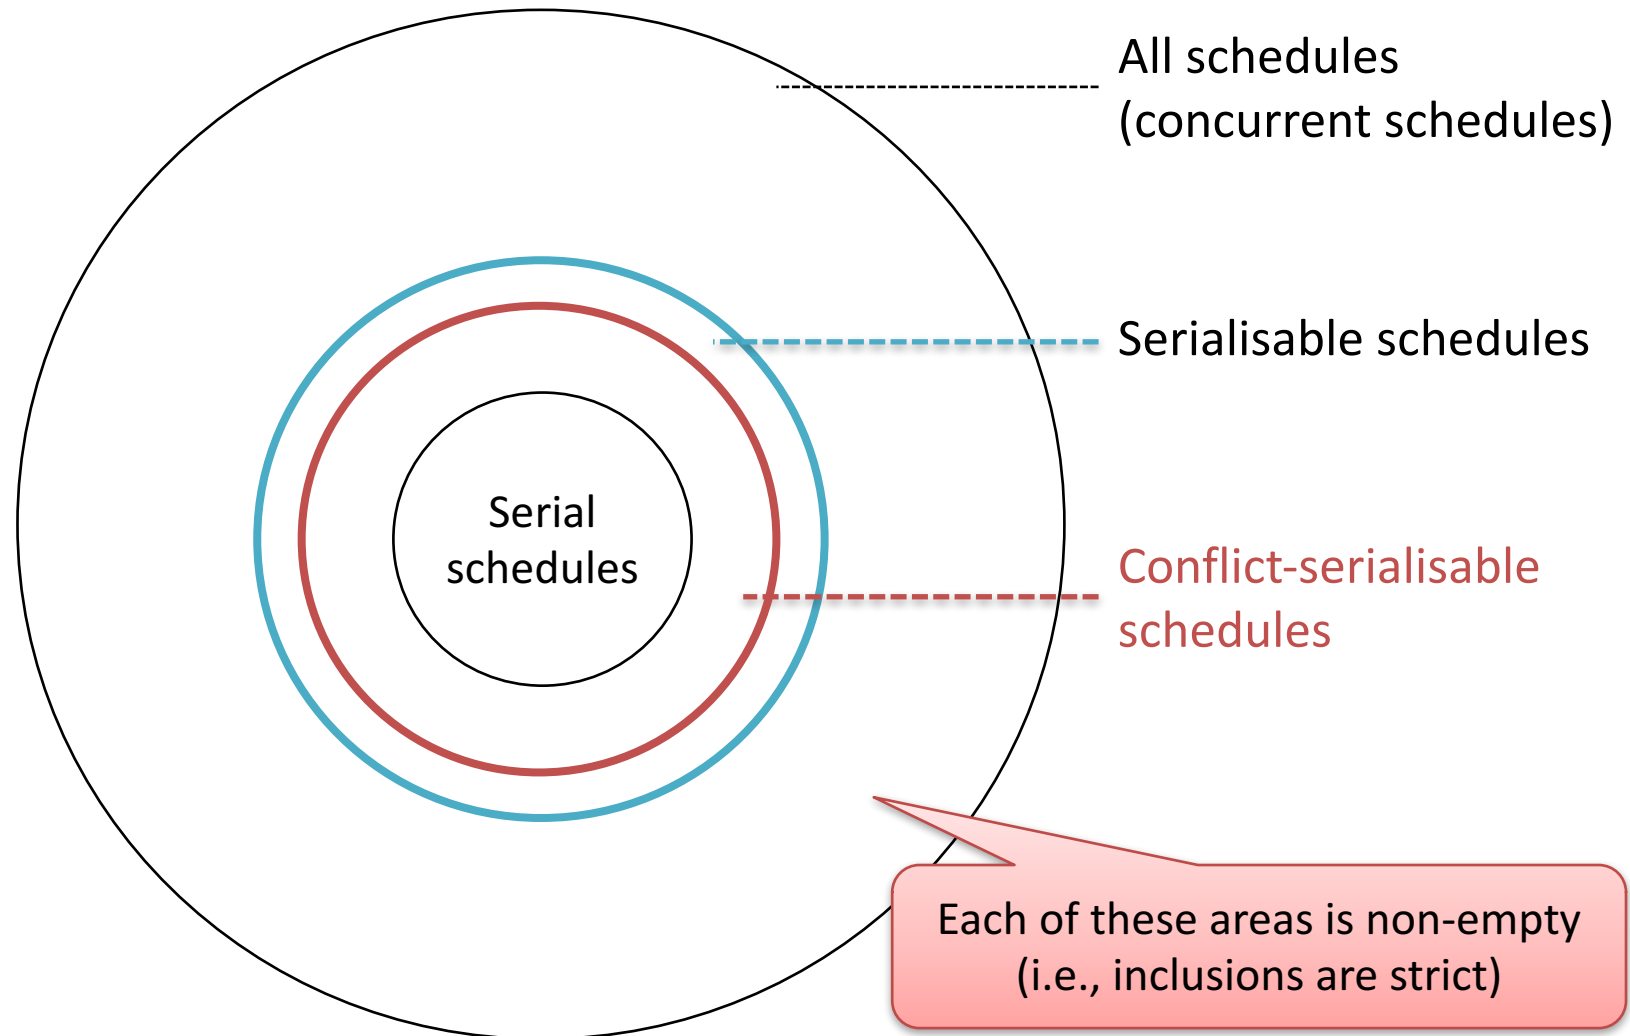
S: $r_2(X); r_1(Y); w_2(X); r_2(Y); r_3(X); w_1(Y); w_3(X); w_2(Y)$

Example 2

What about this one?

S: $r_2(X); r_1(Y); w_2(X); r_2(Y); r_3(X); w_1(Y); w_3(X); w_2(Y)$

Schedules



Example

A schedule that is **not** conflict-serialisable, but serialisable:

S: $w_2(X); w_1(X); w_1(Y); w_2(Y); w_3(X);$

S': $w_1(X); w_1(Y); w_2(X); w_2(Y); w_3(X);$

Why is it serialisable?

- The change in the database is given by $w_2(Y); w_3(X);$ and neither transaction reads anything


Why is it **not** conflict-serialisable?

- In a bit

How can we *test* if a schedule
is *conflict-serialisable*?

Example

A schedule that is **not** conflict-serialisable:



How do we know?

- Identify the **transactions** in S
- Identify the **conflicts** in S
- Conflicts **impose constraints** on the order of the transactions in any conflict-equivalent serial schedule


The cycle indicates that no serial schedule is conflict-equivalent to S

Precedence Graph

- The **precedence graph** for a schedule S is defined as follows:
 - It is a **directed graph**.
 - Its **nodes** are the transactions that occur in S.
 - It has an **edge** from transaction T_i to transaction T_j if there is a conflicting pair of operations op_1 and op_2 in S such that
 - op_1 appears before op_2 in S
 - op_1 belongs to transaction T_i
 - op_2 belongs to transaction T_j .
- Example:

S: $r_2(X); r_1(Y); w_2(X); r_2(Y); r_3(X); w_1(Y); w_3(X); w_2(Y)$

Precedence graph for S:



Testing Conflict-Serialisability

- To test if a schedule S is **conflict-serialisable**:
 - Construct the precedence graph for S .
 - If the precedence graph is **acyclic**, then S is conflict-serialisable. Otherwise not.


Acyclic graph: graph without a directed cycle

- Example 1: $S: r_1(X); w_1(X); r_2(X); w_2(X); r_1(Y); w_1(Y); r_2(Y); w_2(Y)$
Precedence graph for S :



has no cycle →
 S is conflict-serialisable

- Example 2: $S: r_2(X); r_1(Y); w_2(X); r_2(Y); r_3(X); w_1(Y); w_3(X); w_2(Y)$
Precedence graph for S :



contains a cycle →
 S is not conflict-serialisable

Why does this work?



This says:

- There is a conflict between an operation in T_1 (that appears first) and an operation in T_2

All conflict-equivalent schedulers: operation in T_1 is before operation in T_2

Why does this work?



All conflict-equivalent schedulers: operation x in T_1 is before operation y in T_2

- Proof by contradiction: Assume conflict-equivalent scheduler S' where this is not so
- Consider first *consecutive* swap between S and S' where x goes from being before y to being after y
- In that swap, either:
 - We swap x and y (not legal since they conflict)
 - Or we swap at most 1 of them (contradicts choice of swap)

Implication of a cycle

- A cycle in the precedence graph



- In serial scheduler: No transaction can be first among those in the cycle (since another must be before)

Exercise (3 min)


- Verify that the schedule below is **not** conflict-serialisable (using a precedence graph).

$S: w_2(X); w_1(X); w_1(Y); w_2(Y); w_3(X);$

Exercise (3 min)

- Verify that the schedule below is **not** conflict-serialisable (using a precedence graph).

$S: w_2(X); w_1(X); w_1(Y); w_2(Y); w_3(X);$



Exercise 2 (5 min)

- Verify that the schedule below is conflict-serialisable (using a precedence graph).


S: $r_1(Y)$, $r_3(Y)$, $r_1(X)$, $r_2(X)$, $w_2(X)$, $r_3(Z)$, $w_3(Z)$, $r_1(Z)$, $w_1(Y)$, $r_2(Z)$

- Can you find a conflict-equivalent serial schedule (using your precedence graph)?

Exercise 2 (5 min)

- Verify that the schedule below is conflict-serialisable (using a precedence graph).

S: $r_1(Y)$, $r_3(Y)$, $r_1(X)$, $r_2(X)$, $w_2(X)$, $r_3(Z)$, $w_3(Z)$, $r_1(Z)$, $w_1(Y)$, $r_2(Z)$



- Can you find a conflict-equivalent serial schedule (using your precedence graph)?

$r_3(Y)$, $r_3(Z)$, $w_3(Z)$

Exercise 2 (5 min)

- Verify that the schedule below is conflict-serialisable (using a precedence graph).

S: $r_1(Y)$, $r_3(Y)$, $r_1(X)$, $r_2(X)$, $w_2(X)$, $r_3(Z)$, $w_3(Z)$, $r_1(Z)$, $w_1(Y)$, $r_2(Z)$



has no cycle →
S is conflict-serialisable

Any topological ordering yields
a conflict-equivalent serial schedule

- Can you find a conflict-equivalent serial schedule (using your precedence graph)?

$r_3(Y)$, $r_3(Z)$, $w_3(Z)$

Exercise 2 (5 min)

- Verify that the schedule below is conflict-serialisable (using a precedence graph).

S: $r_1(Y)$, $r_3(Y)$, $r_1(X)$, $r_2(X)$, $w_2(X)$, $r_3(Z)$, $w_3(Z)$, $r_1(Z)$, $w_1(Y)$, $r_2(Z)$

T_2

- Can you find a conflict-equivalent serial schedule (using your precedence graph)?

$r_3(Y)$, $r_3(Z)$, $w_3(Z)$, $r_1(Y)$, $r_1(X)$, $r_1(Z)$, $w_1(Y)$,

Exercise 2 (5 min)

- Verify that the schedule below is conflict-serialisable (using a precedence graph).

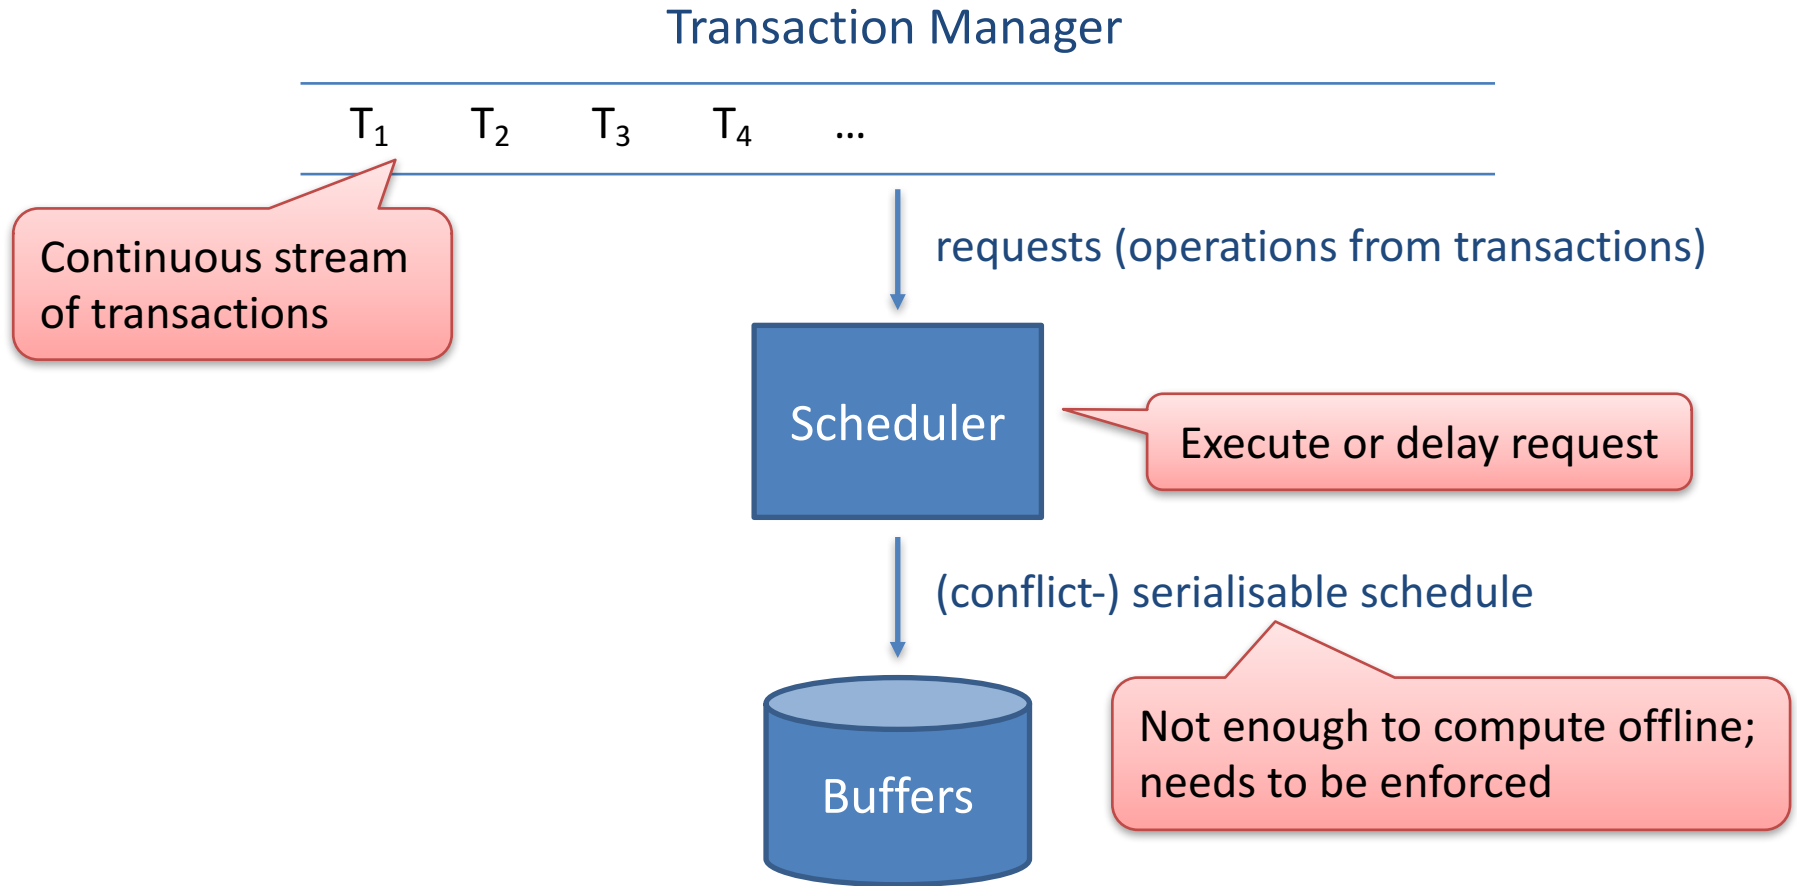
S: $r_1(Y), r_3(Y), r_1(X), r_2(X), w_2(X), r_3(Z), w_3(Z), r_1(Z), w_1(Y), r_2(Z)$

- Can you find a conflict-equivalent serial schedule (using your precedence graph)?

$r_3(Y), r_3(Z), w_3(Z), r_1(Y), r_1(X), r_1(Z), w_1(Y), r_2(X), w_2(X), r_2(Z)$

Are we done?

Transaction Scheduling in a DBMS



Enforcing Conflict-Serialisability Using Locks



Summary

- What schedules are “good”?
- First approximation: serialisable schedules
 - Equivalent to serial schedules
 - Guarantee correctness & consistency
 - But: difficult to check
- A good compromise: conflict-serialisable schedules
 - Conflict-equivalent to serial schedules
 - Imply serialisability (so inherit all the nice properties)