

COMP207

Database Development

Lecture 8

Transaction Management:
Dealing With Transaction Aborts &
System Failures

No lecture later today!


- Lecture today 17-18 is canceled
 - People can go to prof Paul Spirakis Inaugural lecture (if they want). That is 17:30-20 in ELT

Review of Undo Logging


- Logs activities with the goal of restoring (“undoing”) a previous database state.
- Log records (or log entries):
 - **<START T>**: Transaction T has started.
 - **<COMMIT T>**: Transaction T has committed.
 - **<ABORT T>**: Transaction T was aborted.
 - **<T, X, v>**: Transaction T has updated the value of database item X, and the old value of X was v.
 - Response to **write_item(X)**
 - If this entry occurs in the log, then the new value of X might not have been written to the database yet.
- Slightly different records for redo logging (later...)

Undo Logging: Procedure

1. If transaction T updates database item X and the old value was v, then $\langle T, X, v \rangle$ must be written to the log on disk **before** X is written to disk.



2. If transaction T commits, then $\langle \text{COMMIT } T \rangle$ must be written to disk **as soon as** all database elements changed by T have been written to disk



Review of Undo Logging

| Time | Transaction T ₁ | Transaction T ₂ |
|------|----------------------------|----------------------------|
| 1 | read_item(X) | |
| 2 | X := X * 2 | |
| 3 | write_item(X) | |
| 4 | | read_item(X) |
| 5 | read_item(Y) | |
| 6 | | X := X * 3 |
| 7 | | write_item(X) |
| 8 | Y := X + Y | |
| 9 | write_item(Y) | |

X = 1
Y = 2

- How does undo logging work on this schedule?
 - Which **log entries** are written to buffer/disk & when?
 - Which **other operations** must be executed & when?

| Time | Transaction T ₁ | Transaction T ₂ | Log (buffer) | Log (disk) |
|------|----------------------------|----------------------------|--------------------------|------------|
| 0 | | | <START T ₁ > | |
| 1 | read_item(X) | | | |
| 2 | X := X * 2 | | | |
| 3 | write_item(X) | | <T ₁ , X, 1> | |
| 4 | | | <START T ₂ > | |
| 5 | | read_item(X) | | |
| 6 | read_item(Y) | | | |
| 7 | | X := X * 3 | | |
| 8 | | write_item(X) | <T ₂ , X, 2> | |
| 9 | Y := X + Y | | | |
| 10 | write_item(Y) | | <T ₁ , Y, 2> | |
| 11 | flush_log | | | |
| 12 | output(X) | | | |
| 13 | output(Y) | | | |
| 14 | | | <COMMIT T ₁ > | |
| 15 | flush_log | | | |
| 16 | | flush_log | | |
| 17 | | output(X) | | |
| 18 | | | <COMMIT T ₂ > | |
| 19 | | flush_log | | |

X = 1
Y = 2

What if a transaction aborts?

If a Transaction Aborts...

- Use the undo log to undo all changes made by the transaction.
 - Similar to recovery with undo logs
 - But focuses on a single transaction
- Procedure:
 - Assume T aborts
 - Traverse the undo log from the last to the first item
 - If we see $\langle T, X, v \rangle$, change the value of X on disk back to v.

Other Logging Variants

Redo Logging
Undo/Redo Logging

Redo Logging

- Logs activities with the goal of *restoring* committed transactions (ignores incomplete transactions).
- Log records:
 - Same as before, but...
 - New meaning of $\langle T, X, v \rangle$: “Transaction T has updated the value of database item X & the new value of X is v.”
 - Direct response to **write_item(X)**
 - Haven’t changed X on disk yet!
- Have to modify the logging procedure...

Redo Logging: Procedure

1. T first writes all log records for all updates to disk
2. T writes $\langle \text{COMMIT } T \rangle$ to the log on disk
3. T writes all updates to disk



Example

| | | Local | | Buffer | | Database | | | |
|------|---------------|-------|----|--------|----|----------|----|--------------|------------|
| Time | Transaction | X | Y | X | Y | X | Y | Log (buffer) | Log (disk) |
| 0 | | | | | | 1 | 10 | <START T> | |
| 1 | read_item(X) | 1 | | 1 | | 1 | 10 | | |
| 2 | X := X*2 | 2 | | 1 | | 1 | 10 | | |
| 3 | write_item(X) | 2 | | 2 | | 1 | 10 | <T, X, 2> | |
| 4 | read_item(Y) | 2 | 10 | 2 | 10 | 1 | 10 | | |
| 5 | Y := Y*2 | 2 | 20 | 2 | 10 | 1 | 10 | | |
| 6 | write_item(Y) | 2 | 20 | 2 | 20 | 1 | 10 | <T, Y, 20> | |
| 7 | | | | | | | | <COMMIT T> | |
| 8 | flush_log | | | | | | | | |
| 9 | output(X) | 2 | 20 | 2 | 20 | 2 | 10 | | |
| 10 | output(Y) | 2 | 20 | 2 | 20 | 2 | 20 | | |

Redo Logging: Procedure

1. T first writes all log records for all updates to disk
2. T writes <COMMIT T> to the log on disk
3. T writes all updates to disk



Fundamental property of redo logs:

- <COMMIT T> occurs in log → log contains complete information on T
- <COMMIT T> doesn't occur in log → T hasn't written anything to disk

Recovery With Redo Logs

- Essentially: reverse of undo logging
- Procedure:
 - Identify all the transactions with **commit** log records.
 - Traverse the log from first to last.
 - If we see $\langle T, X, v \rangle$ and T has a **commit** log record, then change the value of X on disk to v.
 - For each *incomplete* transaction T, write **<ABORT T>** into the log on disk.

What could go wrong if we wrote to disk **before** flushing **<commit T>**?

Exercise with Redo Logging

| Time | Transaction T ₁ | Transaction T ₂ |
|------|----------------------------|----------------------------|
| 1 | read_item(X) | |
| 2 | X := X * 2 | |
| 3 | write_item(X) | |
| 4 | | read_item(X) |
| 5 | read_item(Y) | |
| 6 | | X := X * 3 |
| 7 | | write_item(X) |
| 8 | Y := X + Y | |
| 9 | write_item(Y) | |

X = 1
Y = 2

- How does redo logging work on this schedule?
 - Which **log entries** are written to buffer/disk & when?
 - Which **other operations** must be executed & when?

| Time | Transaction T ₁ | Transaction T ₂ | Log (buffer) | Log (disk) |
|------|----------------------------|----------------------------|--------------------------|------------|
| 0 | | | <START T ₁ > | |
| 1 | read_item(X) | | | |
| 2 | X := X * 2 | | | |
| 3 | write_item(X) | | <T ₁ , X, 2> | |
| 4 | | | <START T ₂ > | |
| 5 | | read_item(X) | | |
| 6 | read_item(Y) | | | |
| 7 | | X := X * 3 | | |
| 8 | | write_item(X) | <T ₂ , X, 6> | |
| 9 | Y := X + Y | | | |
| 10 | write_item(Y) | | <T ₁ , Y, 4> | |
| 11 | | | <COMMIT T ₁ > | |
| 12 | flush_log | | | |
| 13 | output(X) | | | |
| 14 | output(Y) | | | |
| 15 | | | <COMMIT T ₂ > | |
| 16 | | flush_log | | |
| 17 | | output(X) | | |
| 18 | | | | |
| 19 | | | | |

X = 1
Y = 2

Summary

- Aborted transactions & system failures can be dealt with using careful logging & restoring data using logs
- Undo logging:
 - Maintains atomicity
 - Supports undoing aborted transactions & recovers last consistent database state in case of system failure
- Other variants:
 - Redo logging: for Durability
 - Undo/Redo: Best variant
- Checkpoints speed up recovery
- Dirty reads may cause certain problems → problem and solutions tomorrow