

COMP207

Database Development

Lecture 9

Transaction Management: Reconciling Conflict-Serialisability & Recovery

- Nonquiescent checkpoints will not be part of the exam

Canceled lecture

- Tuesday 17-18 is canceled again
 - This time because of a networking event for second years in Sensor City that starts at 17 as well (no more slots though)
- Tuesday 13-14 is still going on though

Review of Undo & Redo Logging

- Logs activities with the goal of:
 - (For Undo): “undoing” to a previous database state.
 - (For Redo): “redoing” a database state that has been lost.
- Log records (or log entries):
 - **<START T>**: Transaction T has started.
 - **<COMMIT T>**: Transaction T has committed.
 - **<ABORT T>**: Transaction T was aborted.
 - **<T, X, v>**: Transaction T has updated the value of database item X, and the (**old** for **undo** and **new** for **redo**) value of X was v.
 - Response to **write_item(X)**
 - If this entry occurs in the log, then the new value of X might not have been written to the database yet

Undo/Redo Logging

- Good properties of undo logging and redo logging
- Log records:
 - Same as before, but replace $\langle T, X, v \rangle$
 - $\langle T, X, v, w \rangle$: “Transaction T has updated the value of database item X, and the **old/new** value of X is v/w.”
- Procedure:
 - Write all log records for all updates to database items first
 - Then write updates to disk
 - $\langle \text{COMMIT } T \rangle$ can be written to disk before or after all changes have been written to disk
- Recovery needs to process log in both directions

Review of Undo/Redo Logging

5
min

| Time | Transaction T ₁ | Transaction T ₂ |
|------|----------------------------|----------------------------|
| 1 | read_item(X) | |
| 2 | X := X * 2 | |
| 3 | write_item(X) | |
| 4 | | read_item(X) |
| 5 | read_item(Y) | |
| 6 | | X := X * 3 |
| 7 | | write_item(X) |
| 8 | Y := X + Y | |
| 9 | write_item(Y) | |

X = 1
Y = 2

- How does undo/redo logging work on this schedule?
 - Which **log entries** are written to buffer/disk & when?
 - Which **other operations** must be executed & when?

| Time | Transaction T ₁ | Transaction T ₂ | Local T ₁ | | Local T ₂ | | Buffer | | Disk | | Buffer log |
|------|----------------------------|----------------------------|----------------------|---|----------------------|---|--------|---|------|---|----------------------------|
| | | | X | Y | X | Y | X | Y | X | Y | |
| 0 | | | | | | | | | 1 | 2 | <START T ₁ > |
| 1 | read_item(X) | | 1 | | | | 1 | | 1 | 2 | |
| 2 | X := X * 2 | | 2 | | | | 1 | | 1 | 2 | |
| 3 | write_item(X) | | 2 | | | | 2 | | 1 | 2 | <T ₁ , X, 1, 2> |
| 4 | | | 2 | | | | 2 | | 1 | 2 | <START T ₂ > |
| 5 | | read_item(X) | 2 | | 2 | | 2 | | 1 | 2 | |
| 6 | read_item(Y) | | 2 | 2 | 2 | | 2 | 2 | 1 | 2 | |
| 7 | | X := X * 3 | 2 | 2 | 6 | | 2 | 2 | 1 | 2 | |
| 8 | | write_item(X) | 2 | 2 | 6 | | 6 | 2 | 1 | 2 | <T ₂ , X, 2, 6> |
| 9 | Y := X + Y | | 2 | 4 | 6 | | 6 | 2 | 1 | 2 | |
| 10 | write_item(Y) | | 2 | 4 | 6 | | 6 | 4 | 1 | 2 | <T ₁ , Y, 2, 4> |
| 11 | | | 2 | 4 | 6 | | 6 | 4 | 1 | 2 | <COMMIT T ₁ > |
| 12 | flush_log | | 2 | 4 | 6 | | 6 | 4 | 1 | 2 | |
| 13 | output(X) | | 2 | 4 | 6 | | 6 | 4 | 6 | 2 | |
| 14 | output(Y) | | 2 | 4 | 6 | | 6 | 4 | 6 | 4 | |
| 15 | | | 2 | 4 | 6 | | 6 | 4 | 6 | 4 | <COMMIT T ₂ > |
| 16 | | flush_log | 2 | 4 | 6 | | 6 | 4 | 6 | 4 | |
| 17 | | output(X) | 2 | 4 | 6 | | 6 | 4 | 6 | 4 | |

Why are DBMS using Undo/Redo?

Undo without Redo

- Undo essentially ensures Atomicity
- Can ensure durability using Force
 - **Force** the writing of updates to disk before commit
 - (**No Force** is not to require this)
 - Force is expensive in disk operations

Redo without Undo

- Redo essentially ensures Durability
- Can ensure atomicity using No Steal
 - **No Steal** means that uncommitted data may not overwrite committed data on disk
 - (**Steal** is not to require this)
 - **No Steal** is expensive to ensure

Ensuring A and D

- Could ensure Atomicity and Durability without log using No Steal/Force
 - Very hard and expensive to ensure
 - (Must commit **and** write every change to disk at once)
- In practice:
 - Want Steal/No Force (cheapest in time) → Use Undo/Redo

More Efficient Recovery via Checkpoints


Simple Checkpointing

- Idea: *checkpoint* the log periodically
 - Every m min., after t transactions since last checkpoint, ...
 - No need to undo transactions before t
- Procedure:
 1. Stop accepting new transactions
 2. Wait until all active transactions finish and have written COMMIT or ABORT record to the log.
 3. Flush the log to disk.
 4. Write a log record **<CHECKPOINT>**.
 5. Flush the log to disk again.
 6. Resume accepting transactions.

There are variants of checkpointing that avoid this See later!

Recovery via Simple Checkpoints

Recovery With Simple Checkpoints



| Time | Transaction T ₁ | Transaction T ₂ | Log (buffer) | Log (disk) |
|------|----------------------------|----------------------------|--------------------------|------------|
| 0 | | | <START T ₁ > | |
| 1 | read_item(X) | | | |
| 2 | X := X * 2 | | | |
| 3 | write_item(X) | | <T ₁ , X, 1> | |
| 4 | | | <START T ₂ > | |
| 5 | | read_item(X) | | |
| 6 | read_item(Y) | | | |
| 7 | | X := X * 3 | | |
| 8 | | write_item(X) | <T ₂ , X, 2> | |
| 9 | Y := X + Y | | | |
| 10 | write_item(Y) | | <T ₁ , Y, 2> | |
| 11 | flush_log | | | |
| 12 | output(X) | | | |
| 13 | output(Y) | | | |
| 14 | | | <COMMIT T ₁ > | |
| 15 | flush_log | | | |
| 16 | | flush_log | | |
| | | output(X) | | |
| | | | <COMMIT T ₂ > | |
| 19 | | flush_log | | |

Checkpointing starts

Transaction T₃ is submitted

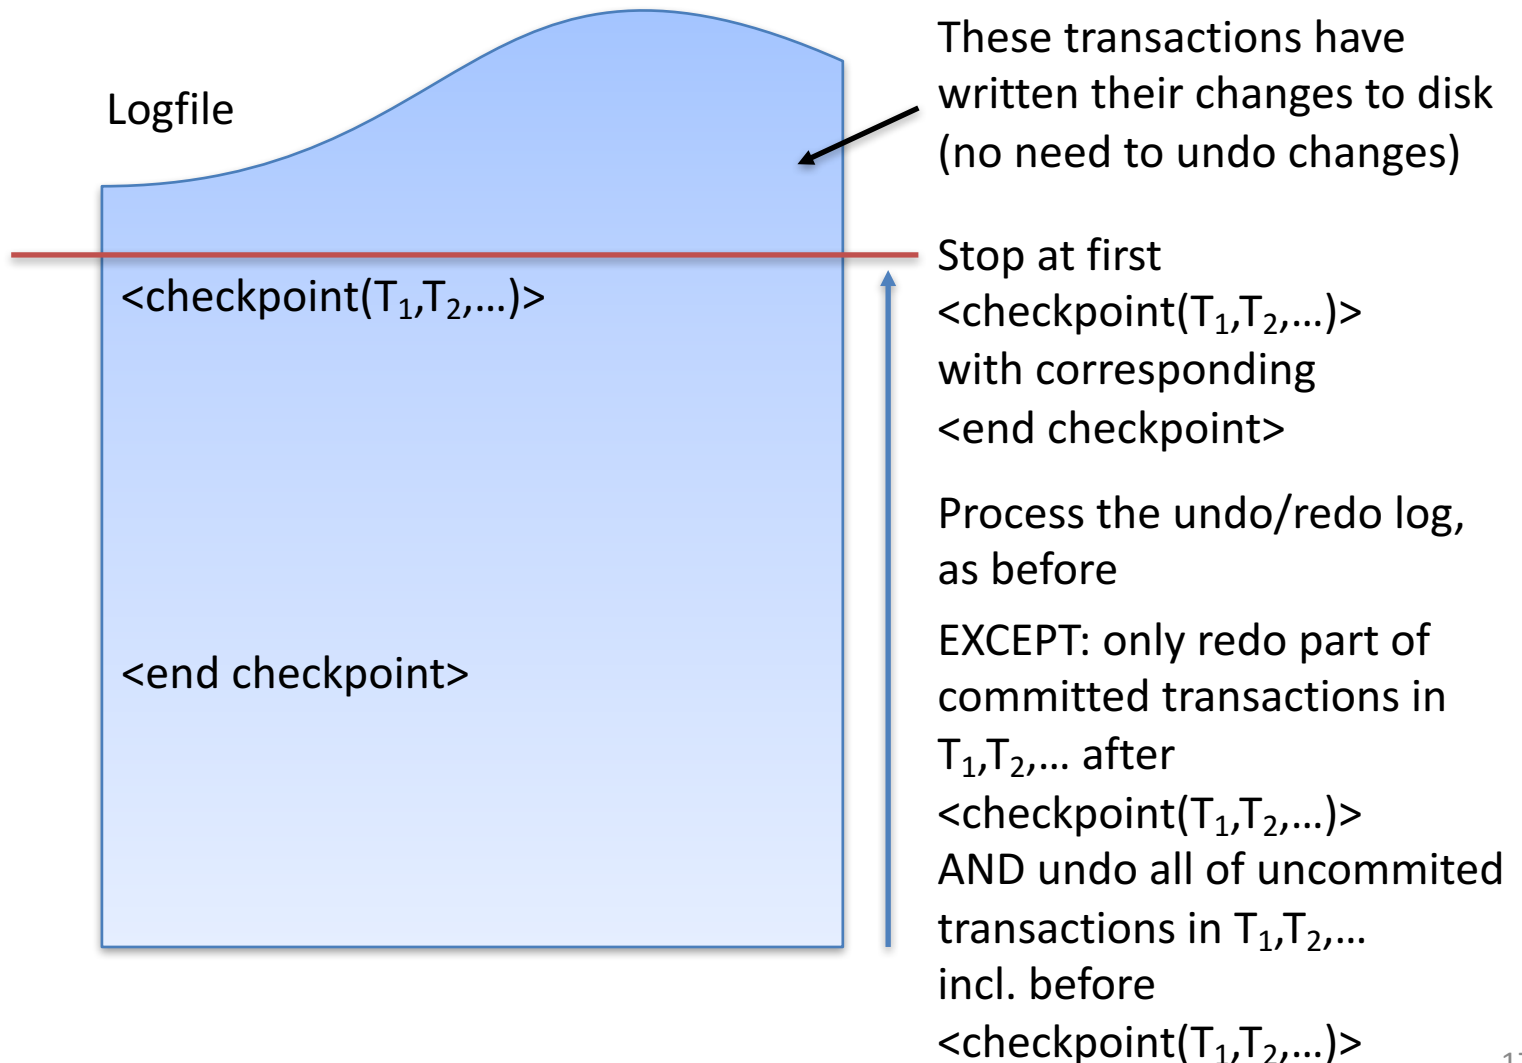
What will happen?

Better checkpoints!

Nonquiescent Checkpoints

- Requirements:
 - Undo/Redo logging
 - Transactions do not write to buffers(!) before they are sure they want to commit
- Procedure:
 - Write $\langle \text{Checkpoint}(T_1, T_2, \dots) \rangle$ in log and flush it
 - » T_1, T_2, \dots are the transaction in progress (i.e. not committed and not aborted)
 - Write the content of the **changed** buffer to disk (i.e. **output**)
 - Write $\langle \text{End Checkpoint} \rangle$ in log and flush it

Recovery via Nonquiescent Checkpoints



Nonquiescent advantages

Nonquiescent checkpoints...

- does not require delaying transactions
- can be forcefully finished

Dirty reads

“Dirty Reads”

Read something written by an uncommitted transaction

- In practice, the isolation property is often not fully enforced (→ “dirty reads” may occur)
- Reason: efficiency!
 - Spend less time on preventing “dirty reads”
 - Gain “more parallelism” by executing some transactions that would have to wait to prevent “dirty reads”

- You can decide:

Other option:
READ ONLY

Other levels in SQL:
READ COMMITTED,
REPEATABLE READ,
SERIALIZABLE

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

- “Dirty reads” can slow down the system when transactions have to abort

Dirty Reads and Rollbacks

| Time | Transaction T ₁ | Transaction T ₂ | X | Y |
|------|----------------------------|----------------------------|-----|---|
| 0 | lock(X) | | 1 | 2 |
| 1 | read_item(X) | | | |
| 2 | X := X + 100 | | 101 | |
| 3 | write_item(X) | | | |
| 4 | lock(Y) | | | |
| 5 | unlock(X) | | | |
| 6 | | lock(X) | | |
| 7 | | read_item(X) | | |
| 8 | | X := X * 2 | 202 | |
| 9 | | write_item(X) | | |
| 10 | | lock(Y) | | |
| 11 | read_item(Y) | | | |
| 12 | abort | | | |
| 13 | | lock(Y) | | |
| 14 | | ... | | |
| 15 | | commit | | |

denied

scheduler automatically unlocks Y

T₂ depends on “dirty data” from T₁
→ must abort T₂, too!

Cascading Rollback

If a transaction T aborts:

- Find all transactions that have read items written by T.
- Recursively abort all transactions that have read items written by an aborted transaction.

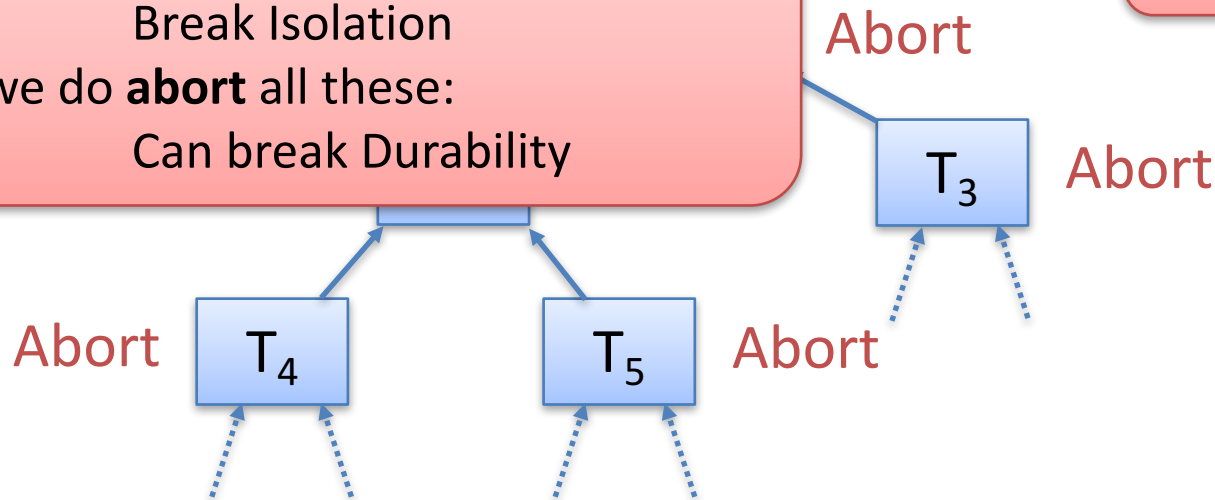
If we do **not abort** all these:

Break Isolation

If we do **abort** all these:

Can break Durability

Very slow →
want to avoid this



Isolation vs Durability

| Time | Transaction T ₁ | Transaction T ₂ | X | Y |
|------|----------------------------|----------------------------|-----|---|
| 0 | lock(X) | | 1 | 2 |
| 1 | read_item(X) | | | |
| 2 | X := X + 100 | | 101 | |
| 3 | write_item(X) | | | |
| 4 | lock(Y) | | | |
| 5 | unlock(X) | | | |
| | | lock(X) | | |
| | | read_item(X) | | |
| | | X := X * 2 | 202 | |
| | | write_item(X) | | |
| | | commit | | |
| 11 | read_item(Y) | | | |
| 12 | abort | | | |

If we do **not abort** T₂:
Break Isolation

If we do **abort** T₂:
Break Durability

Conflict-Serialisability vs Recovery

Conflict-Serialisability

- Many nice properties:
 - Equivalent to serial schedules
 - Ensure consistency / correctness
- Can be enforced by two-phase locking (2PL)


Logging and Recovery

- Suitable logging techniques ensure that we can restore desired database states
 - Undo incomplete transactions
 - Redo committed transactions
 - Undo a single or a selected number of transactions
- Robust: works even after system failures

Problem: cascading rollbacks may be necessary!

Recoverable Schedules

- The problem for Durability in regards to cascading rollbacks occur because a transaction T_1 reads data from some transaction T_2 , then T_1 commits and afterwards T_2 aborts.
- A schedule S is **recoverable** if the following is true:
 - if a transaction T_1 commits and has read an item X that was written before by a different transaction T_2 , ...
 - then **T_2 must commit before T_1 commits.**




Can still do cascading rollbacks, but only active transactions can be forced to abort

Example


- A **recoverable** schedule:

$S_1: w_2(X); w_1(Y); w_1(X); r_2(Y); w_2(Y); c_1; c_2$

T_2 reads data
that was written
before by T_1




T_1 must commit
before T_2 can commit



- A **non-recoverable** schedule:

$S_2: w_1(X); w_1(Y); w_2(X); r_2(Y); w_2(Y); c_2; c_1$

T_2 reads data
that was written
before by T_1



But: T_2 commits first



- Note:
 - S_1 is *not* serialisable.
 - S_2 is serialisable.

Summary

- Undo/Redo logging and why it is used
- Reconciliation of conflict-serialisability and recovery
 - Can lead to problems (**cascading rollbacks**) if done naively
 - Avoiding cascading rollbacks requires a smarter way of scheduling transactions
- Ideas:
 - **Recoverable schedules**: T commits only if all transactions that T has read from have committed