

COMP207

Database Development

Lecture 10

Transaction Management:
More on Combining Concurrency and Recovery
and how to Prevent Deadlocks

Conflict-Serialisability vs Recovery

Conflict-Serialisability


- Many nice properties:
 - Equivalent to serial schedules
 - Ensure consistency / correctness
- Can be enforced by two-phase locking (2PL)

Logging and Recovery

- Suitable logging techniques ensure that we can restore desired database states
 - Undo incomplete transactions
 - Redo committed transactions
 - Undo a single or a selected number of transactions
- Robust: works even after system failures

Recoverable Schedules

- The problem for Durability in regards to cascading rollbacks occur because a transaction T_1 reads data from some transaction T_2 , then T_1 commits and afterwards T_2 aborts.
- A schedule S is **recoverable** if the following is true:
 - if a transaction T_1 commits and has read an item X that was written before by a different transaction T_2 , ...
 - then **T_2 must commit before T_1 commits.**



Can still do cascading rollbacks, but only active transactions can be forced to abort

Recoverable Schedules

- Additional implicit requirement:
All log records have to reach disk in the order in which they are written.
- Compare:
 - Recoverable: $S_1: w_2(X); w_1(Y); w_1(X); r_2(Y); w_2(Y); c_1; c_2$
 - Not recoverable: $S_3: w_2(X); w_1(Y); w_1(X); r_2(Y); w_2(Y); c_2; c_1$

If in S_1 the commit record for T_2 would reach disk earlier than the commit record for T_1 , then T_1 could in principle abort → cascading rollback

Cascading Rollbacks Again

- A recoverable schedule:

$S_4: w_1(X); w_1(Y); w_2(X); r_2(Y); w_2(Y); c_1; c_2$




Suppose T_1 needs to be rolled back here

- T_1 rolls back $\rightarrow T_2$ has to be rolled back

Cascadeless Schedules

- A schedule is **cascadeless** if each transaction in it reads only values that were written by transactions that have already committed.

No reading of “dirty data”.
No cascading rollbacks.



Cascadeless Schedules

- A schedule is **cascadeless** if each transaction in it reads only values that were written by transactions that have already committed.

No reading of “dirty data”.
No cascading rollbacks.



- As for recoverable schedules:
Log records have to reach disk in the right order.

Example

- The schedules S_1 - S_4 on the previous slides are **not cascadeless**:

S_1 : $w_2(X)$; $w_1(Y)$; $w_1(X)$; $r_2(Y)$; $w_2(Y)$; c_1 ; c_2

reads uncommitted data from T_1

S_2 : $w_1(X)$; $w_1(Y)$; $w_2(X)$; $r_2(Y)$; $w_2(Y)$; c_2 ; c_1

S_3 : $w_2(X)$; $w_1(Y)$; $w_1(X)$; $r_2(Y)$; $w_2(Y)$; c_2 ; c_1

S_4 : $w_1(X)$; $w_1(Y)$; $w_2(X)$; $r_2(Y)$; $w_2(Y)$; c_1 ; c_2

- This variant of S_1 is **cascadeless**:

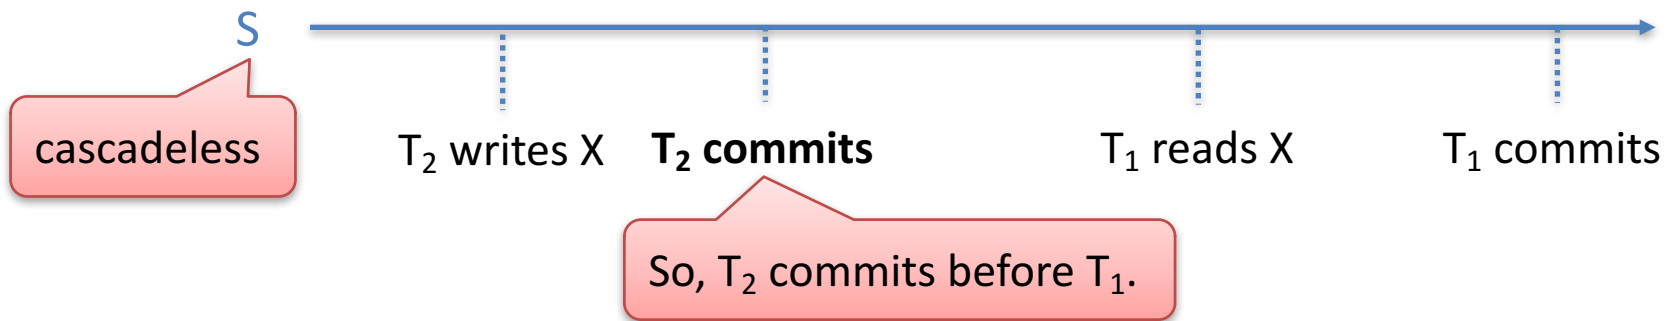
S_5 : $w_2(X)$; $w_1(Y)$; $w_1(X)$; c_1 ; $r_2(Y)$; $w_2(Y)$; c_2

reads committed data from T_1

- Note: S_5 is **not serialisable**.

Cascadeless Schedules: Properties

- Cascadeless schedules are **recoverable**:

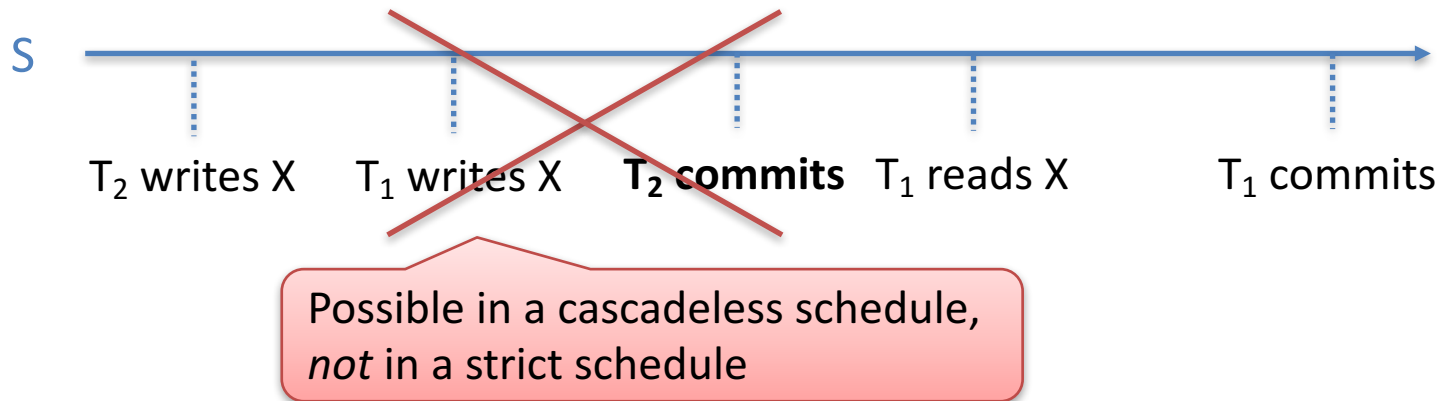


- Cascadeless schedules are in general **not serialisable**.
(recall example on previous slide)

Can We Have Both?
No Cascading Aborts & Serialisability?

Strict Schedules

- A schedule is **strict** if each transaction in it reads and writes only values that were written by transactions that have already committed.



Strict Schedules

- A schedule is **strict** if each transaction in it reads and writes only values that were written by transactions that have already committed.



- Of course, log records have to reach disk in order.

Strict Two-Phase Locking (Strict 2PL)

- Most popular variant of two-phase locking (2PL)
- Enforces both:
 - Conflict-serialisability
 - Strict schedules
- **Strict locking condition**
(in addition to 2PL condition):
 - with simple locking: **any lock**
 - with shared/exclusive locks: **just exclusive locks**

A transaction T **must not release any lock (that allows T to write data)** until:

- T has committed or aborted, and
- the commit/abort log record has been written to disk.

Example 1

2PL transaction

Not strict 2PL

Transaction T

lock(X)

read_item(X)

$X := X + 100$

write_item(X)

lock(Y)

unlock(X)

read_item(Y)

$Y := Y + 100$

write_item(Y)

unlock(Y)

commit

For undo logging, we assume that **commit...**

1. Writes all log records to disk
2. Writes all modified database items to disk
3. Writes the commit record to disk

Example 2

2PL transaction

Transaction T
lock(X)
read_item(X)
$X := X + 100$
write_item(X)
lock(Y)
unlock(X)
read_item(Y)
$Y := Y + 100$
write_item(Y)
unlock(Y)
commit

Strict 2PL transaction

New transaction T'
lock(X)
read_item(X)
$X := X + 100$
write_item(X)
lock(Y)
read_item(Y)
$Y := Y + 100$
write_item(Y)
commit
unlock(X)
unlock(Y)

Locks released only after fully committed,
and all log records written to disk

Example With Shared/Exclusive Locks

Transaction T

s-lock(X)

read_item(X)

x-lock(Y)

unlock(X)

read_item(Y)

$Y := X + Y$

write_item(Y)

commit

unlock(Y)

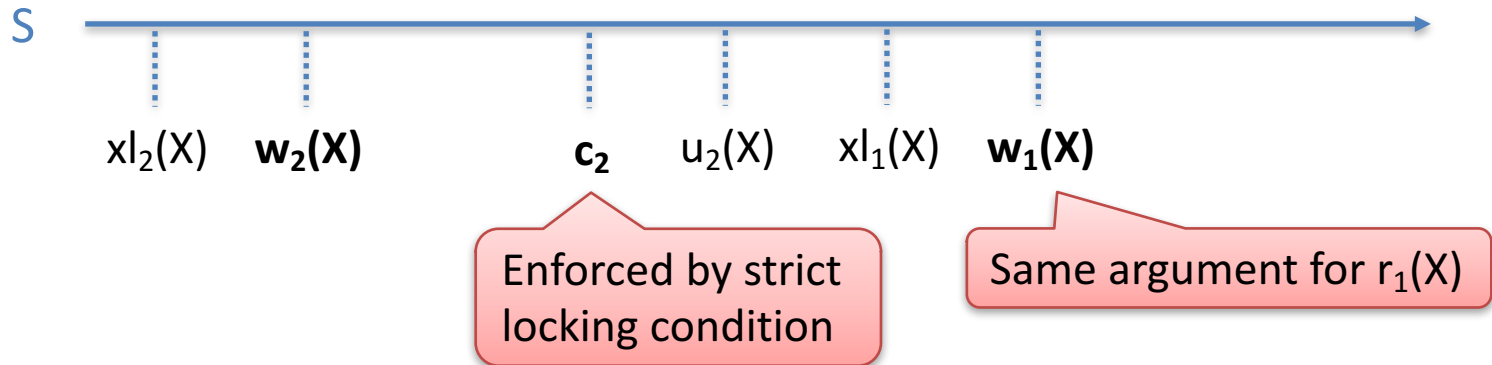
Strict 2PL

T can release the shared lock on X here (shared locks do not allow T to write data)

T is allowed to release the exclusive lock on Y only here.

Strict Two-Phase Locking *Enforces* Conflict-Serialisable & Strict Schedules

- If S is a schedule consisting of strict 2PL transactions:
 - S is conflict-serialisable.
 - S is strict.
- Strictness:



Still... Risk of Deadlocks

T ₁
lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
commit
unlock(X)
unlock(Y)

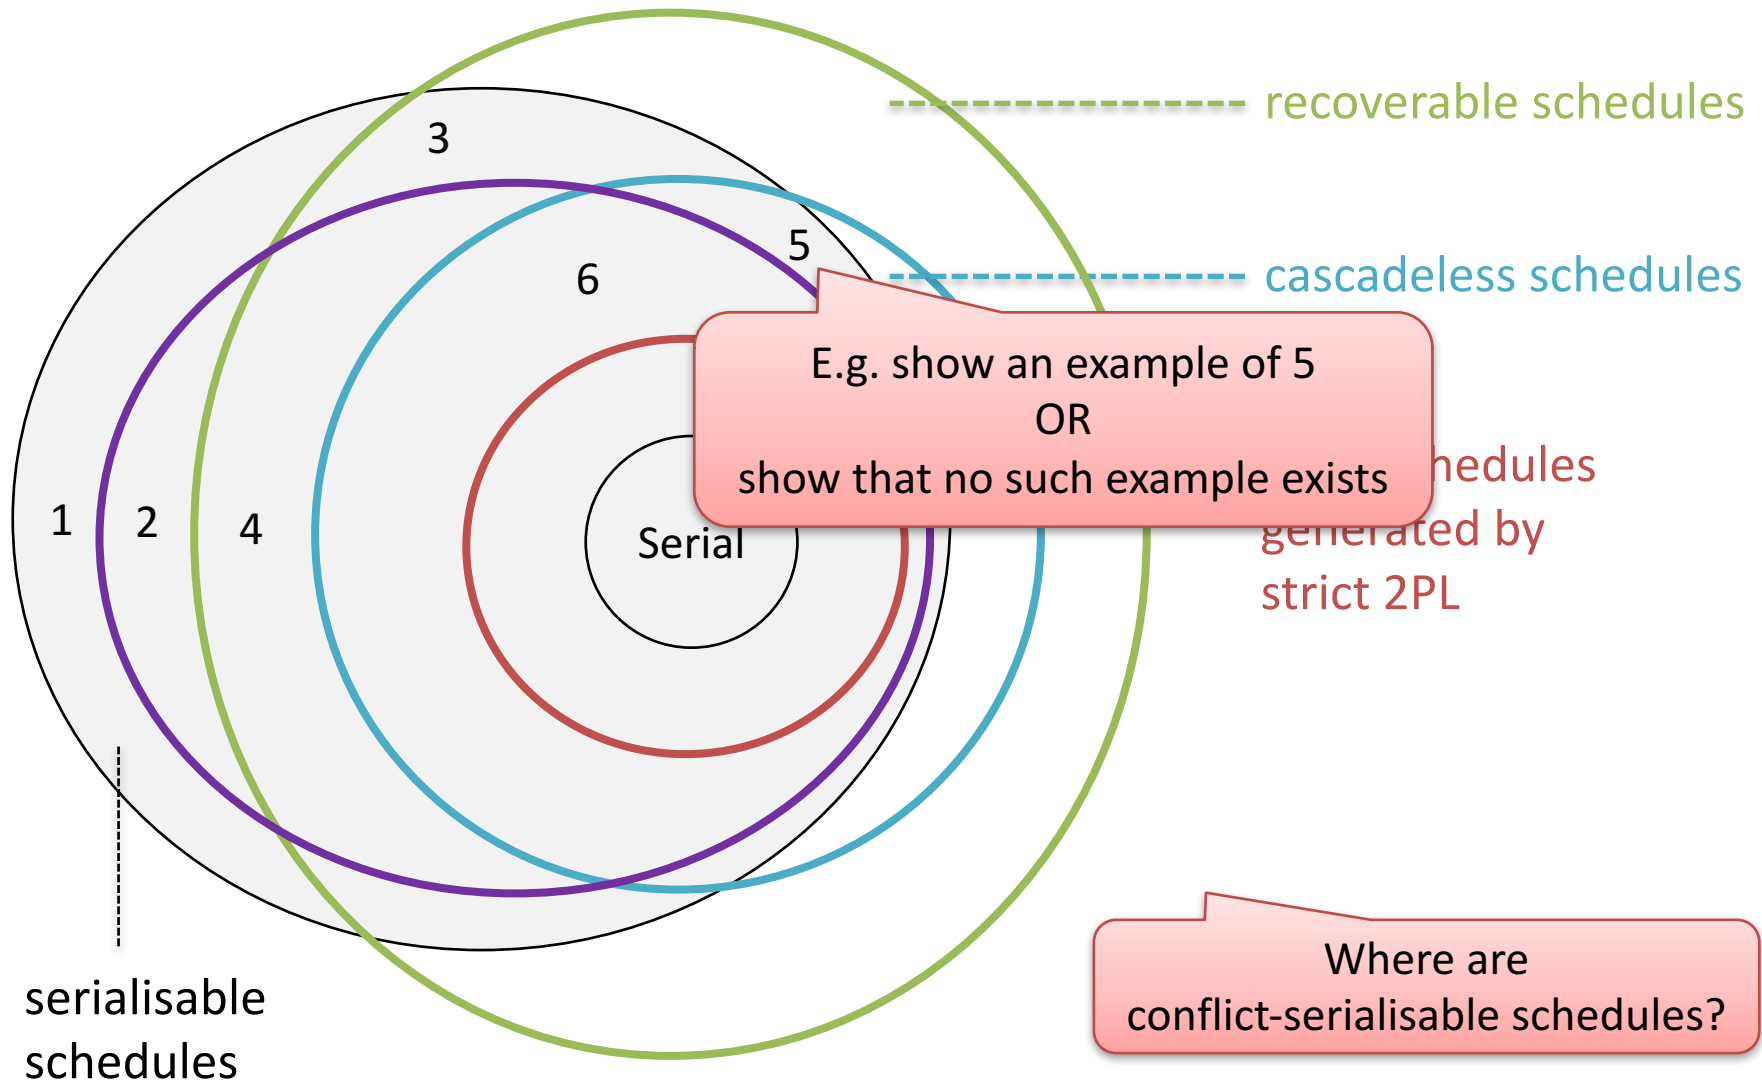
T ₂
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
lock(X)
read_item(X)
X := X + 100
write_item(X)
commit
unlock(X)
unlock(Y)

$l_1(X); r_1(X); w_1(X); l_2(Y); r_2(Y); w_2(Y); \underline{\quad?}$

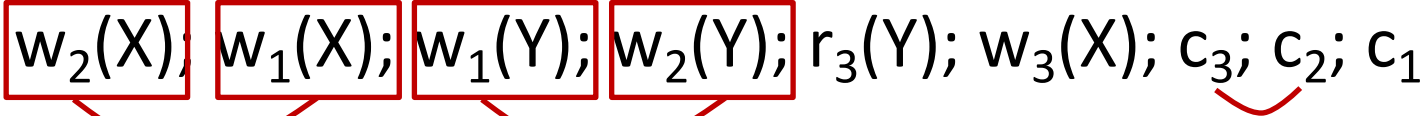
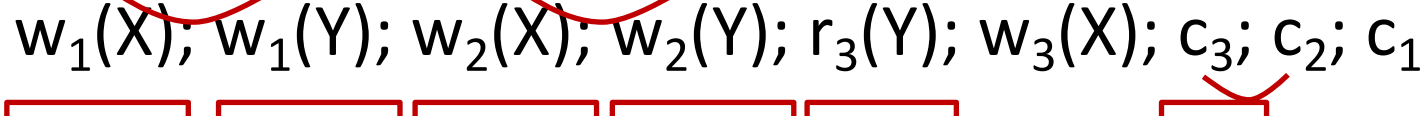
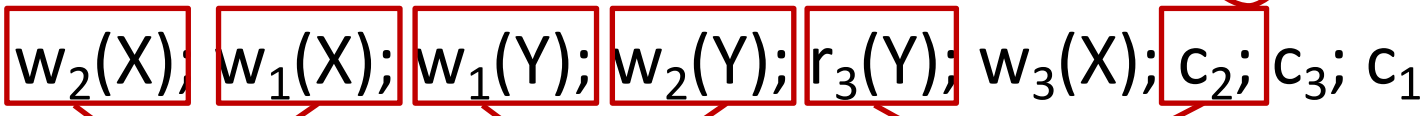
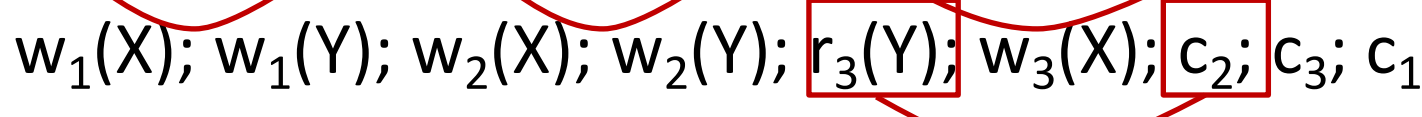
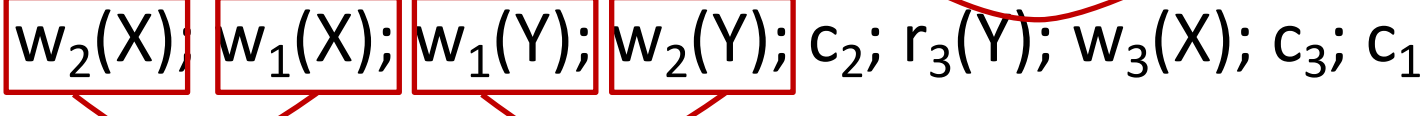
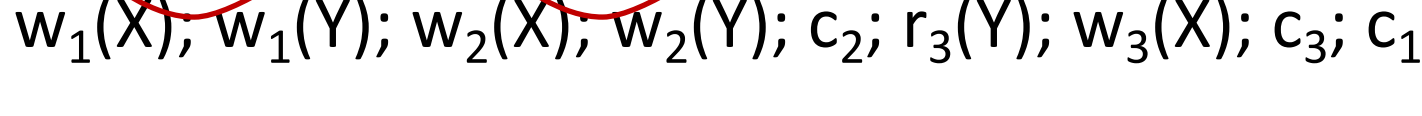
T₂'s request for lock on X denied

T₁'s request for lock on Y denied

How the Types of Schedules are Related



Example or proof

1. Example: $w_2(X); w_1(X); w_1(Y); w_2(Y); r_3(Y); w_3(X); c_3; c_2; c_1$ 
2. Example: $w_1(X); w_1(Y); w_2(X); w_2(Y); r_3(Y); w_3(X); c_3; c_2; c_1$ 
3. Example: $w_2(X); w_1(X); w_1(Y); w_2(Y); r_3(Y); w_3(X); c_2; c_3; c_1$ 
4. Example: $w_1(X); w_1(Y); w_2(X); w_2(Y); r_3(Y); w_3(X); c_2; c_3; c_1$ 
5. Example: $w_2(X); w_1(X); w_1(Y); w_2(Y); c_2; r_3(Y); w_3(X); c_3; c_1$ 
6. Example: $w_1(X); w_1(Y); w_2(X); w_2(Y); c_2; r_3(Y); w_3(X); c_3; c_1$ 

Strict 2PL and Deadlocks

- Strict 2PL yields **conflict-serialisable, strict schedules**
- Problem: **deadlocks**

T ₁
lock(X)
read_item(X)
X := X + 100
write_item(X)
lock(Y)
...

T ₂
lock(Y)
read_item(Y)
Y := Y + 100
write_item(Y)
lock(X)
...

$l_1(X); r_1(X); w_1(X);$
 $l_2(Y); r_2(Y); w_2(Y); \underline{\quad ? \quad}$

Roll back (and restart)
one of the transactions

- Two approaches for deadlock prevention:
 - Detect deadlocks & fix them
 - Enforce deadlock-free schedules


Not based on (strict) 2PL

Deadlock Detection

Deadlock Detection: Approaches


- **Timeouts**

- Assume a transaction is in a deadlock if it exceeds a given time limit



- **Wait-for graphs**

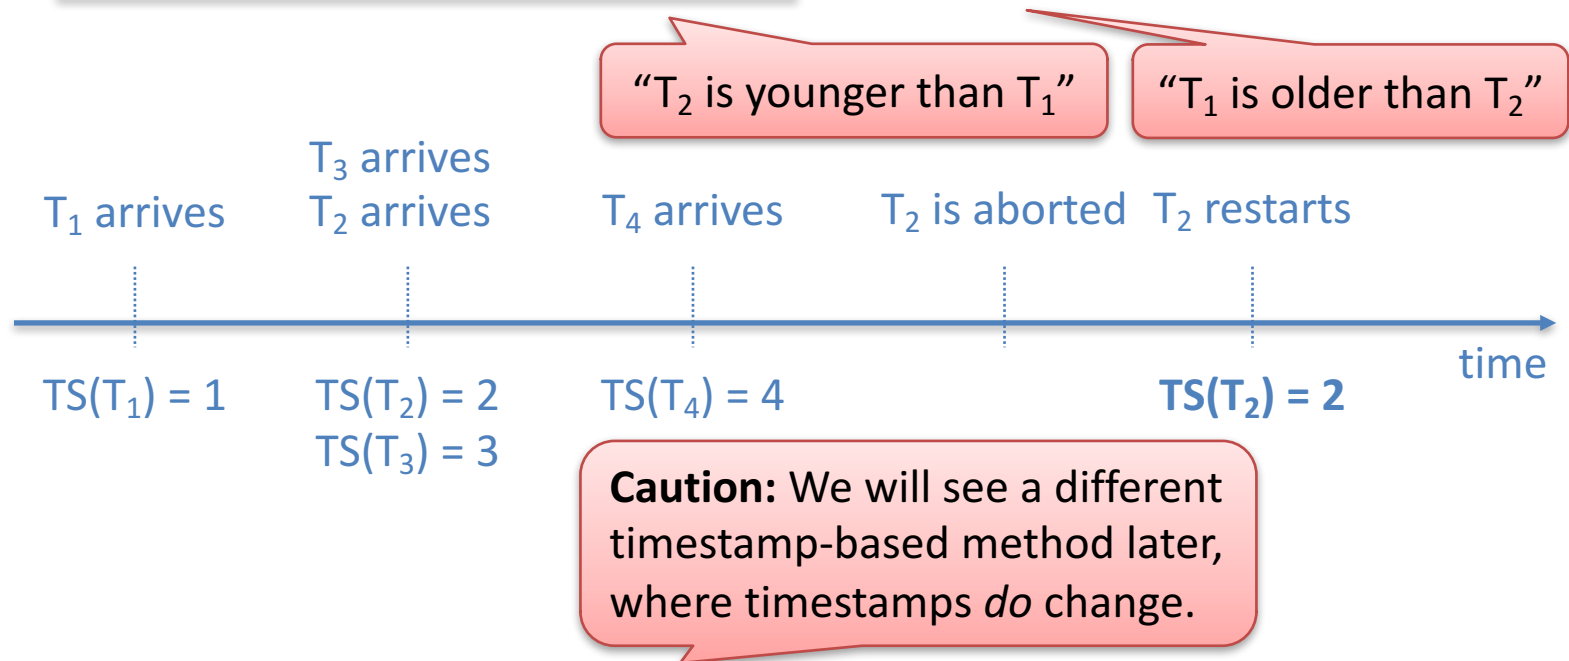
- Nodes: transactions
- Edge from T_1 to T_2 if T_1 waits for T_2 to release a lock
- Deadlocks correspond to cycles



- **Timestamp-based**

Timestamps for Deadlock Detection

- Each transaction T is assigned a unique integer $TS(T)$ upon arrival (the **timestamp of T**).
- If T_1 arrived earlier than T_2 , we require $TS(T_1) < TS(T_2)$



- Timestamps do not change** even after a restart!

How Are Timestamps Used?

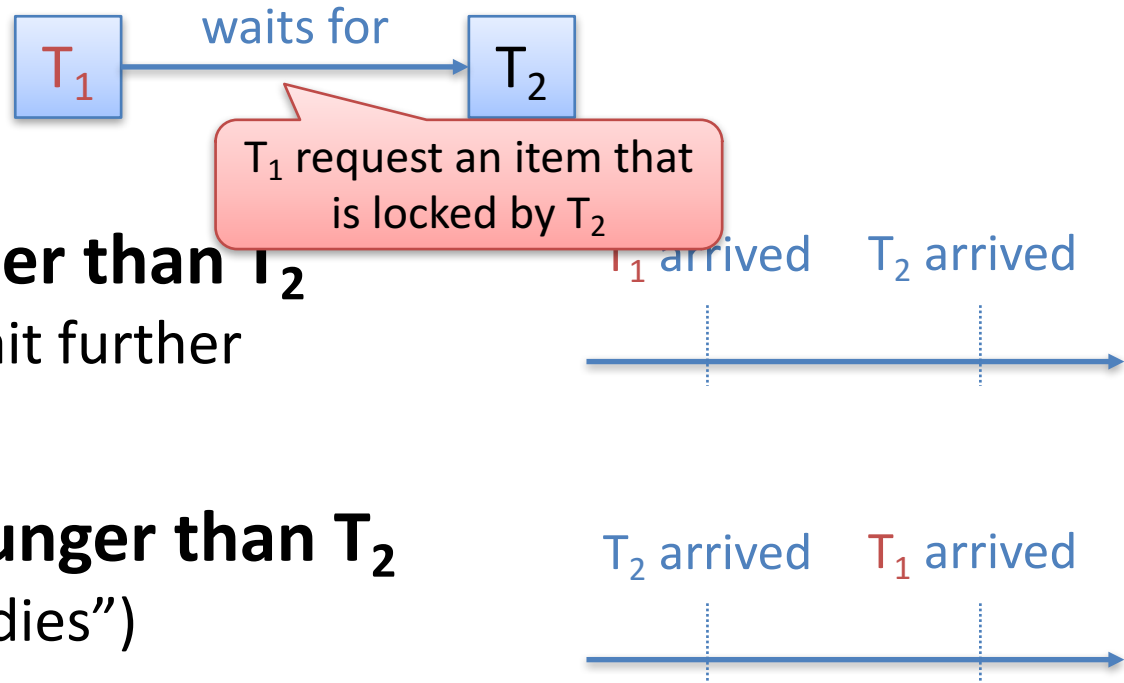
- Want to prevent cyclic dependencies such as
 - T_1 holds a lock on X_1 and **waits for a lock on X_2**
 - T_2 holds a lock on X_2 and **waits for a lock on X_3**
 - ...
 - T_n holds a lock on X_n and **waits for a lock on X_1**



- Use timestamps to decide which transaction can wait further and which must abort to prevent deadlock

Wait-Die Scheme

(“older transactions *always* wait for unlocks”)



- **Case 1: T_1 is older than T_2**

T_1 is allowed to wait further for T_2 to unlock

- **Case 2: T_1 is younger than T_2**

T_1 is rolled back (“dies”)

- Note: only older transactions are allowed to wait, so no cyclic dependencies created


Wound-Wait Scheme

(“older transactions *never* wait for unlocks”)




- **Case 1: T_1 is older than T_2**

T_2 is rolled back unless it has finished (it is “wounded”)



- **Case 2: T_1 is younger than T_2**

T_1 is allowed to wait further for T_2 to unlock



- Note: only younger transactions are allowed to wait, so no cyclic dependencies created

To be continued...