

COMP207

Database Development

Lecture 11

Transaction Management:
Timestamp-based Scheduling &
Concluding Remarks


Deadlock prevention

- Two approaches for deadlock prevention:
 - **Detect deadlocks & fix them**
 - **Enforce deadlock-free schedules**

Deadlock Detection: Approaches


- **Timeouts**

- Assume a transaction is in a deadlock if it exceeds a given time limit



- **Wait-for graphs**

- Nodes: transactions
- Edge from T_1 to T_2 if T_1 waits for T_2 to release a lock
- Deadlocks correspond to cycles



- **Timestamp-based**

Timestamps for Deadlock Detection


- Each transaction T is assigned a unique integer **$TS(T)$** upon **arrival** (the **timestamp of T**), such that later arrival means higher number
- **Timestamps for deadlock detection do not change** even after a restart!
- Two schemes, each based on some times doing rollback when a transaction T_1 requests a lock hold by another transaction T_2 :
 - **Wait-Die** – T_1 is rolled back, if it is younger than T_2
 - **Wound-Wait** – T_2 is rolled back, if it is younger than T_1

Why Wound-Wait Works

Eventually, any finite number of transactions finishes under Wound-Wait

- At all times, the oldest transaction can move
- Hence, eventually it finishes and there is one less transaction left and we are still doing Wound-Wait!
- Wait-Die is similar, but we look at the oldest transaction or the transaction it (recursively) waits for

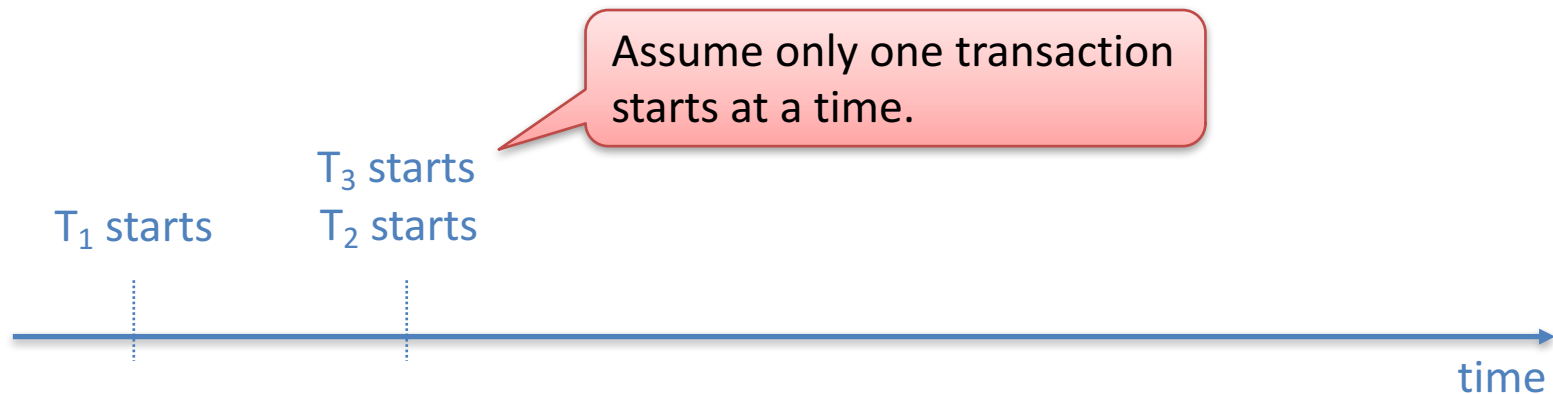
Deadlock prevention

- Two approaches for deadlock prevention:
 - **Detect deadlocks & fix them** 
 - **Enforce deadlock-free schedules**

Uses timestamps too!

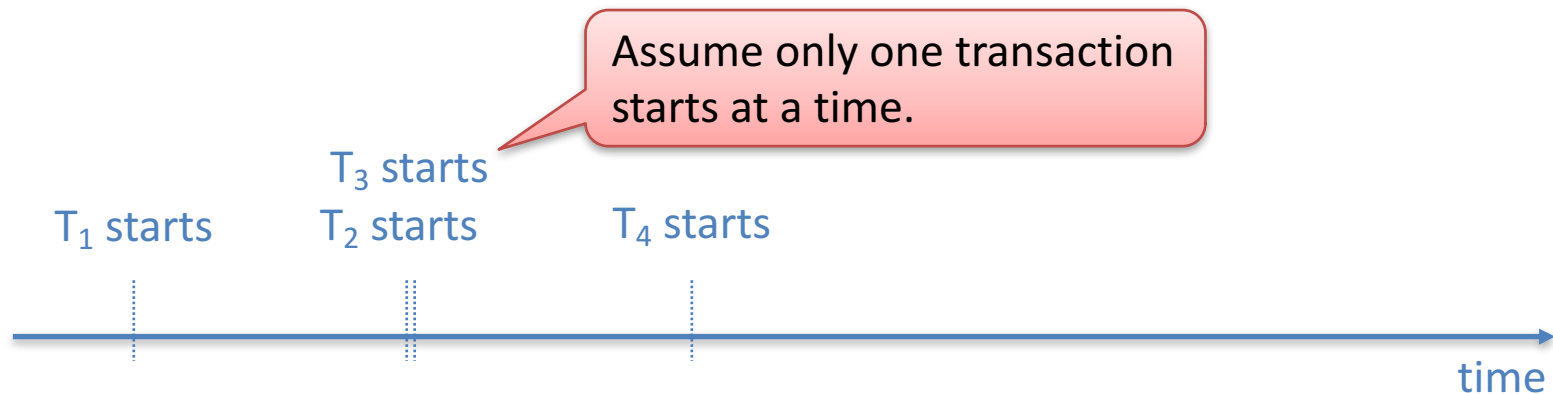
Basic Idea

- Schedule transactions so that the effect is the same as **executing each transaction instantaneously** when it is started.



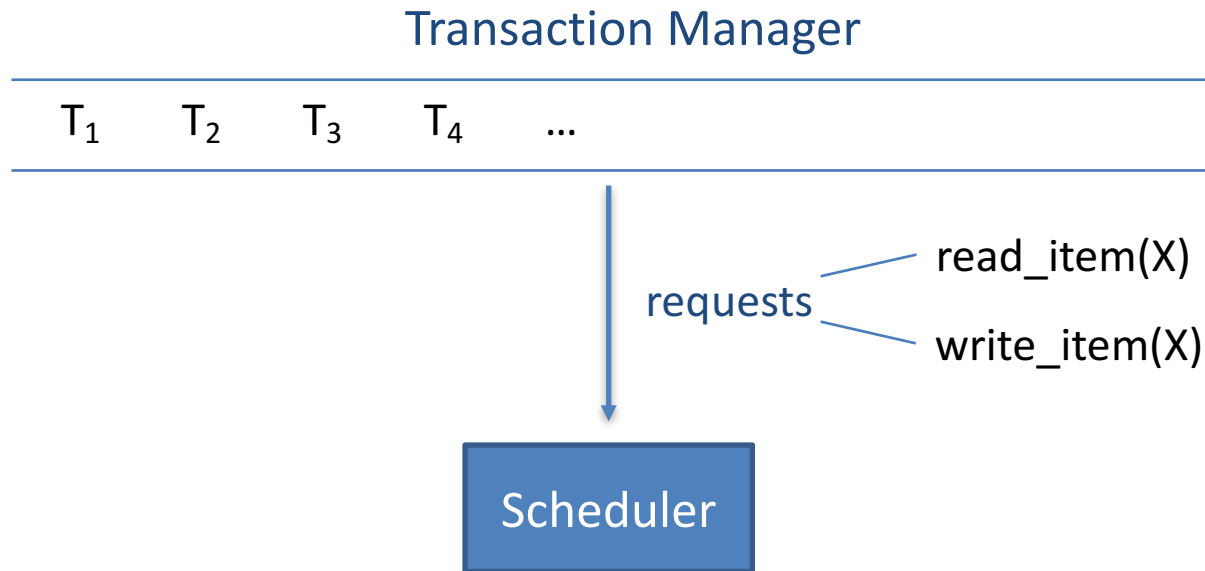
Basic Idea

- Schedule transactions so that the effect is the same as **executing each transaction instantaneously** when it is started.



- Equivalent to serial schedule that has all transactions in the order of their start time.

Timestamp-Based Schedulers



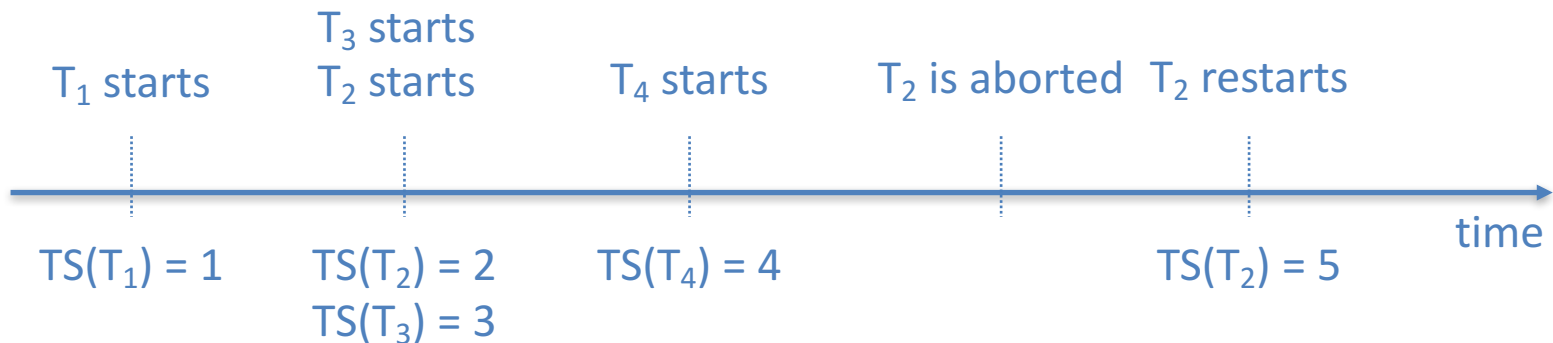
Possible actions of the scheduler:

- Grant request
- Abort transaction
- Delay transaction

Timestamps

- Each transaction **T** is assigned a unique integer **TS(T)** when it starts (the **timestamp of T**).
- If T_1 started earlier than T_2 , we require **$TS(T_1) < TS(T_2)$**

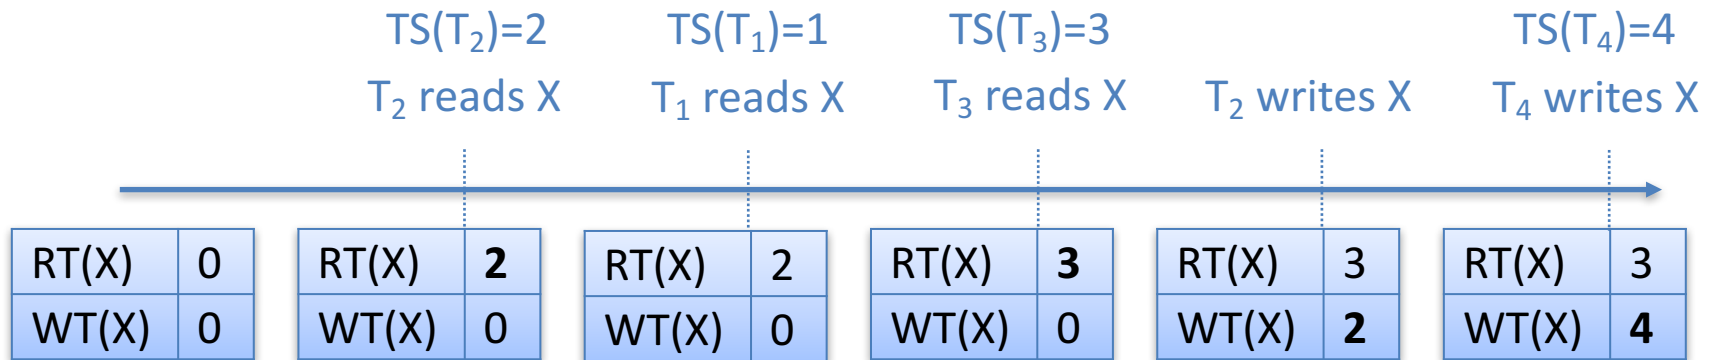
" T_2 is younger than T_1 "



- We assign a **new timestamp even after a restart!**

Additional Bookkeeping

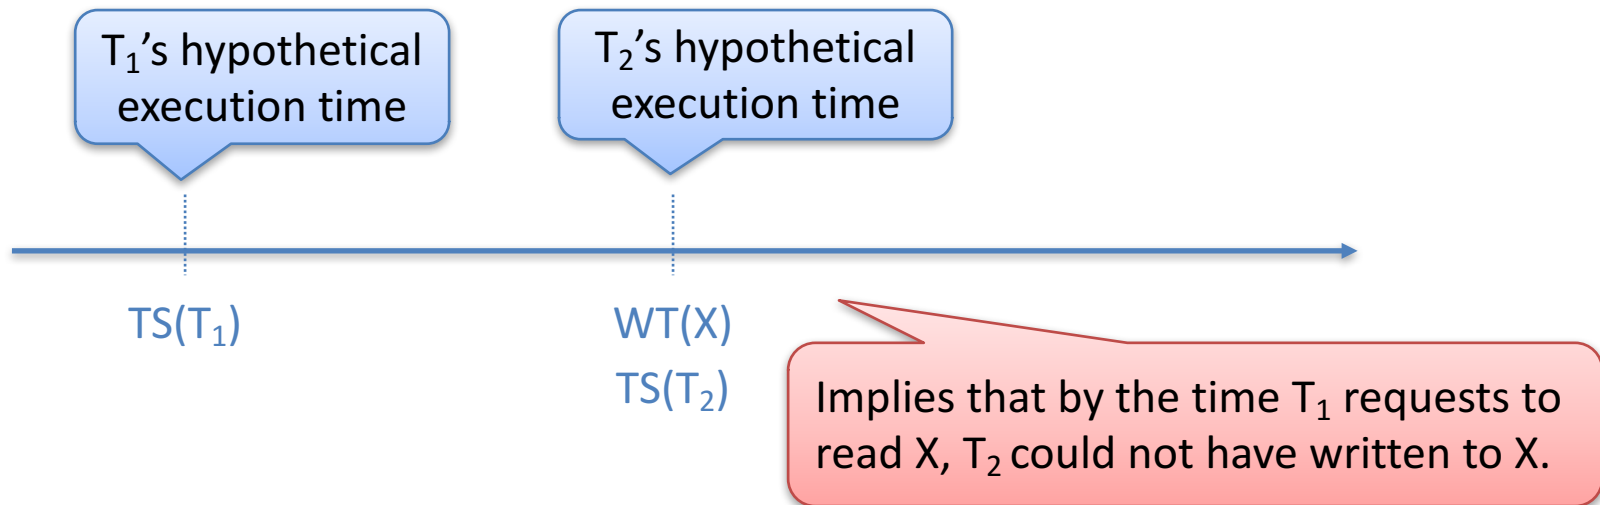
- For each database item X , maintain:
 - Read Time of X : $RT(X)$**
Timestamp of youngest transaction that read X
 - Write Time of X : $WT(X)$**
Timestamp of youngest transaction that wrote X



Read Requests

If T_1 requests to **read** X:

- **Abort & restart T_1** if $WT(X) > TS(T_1)$

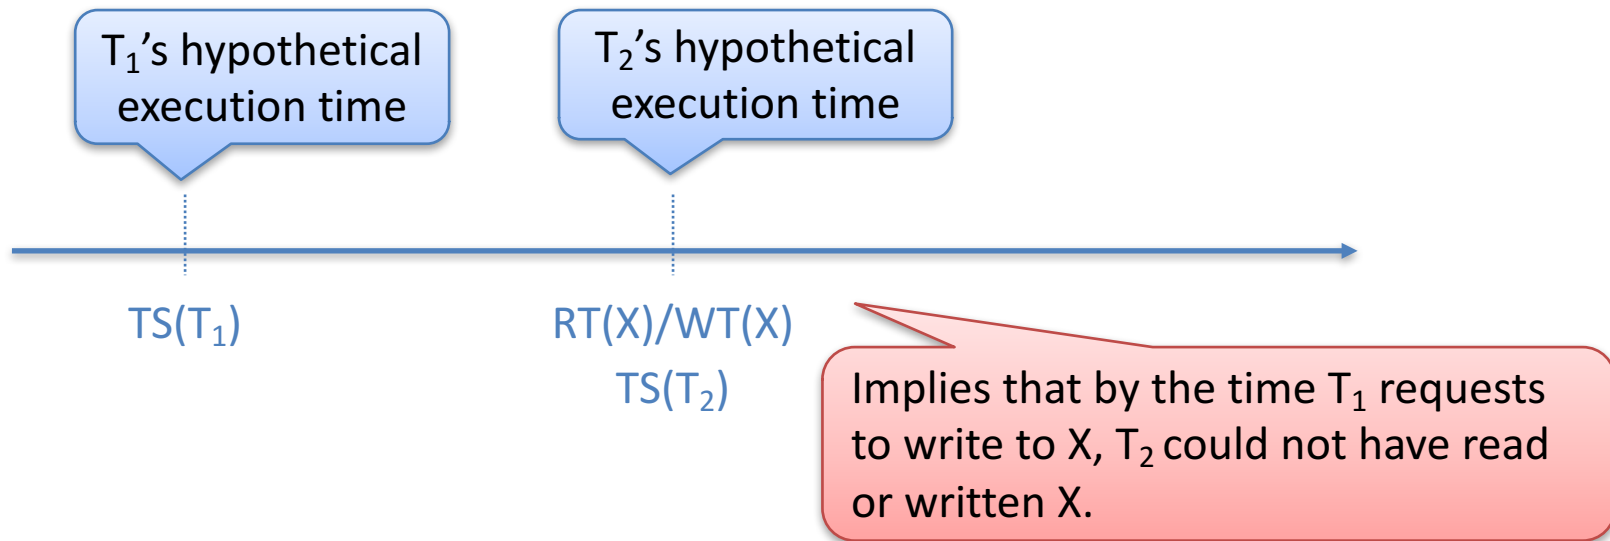


- **Grant request** otherwise

Write Requests

If T_1 requests to **write** X :

- **Abort & restart T_1** if $RT(X) > TS(T_1)$ or $WT(X) > TS(T_1)$



- **Grant request** otherwise

Example 1

| Time | T_1 (TS = 1) | T_2 (TS = 2) | X | | Y | |
|------|----------------|----------------|----|----|----|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | | | | | | |
| 2 | | | | | | |
| 3 | | | | | | |
| 4 | | | | | | |
| 5 | | | | | | |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |

Example 1

| Time | T_1 (TS = 1) | T_2 (TS = 2) | X | | Y | |
|------|----------------|----------------|----|----|----|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | | | | | | |
| 3 | | | | | | |
| 4 | | | | | | |
| 5 | | | | | | |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |


X last
read



Example 1


| Time | T_1 (TS = 1) | T_2 (TS = 2) | X | | Y | |
|------|----------------|----------------|----|----|----|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | $X := X + 100$ | | 1 | 0 | 0 | 0 |
| 3 | | | | | | |
| 4 | | | | | | |
| 5 | | | | | | |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |

X last
read




Example 1

| Time | T_1 (TS = 1) | T_2 (TS = 2) | X | | Y | |
|------|----------------|----------------|----------|----|----------|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | $X := X + 100$ | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | | | | | |
| 5 | | | | | | |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |




Example 1

| Time | T ₁ (TS = 1) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | | | | | |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |



Example 1


| Time | T ₁ (TS = 1) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |



Example 1

| Time | T ₁ (TS = 1) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | | | | |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |


Must abort
& restart T₁



Example 1

New timestamp


| Time | T_1 (TS = 3) | T_2 (TS = 2) | X | | Y | |
|------|----------------|----------------|----|----|----|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | $X := X + 100$ | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | $Y := Y * 2$ | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | | | | | | |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |



Example 1


New timestamp

| Time | T_1 (TS = 3) | T_2 (TS = 2) | X | | Y | |
|------|-------------------------|----------------|----|----|----|----|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | $X := X + 100$ | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | $Y := Y * 2$ | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | read_item(X) | | 3 | 0 | 2 | 2 |
| 8 | | | | | | |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |




Example 1

| Time | T ₁ (TS = 3) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | read_item(X) | | 3 | 0 | 2 | 2 |
| 8 | X := X + 100 | | 3 | 0 | 2 | 2 |
| 9 | | | | | | |
| 10 | | | | | | |
| 11 | | | | | | |




Example 1

| Time | T ₁ (TS = 3) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | read_item(X) | | 3 | 0 | 2 | 2 |
| 8 | X := X + 100 | | 3 | 0 | 2 | 2 |
| 9 | read_item(Y) | | 3 | 0 | 3 | 2 |
| 10 | | | | | | |
| 11 | | | | | | |




Example 1

| Time | T ₁ (TS = 3) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | read_item(X) | | 3 | 0 | 2 | 2 |
| 8 | X := X + 100 | | 3 | 0 | 2 | 2 |
| 9 | read_item(Y) | | 3 | 0 | 3 | 2 |
| 10 | Y := Y * 3 | | 3 | 0 | 3 | 2 |
| 11 | | | | | | |



Example 1

| Time | T ₁ (TS = 3) | T ₂ (TS = 2) | X | | Y | |
|------|-------------------------|-------------------------|----------|----|----------|----------|
| | | | RT | WT | RT | WT |
| 0 | | | 0 | 0 | 0 | 0 |
| 1 | read_item(X) | | 1 | 0 | 0 | 0 |
| 2 | X := X + 100 | | 1 | 0 | 0 | 0 |
| 3 | | read_item(Y) | 1 | 0 | 2 | 0 |
| 4 | | Y := Y * 2 | 1 | 0 | 2 | 0 |
| 5 | | write_item(Y) | 1 | 0 | 2 | 2 |
| 6 | read_item(Y) | | 0 | 0 | 2 | 2 |
| 7 | read_item(X) | | 3 | 0 | 2 | 2 |
| 8 | X := X + 100 | | 3 | 0 | 2 | 2 |
| 9 | read_item(Y) | | 3 | 0 | 3 | 2 |
| 10 | Y := Y * 3 | | 3 | 0 | 3 | 2 |
| 11 | write_item(Y) | | 3 | 0 | 3 | 3 |

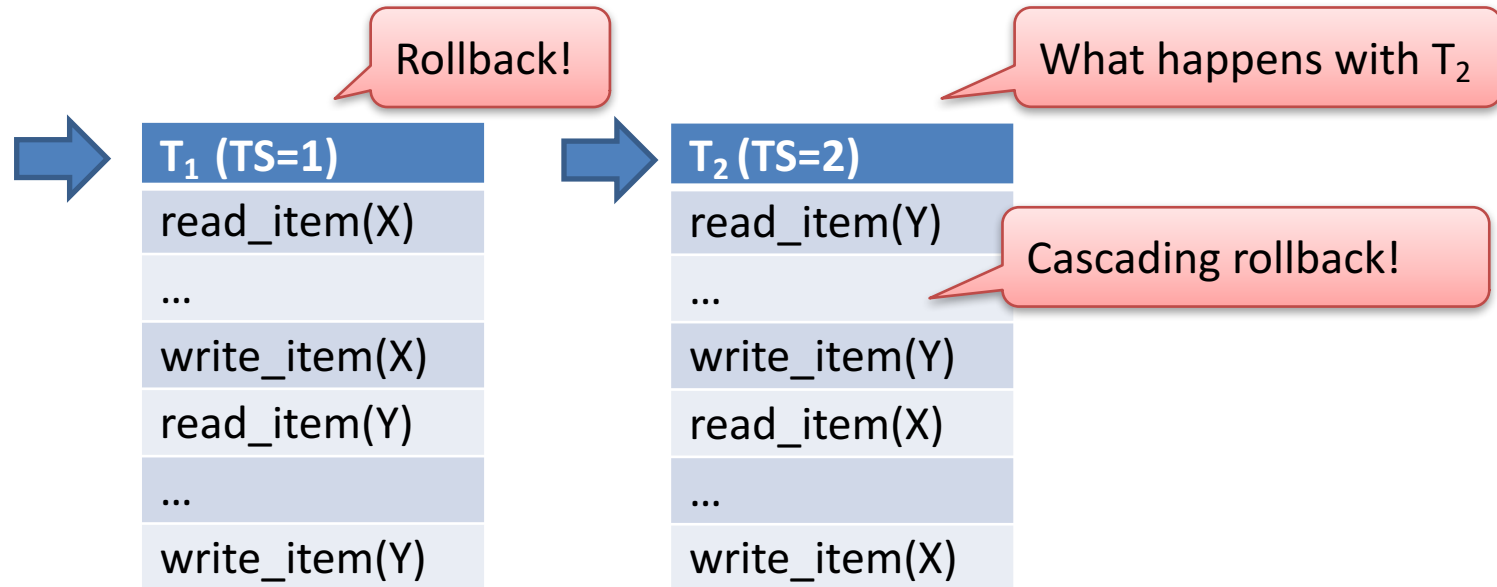


Example 2

| T_1 | T_2 |
|---------------|---------------|
| read_item(X) | read_item(Y) |
| ... | ... |
| write_item(X) | write_item(Y) |
| read_item(Y) | read_item(X) |
| ... | ... |
| write_item(Y) | write_item(X) |

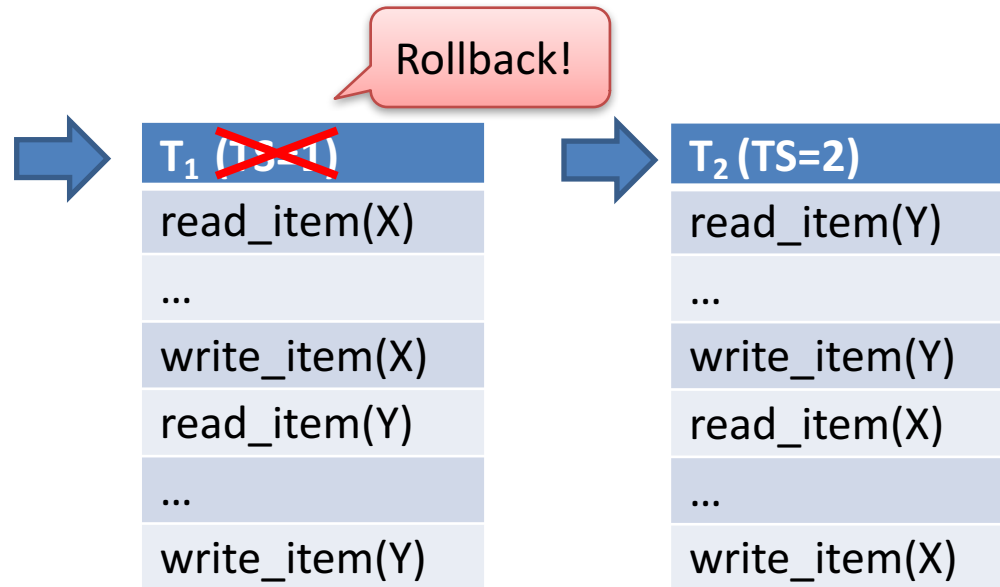
- Assume a timestamp-based scheduler
 - T_1 starts first
 - Lines 1-3 of T_1 are executed first, then lines 1-3 of T_2
- Which operations could be executed next?

Example 2



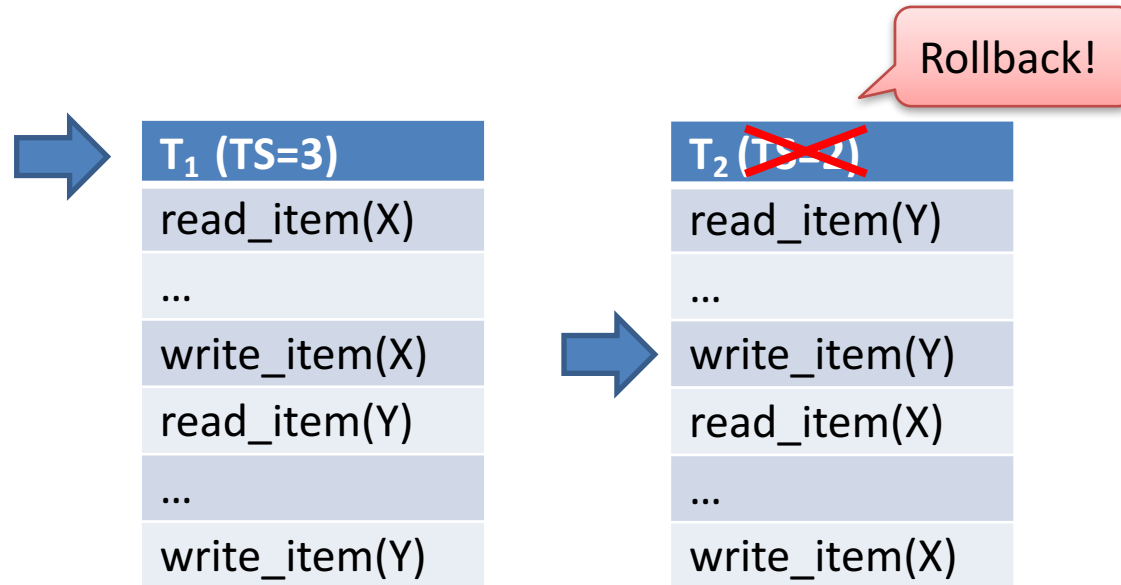
- Assume a timestamp-based scheduler
 - T_1 starts first
 - Lines 1-3 of T_1 are executed first, then lines 1-3 of T_2
- Which operations could be executed next?

Example 2



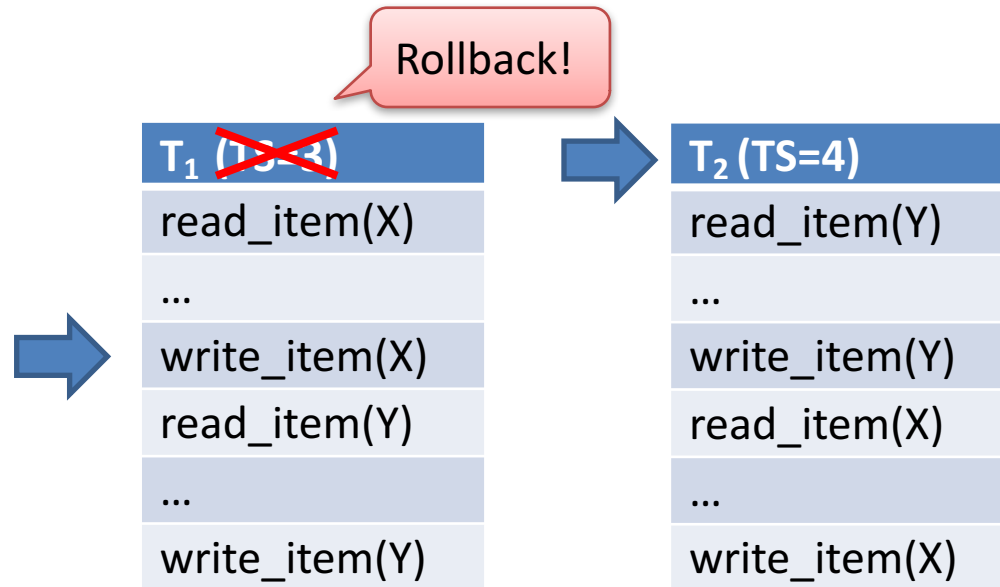
- Assume a timestamp-based scheduler
 - T₁ starts first
 - Lines 1-3 of T₁ are executed first, then lines 1-3 of T₂
- Which operations could be executed next?

Example 2



- Assume a timestamp-based scheduler
 - T_1 starts first
 - Lines 1-3 of T_1 are executed first, then lines 1-3 of T_2
- Which operations could be executed next?

Example 2



- Assume a timestamp-based scheduler
 - T₁ starts first
 - Lines 1-3 of T₁ are executed first, then lines 1-3 of T₂
- Which operations could be executed next?

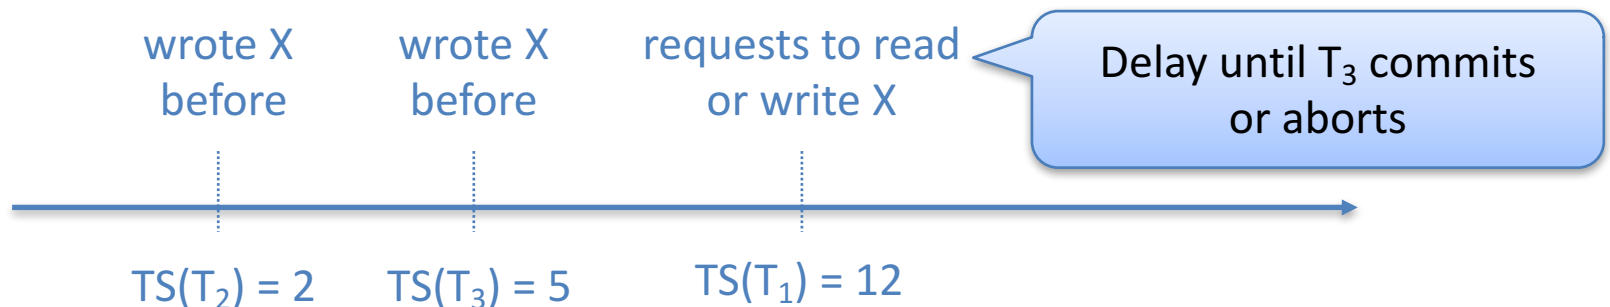
Timestamp-based Scheduling

- Nice properties:
 - Enforces conflict-serialisable schedules
 - Deadlocks don't occur
- Bad properties:
 - **Cascading rollbacks**
 - **Starvation can occur** (cyclic aborts & restarts of transactions)
- Starvation can be prevented using appropriate techniques (not in this module)

Ensuring Strictness

- Schedules enforced by timestamp-based schedulers are not strict.
- Additional condition to enforce a **strict schedule**:

Delay read or write requests until the youngest transaction who wrote X before has committed or aborted.



Locking or Timestamping?

Timestamping vs Locking

- Space usage roughly the same
 - Timestamping: manage timestamps & read/write times
 - Locking: manage locks
- Optimistic vs conservative
 - **Timestamping is optimistic**: “unserialisable behaviour will not occur, or occurs rarely and can then be fixed”
 - **Locking is conservative**: “unserialisable behaviour will occur unless we take precautions”
- Both have issues
 - Locking may cause **deadlocks** (but can be fixed)
 - Timestamping may cause **starvation** (but can be fixed)

Timestamping or Locking?

Why might this be the case?

- Timestamping might be preferable if there is little interaction between transactions, e.g.,
 - Transactions are read-only (SELECT-FROM-WHERE)
 - Transactions access distinct database items
- Some DBMSs use a combination of locking and timestamping
 - Locking for transactions that read & write
 - Timestamping for read-only transactions