

# COMP207

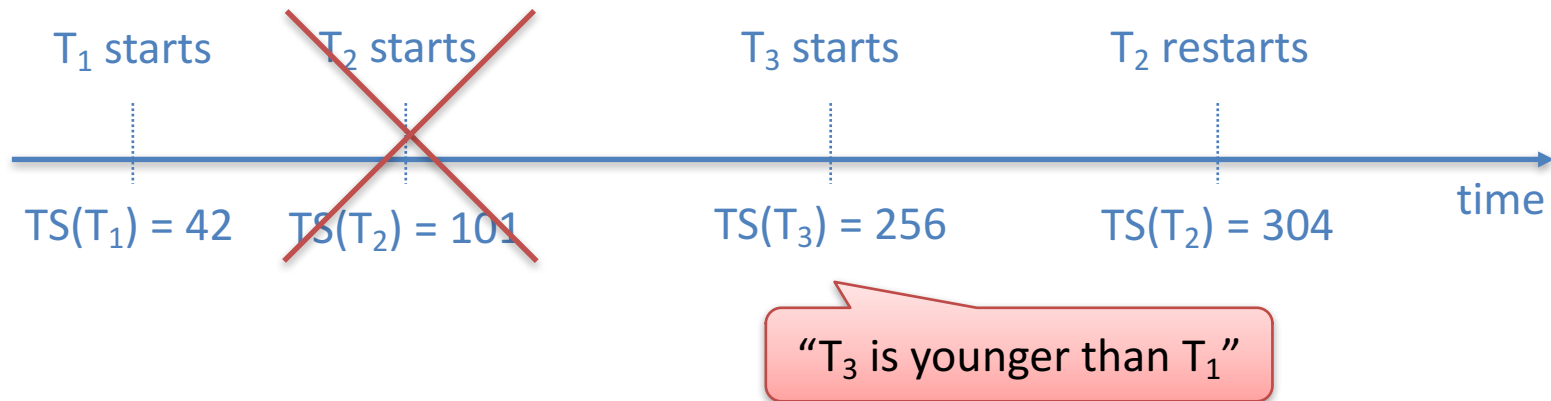
# Database Development

## Lecture 12

Review of Chapter 2 and  
the start of Chapter 3: Query Processing:  
From SQL to Relational Algebra

# Timestamp-based Scheduling

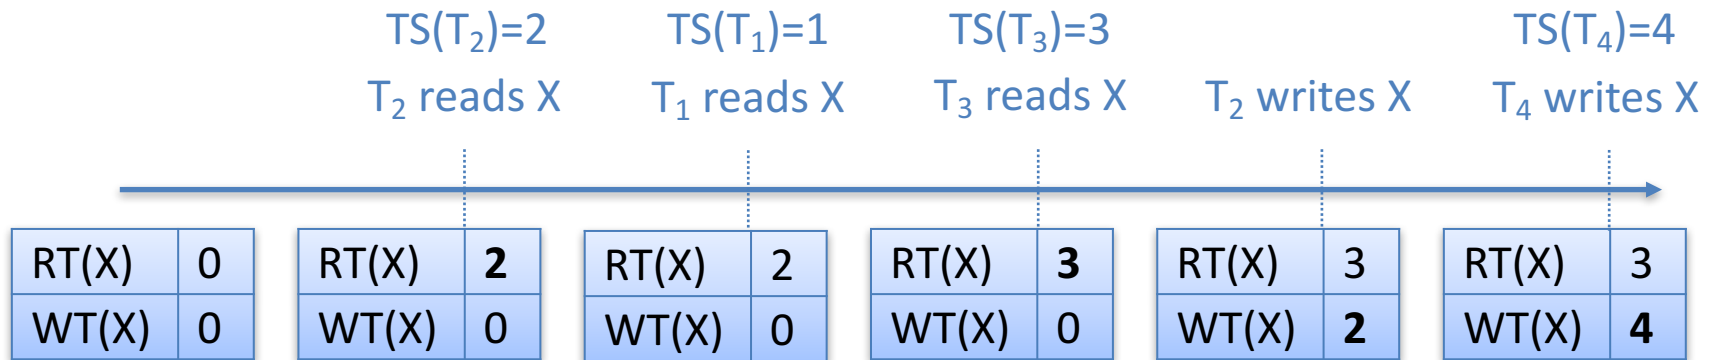
- Enforces a schedule that is *equivalent to the serial schedule* that contains all transactions ordered by start time.



- Uses timestamps
  - If a transaction  $T$  starts, assigns a new timestamp  $TS(T)$  to it
  - “Younger” transactions are those with higher timestamps
  - Timestamps are used to decide whether to grant read/write requests

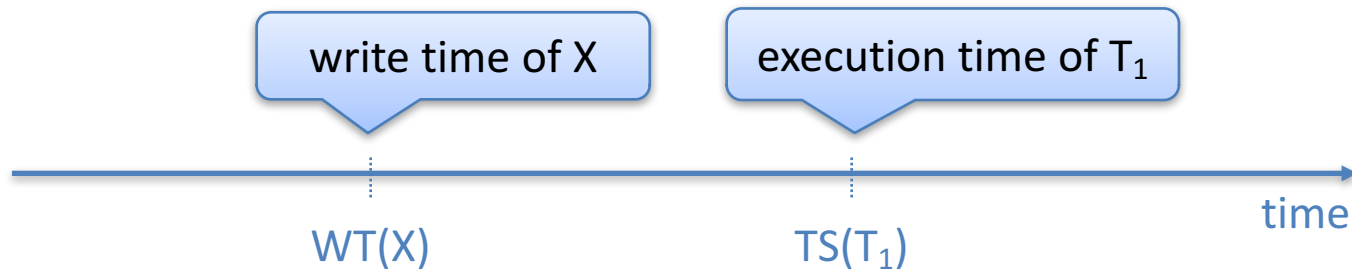
# Additional Bookkeeping

- For each database item  $X$ , maintain:
  - Read Time of  $X$ :  $RT(X)$**   
Timestamp of youngest transaction that read  $X$
  - Write Time of  $X$ :  $WT(X)$**   
Timestamp of youngest transaction that wrote  $X$



# Read Requests

$T_1$  requests to **read** X ...

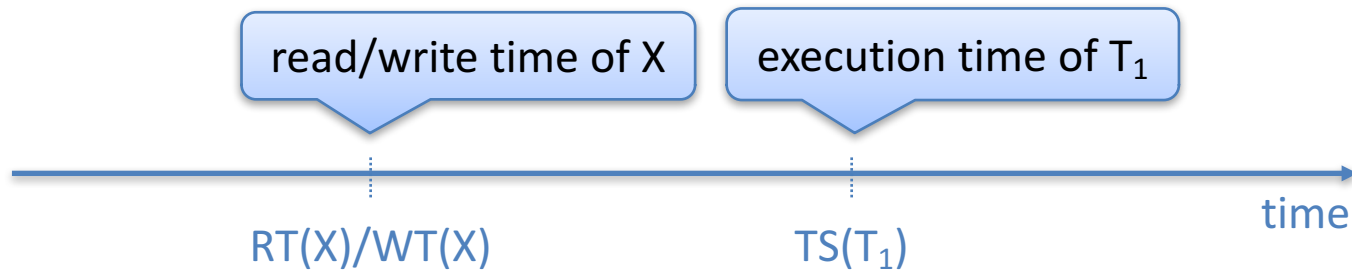


"X was last written before  $T_1$  started"

- **Grant request if  $WT(X) \leq TS(T_1)$**
- **Abort & restart  $T_1$  otherwise**

# Write Requests

$T_1$  requests to **write** X ...



"X was last read or written before  $T_1$  started"

- **Grant request** if  $RT(X) \leq TS(T_1)$  and  $WT(X) \leq TS(T_1)$
- **Abort & restart  $T_1$**  otherwise

# MySQL (with InnoDB)

- Uses Wait-For graphs
- Except: If transaction has line of code that is rolled back it is
- E.g.  $T_1$  is rolled back in this case:



Also uses a timestamp based approach to ensure that reads do not interfere with writes

- If cycle: rollback the smallest transaction
- Deadlock detection can be switched off, in which case time-out is used (on locks)

# Other DBMS

## PostgreSQL

- Uses timeout on locks followed by Wait-For graphs
- Like MySQL: Uses a timestamp based approach for ensuring that reads do not interfere with writes

## Oracle DB

- Uses timeout directly or timeout followed by Wait-For graphs
- Does not use locks on read

## IBM DB2

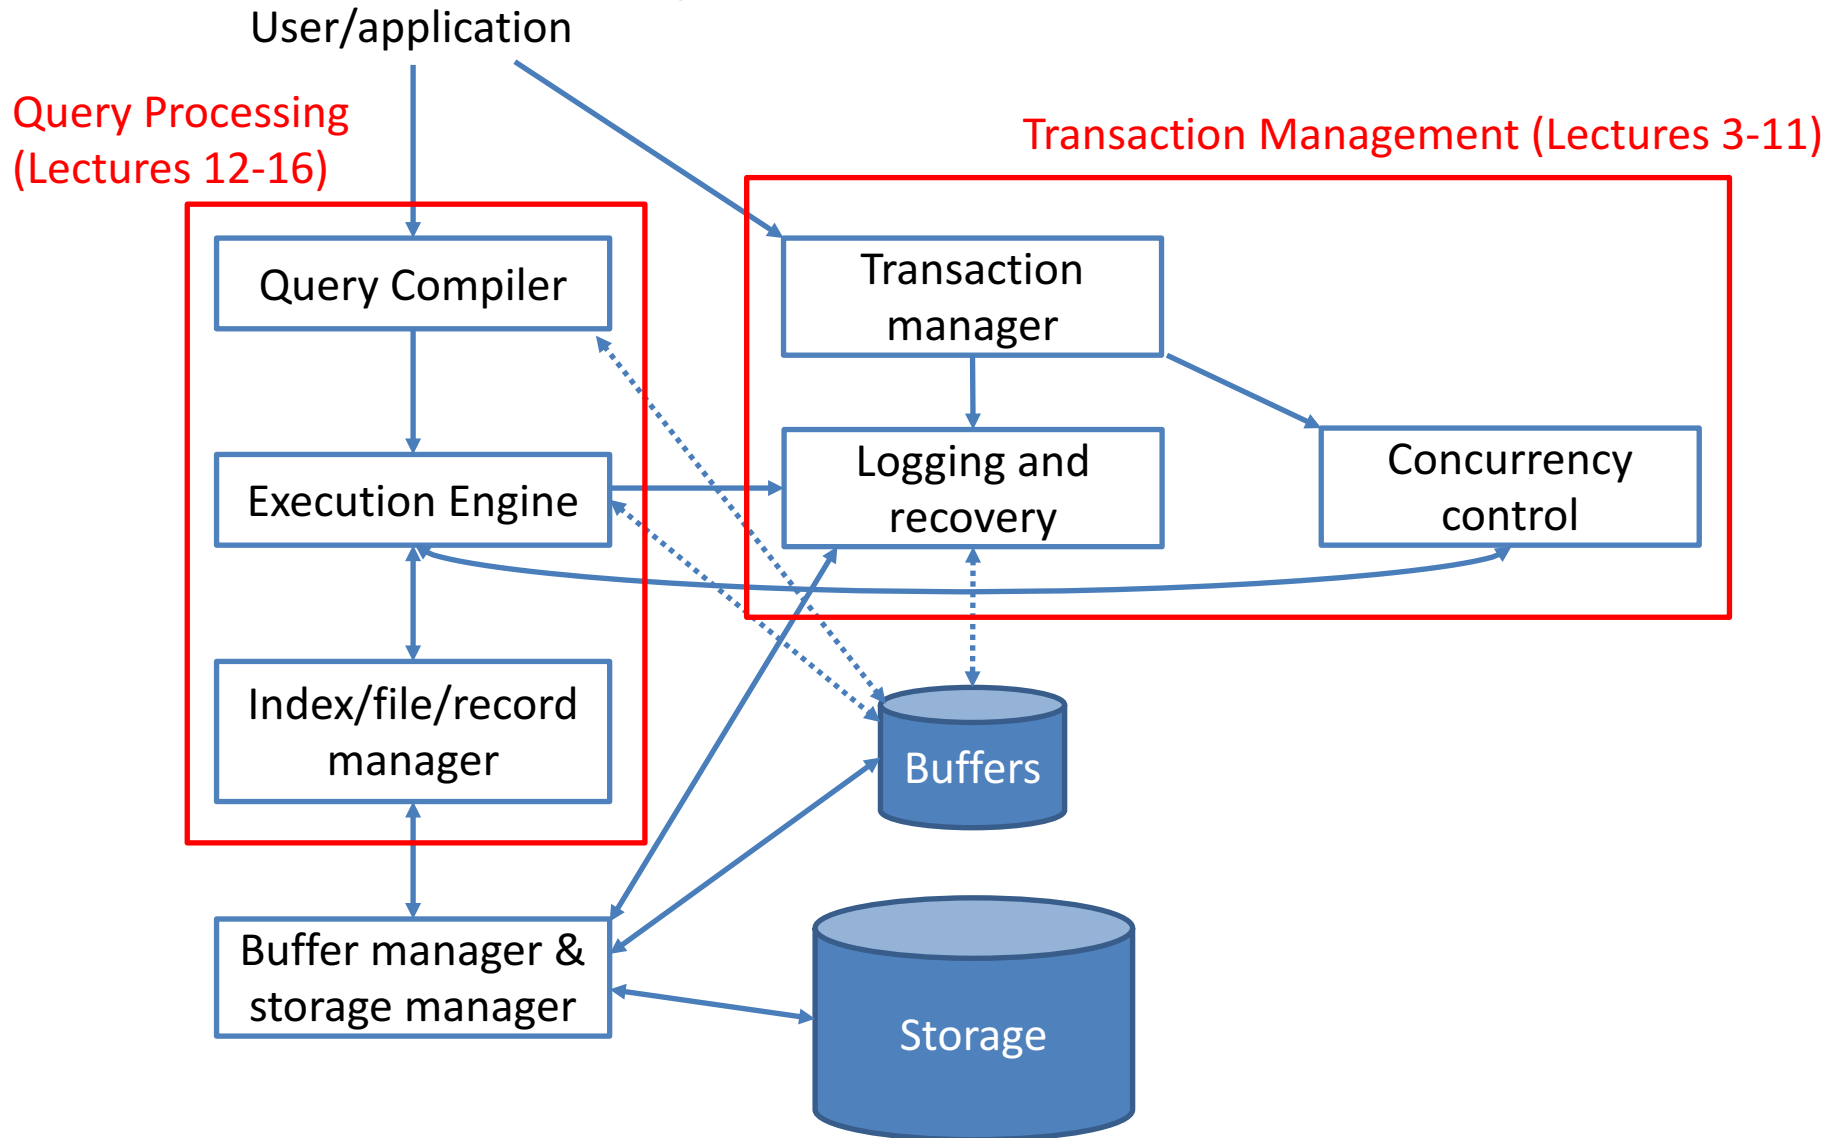
- Uses lock timeout or global time-out followed by Wait-For graphs
- Uses update-locks

## Concluding Remarks



# Relational DBMS Components

(Simplified, from Lecture 1)



# Transaction Management Review

- Dealing with **transactions** is a core task of DBMS
  - Many things can go wrong when processing transactions, even when executing single SQL statements.
  - Need to ensure **ACID properties**
- Requires careful **scheduling** of transactions and **logging** of relevant information
  - Schedules should be **conflict-serialisable**
  - Schedules should be **strict**
- Methods for enforcing conflict-serialisability & strictness:
  - **Strict two-phase locking & deadlock prevention** methods
  - **Timestamping**

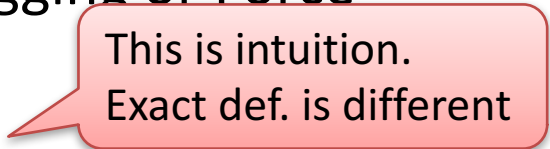
# ACID

- **A**tomicity

- Transactions are fully executed or not at all
- Ensured by Undo logging, Undo/Redo logging or Force

- **C**onsistency

- Schedule executes transactions equivalent to a serial schedule
- (needs two assumptions for this: non-database operations can be ignored and if a schedule is serial, then it is consistent)
- Ensured by Serializability, Conflict-Serializability, 2PL and Timestamp-based Scheduling (also Strict versions of the last two)



This is intuition.  
Exact def. is different

# ACID continued

- **I**solat**I**on
  - Transactions never reads uncommitted data
  - Ensured by Cascadeless and Strict schedules (incl. Strict 2PL and Strict Timestamp-based schedules)
- **D**urability
  - If a transaction is committed, it does not disappear
  - Ensured by Redo logging, Undo/Redo logging or No Steal
  - Recoverable schedules are also required

# Transaction Support in DBMS

- Part of SQL:
  - Begin/end transactions, isolation levels, auto commit, ...
  - Need to understand the consequences of these commands to make effective use of DBMS
    - When to combine different SQL statements into a transaction?
    - When do we need (conflict) serialisability? When is a weaker isolation level fine?
- Widespread ACID support in major DBMSs
  - Fully ACID compliant: PostgreSQL, Oracle DB, IBM DB2, ...
  - Partly ACID compliant: MySQL (full compliance requires additional engines like InnoDB)

# Transactions Beyond DBMS

- The techniques covered in this chapter are not confined to DBMS
- Similar issues whenever systems share resources
- Some example scenarios:
  - Processes in an operating system that access the same files, network resources, etc.
  - Users editing the same document online
  - Document versioning systems like subversion, git, etc.

# Try it out...

- `CREATE TABLE Student (id INT NOT NULL, name ...);`
- `INSERT INTO Student VALUES (1, 'Anna', ...);`  
`SELECT * FROM Student;`
- **START TRANSACTION;**  
`INSERT INTO Student VALUES (2, 'Ben', ...);`  
`INSERT INTO Student VALUES (3, 'Chloe', ...);`  
**ROLLBACK;**  
`SELECT * FROM Student;`
- Try out reads, writes, different isolation levels, dirty reads, look up the documentation, ...
- Experiment with more complex scenarios...

Try out with any DBMS:  
MySQL, PostgreSQL, ...  
MySQL and PostgreSQL  
are easy to set up.  
Lots of tutorials online.

Which tuples are returned?