

# COMP207

# Database Development

## Lecture 13

### Query Processing: Query Plans and their Execution

# Semijoin ( $\bowtie$ )

- $R_1 \bowtie R_2$  = tuples from  $R_1$  matching tuples in  $R_2$

Modules

Lecturers

```
SELECT *  
FROM Modules  
WHERE EXISTS (SELECT 1  
              FROM Lecturers  
              WHERE  
                Modules.module =  
                Lecturers.module)
```

Two tuples match  
if they have the same values  
for all common attributes

module	year
COMP105	1

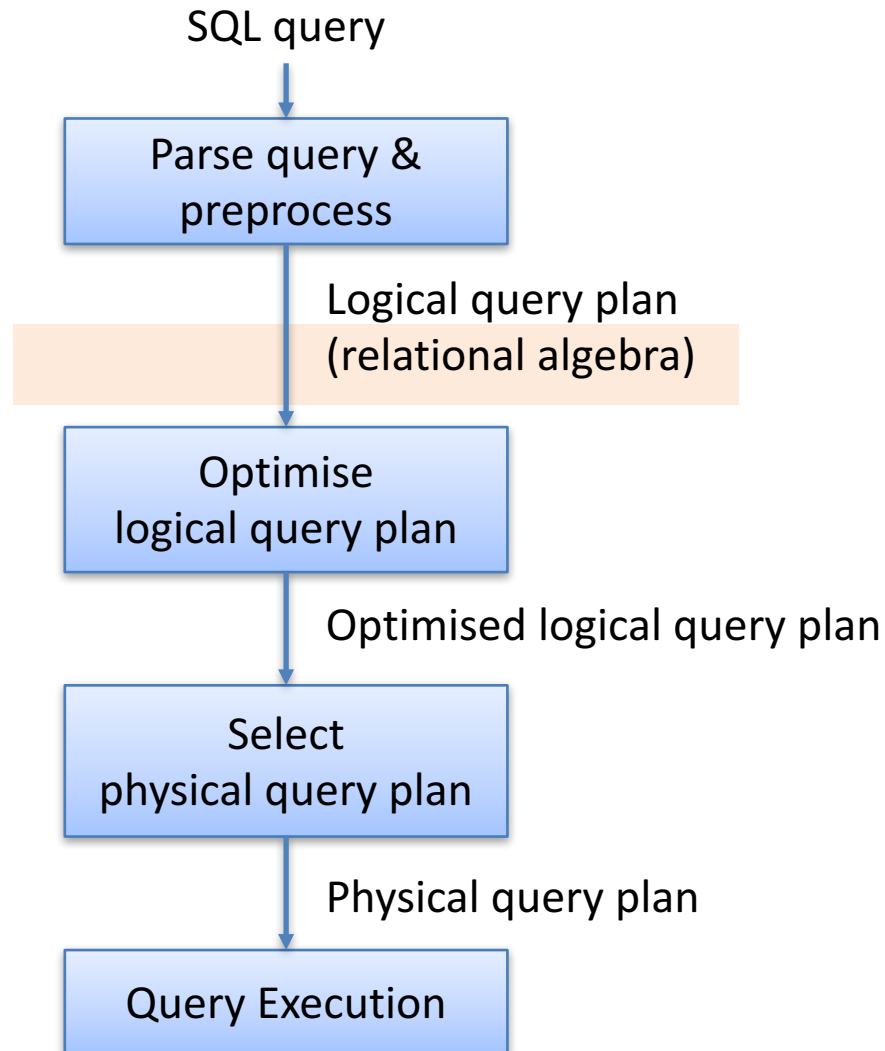
- Can be expressed by the other operators:

$\text{Modules} \bowtie \text{Lecturers} =$

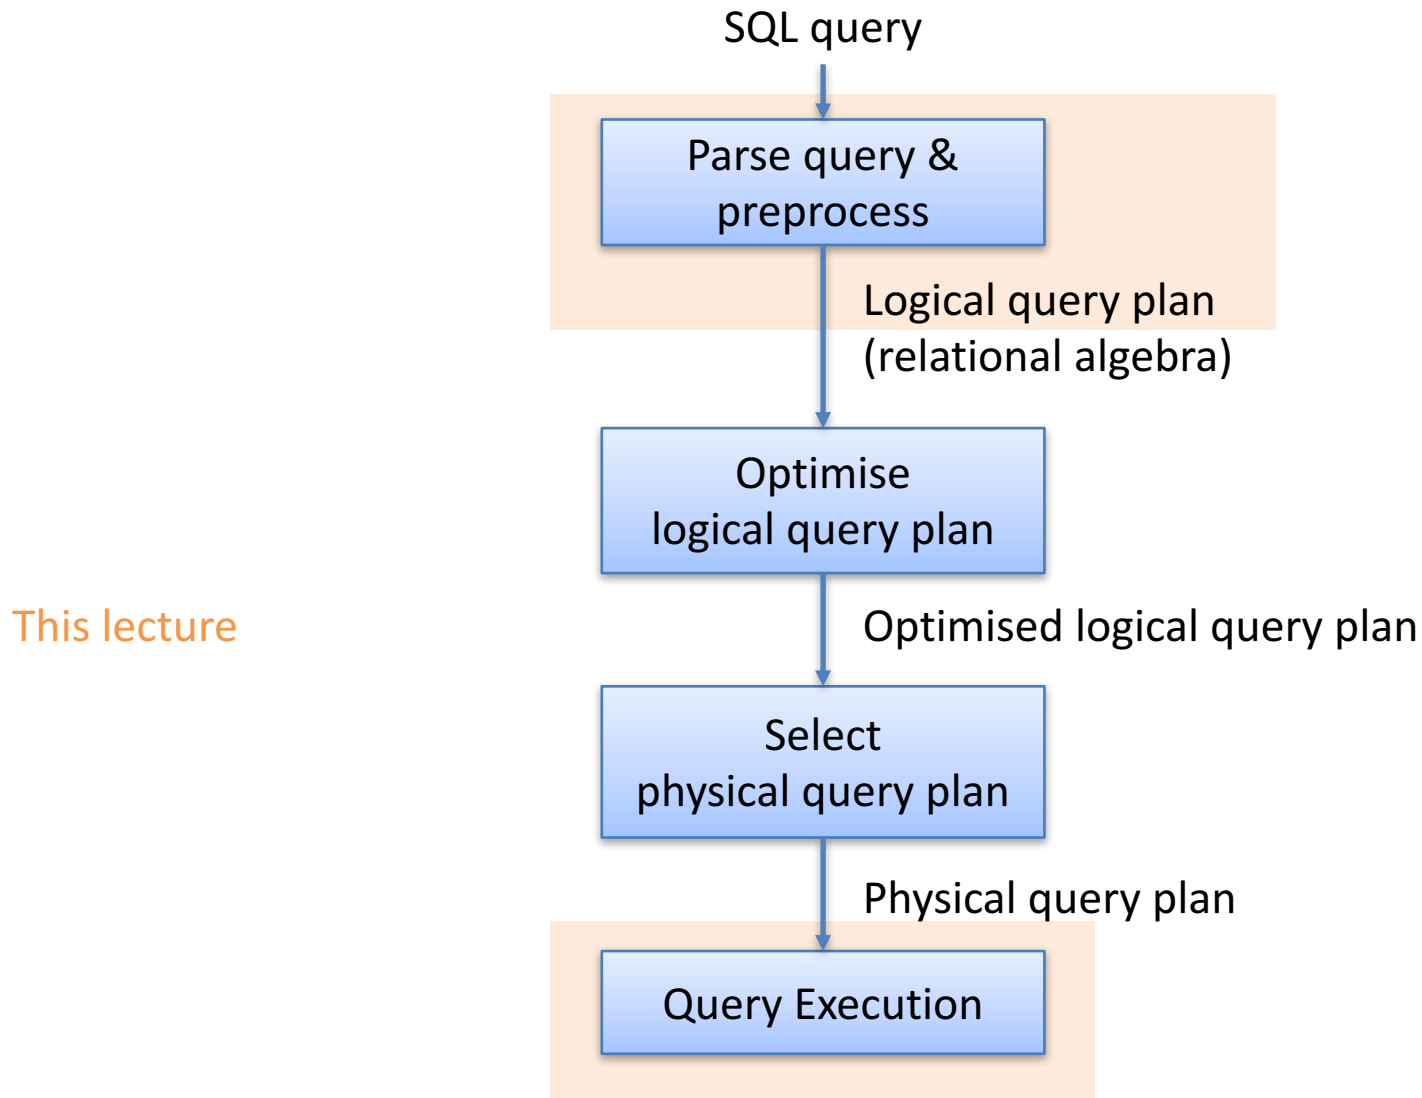
$\pi_{\text{module, year}}(\sigma_{\text{module}=\text{module}'}(\text{Modules} \times \rho_{\text{module} \rightarrow \text{module}'}(\text{Lecturers})))$

# Reminder: Query Processing

Previous lecture



# Reminder: Query Processing



# Query Plans

(a.k.a. Query Trees)

# Query Plans

- A **relational algebra expression** that is obtained from an SQL query is also called a **(logical) query plan**

```
SELECT module, name  
FROM Modules, Lecturers  
WHERE code=module AND year=2;
```

SQL query



```
 $\pi_{\text{module, name}}(\sigma_{\text{code=module AND year=2}}(\text{Modules} \times \text{Lecturers}))$ 
```


Query plan for the query

- Query plans are typically **represented as trees**

# Query Plans As Trees

$\pi_{\text{module,name}}(\sigma_{\text{code=module AND year=2}}(\text{Modules} \times \text{Lecturers}))$

- Tree representation:




- Inner nodes = operators
- Leaves = input relations
- Such trees are evaluated from the leaves to the root

# Exercise

- Represent the following query plan as a tree:

$\sigma_{\text{code}=\text{module}}(\pi_{\text{code}}(\text{Modules}) \times \pi_{\text{module,name}}(\text{Lecturers}))$






# Equivalent Query Plans

- There are typically **many different query plans**
- DBMSs aim to **select a best possible query plan**
- Relational algebra is better suited than SQL for this
  - Can use **equivalence laws** of relational algebra to generate a query plan for the same query that can be executed faster!
  - Example:

Details will come later...


$$\sigma_{\text{condition}}(R_1 \bowtie R_2) = \sigma_{\text{condition}}(R_1) \bowtie \sigma_{\text{condition}}(R_2)$$



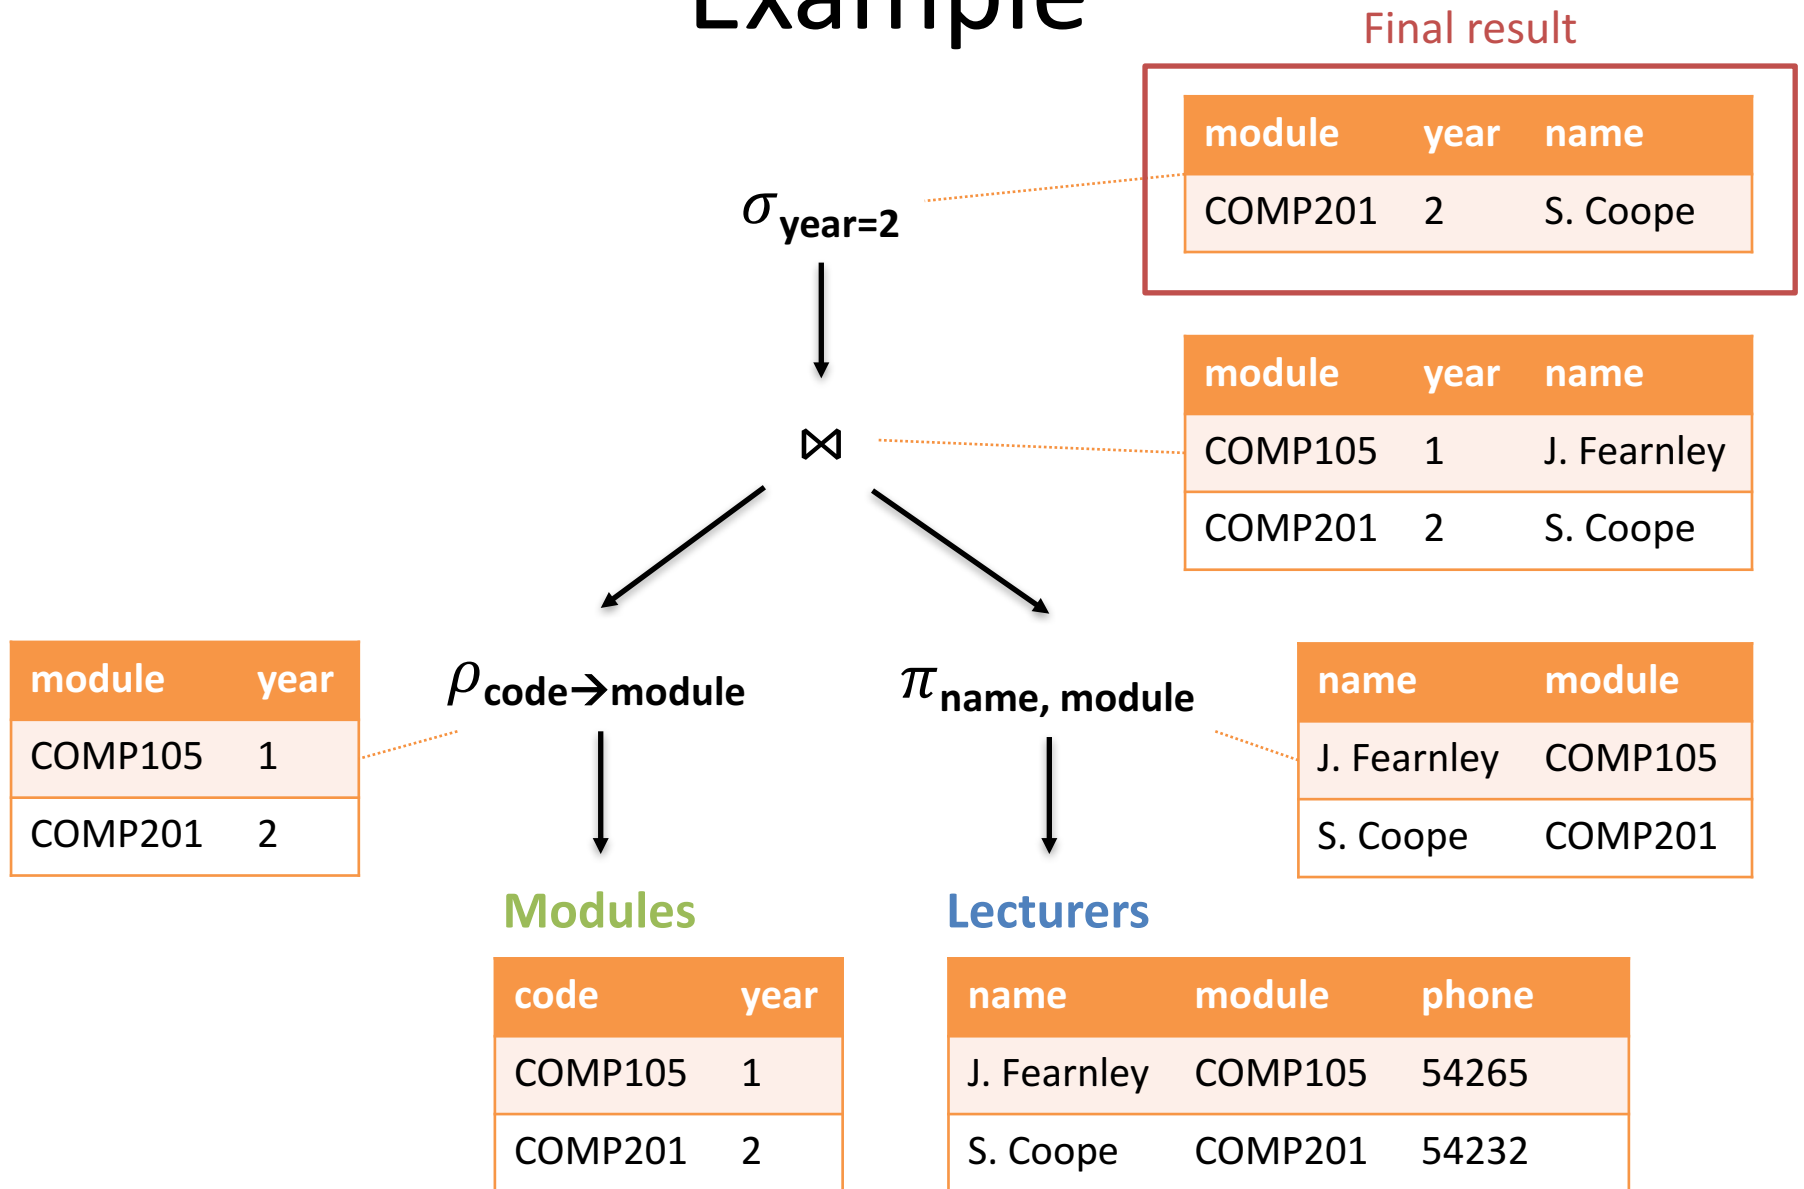
# Execution of Query Plans

# Executing Query Plans

- Query plans tells us exactly how to compute the result to a query
- Proceed from bottom to top:
  - Compute an **intermediate result** for each node
  - For a **leaf** labeled with relation **R**, the intermediate result is **R**.
  - For an **inner node** labeled with operator **op**, get the intermediate result by applying **op** to the childrens' intermediate results.
  - **Result of the query** = intermediate result of the root

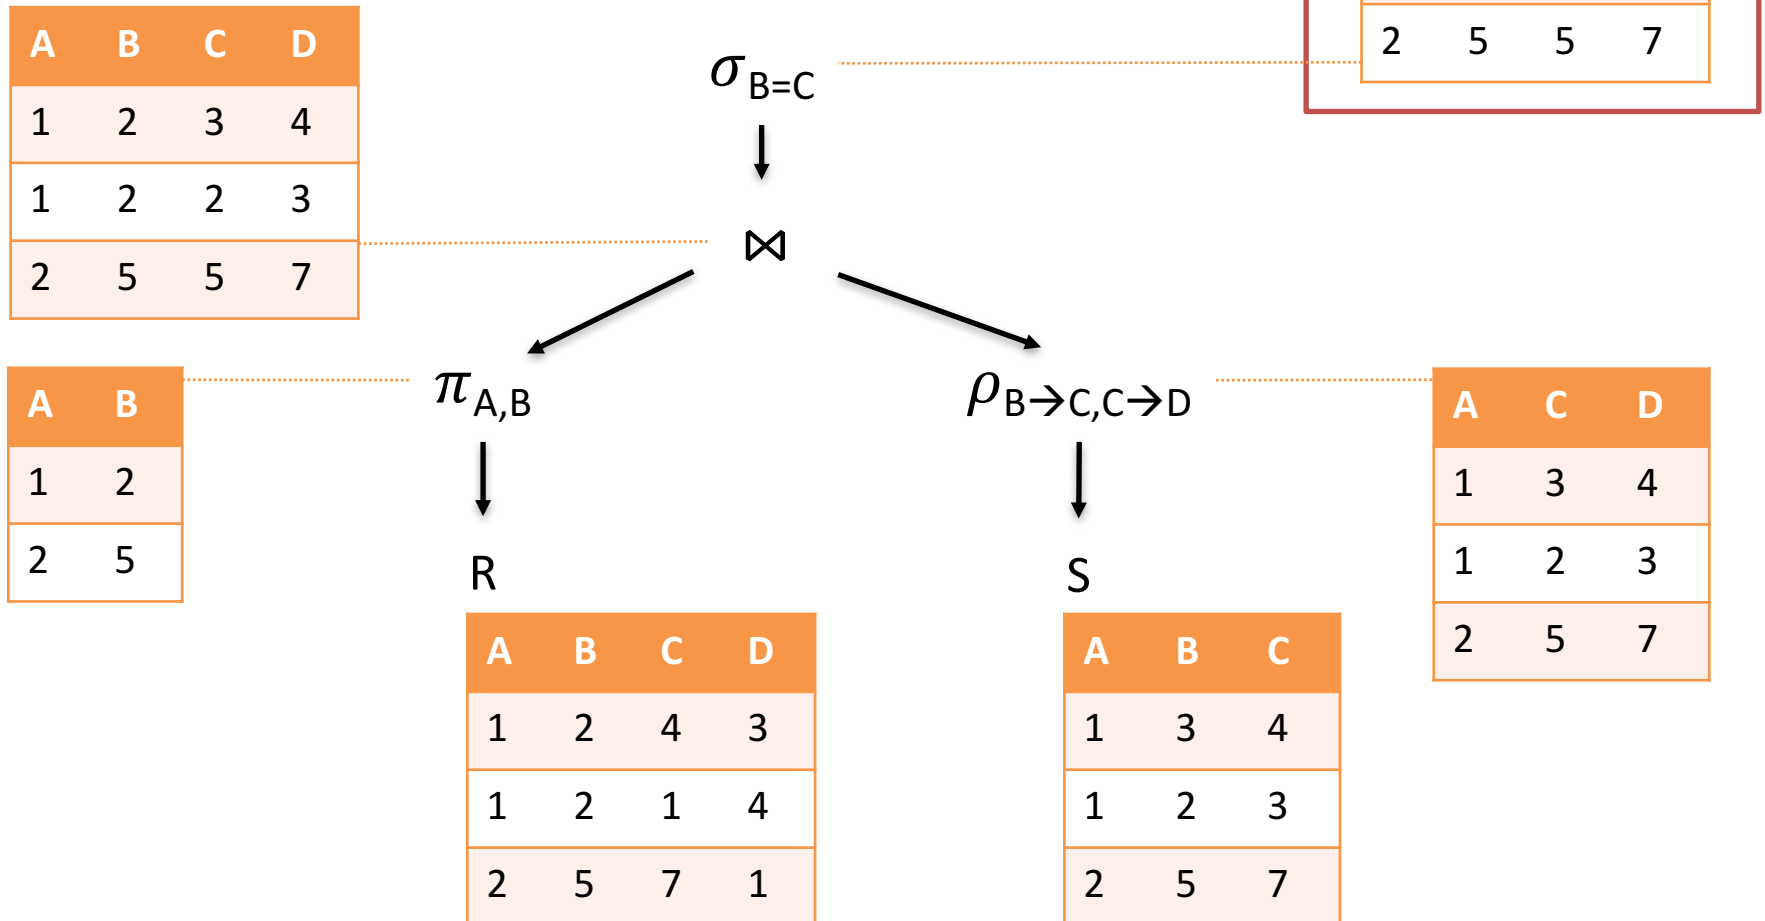


# Example




# Exercise (5 min)

Execute this query plan:



# How to “Apply” An Operator?

- How to compute  $\sigma_{\text{condition}}(R)$ ?



for each tuple **t** in **R**:  
if **t** satisfies condition:  
output **t**

Is there a faster way?

- How to compute  $\pi_{\text{attribute list}}(R)$ ?

Similar! Read **R** only once...

for each tuple **t** in **R**:  
output the restriction  
of **t** to the attributes  
in **attribute list**

- How to compute  $R \bowtie S$ ?

Many different ways...

Is there a faster way?

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

module	year	name
COMP105	1	J. Fearnley

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley



# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

$s$

module	year	name
COMP105	1	J. Fearnley

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley
COMP201	2	S. Coope

$r \bowtie s$

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley
COMP201	2	S. Coope

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

the tuple obtained by joining  $r$  and  $s$

Modules

module	year
COMP105	1
COMP201	2

$\bowtie$

Lecturers

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

=

module	year	name
COMP105	1	J. Fearnley
COMP201	2	S. Coope
COMP201	2	J. Smith

$r \bowtie s$

# Naïve Computation of Joins

- Nested Loop Join Algorithm:

Compute  $R \bowtie S$ :

for each tuple  $r$  in  $R$ :

for each tuple  $s$  in  $S$ :

if  $r$  and  $s$  have the same values for all common attributes:

output  $r \bowtie s$

- Slow: for each tuple  $r$  in  $R$  reads entire relation  $S$
- Running time:  $O(|R| \times |S|)$

Number of tuples in  $R$

Number of tuples in  $S$

*Can We Go Faster?*

*Yes, we can!*



# Equijoins

- Equijoin  $R \bowtie_{A=B} S$  is defined as  $\sigma_{A=B}(R \times S)$

A, B are the **join attributes**

**Modules**

code	year
COMP105	1
COMP201	2

**Lecturers**

name	module
J. Fearnley	COMP105
S. Coope	COMP201
J. Smith	COMP201

**Modules**  $\bowtie_{\text{code}=\text{module}}$  **Lecturers**

code	year	name	module
COMP105	1	J. Fearnley	COMP105
COMP201	2	S. Coope	COMP201
COMP201	2	J. Smith	COMP201


- If  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$ , then  $R \bowtie_{A=B} S$  can be computed with one pass over  $R$  and  $S$

# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

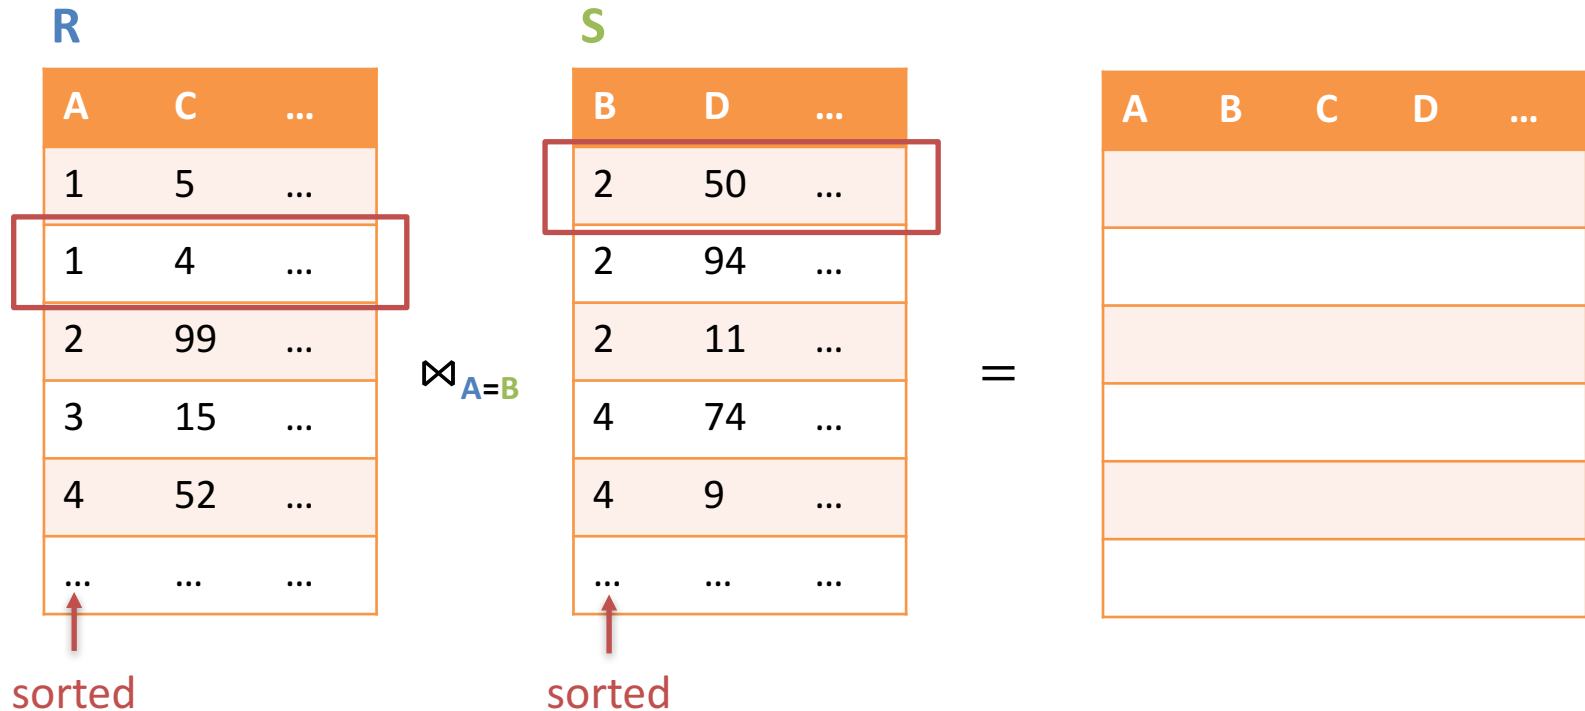


# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

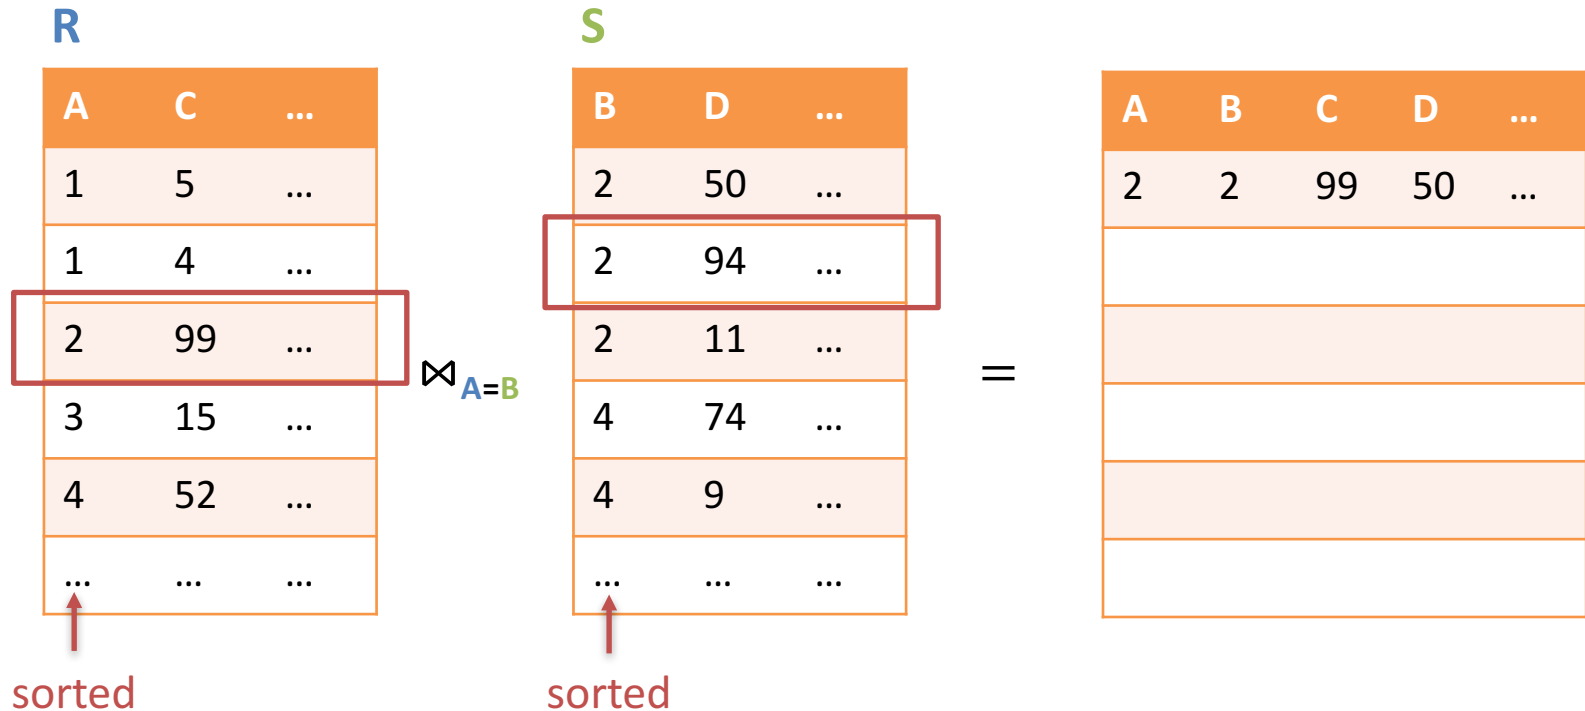


# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$






# Merging as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$




# Merging

as in merge sort

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

Running time:  $O(|R| + |S|)$



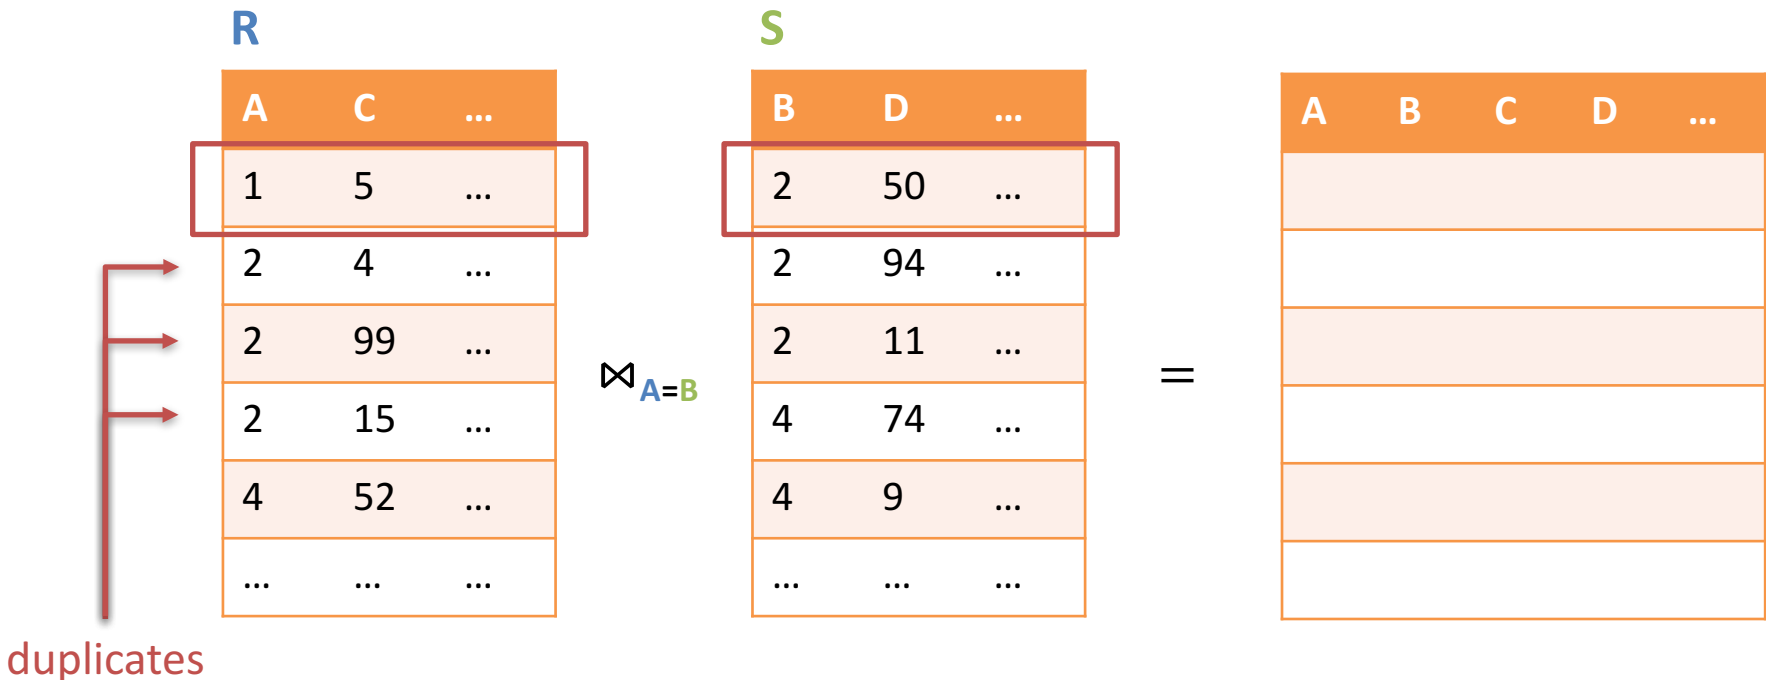
- What if a value occurs twice or more in column  $A$ ?



# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

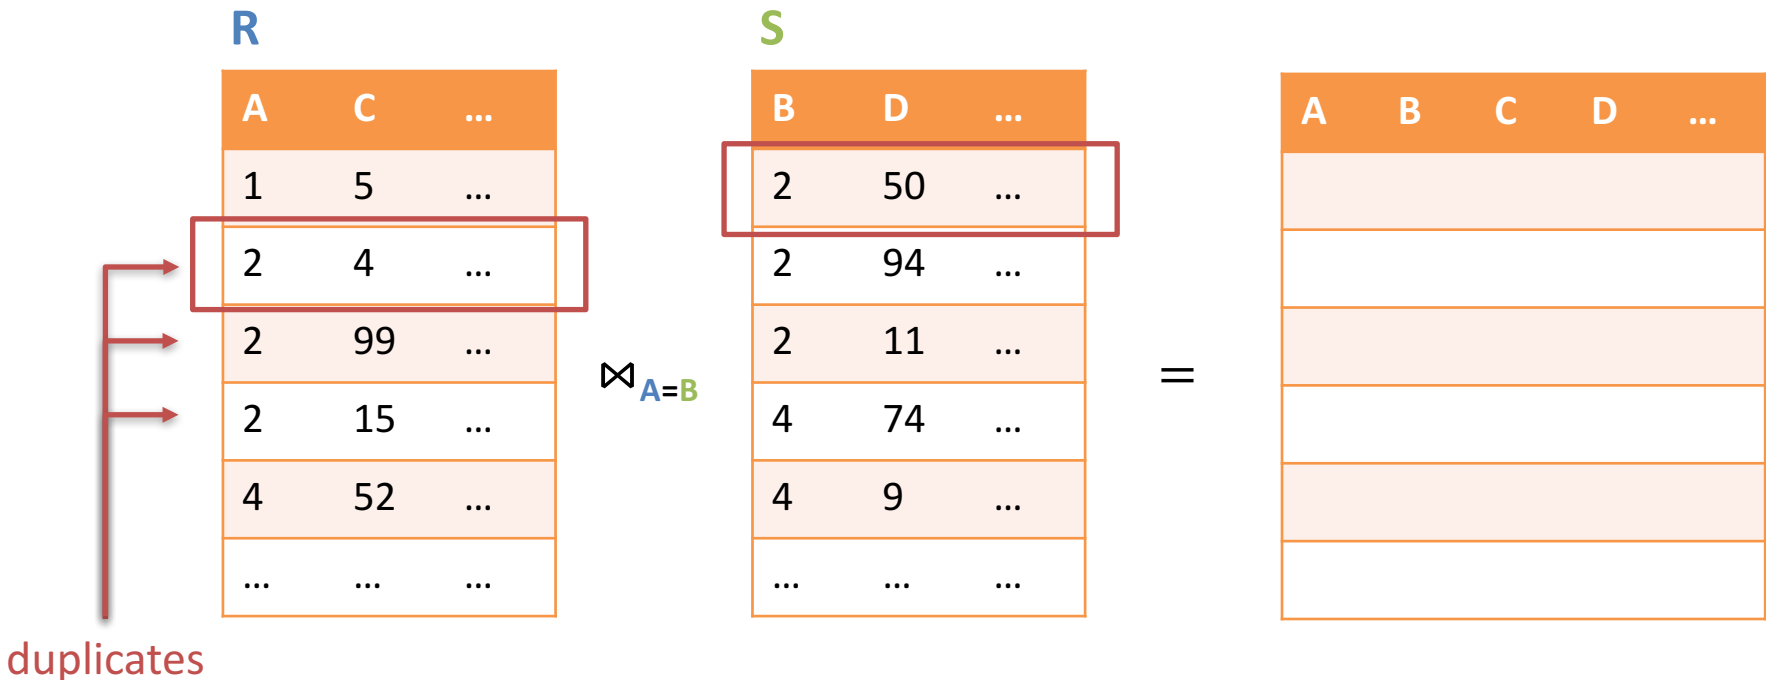


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

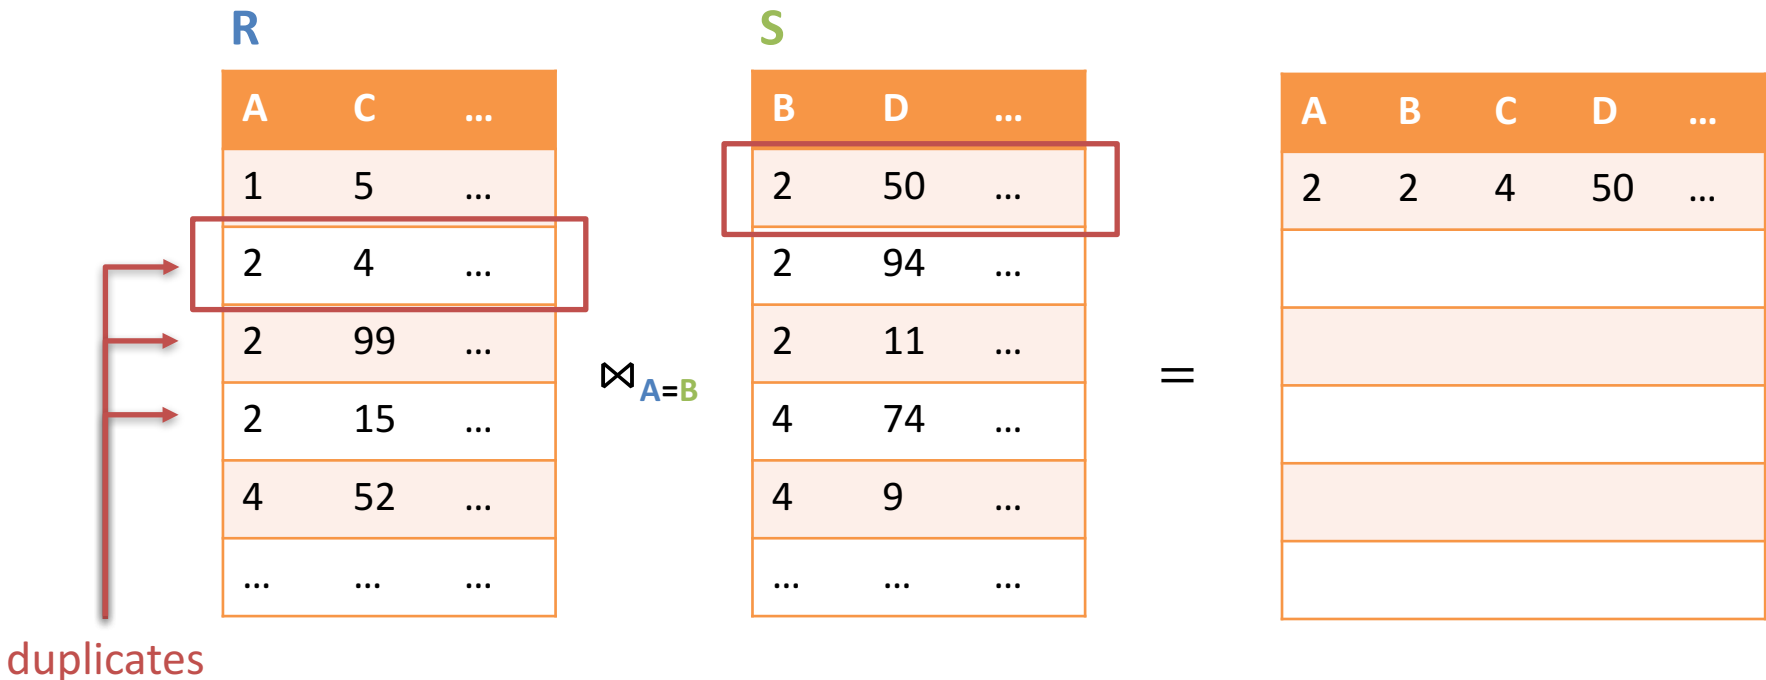


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

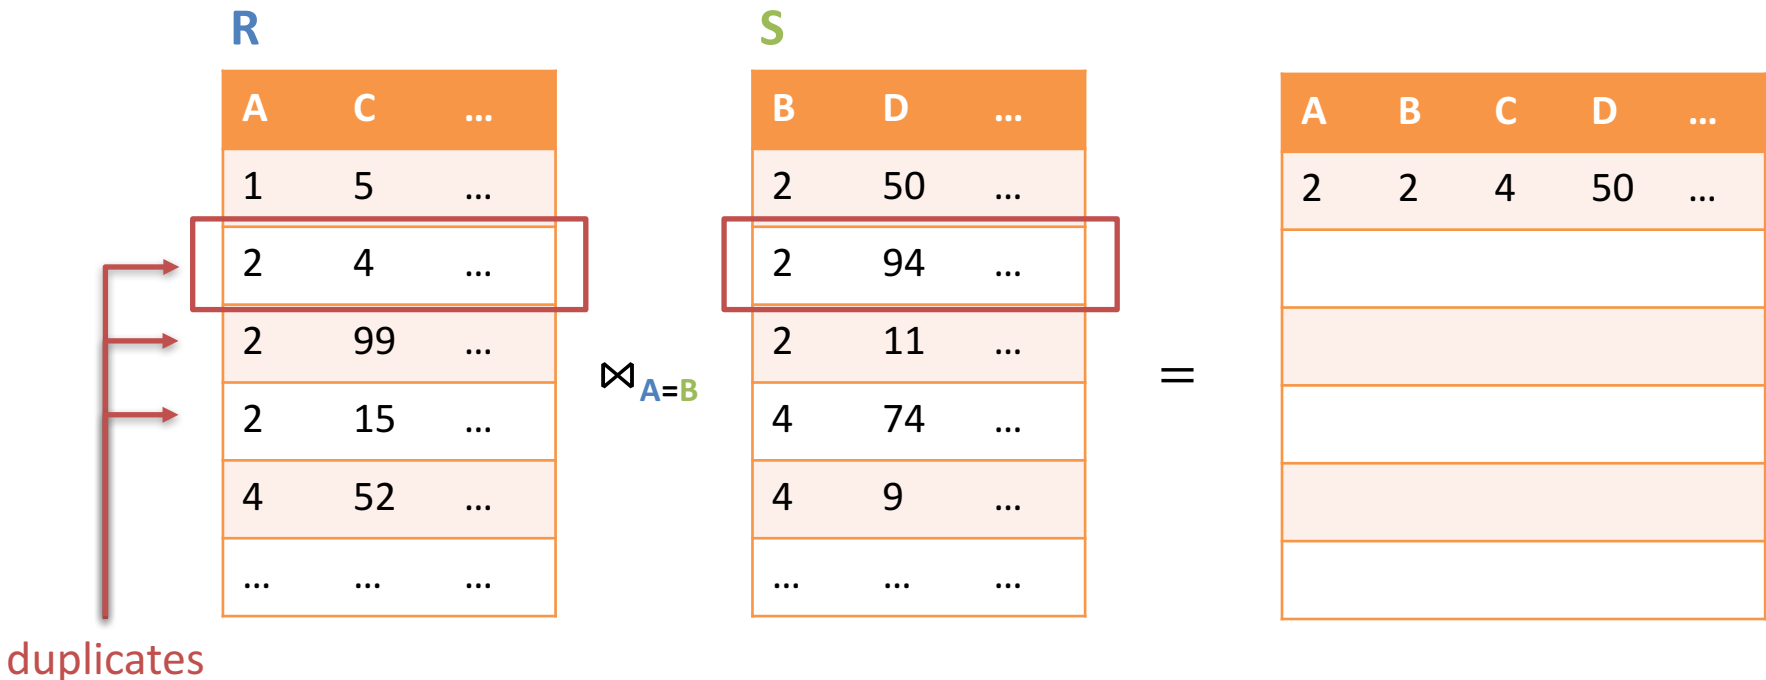


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

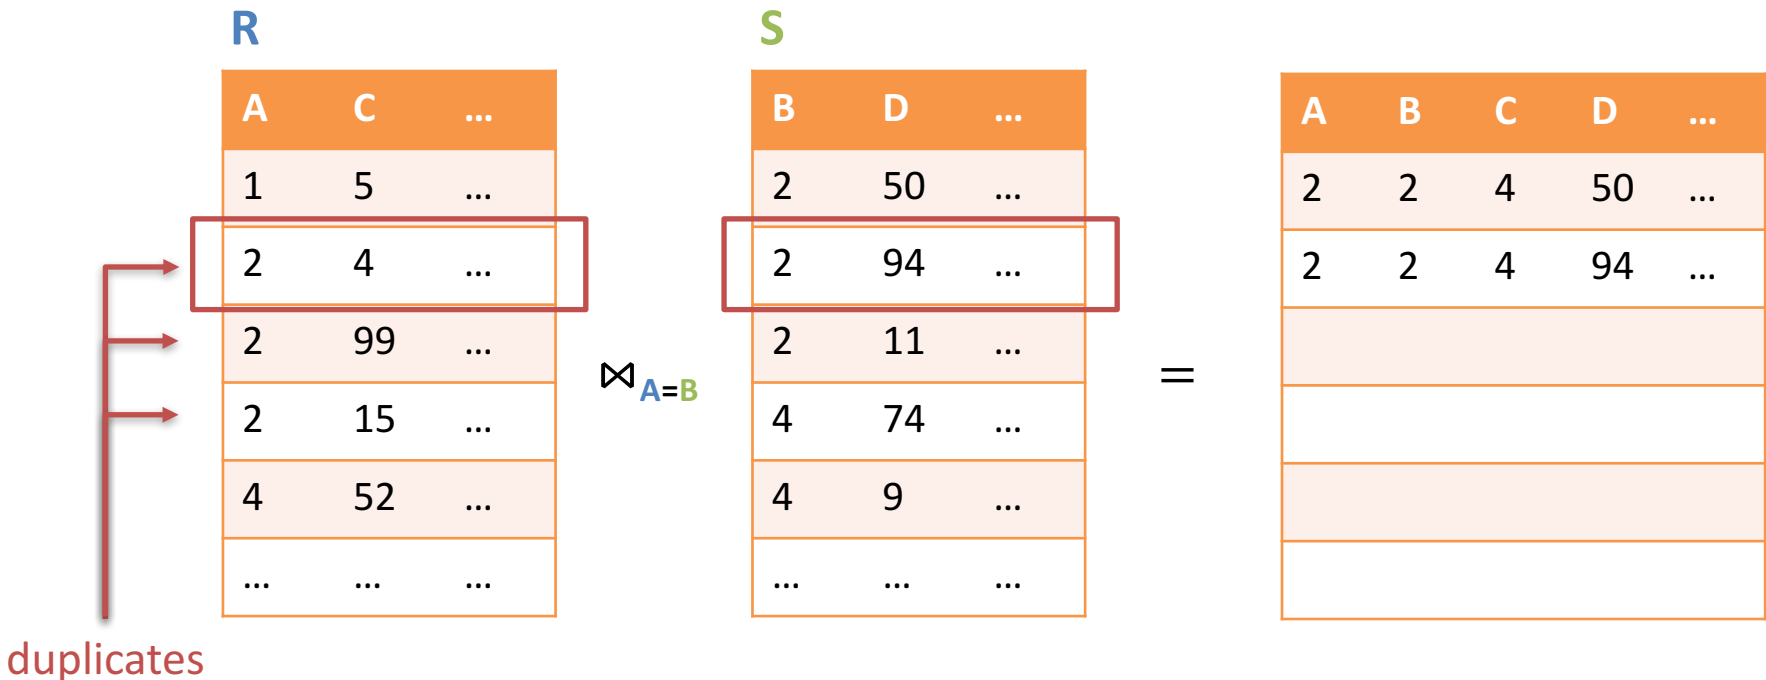


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

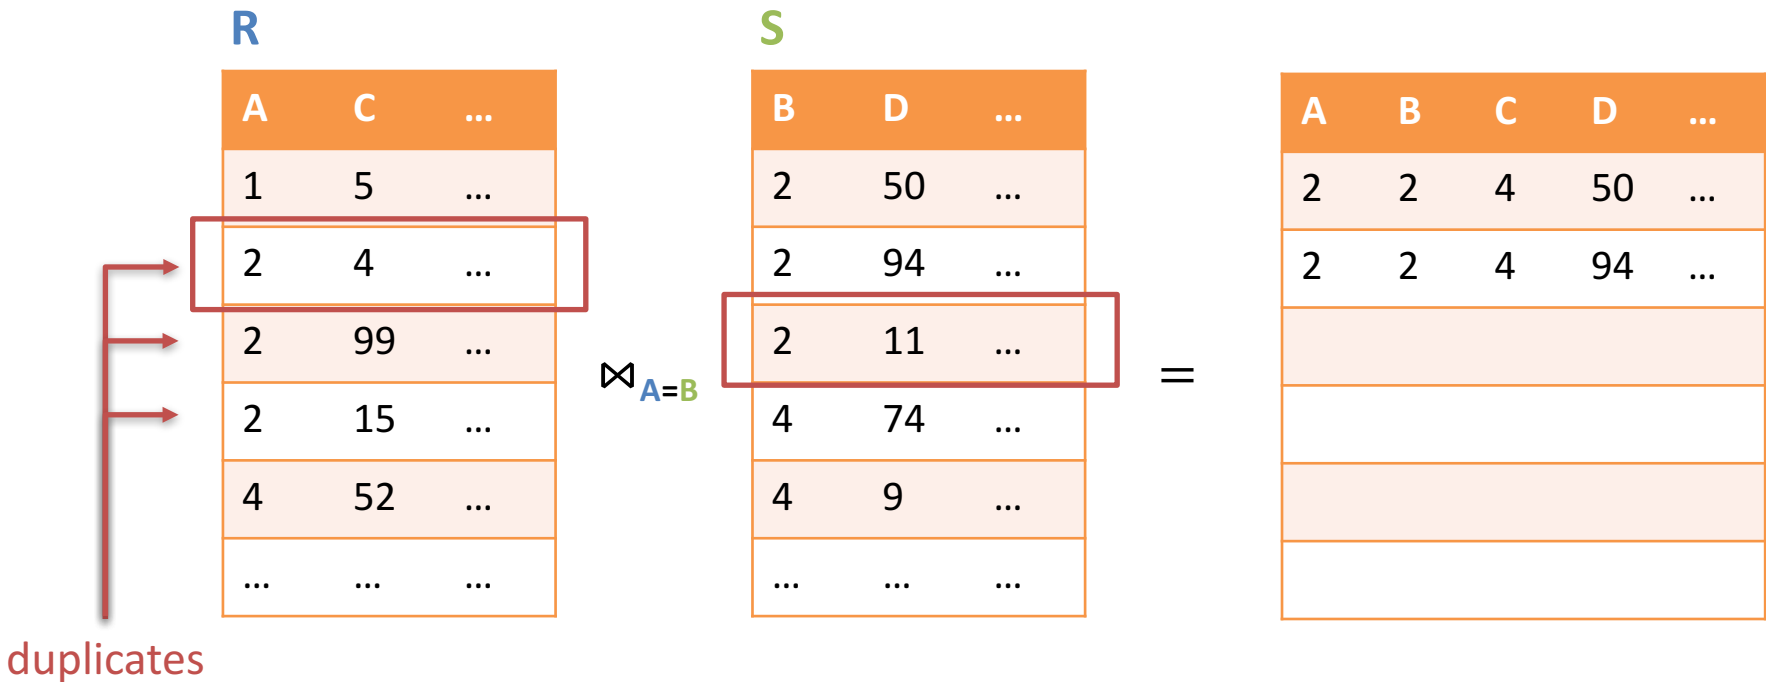


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

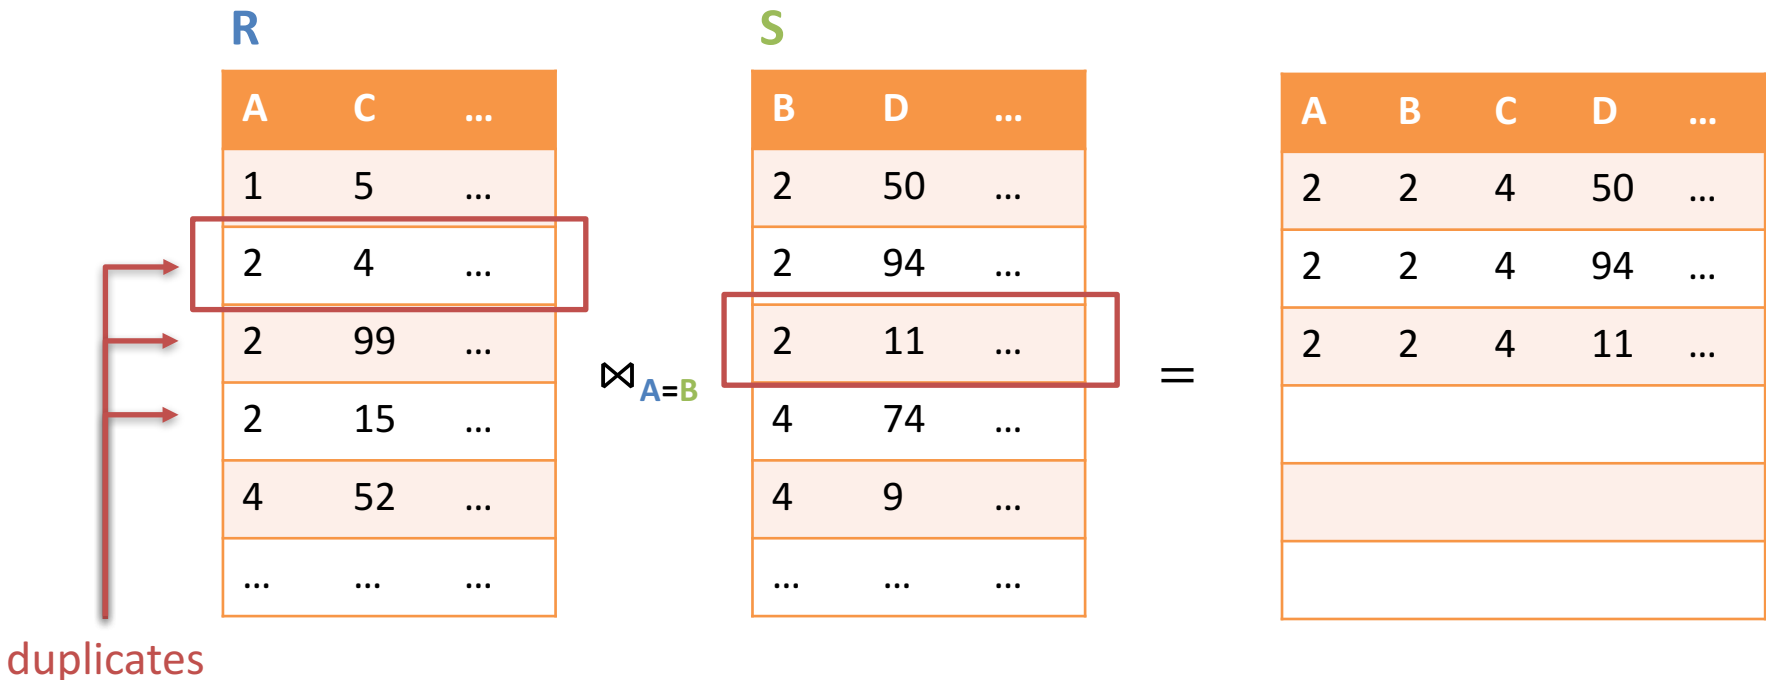


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

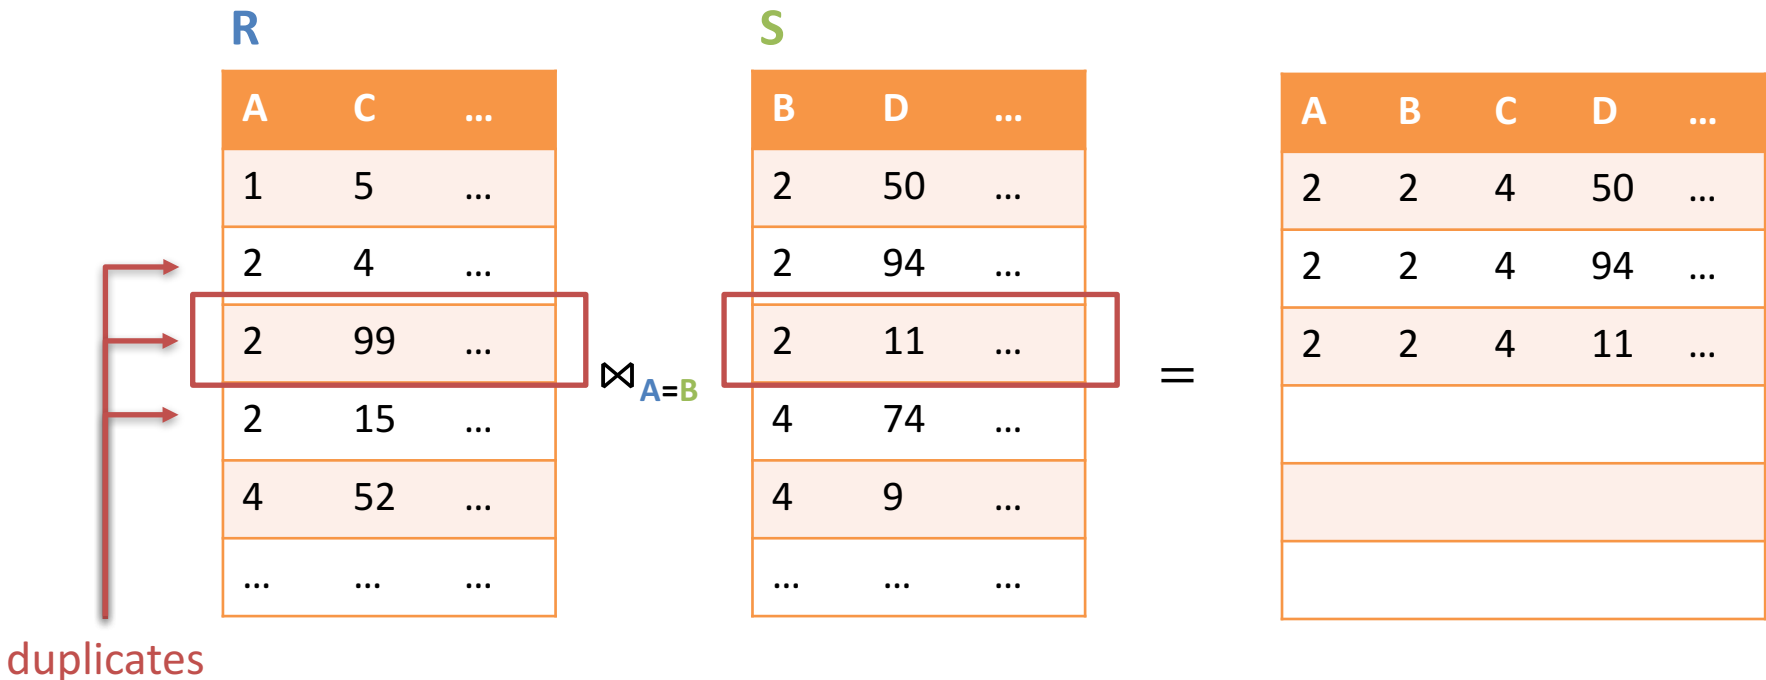


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$



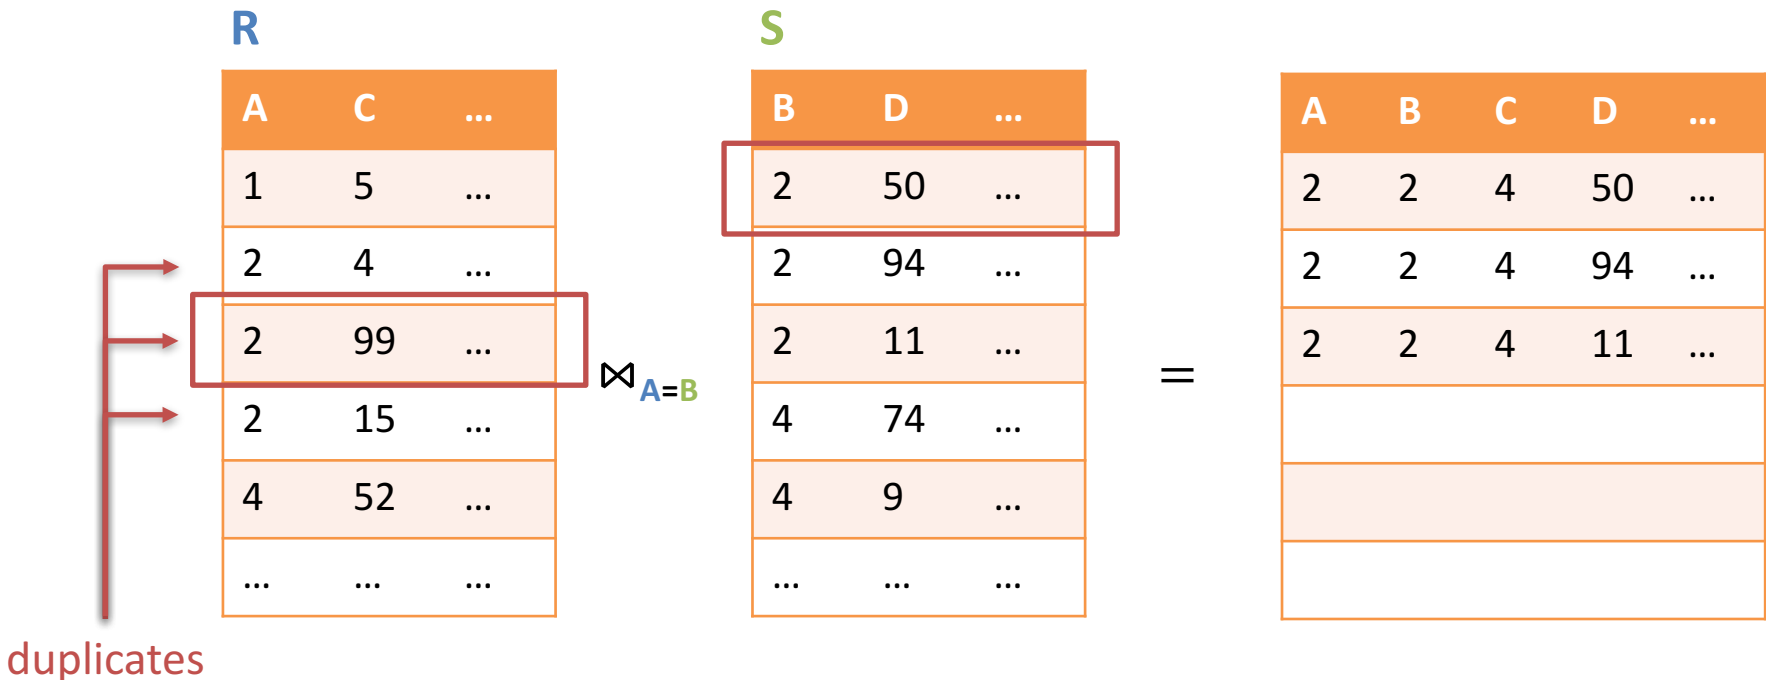
Remember tuples in  $S$  that match with the current value of  $A$



# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

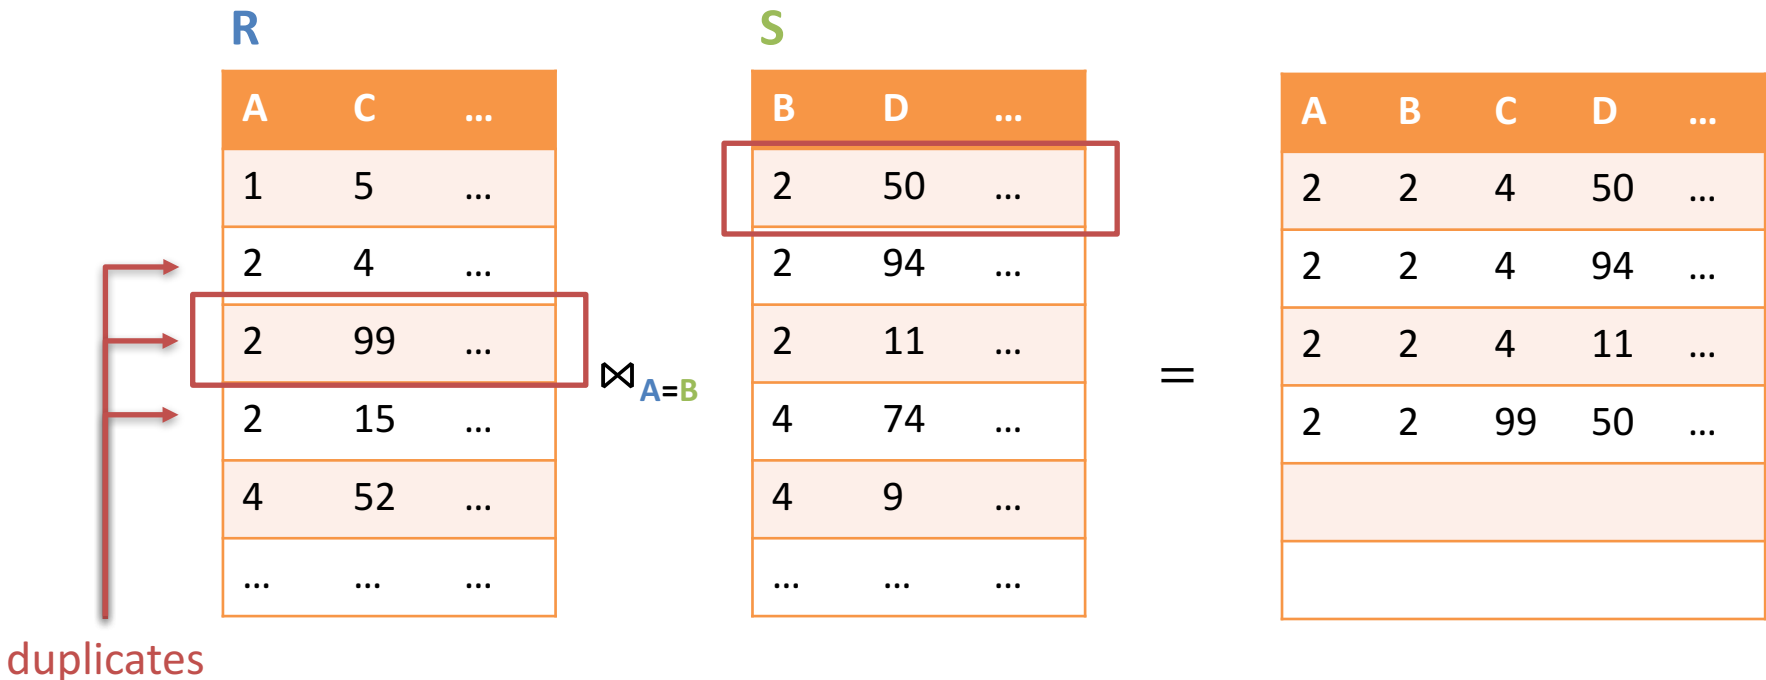


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

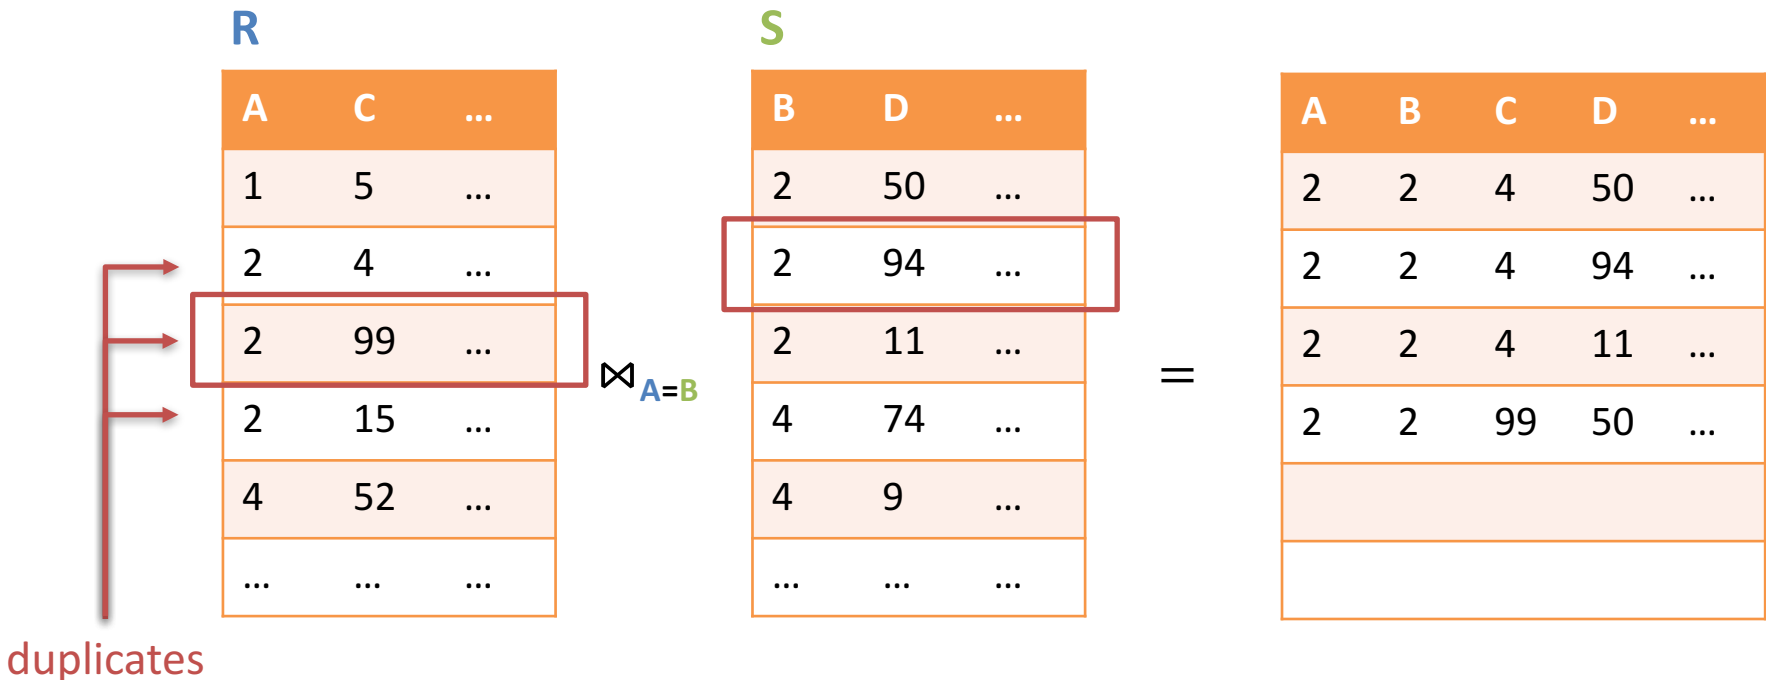


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

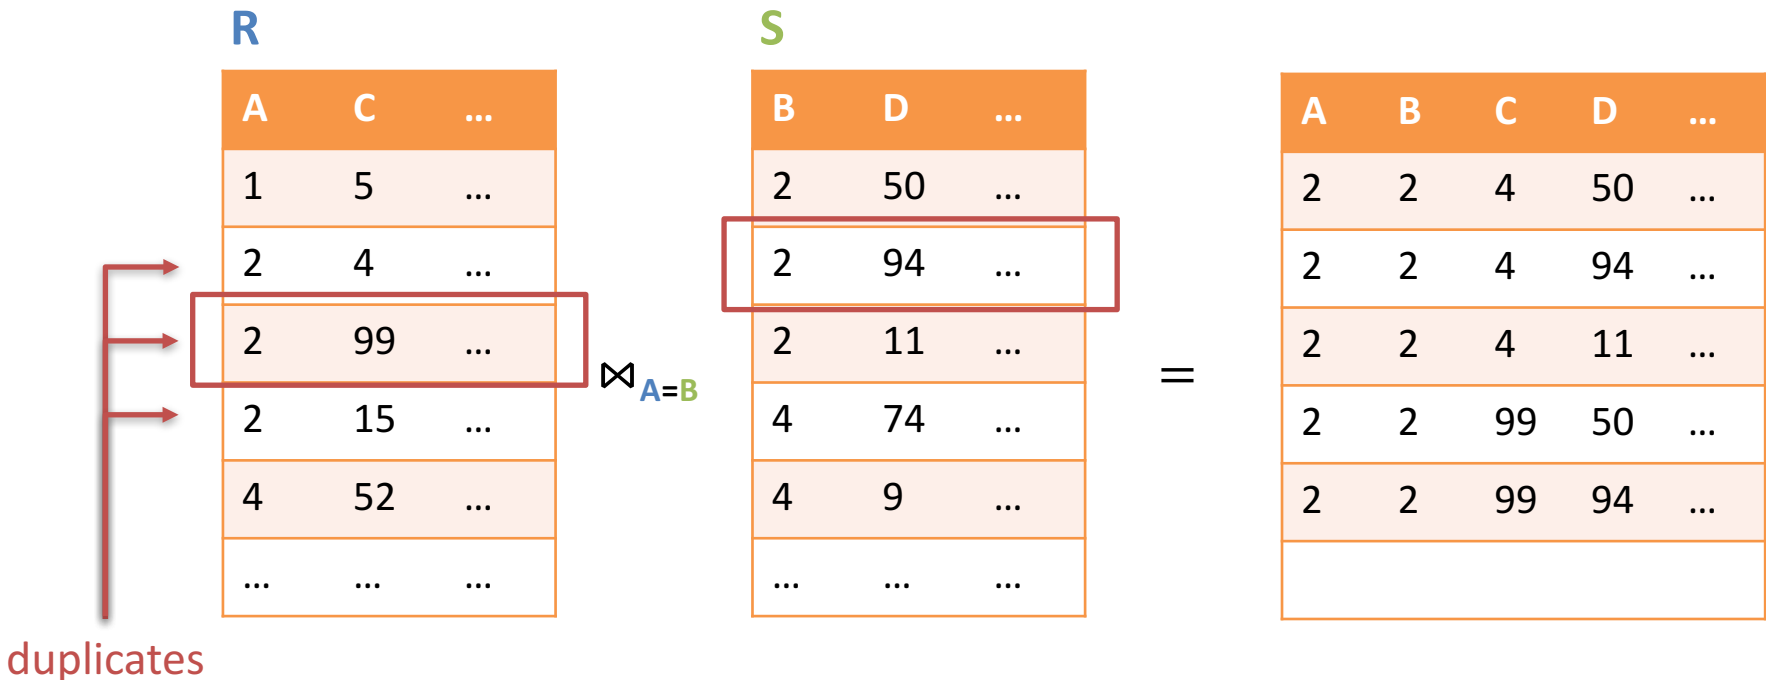


Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$



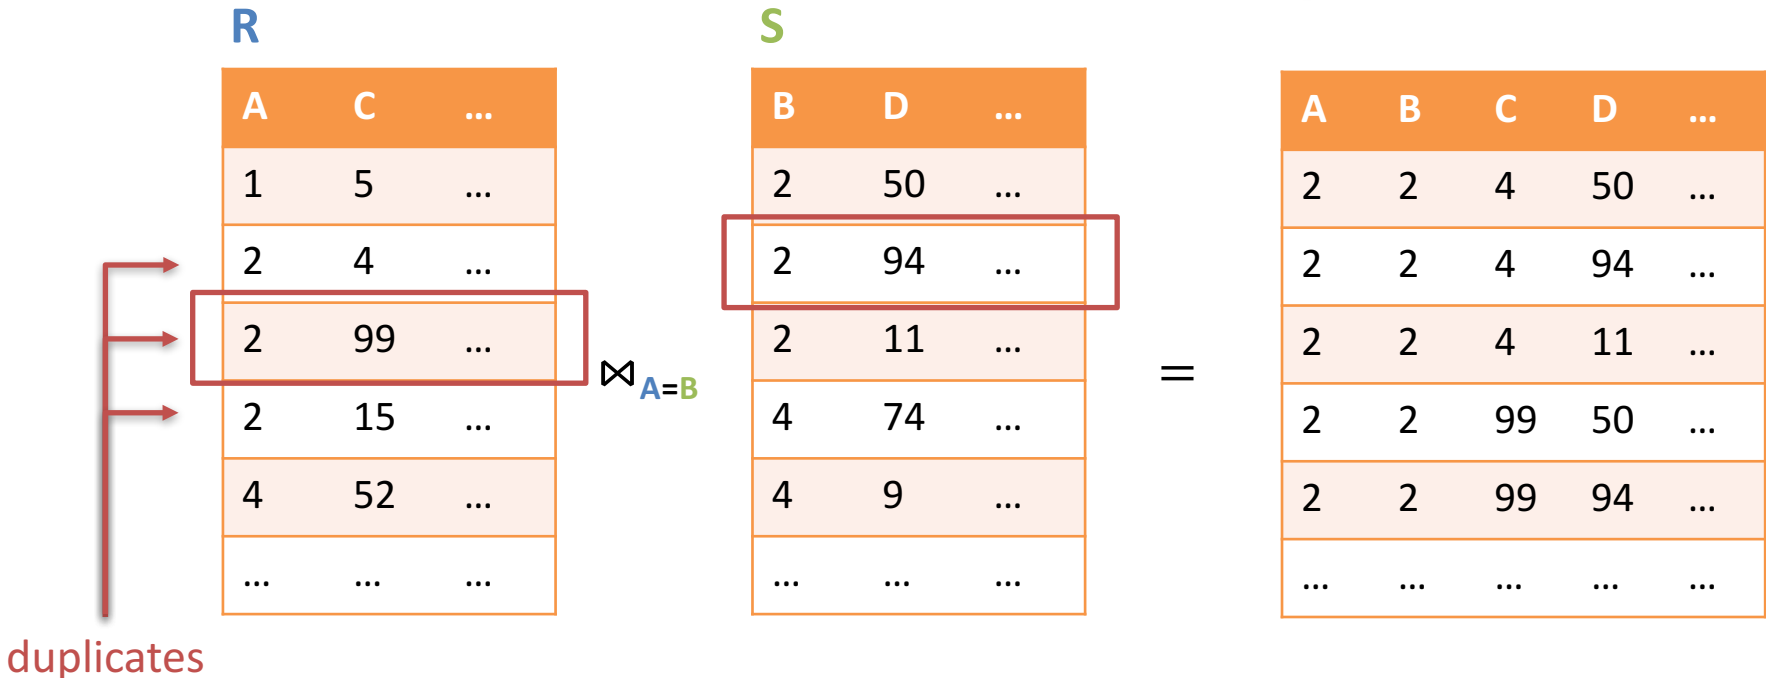
Remember tuples in  $S$  that match with the current value of  $A$

# Merging With Duplicates in Column A

- Goal: compute  $R \bowtie_{A=B} S$

Assume:  $R$  is sorted on  $A$  and  $S$  is sorted on  $B$

What is the running time?



Remember tuples in  $S$  that match with the current value of  $A$

# Faster Joins With Sorting

- Sort Join Algorithm:

**Compute  $R \bowtie_{A=B} S$ :**


1. Sort  $R$  on  $A$  Running time:  $O(|R| \times \log_2 |R|)$
2. Sort  $S$  on  $B$  Running time:  $O(|S| \times \log_2 |S|)$
3. Merge the sorted  $R$  and  $S$  Running time:  $O(|R| + |S|)$

- Typical running time:  $O(|R| \log_2 |R| + |S| \log_2 |S|)$ 
  - If not “too many” values in  $A$  occur multiple times
  - E.g., this is the case if  $A$  is a key
- Typically much faster than Nested Loop Join

# Remarks

- Various **join algorithms** in practice:
  - Index joins
  - Hash joins
  - Multiway joins: join more than two relations at once
- Can compute **other operations of relational algebra** using similar methods as those in this lecture
- **We've neglected that relations are stored on disk**

# Running Time vs Disk Accesses



Relevant parameters:

- **B** = size of a disk block (typically 512→4096 bytes)
- **M** = number of disk blocks that fit into available RAM

Algorithm	No. of elementary operations	No. of disk accesses
Reading a relation <b>R</b>	$O( R )$	$O\left(\frac{ R }{B}\right)$
Sorting <b>R</b> on attribute <b>A</b>	$O( R  \log_2  R )$	$O\left(\frac{ R }{B} \log_M \frac{ R }{B}\right)$

External memory merge sort



# Summary

- Query plans are evaluated bottom-up – from the leaves to the root
- Each operator can be computed in different ways
  - Selection: e.g., linear scans (reading the relation once)
  - Projection: linear scans
  - Joins: Nested Loop Join, Sort Join, Index Join, Hash Join, ...
- Next lecture: faster query processing with indexes



## Can you help us?

Please complete this survey to help with our research. Your answers will help us improve our service for you.

Use this short URL or the QR code

<https://tinyurl.com/y5cot2aj>



Open the camera on your phone to use the QR code