

COMP207


Database Development

Lecture 14

Query Processing: Faster With Indexes

Reminder: Query Plans


- “Recipe” for computing the result to a query



- Executed bottom-up
- Typically several algorithms for computing operators

Reminder: Computing $\sigma_{\text{condition}}(\mathbf{R})$

- Basic procedure:



for each tuple \mathbf{t} in \mathbf{R} :
if \mathbf{t} satisfies condition:
output \mathbf{t}

Is there a faster way?

- Needs to read the entire relation
- Yes, sometimes!

Example

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...

- Selection can be performed faster if we know **where to find the rows for 'G401'**
- Two solutions: **sorting & index**

External sorting (not in exam)

Merge Sort

- (Internal memory) merge sort:

Merge

Divide input in two parts P_1 , P_2

Sort P_1 & P_2 **recursively**

While P_1 or P_2 is not empty:

Add smallest remaining
element from P_1 or P_2
to output

External Merge Sort

- External merge sort:

Number of disk blocks that fit in RAM

Merge

Divide input in **M** parts P_1, P_2, \dots, P_M


Sort P_1, P_2, \dots, P_M **recursively**

While not all P_i are empty:

Add smallest remaining
element from any part P_i to
output

Merge in more details

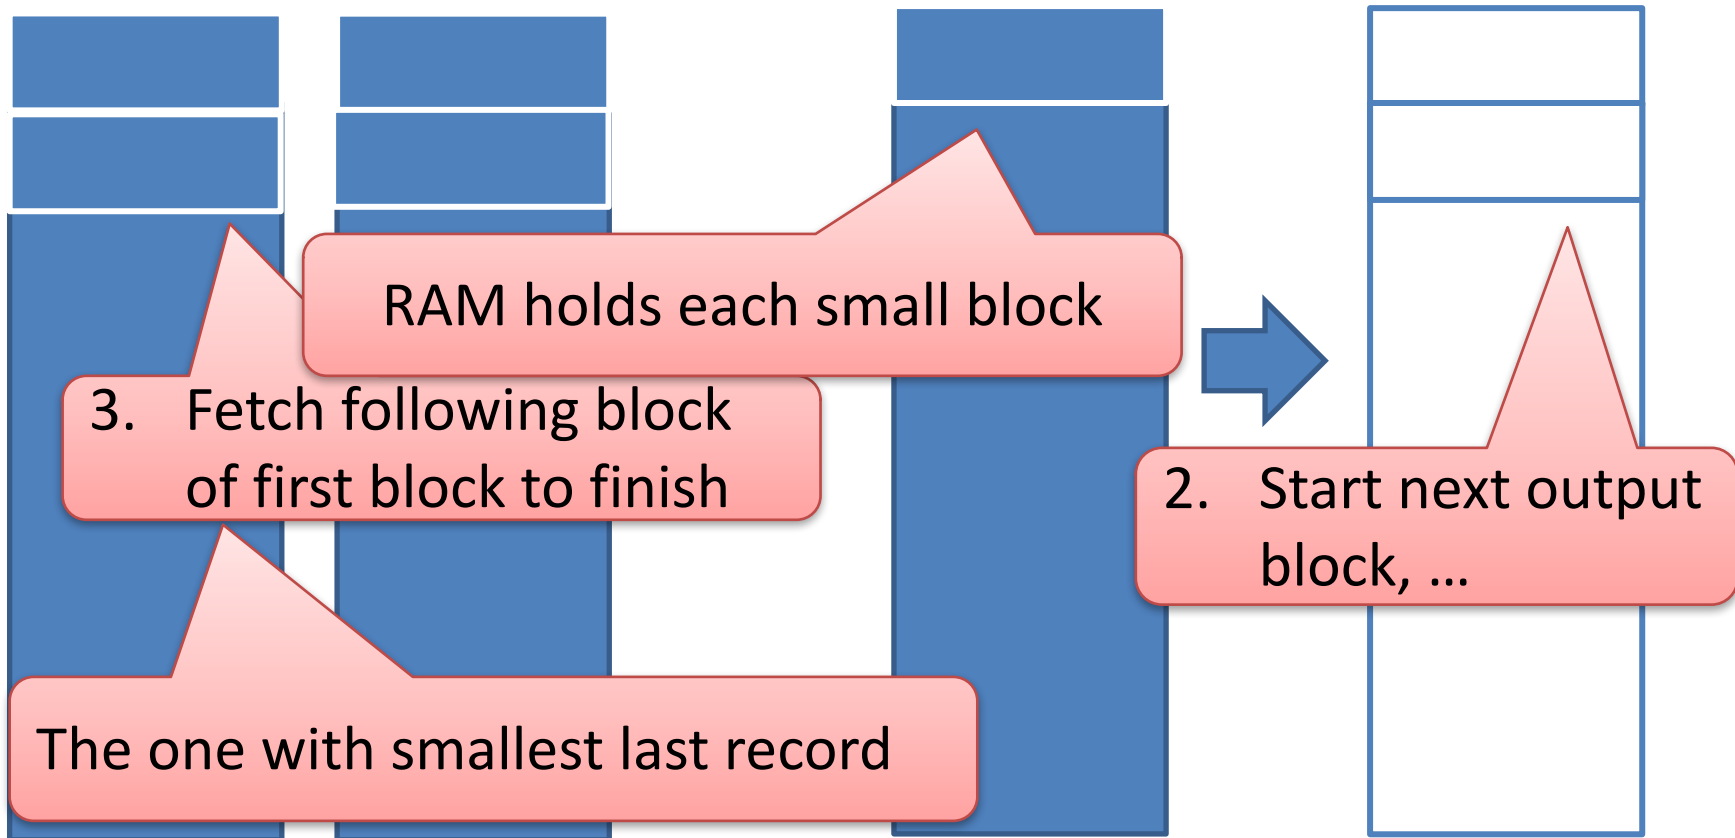
- Goal: Must merge M (perhaps large) parts into one



Merge in more details

- Goal: Must merge M (perhaps large) files

1. Move to disk, ...



Merge Sort

- We split in M buckets in each level of recursion until remainder is below MB (where it fits in RAM and can be sorted directly): $\log_M(N/(MB))$ levels

Divide input in M parts P_1, P_2, \dots, P_M
Sort P_1, P_2, \dots, P_M **recursively**

Total disk operations:
 $O(N/B \log_M(N/(MB)))$

not all P_i are empty:

Add smallest remaining element from any part P_i to output

On each level of recursion we scan through all blocks once
= $O(N/B)$ disk operations,
where B is block size

Comparison

- In practice, since harddisks are slow:
 - Running time is basically time spend on disk operations
 - Thus, time is $O(N/B \log_M(N/(MB)) \times D)$
 - D is time for a disk operation
- Quicksort (or other internal memory sorting algorithms) uses $O(N/B \log_2(N) \times D)$ time!

Numeric example

- Want to sort $N=4$ TB in 4 MB of RAM with $B=4096$

External Merge Sort:

- $M = 4 \text{ MB} / 4096 = 1024$
- $\log_M(N/(M B))$ is then $\approx \log_{1024}(1024^2) = 2$
- We are thus using 2 recursive calls, for 3 scans in total

(Internal) Quick Sort:

- $\log_2(N)=42$
- We are thus using 42 recursive calls, for 43 scans in total

Assuming we are lucky about
pivot elements

Example

$\sigma_{\text{programme}='G401'}(\text{Students})$

Students

id	name	programme
...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...

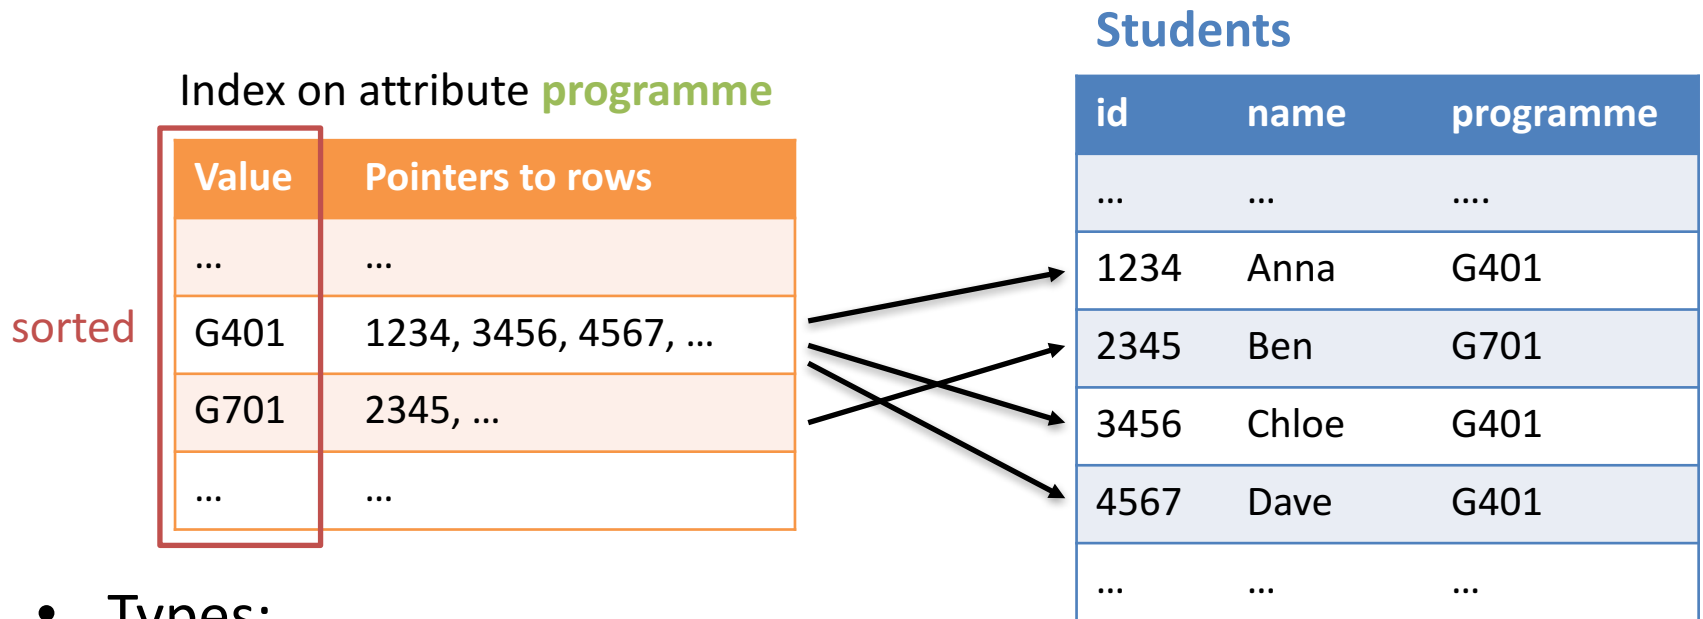
- Selection can be performed faster if we know **where to find the rows for 'G401'**
- Two solutions: **sorting & index**

Indexes

(Basics: COMP102/CSE103)

Index

- Takes the values for one or more attributes of a relation **R**, provides quick access to the tuples with these values



- Types:
 - Secondary: merely points to location of records on disk
 - Primary: in addition, defines how data is sorted on disk

Always dense

Good when attributes
involve primary key

Example Revisited

$\sigma_{\text{programme}=\text{'G401'}}(\text{Students})$

Index on attribute **programme**

Value	Pointers to rows
...	...
G401	1234, 3456, 4567, ...
G701	2345
...	...

Students

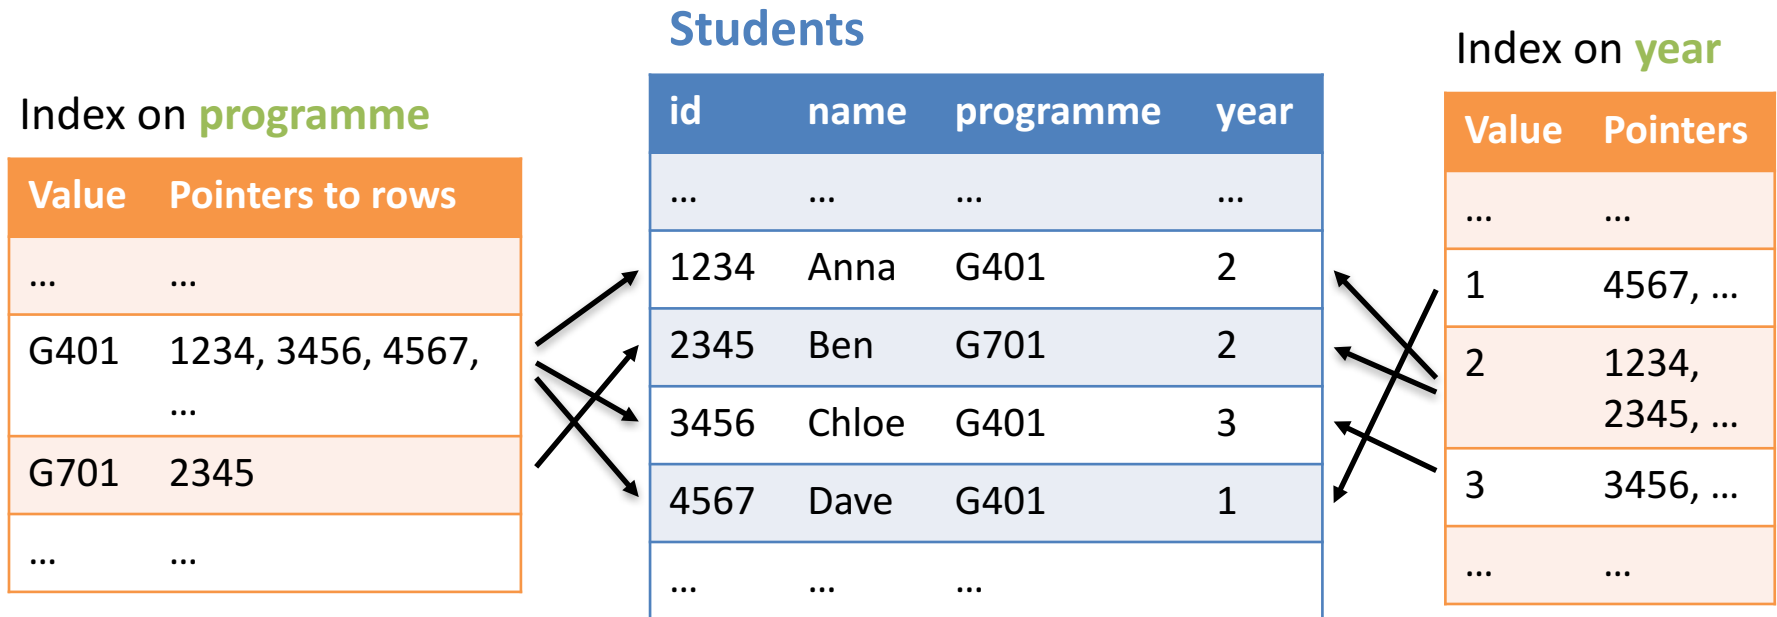
id	name	programme
...
1234	Anna	G401
2345	Ben	G701
3456	Chloe	G401
4567	Dave	G401
...

- Selection with index:
 - Find entry for **'G401'** in index
 - Visit all rows in **Students** whose ids occur in the index entry for **'G401'**

Running time?

Example 2

$\sigma_{\text{programme}='G401' \text{ AND } \text{year}=2}(\text{Students})$



- Selection with two indexes:
 - Find entries for 'G401' & 2 in indexes for programme & year
 - Visit all rows in Students whose ids occur in both index entries

Exercise

Students(id, name, programme, year)

Modules(module_code, module_title, programme, year)

Which indexes could be useful to answer these queries:

- $\pi_{id,name}(\sigma_{year=1}(\text{Students}))$
- $\text{Students} \bowtie \text{Modules}$

Forms of Indexes

- **B+ Trees**



- Good if selection condition specifies a range
- E.g., $\sigma_{\text{programme}='G401' \text{ AND year} > 1}$
- Most widely used

```
CREATE INDEX
  ON Students USING btree
  (programme, year);
```

default

- **Hash tables**

- Good if selection involves equality only
- E.g., $\sigma_{\text{programme}='G401'}$

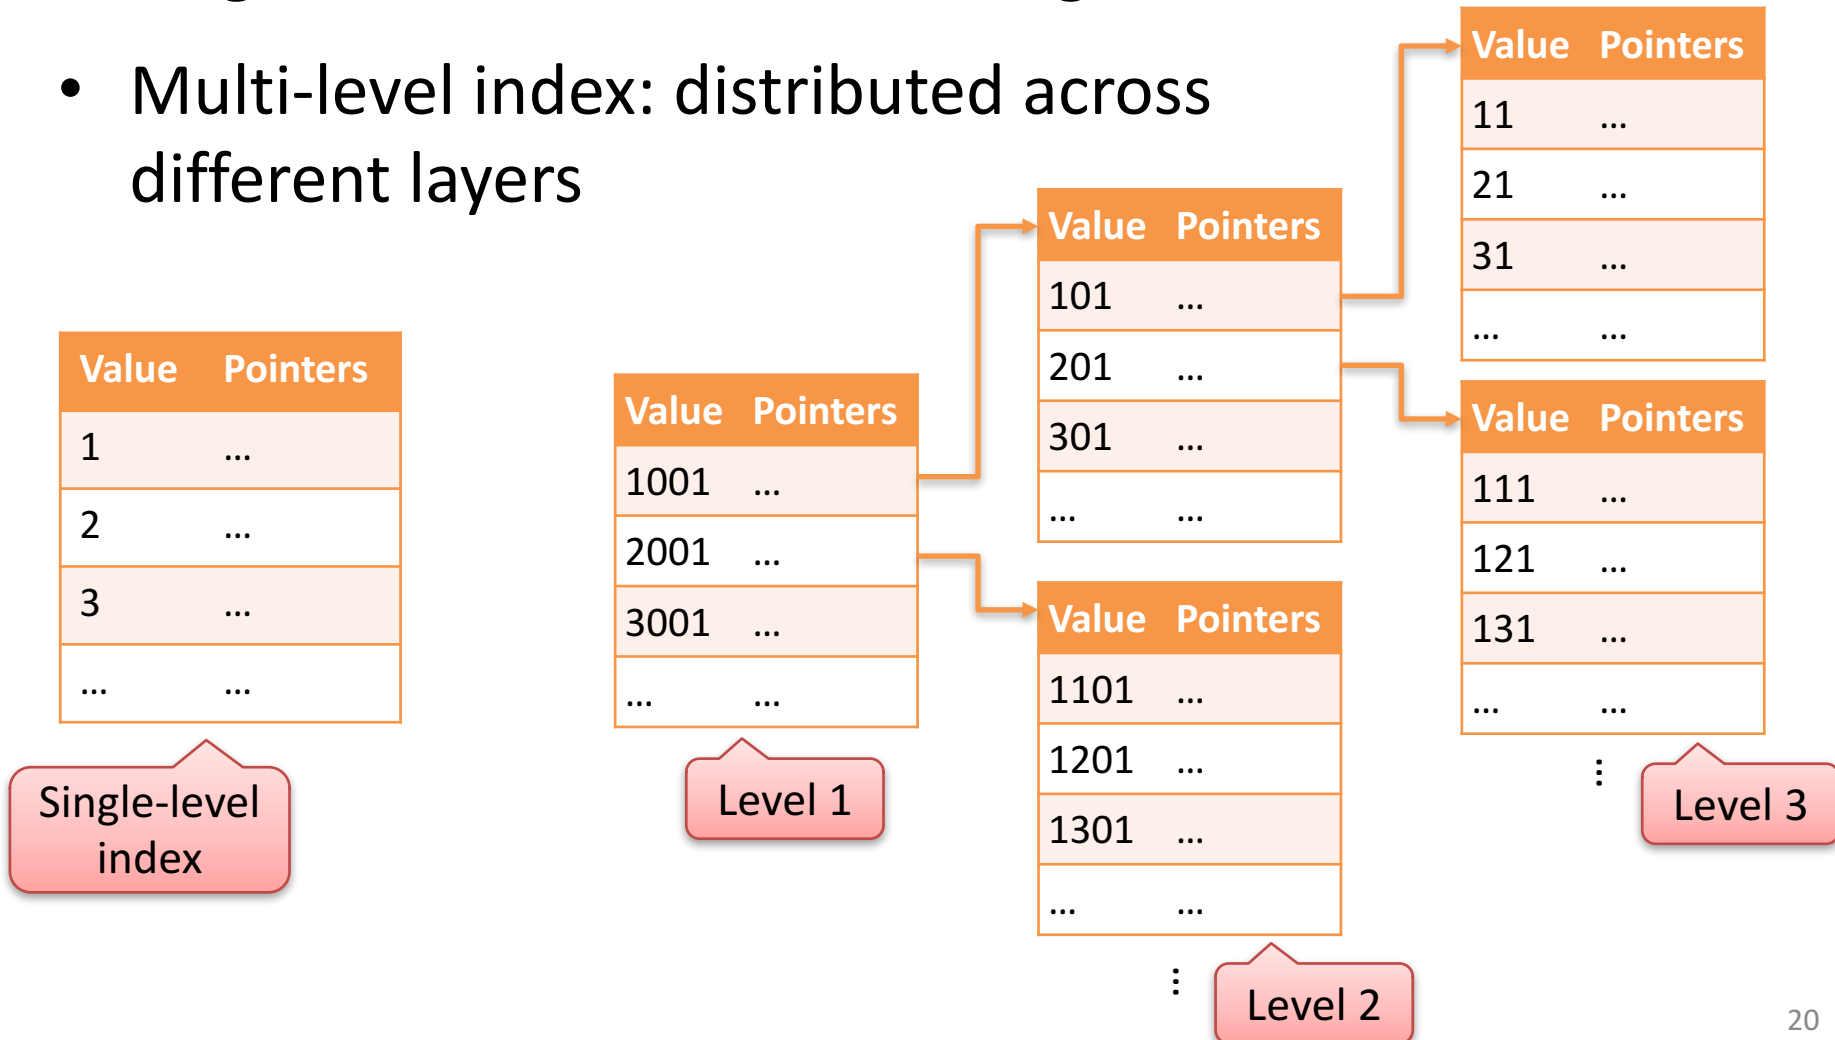
```
CREATE INDEX
  ON Students USING hash
  (programme);
```

```
CREATE INDEX
  ON Students USING hash
  (lower(name));
```

- Many more...

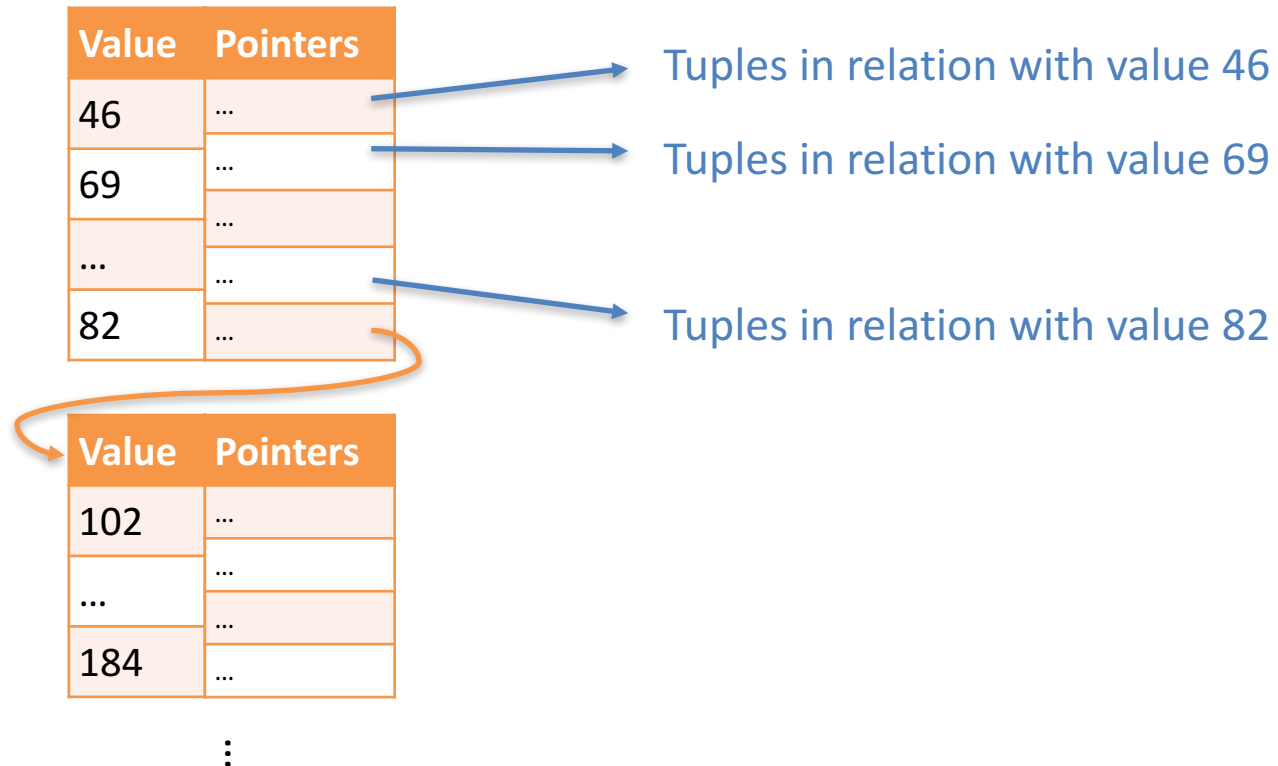
Single Level vs Multi-Level Indexes

- Single level index: stored in single list
- Multi-level index: distributed across different layers

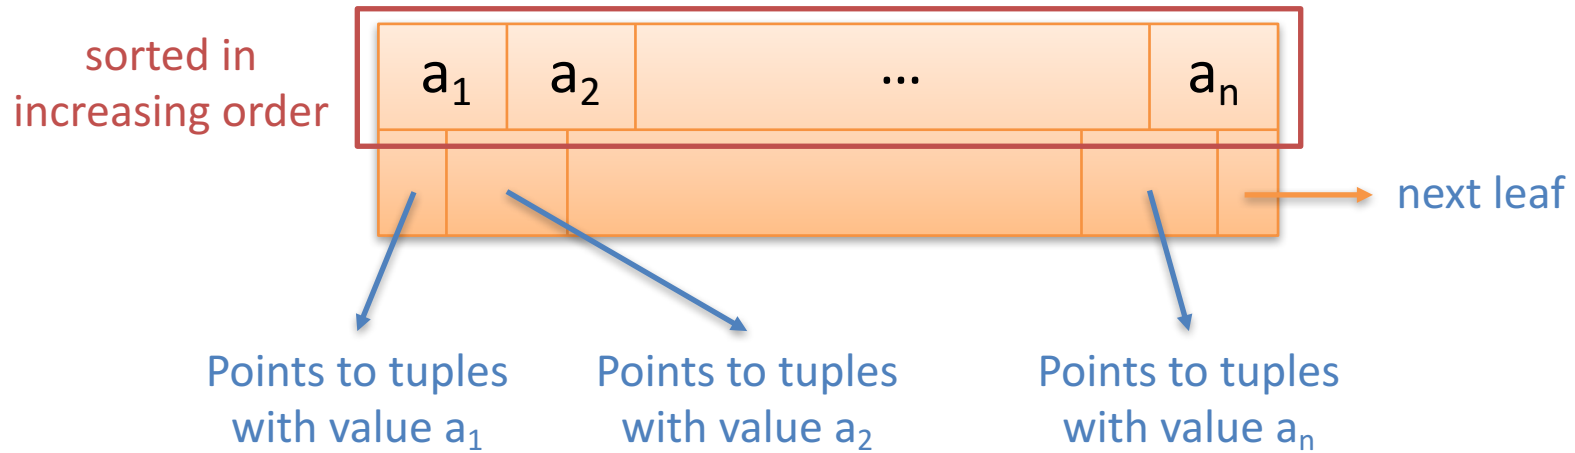


B+ Trees

- A multi-level index, but different shape than the one shown on the previous slide
- E.g., shape of a leaf in a B+ tree:

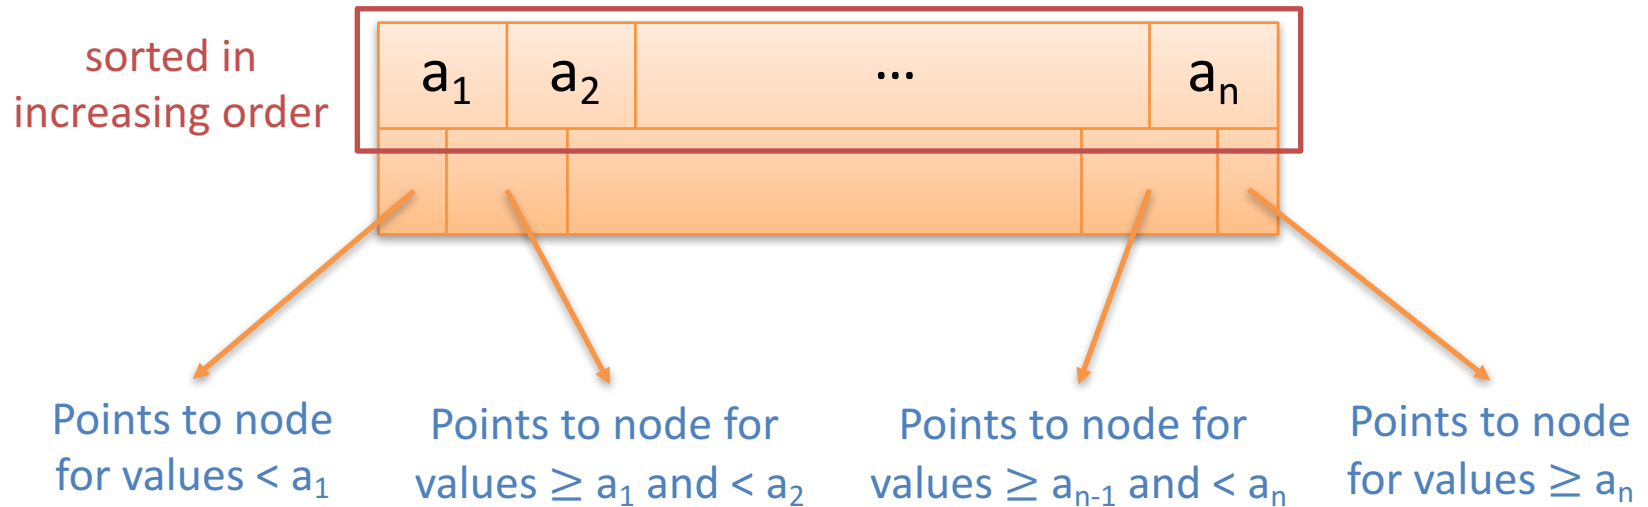


B+ Tree: Leaves (Idea)



- n = chosen such that node fits into a single disk block
 - Example:
 - Disk block size = 512 byte
 - Values: 4 byte integers
 - Pointers: 8 bytes
- } $n = 42$


B+ Tree: Inner Nodes (Idea)



- Pointers point to B+ tree nodes at level below
- n = chosen as before (≥ 1), so there are ≥ 2 pointers

B+ Tree: Leaves (Actually)

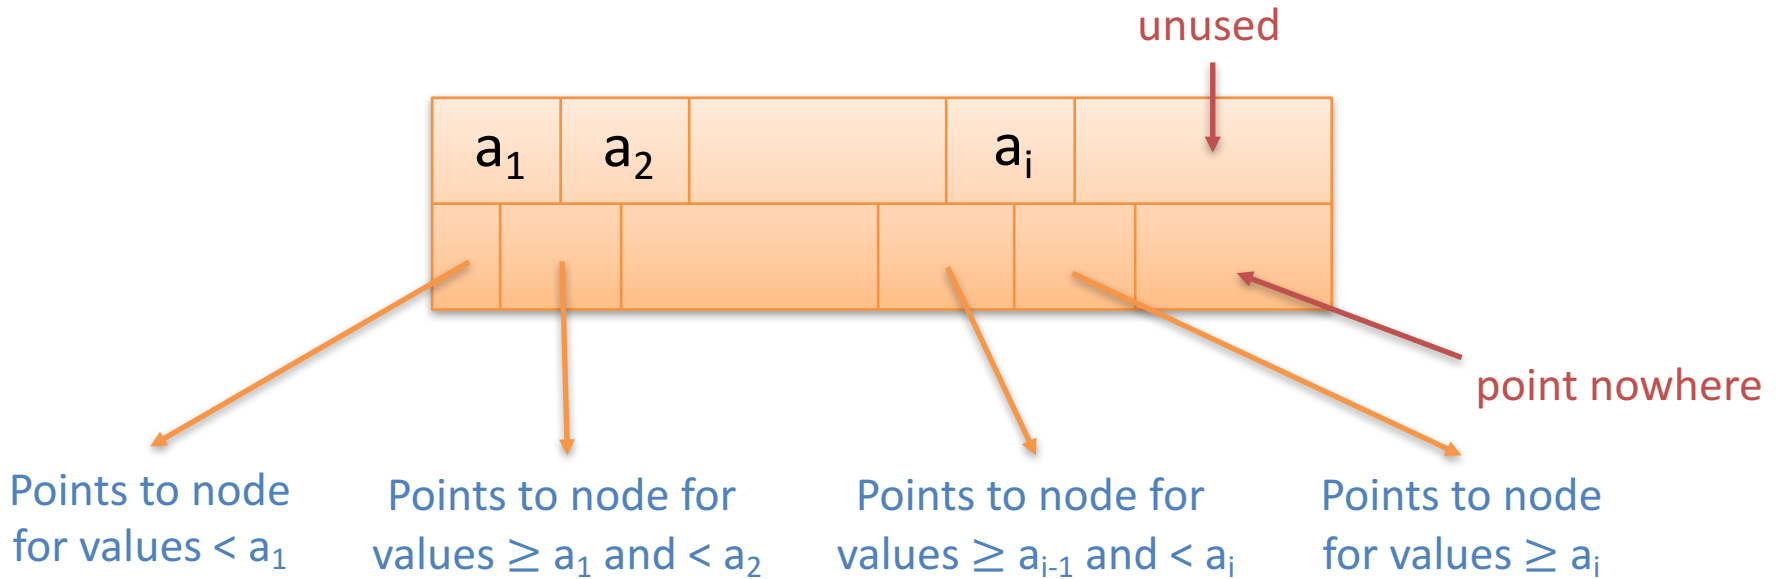
- Not all of the fields have to be used
- Fields are filled from left to right



- Ensure that at least $\left\lfloor \frac{n+1}{2} \right\rfloor$ pointers are used

B+ Tree: Inner Nodes (Actually)

- Similar as for leaves:



- Ensure that at least $\left\lceil \frac{n+1}{2} \right\rceil$ pointers are used
- Exception: root must use ≥ 2 pointers

To be continued...