

COMP207

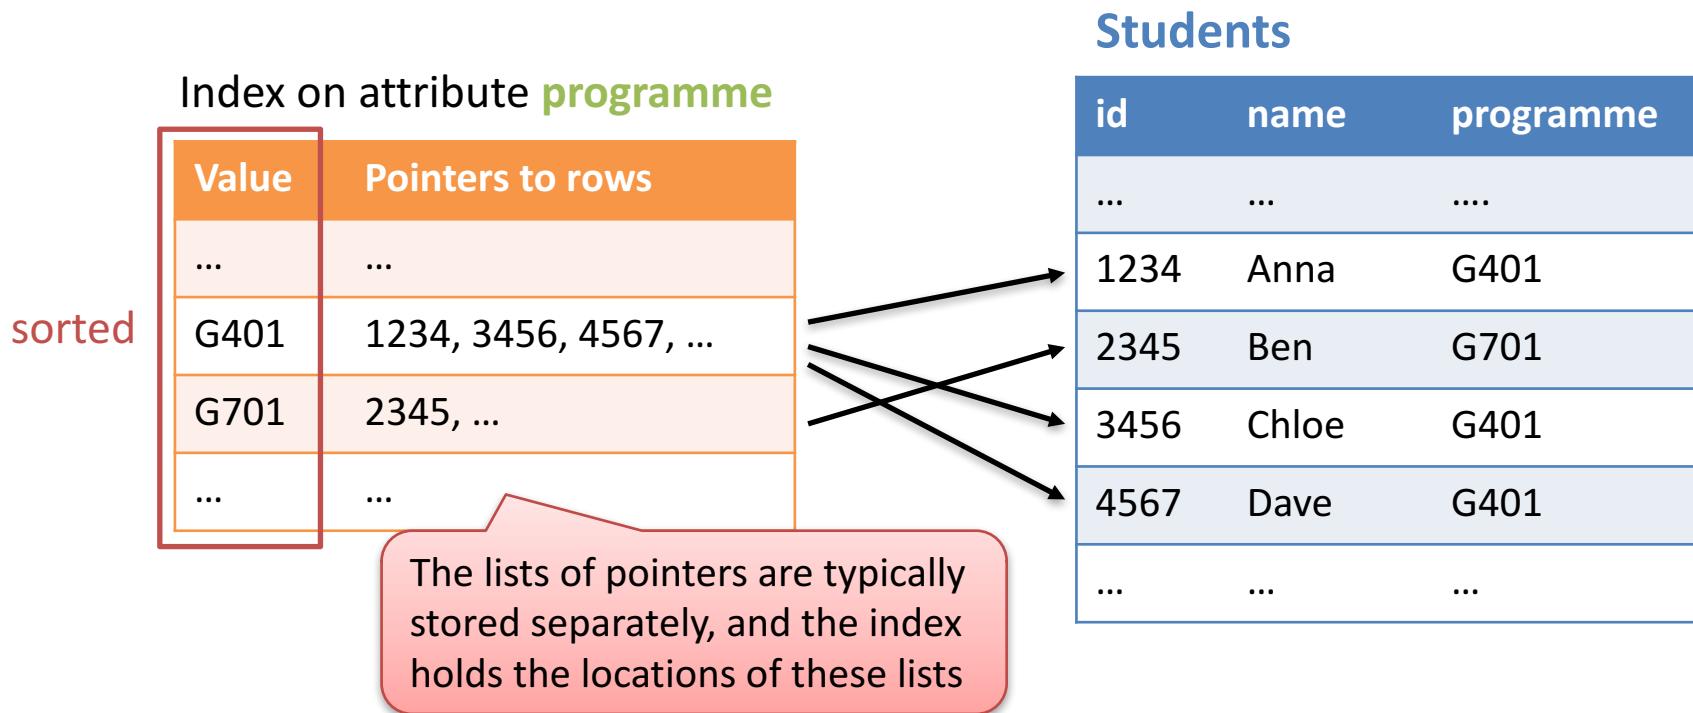
Database Development

Lecture 15

Query Processing:
B+ Tree Indexes

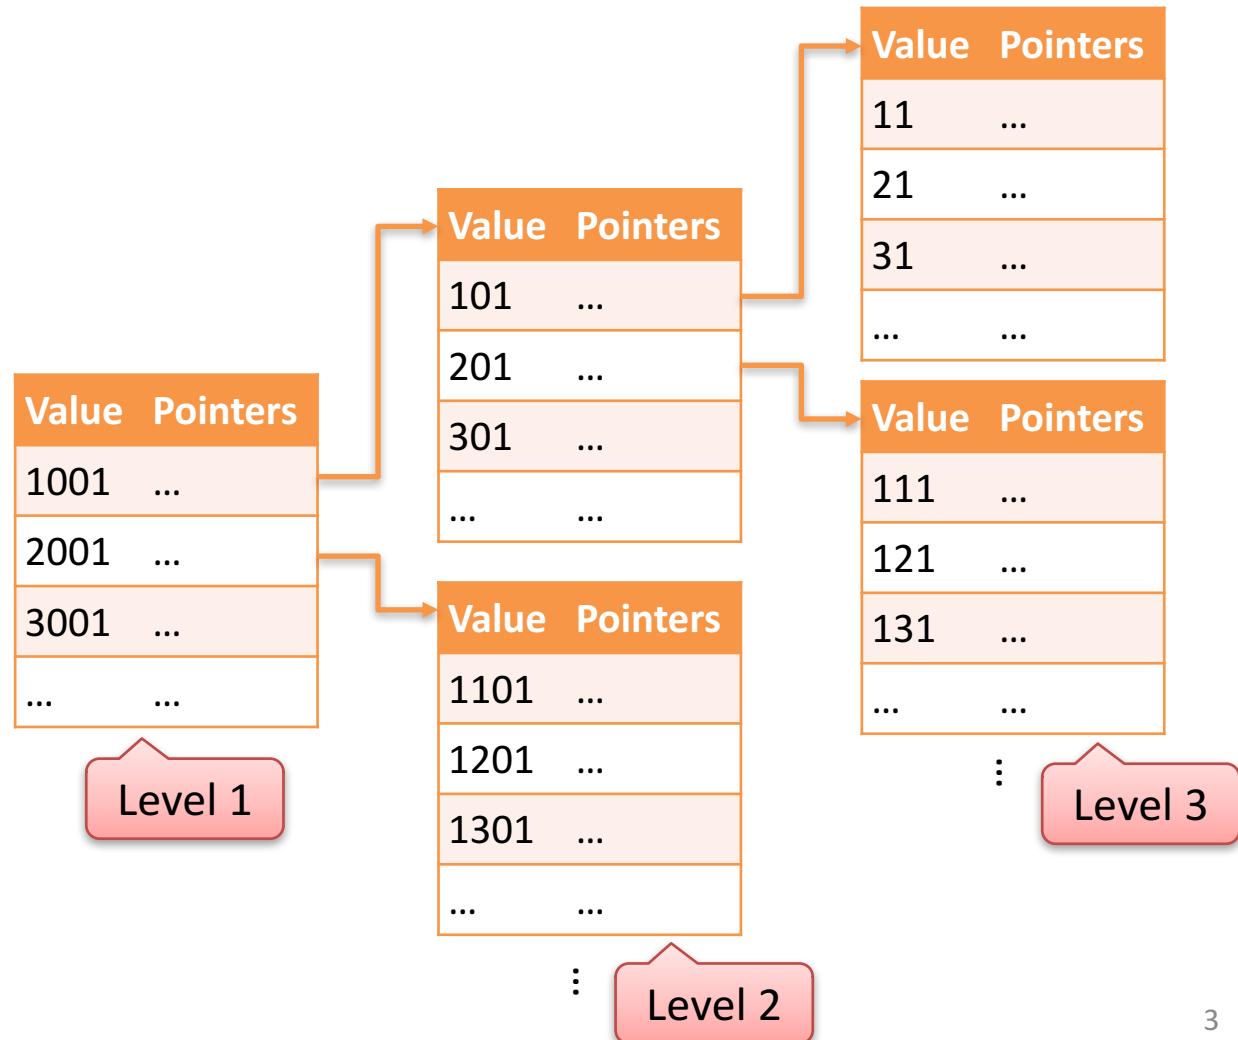
Review: Index

- Takes the values for one or more attributes of a relation **R**, provides quick access to the tuples in **R** with these values



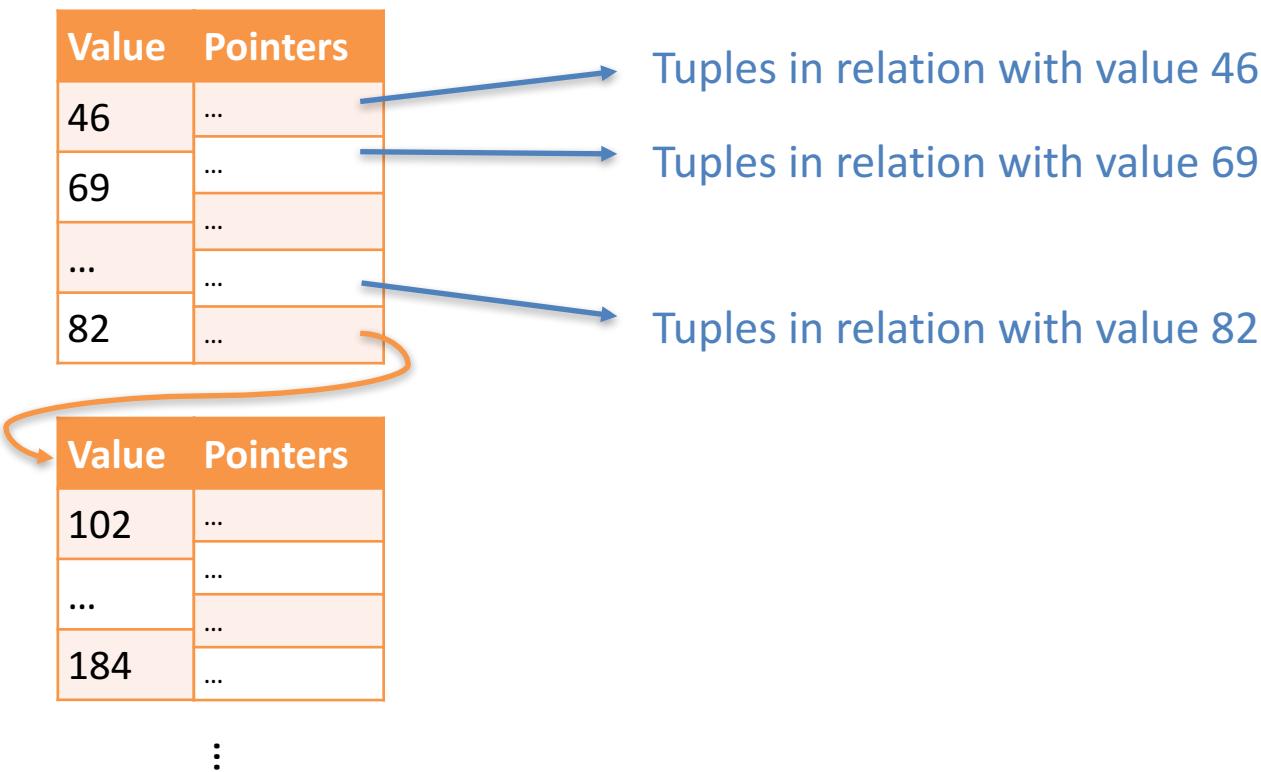
Reminder: Multi-Level Indexes

- Multi-level index: index distributed across different layers

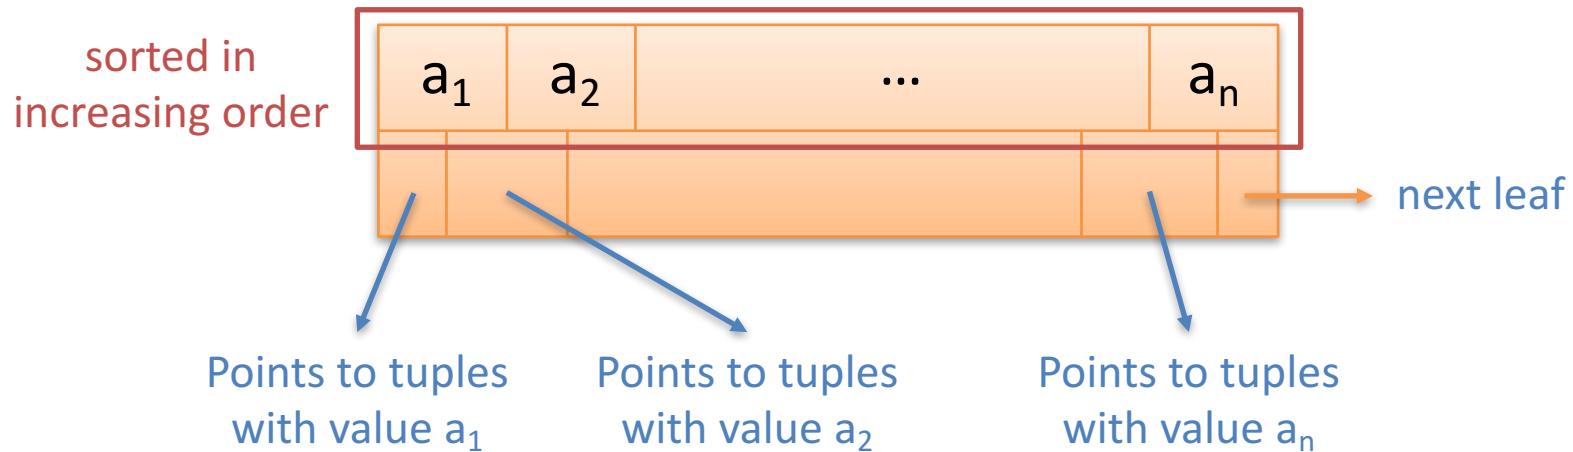


B+ Trees

- A multi-level index, but different shape than the one shown on last slide
- E.g., shape of a leaf in a B+ tree:

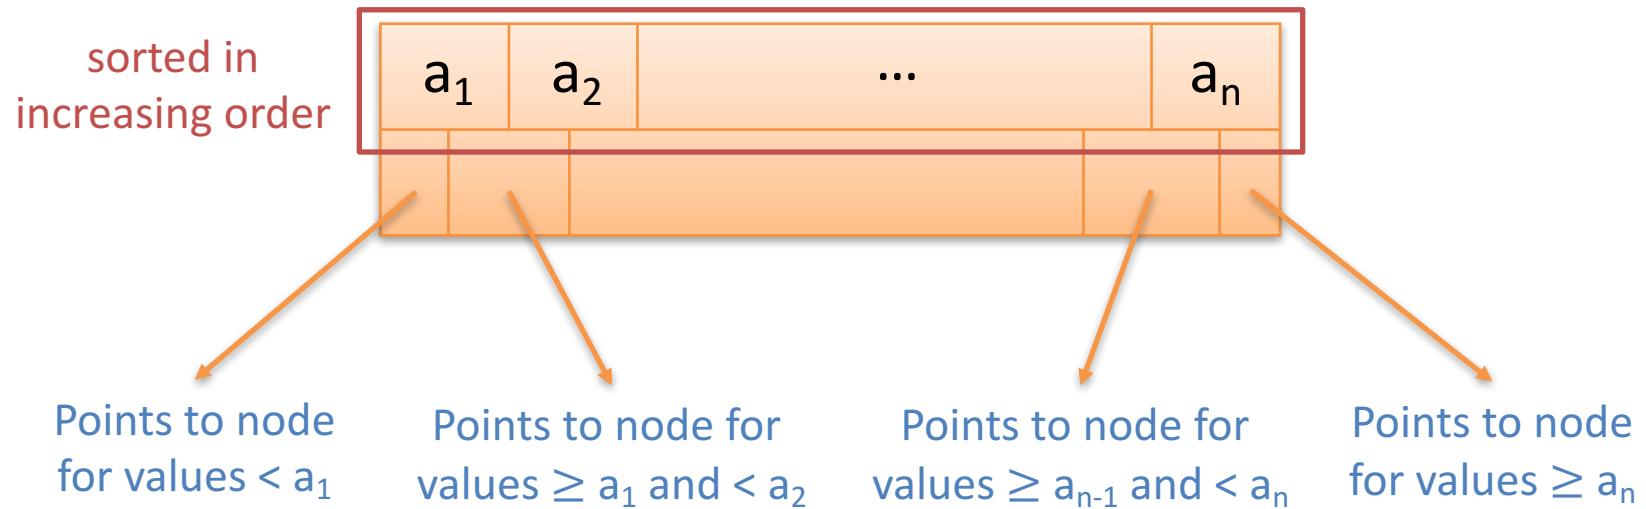


B+ Tree: Leaves (Idea)



- n = chosen such that node fits into a single disk block
 - Example:
 - Disk block size = 512 byte
 - Values: 4 byte integers
 - Pointers: 8 bytes
- $\left. \begin{array}{l} \\ \\ \end{array} \right\} n = 42$
 $(512 \geq 42(8+4)+8)$
 $512 < 43(8+4)+8)$

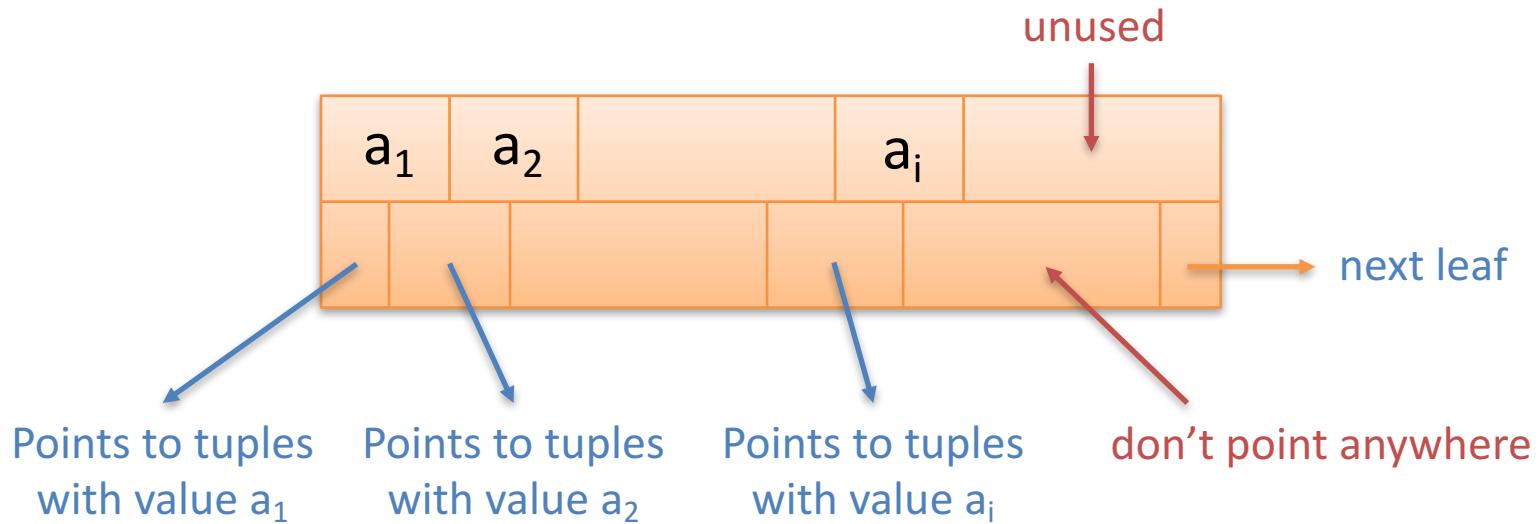
B+ Tree: Inner Nodes (Idea)



- Pointers point to B+ tree nodes at level below
- n = chosen as before (≥ 1), so there are ≥ 2 pointers

B+ Tree: Leaves (Actually)

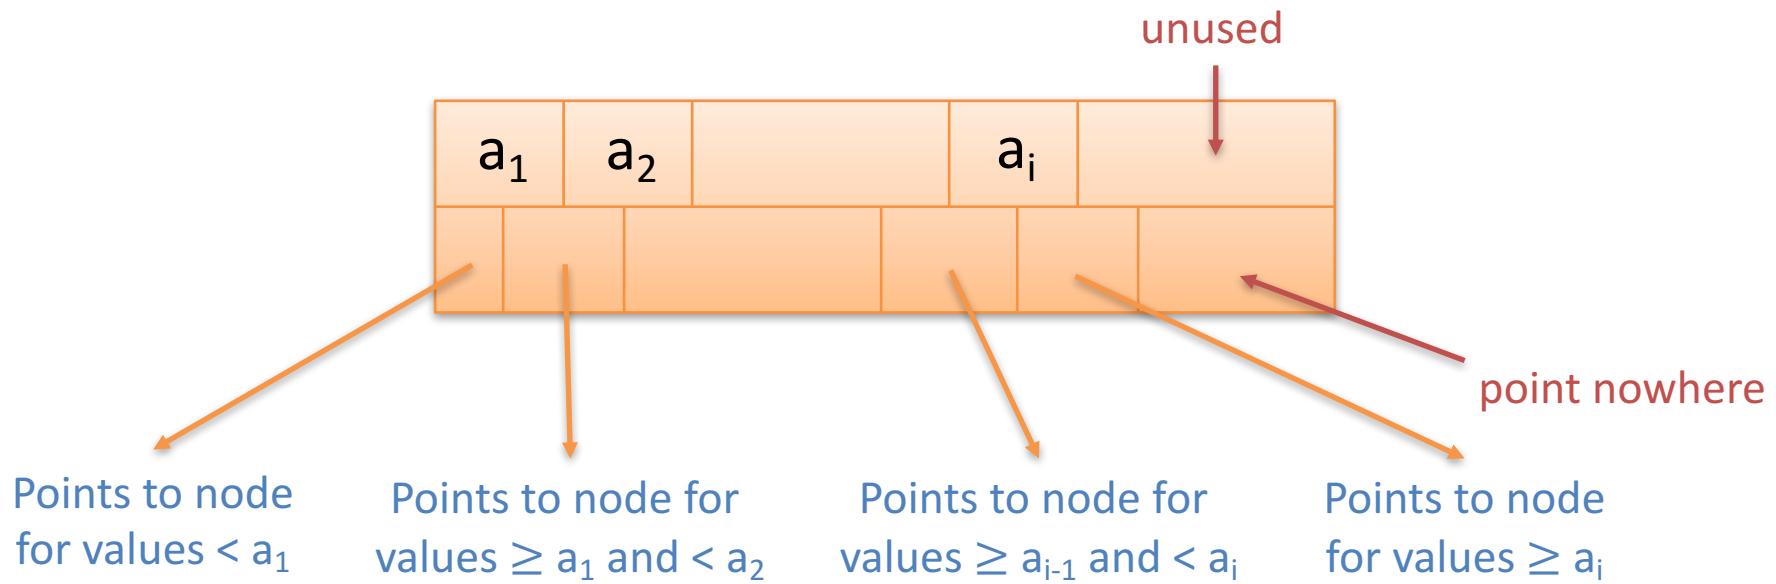
- Not all of the fields have to be used
- Fields are filled from left to right



- Ensure that at least $\left\lfloor \frac{n+1}{2} \right\rfloor$ pointers are used

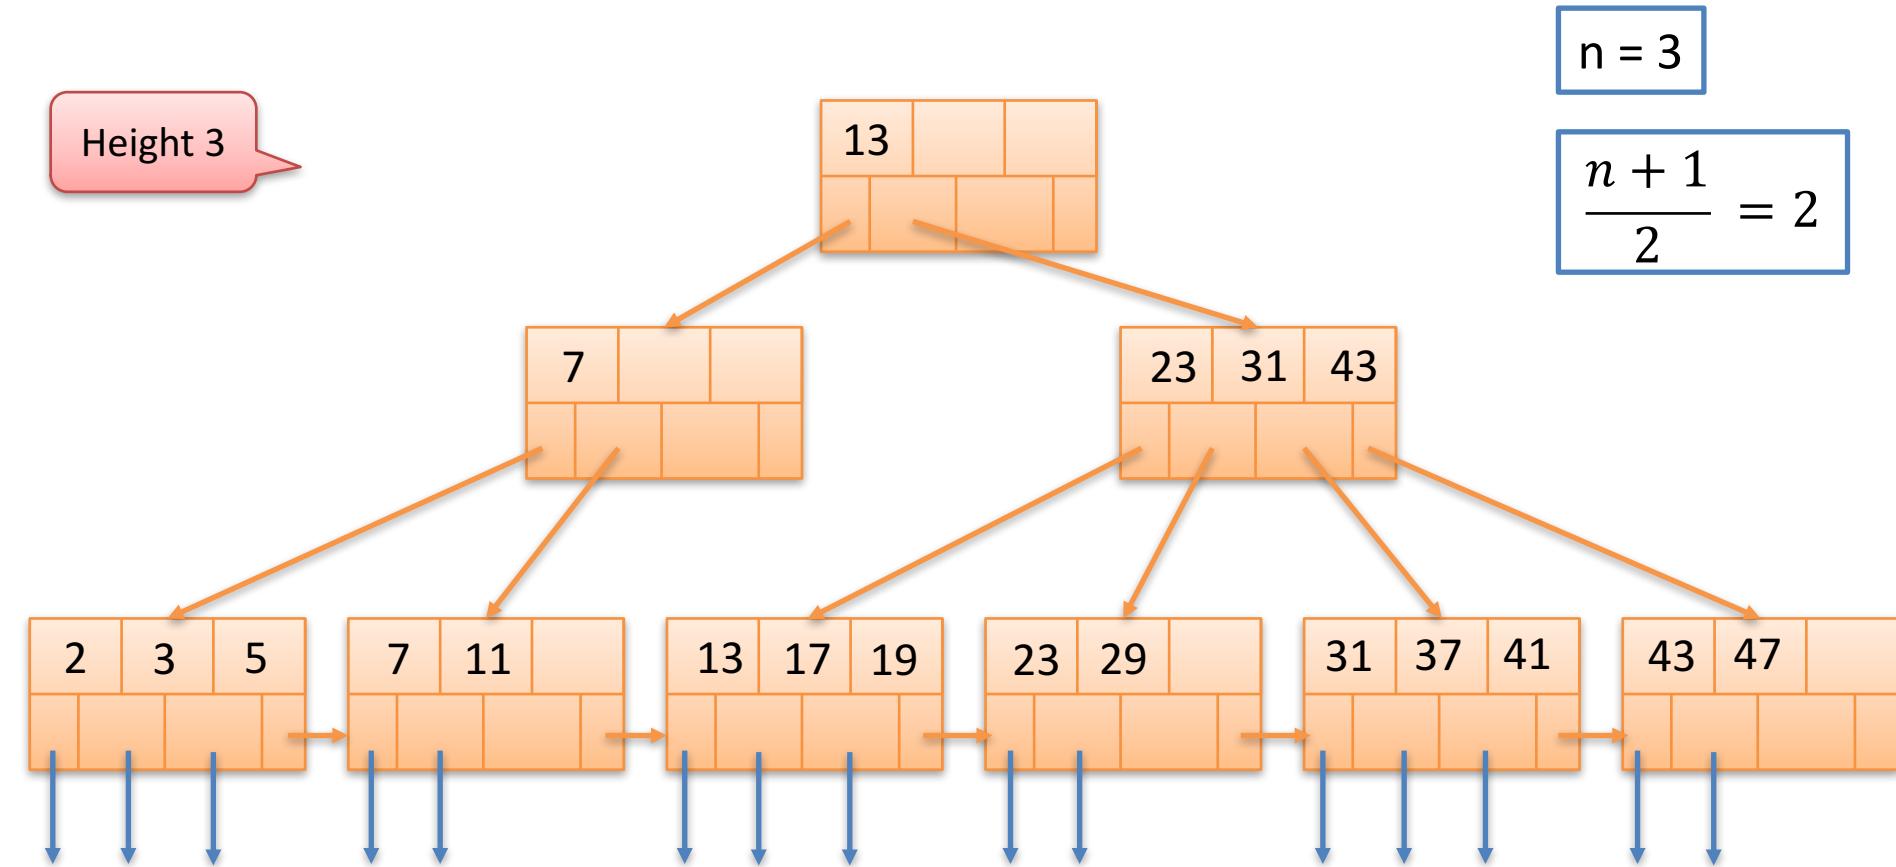
B+ Tree: Inner Nodes (Actually)

- Similar as for leaves:



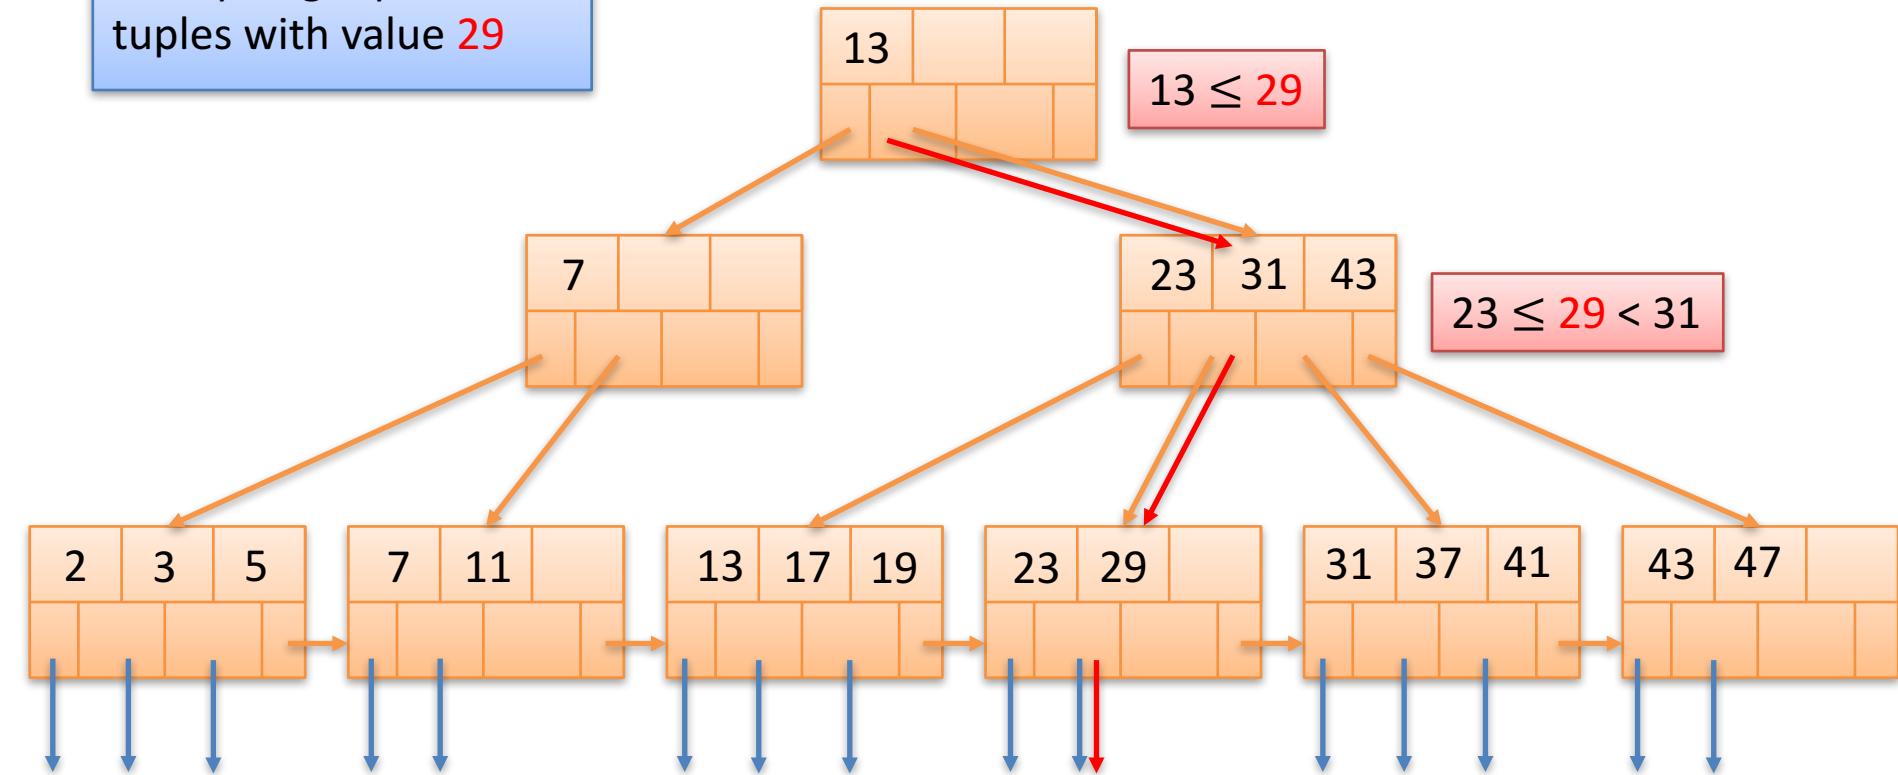
- Ensure that at least $\left\lceil \frac{n+1}{2} \right\rceil$ pointers are used
- Exception: root must use ≥ 2 pointers

A B+ Tree Index



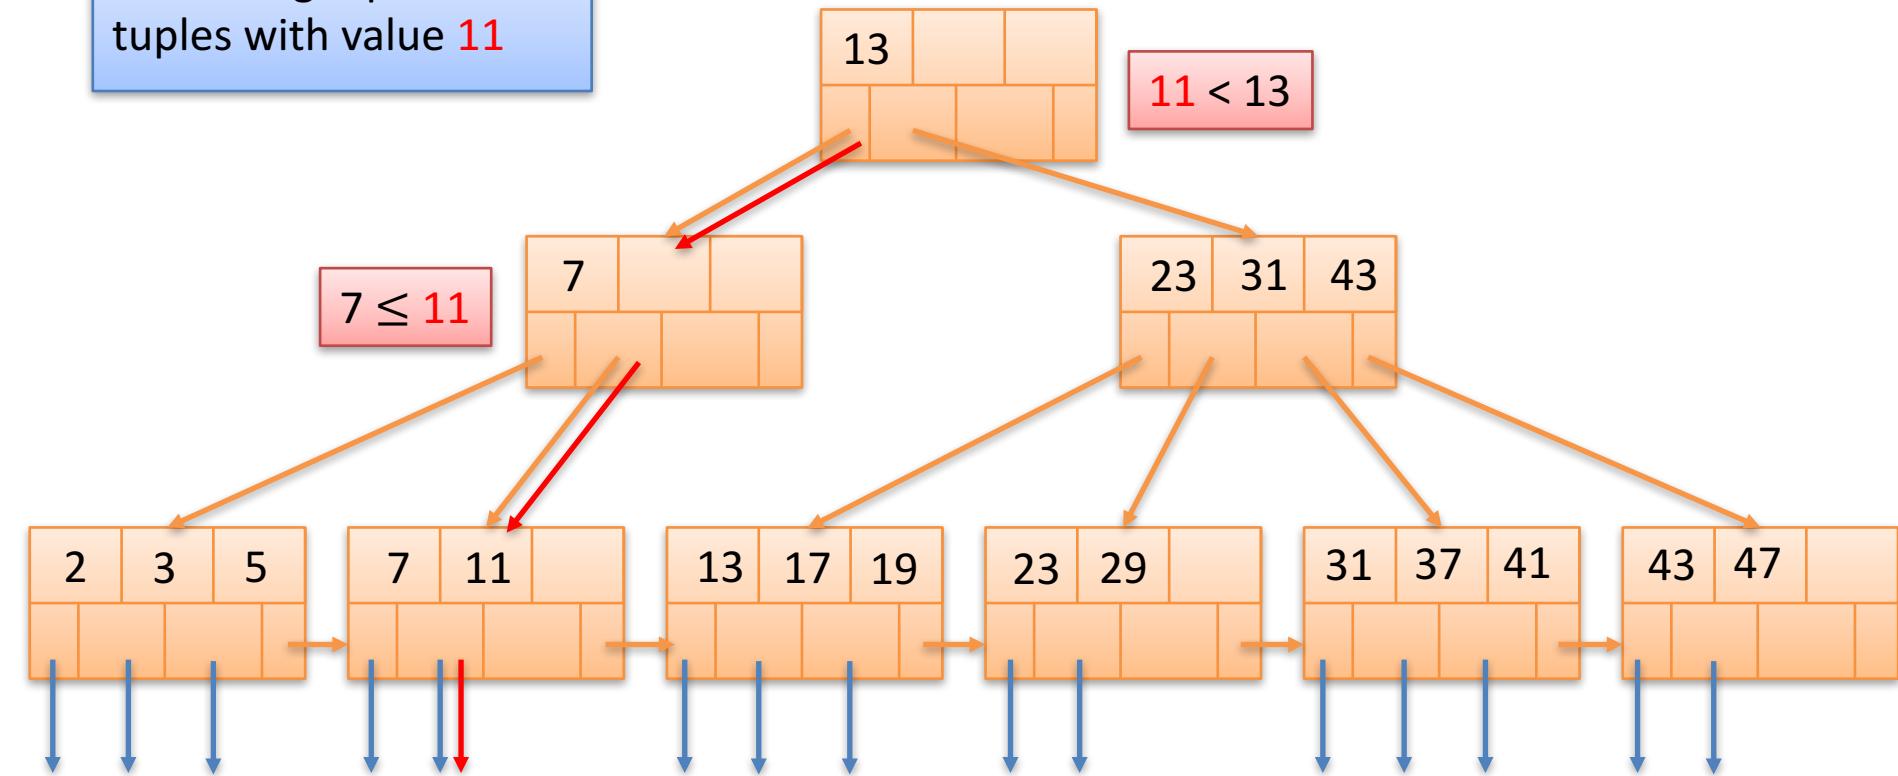
Looking Up Values

Example: get pointer to tuples with value **29**



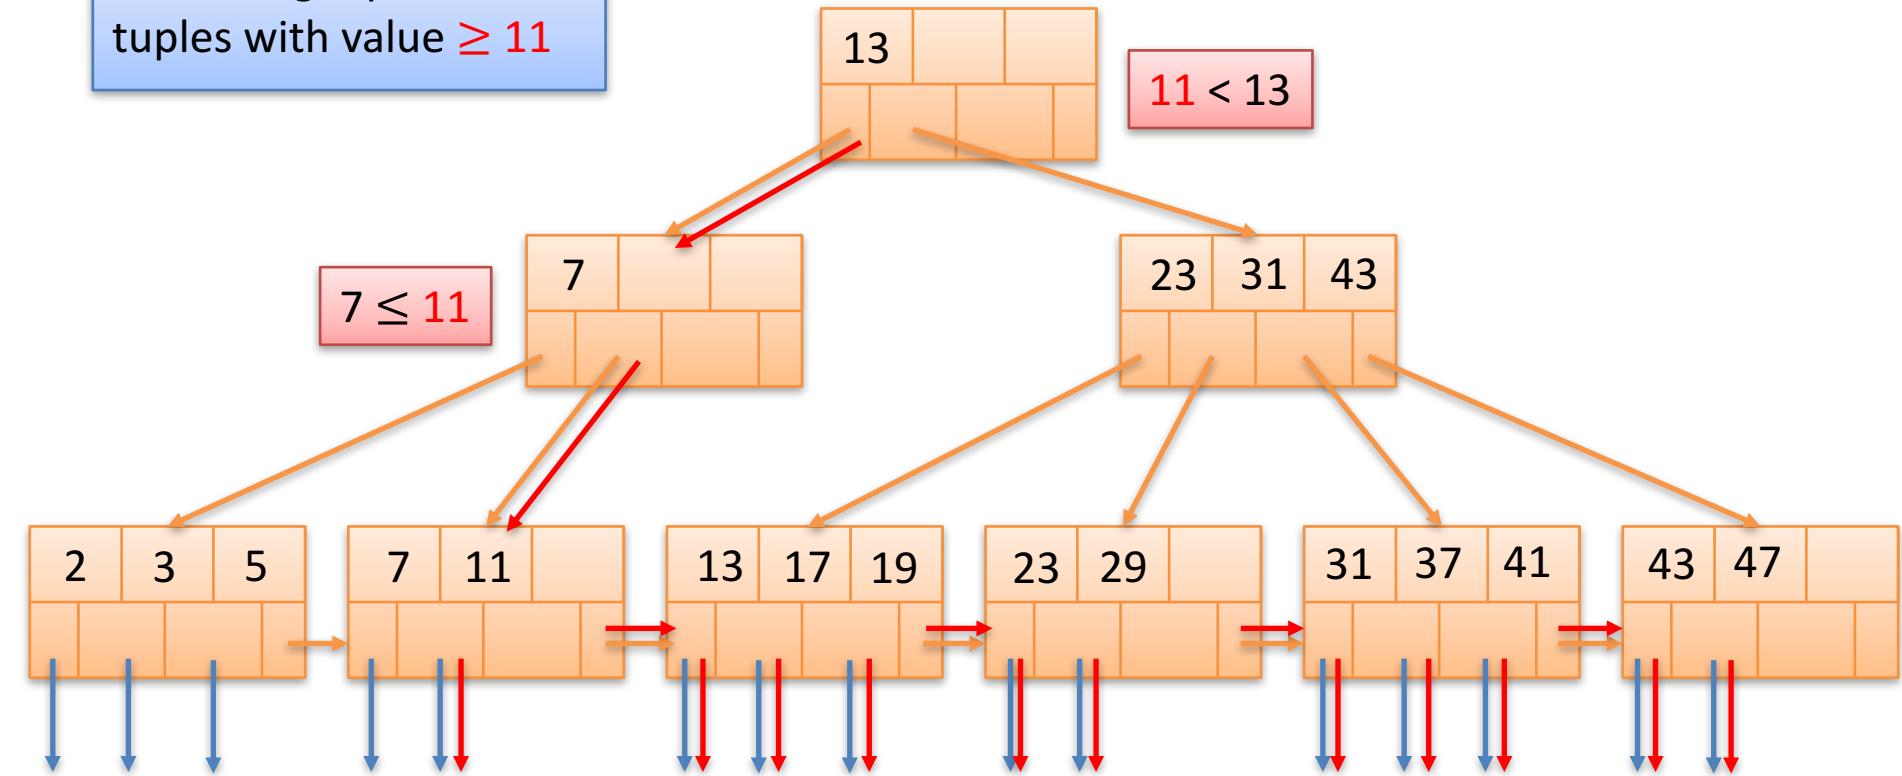
Exercise

Exercise: get pointer to tuples with value 11



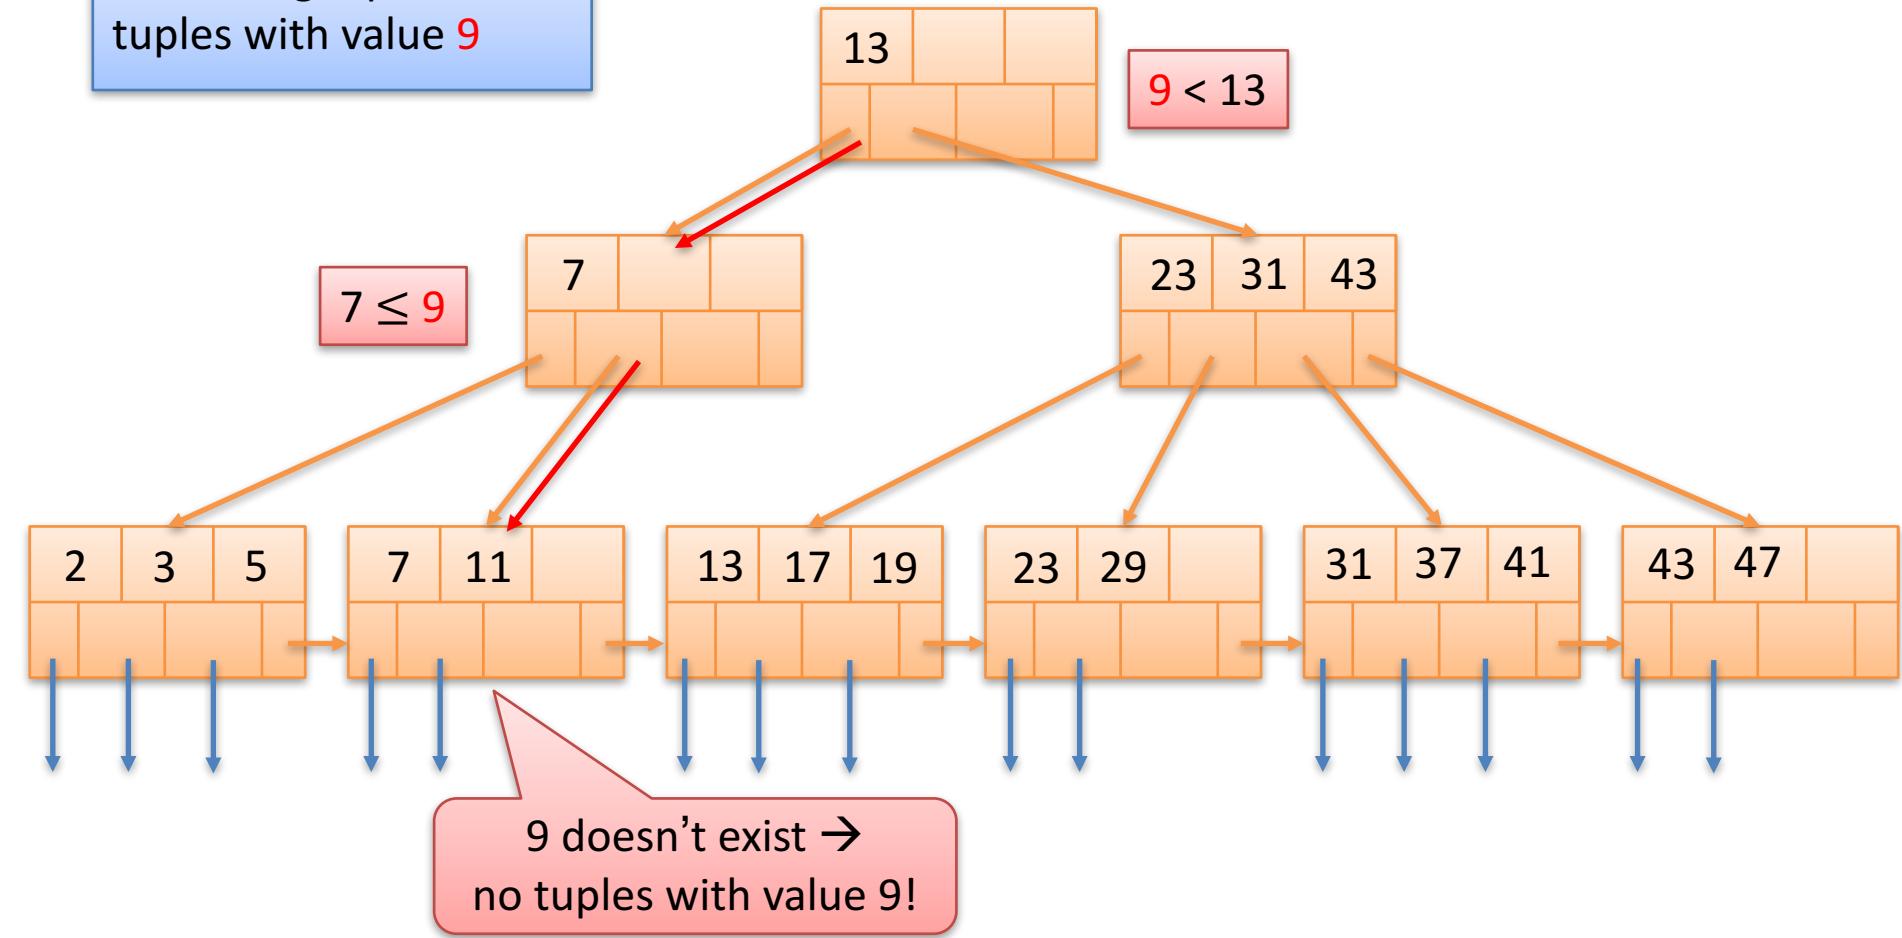
Exercise 2

Exercise: get pointers to tuples with value ≥ 11



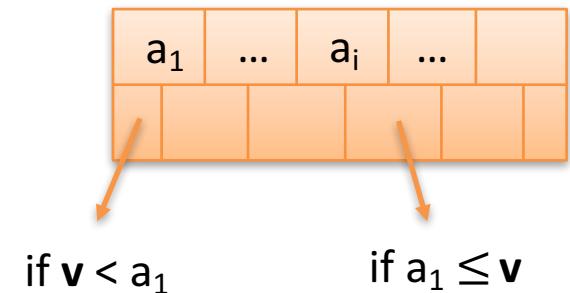
Exercise 3

Exercise: get pointer to tuples with value 9



Looking Up A Value: Summary

- Goal: find the pointer to the rows with value v
- Procedure:
 - Start at the root of the B+ tree
 - While the current node is a non-leaf node:
 - If $v < a_1$, proceed to the first child of the node
 - Otherwise find the **largest i** with $a_i \leq v$ and proceed to the associated child node
 - If the current node is a leaf:
 - If v occurs in the leaf, follow the associated pointer
 - If v does not occur in the leaf, return “ v does not exist in index”
- Running time: $O(h \times \log_2 n)$ “real” running time $O(h \times D)$

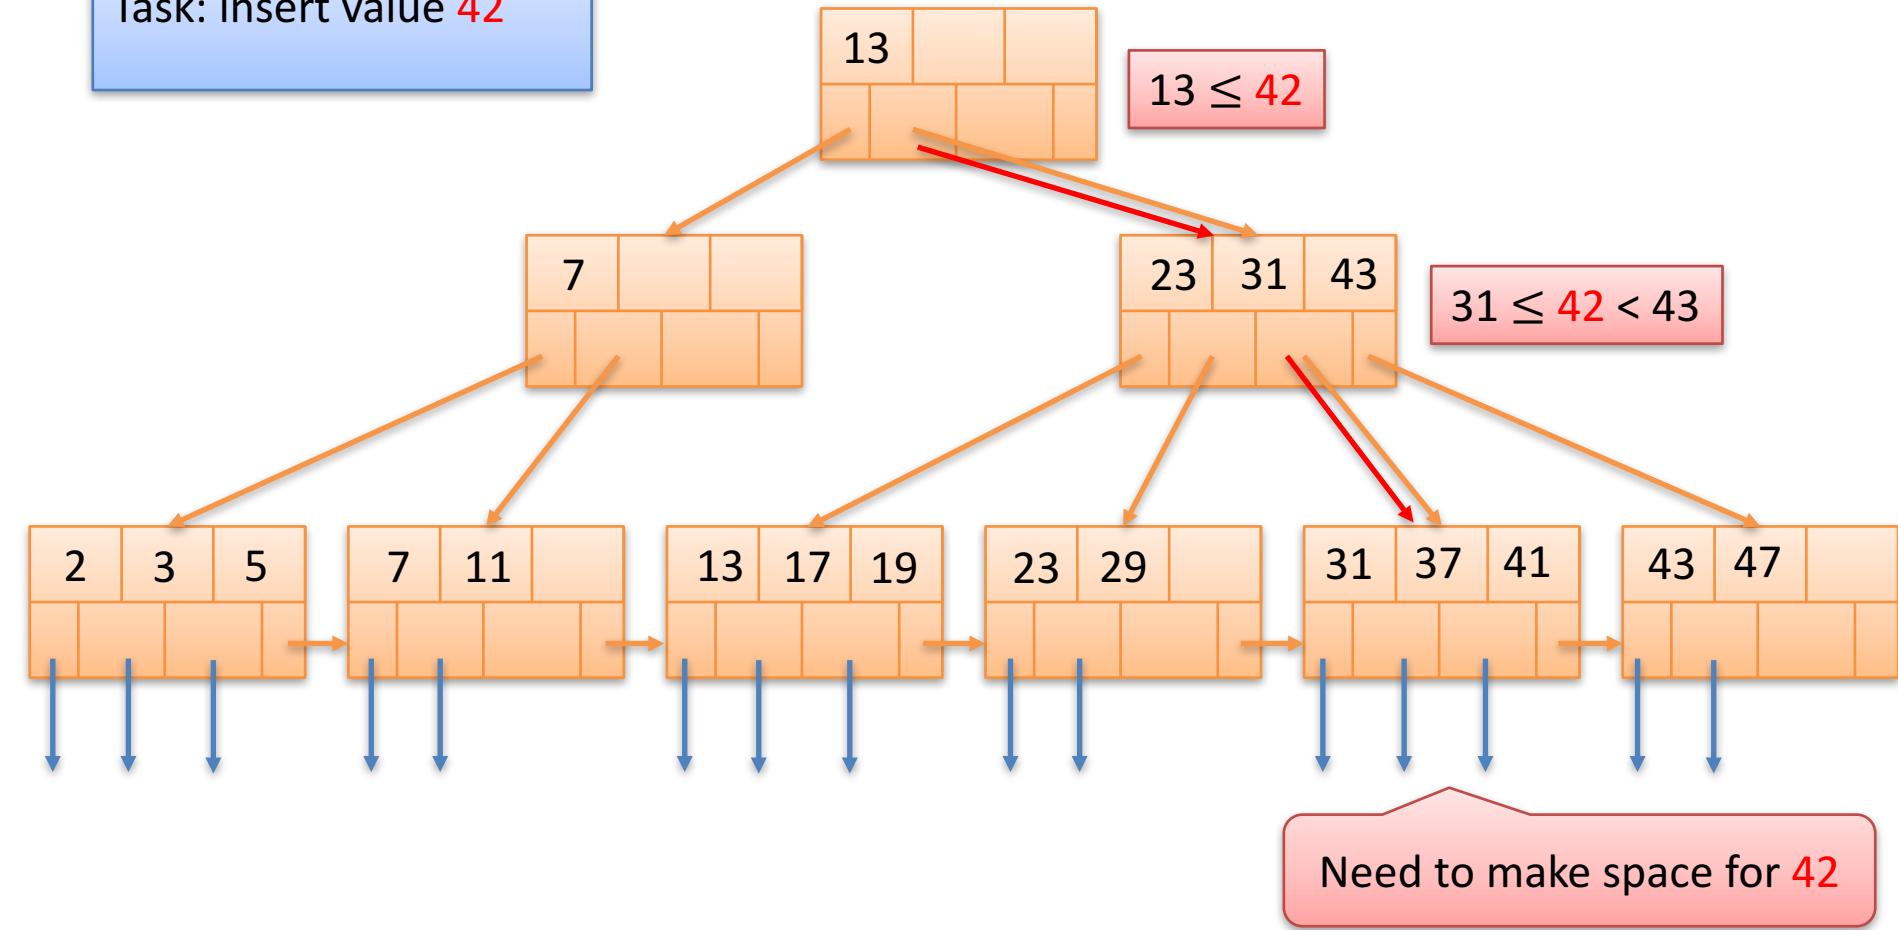


Height of the B+ tree

Time for a disk operation

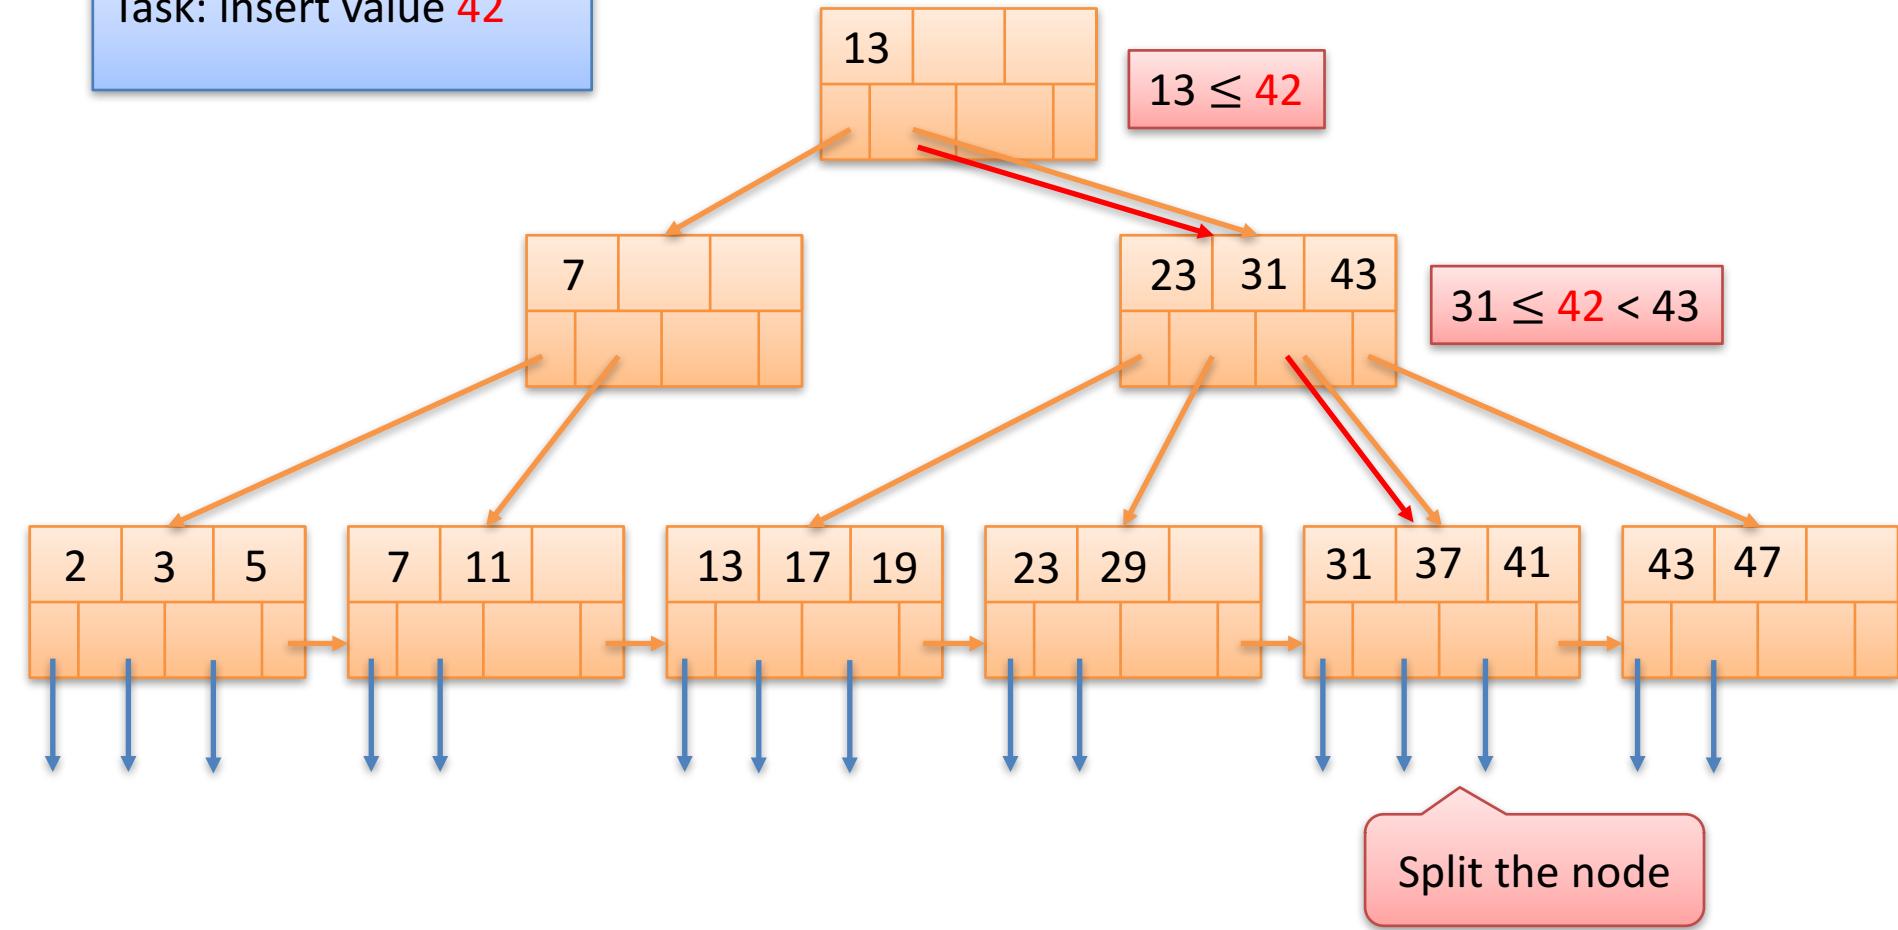
Insertions of Value/Pointer Pairs

Task: insert value **42**



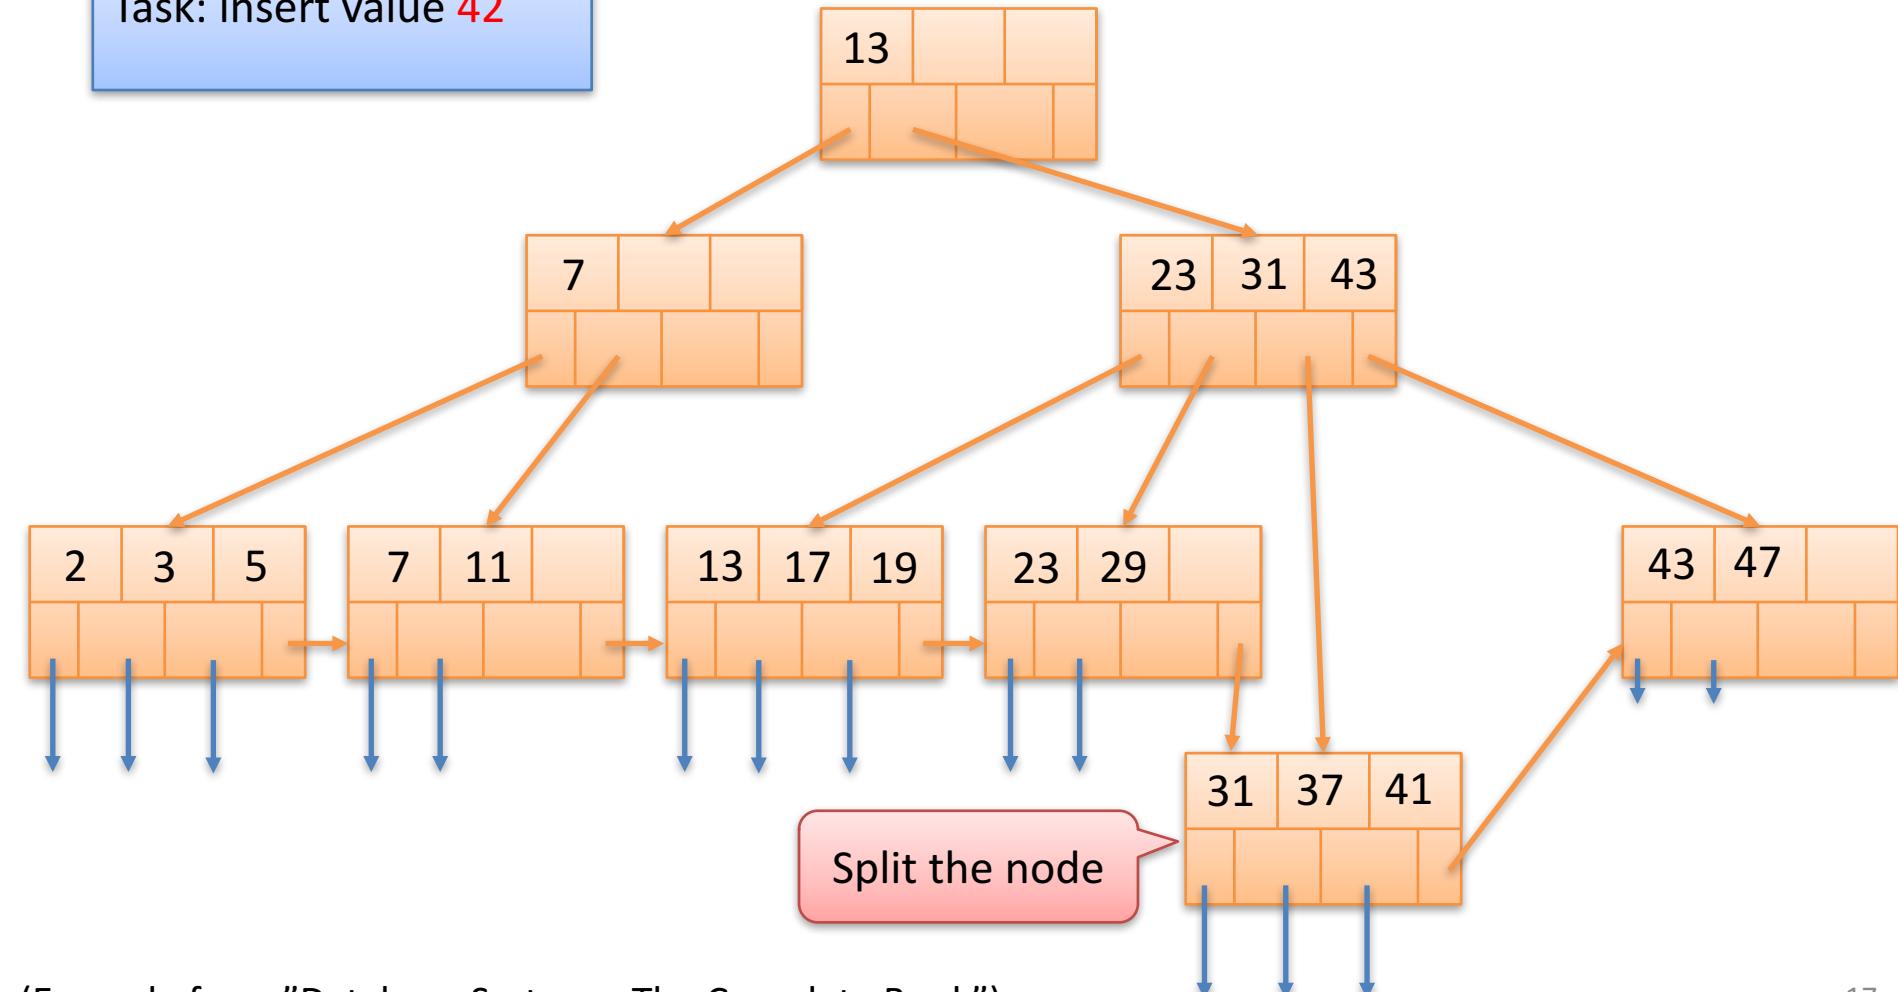
Insertions of Value/Pointer Pairs

Task: insert value **42**



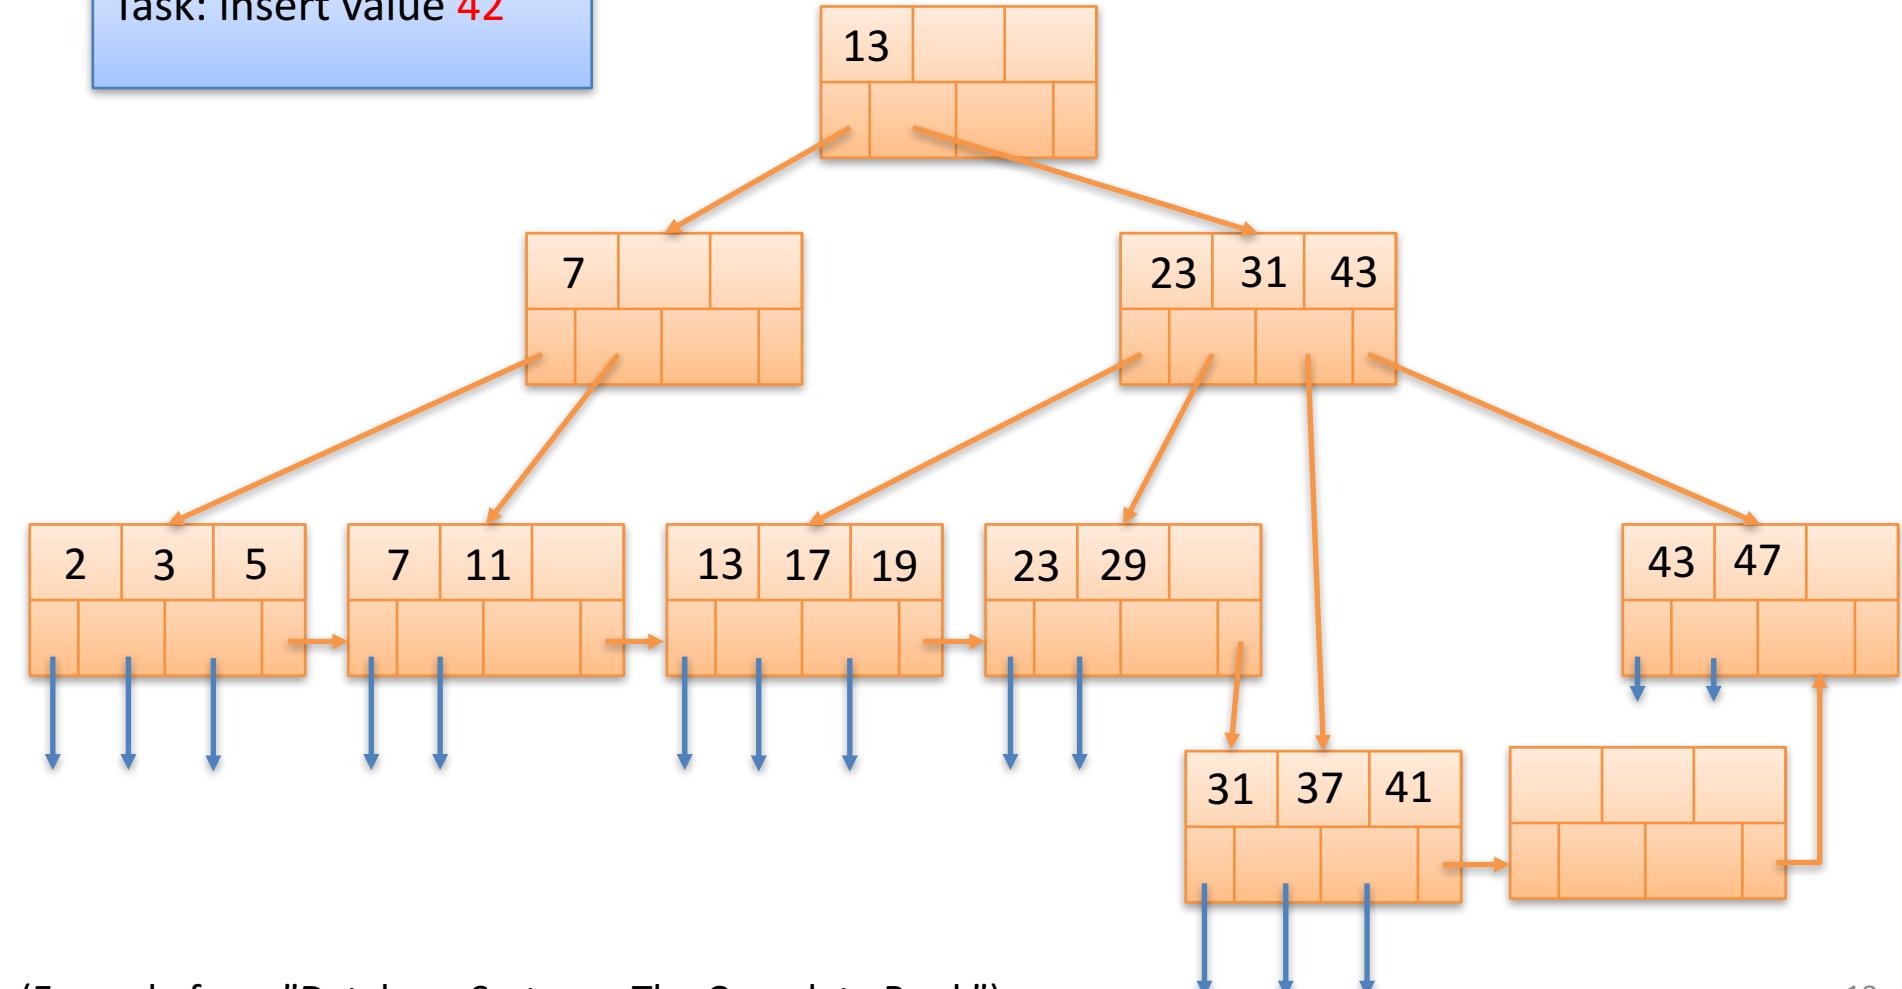
Insertions of Value/Pointer Pairs

Task: insert value **42**



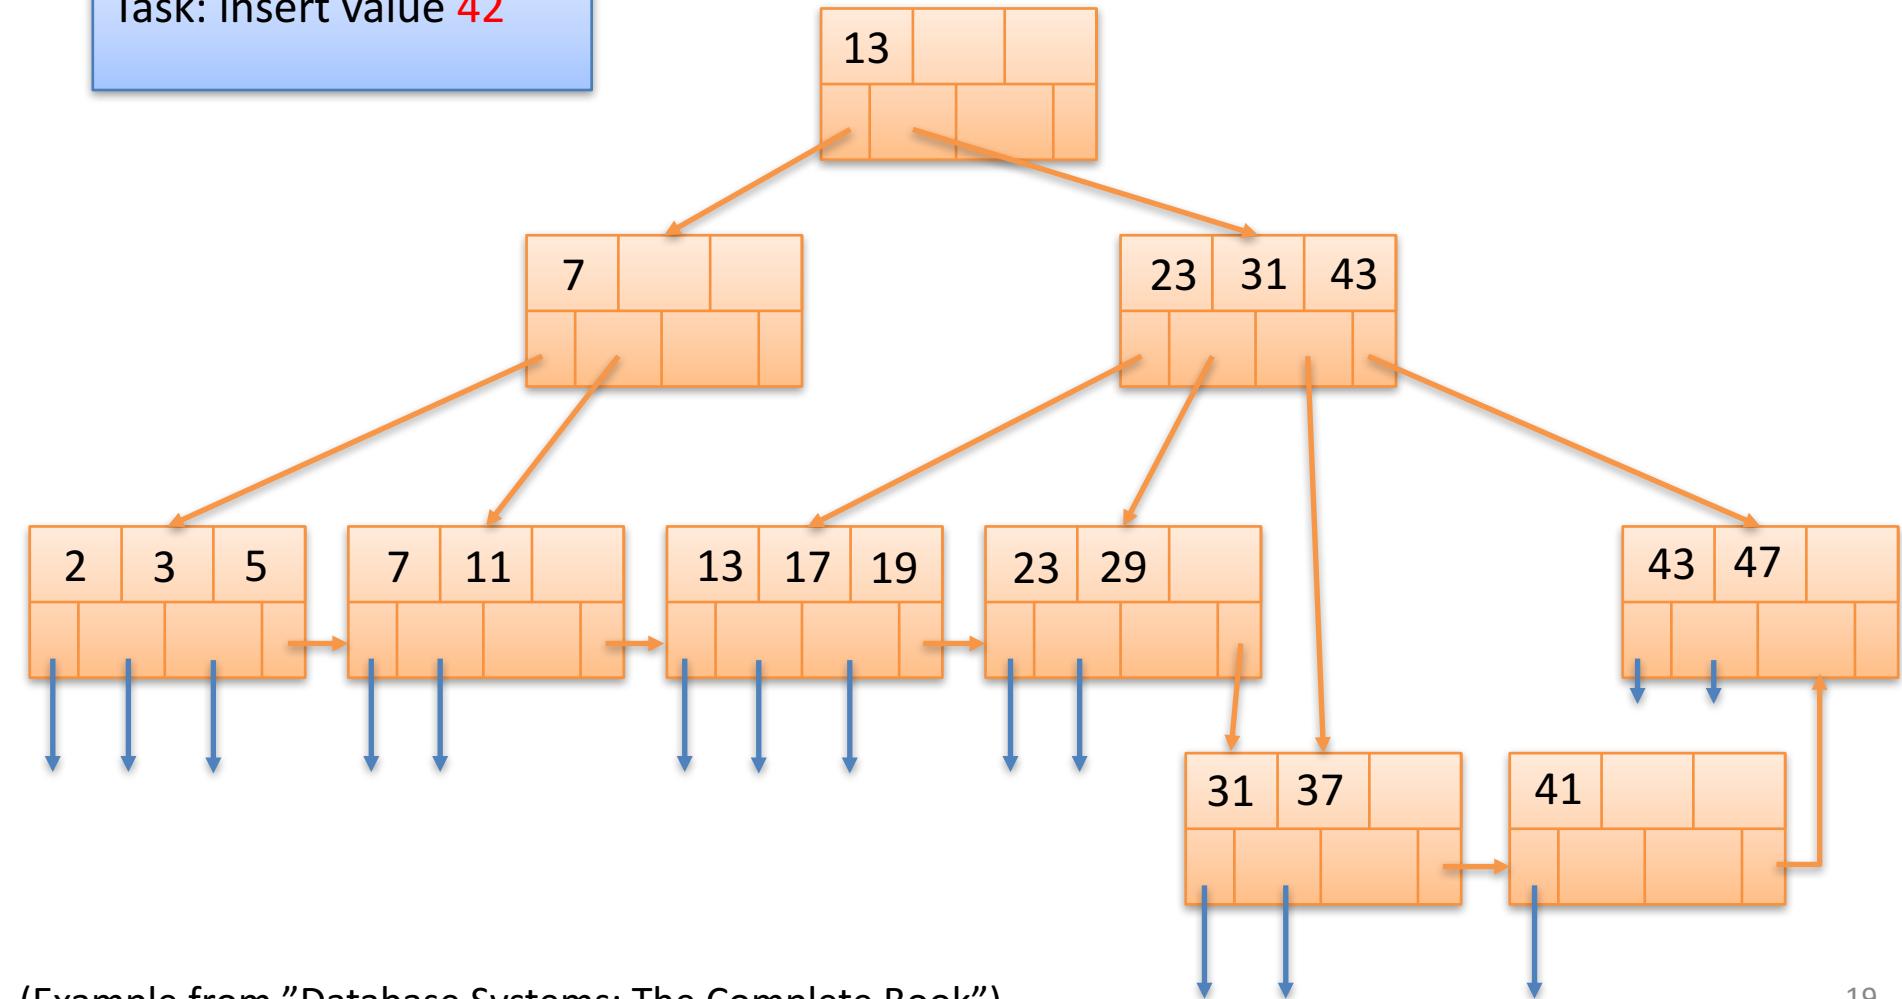
Insertions of Value/Pointer Pairs

Task: insert value **42**



Insertions of Value/Pointer Pairs

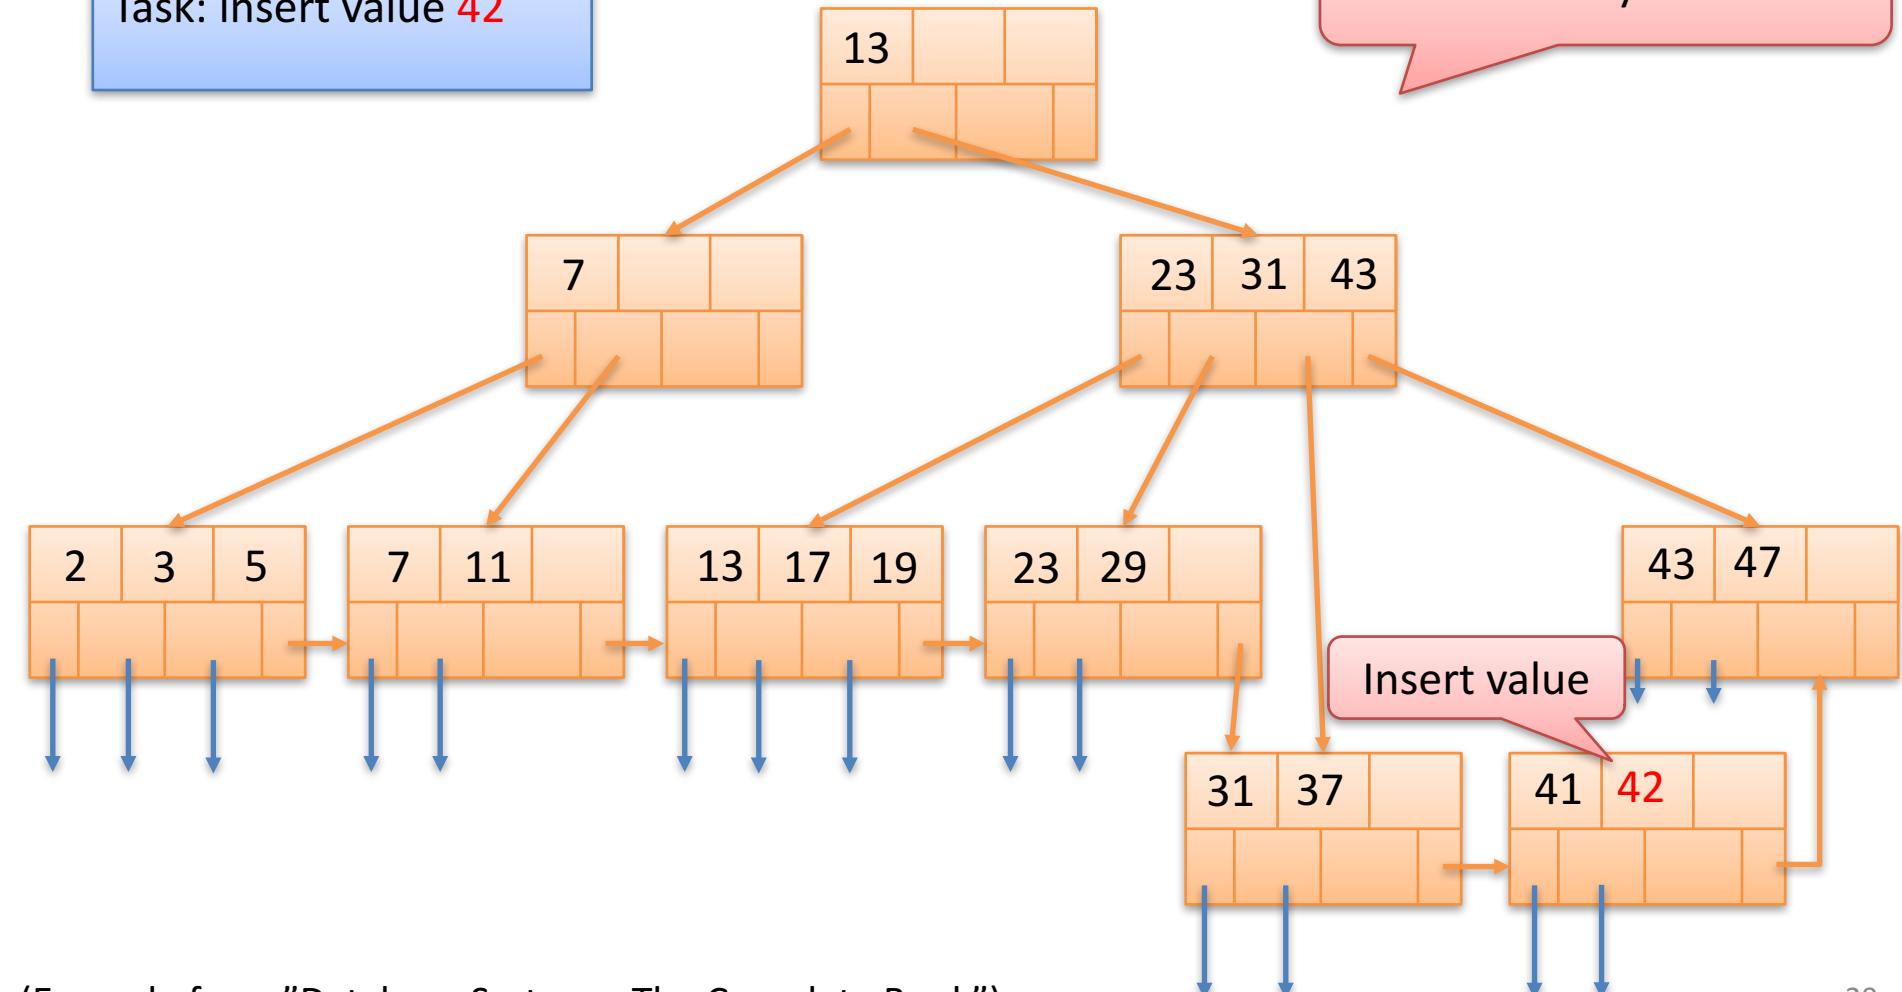
Task: insert value **42**



Insertions of Value/Pointer Pairs

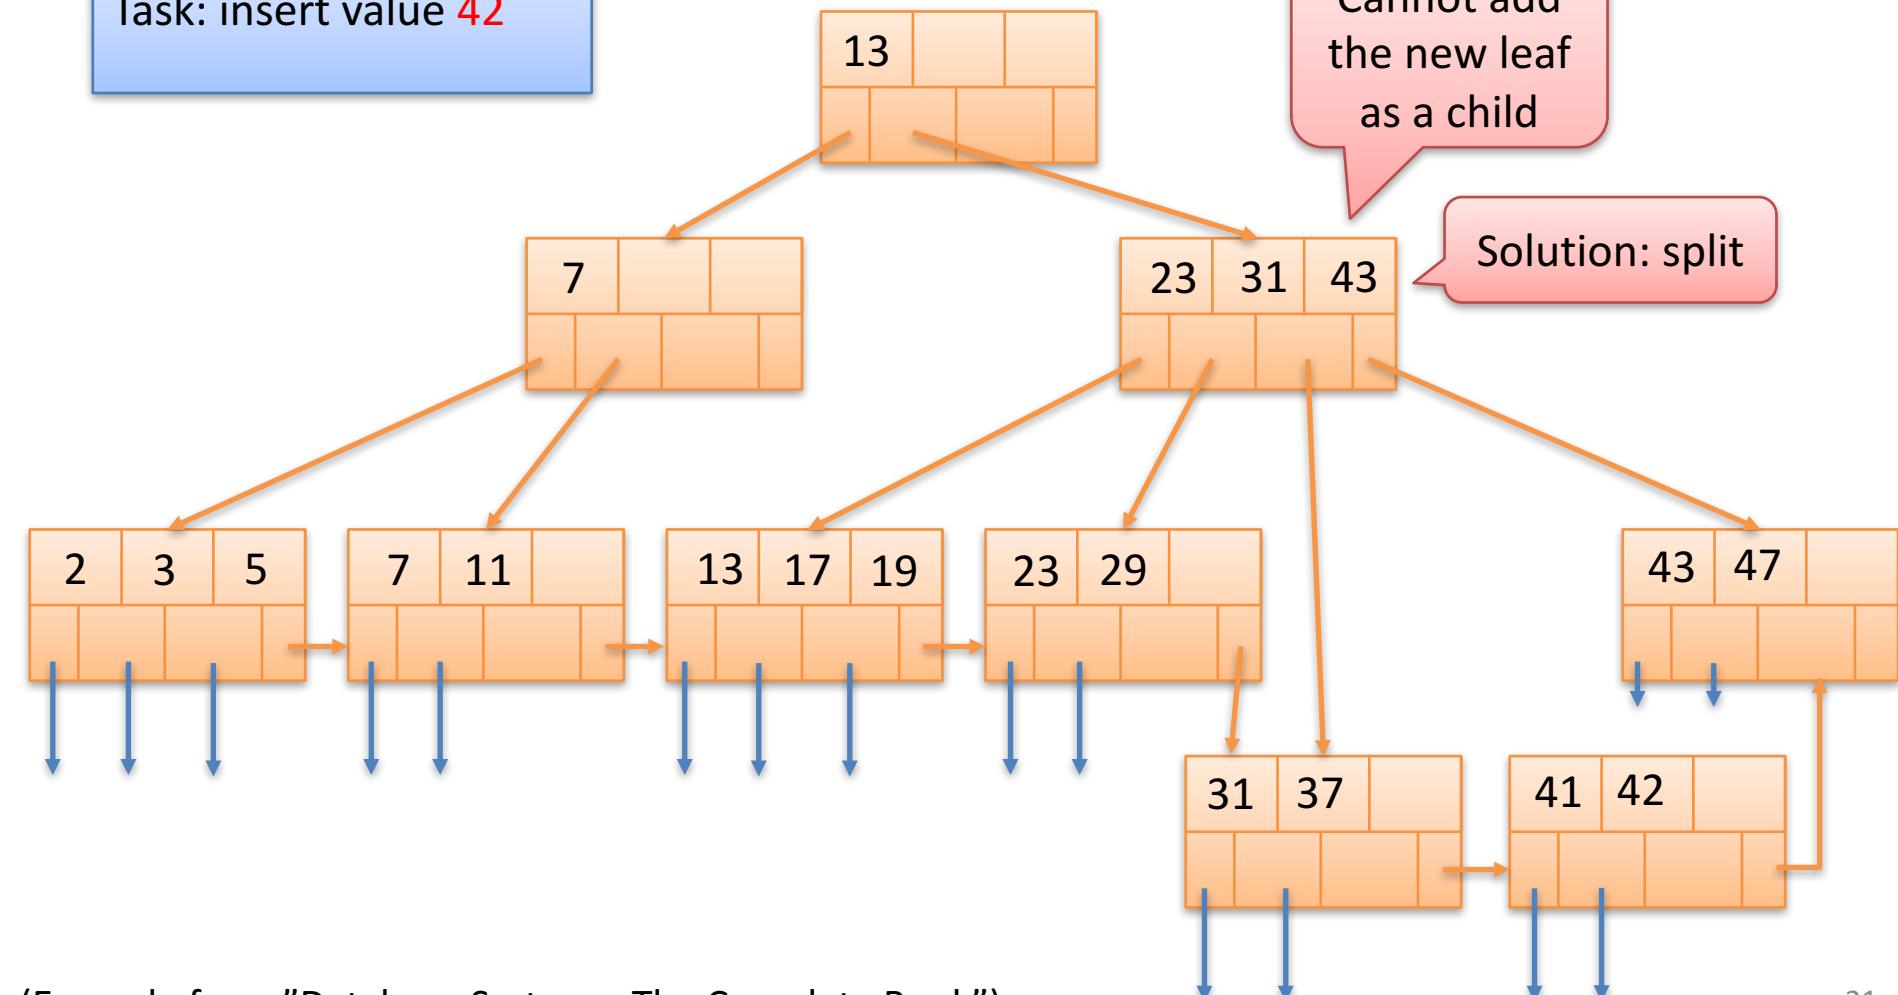
Task: insert value **42**

Problem: not yet a B+ tree



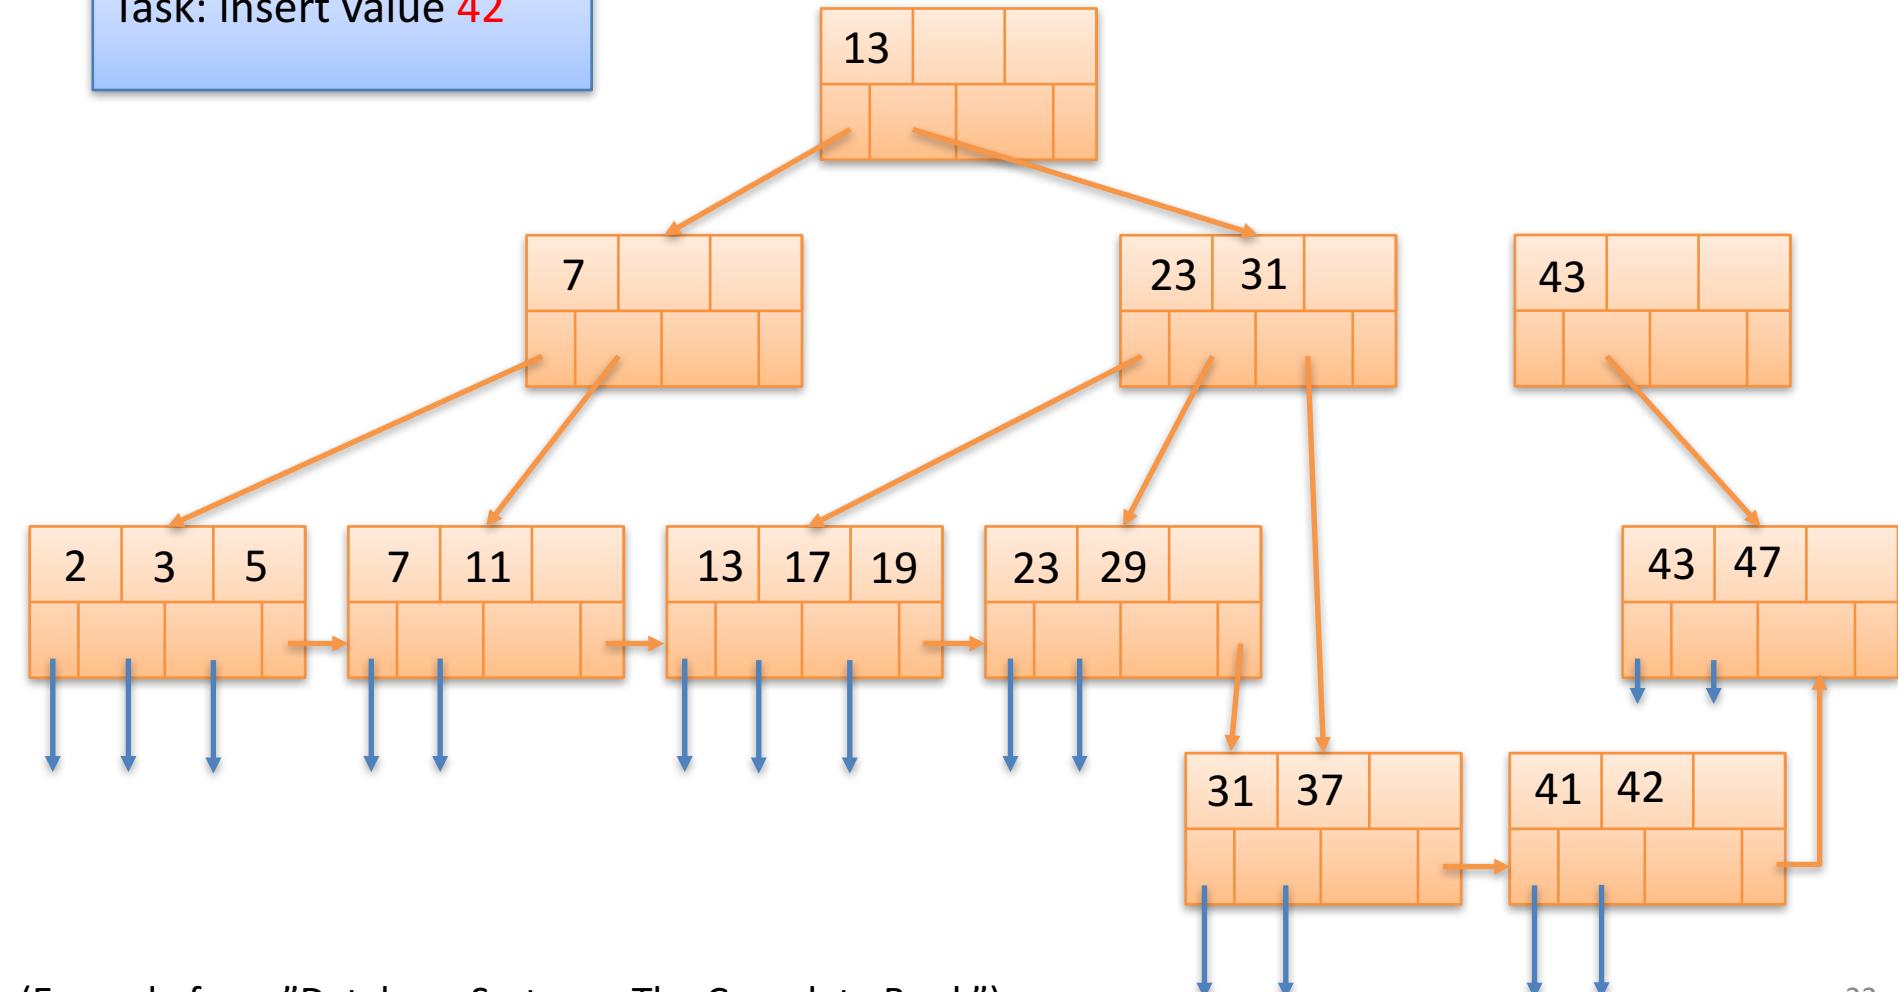
Insertions of Value/Pointer Pairs

Task: insert value **42**



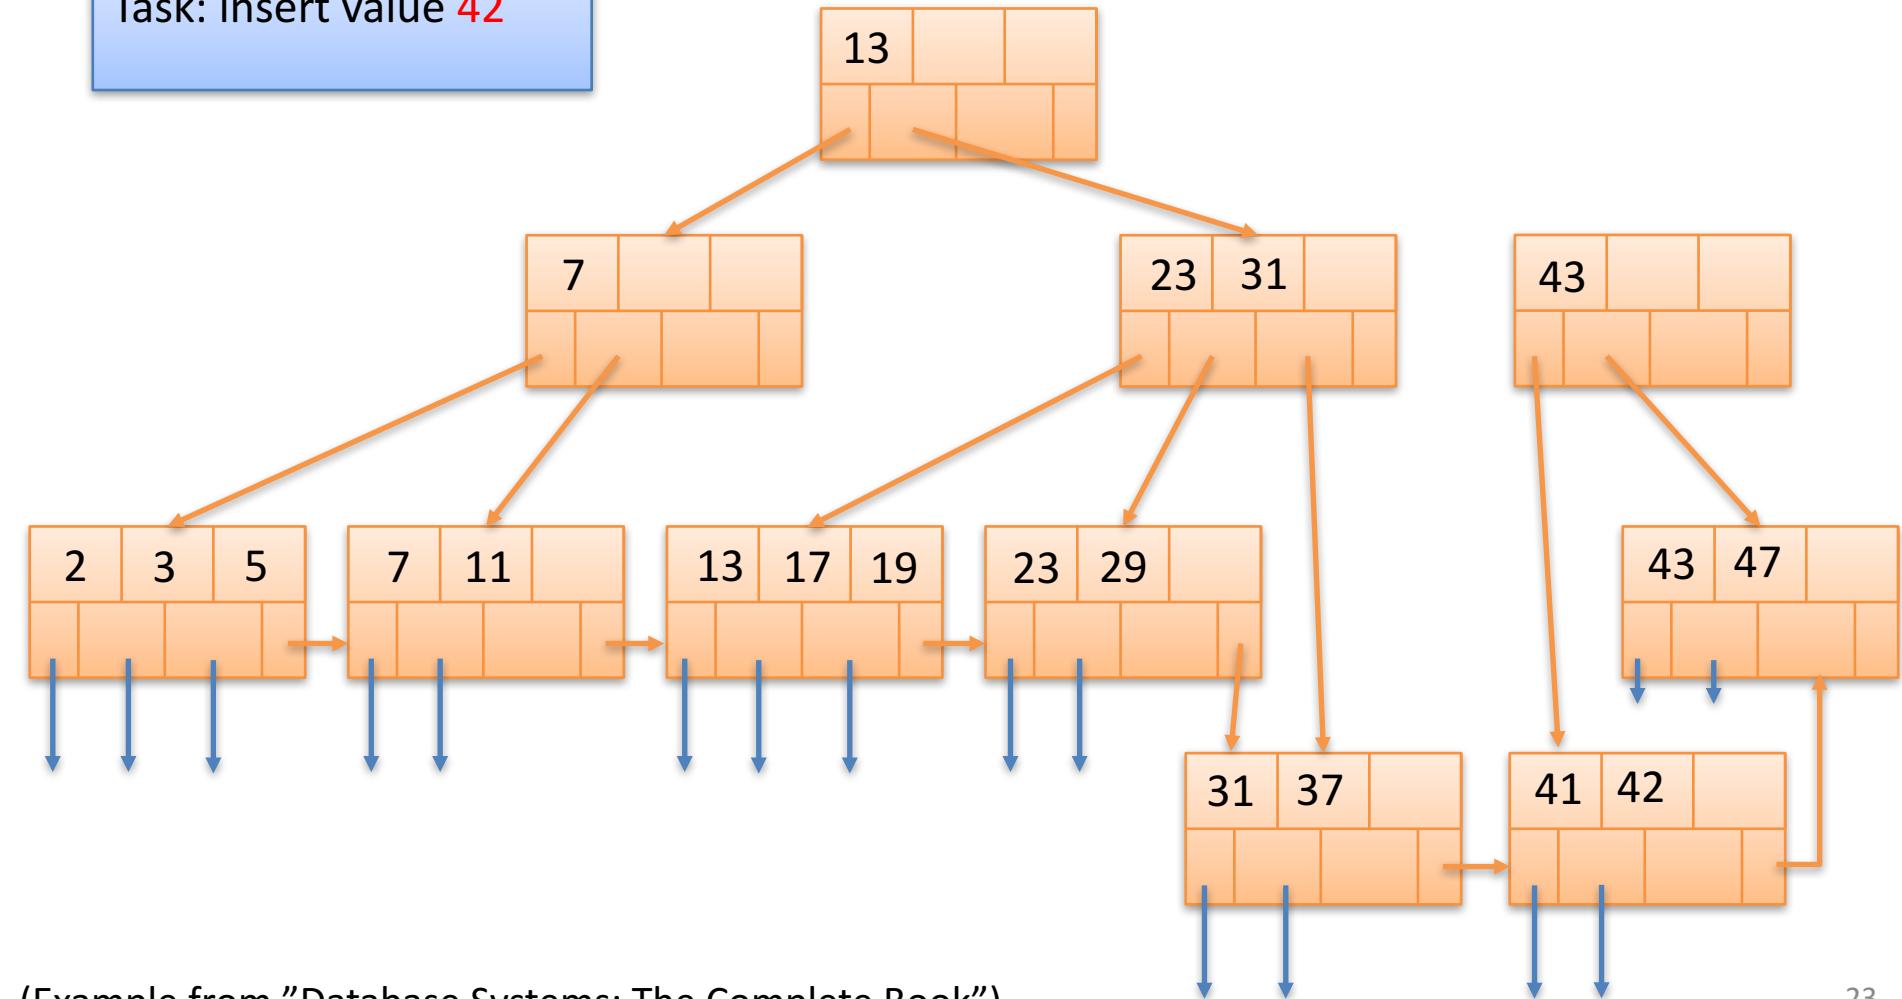
Insertions of Value/Pointer Pairs

Task: insert value **42**



Insertions of Value/Pointer Pairs

Task: insert value **42**

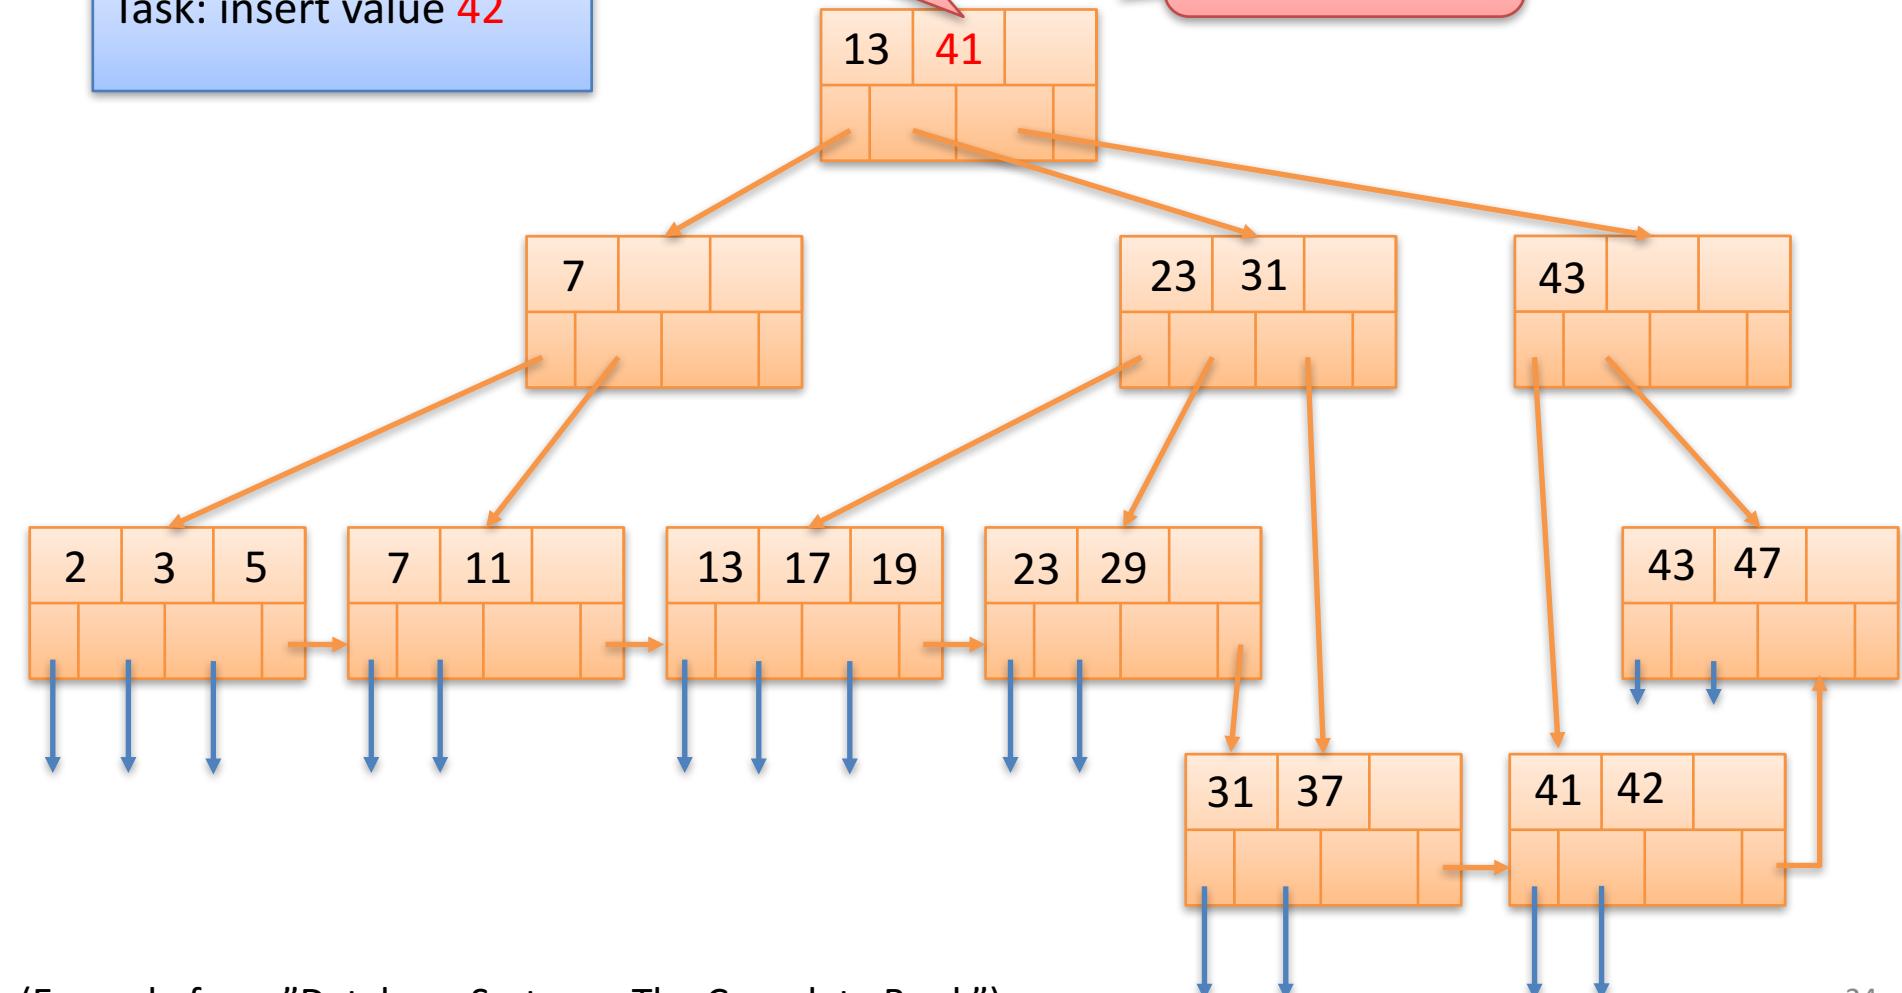


Insertions of Value/Pointer Pairs

Task: insert value **42**

Insert value

Add the new node as a child



Insertion: Summary

- Goal: insert a new value/pointer pair
- Procedure:
 - Find the leaf that should contain the value
 - If the leaf is not full, insert the key value pair at a suitable location
 - If the leaf is full:
 - Split the leaf to make space for the new value/pointer pair
 - Insert the value/pointer pair
 - Connect the leaf to a suitable parent node (which might incur the creation of a new node etc.)
- The B+ tree **remains balanced!**
- Running time: $O(h \times \log_2 n)$
“real” running time $O(h \times D)$

Time for a disk operation

Height of the B+ tree

Deletion

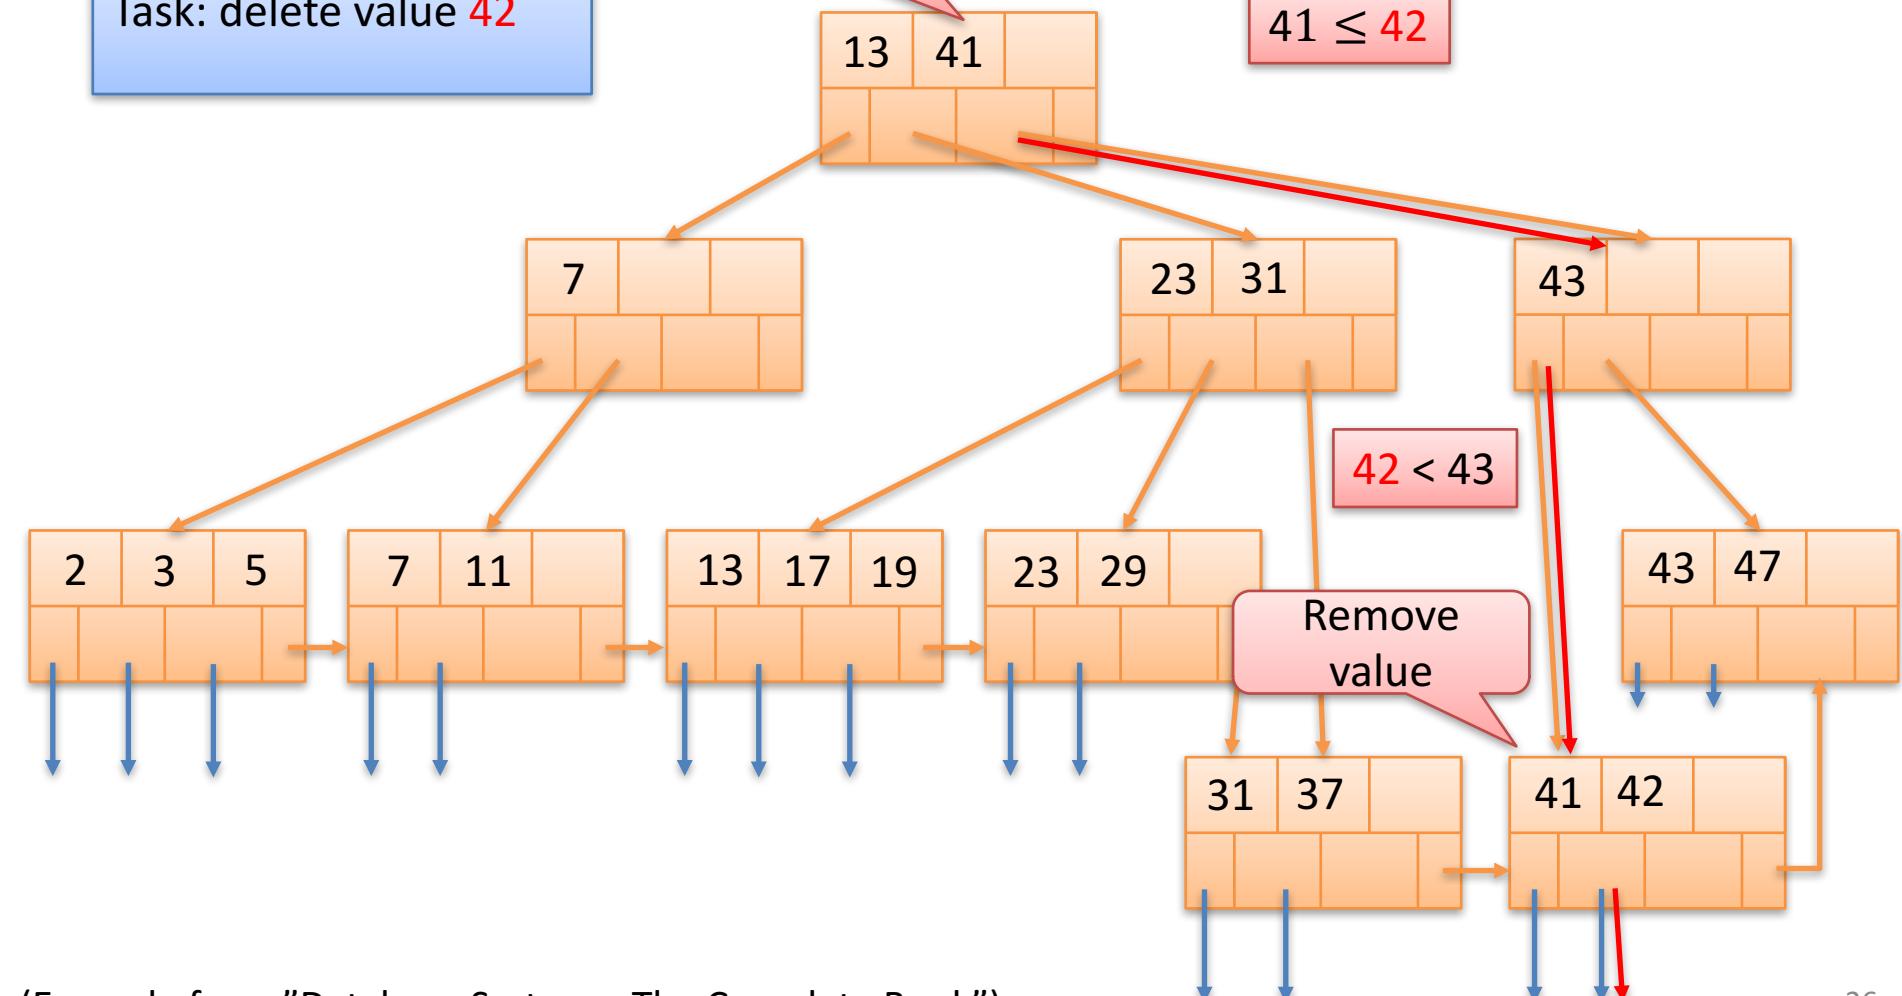
Task: delete value **42**

Find value

$41 \leq 42$

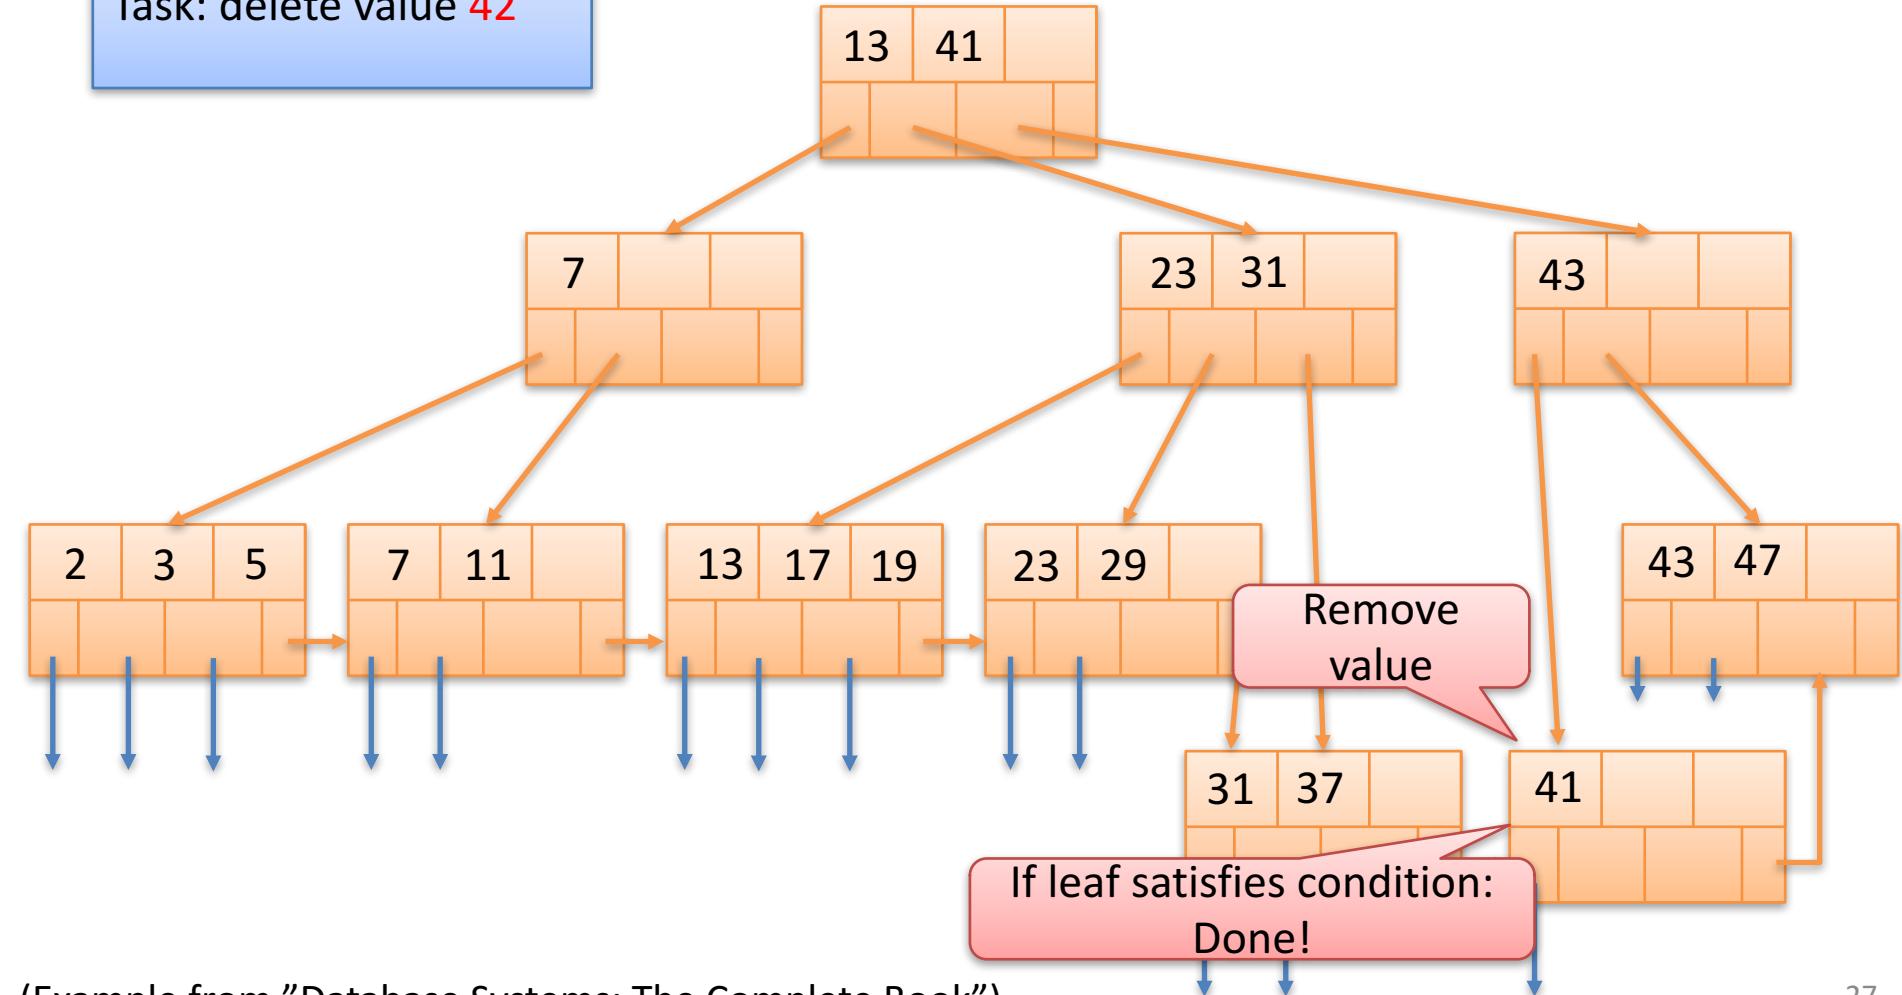
$42 < 43$

Remove value



Deletion

Task: delete value **42**



Deletion

Task: delete value **42**

Clean up!

13 41

7

2 3 5

7 11

13 17 19

22 21

43

Remove
value

31 37

41

If leaf satisfies condition:
Done!

Side note

- There may be more than one B+ tree with the same leaves

Deletion part 2

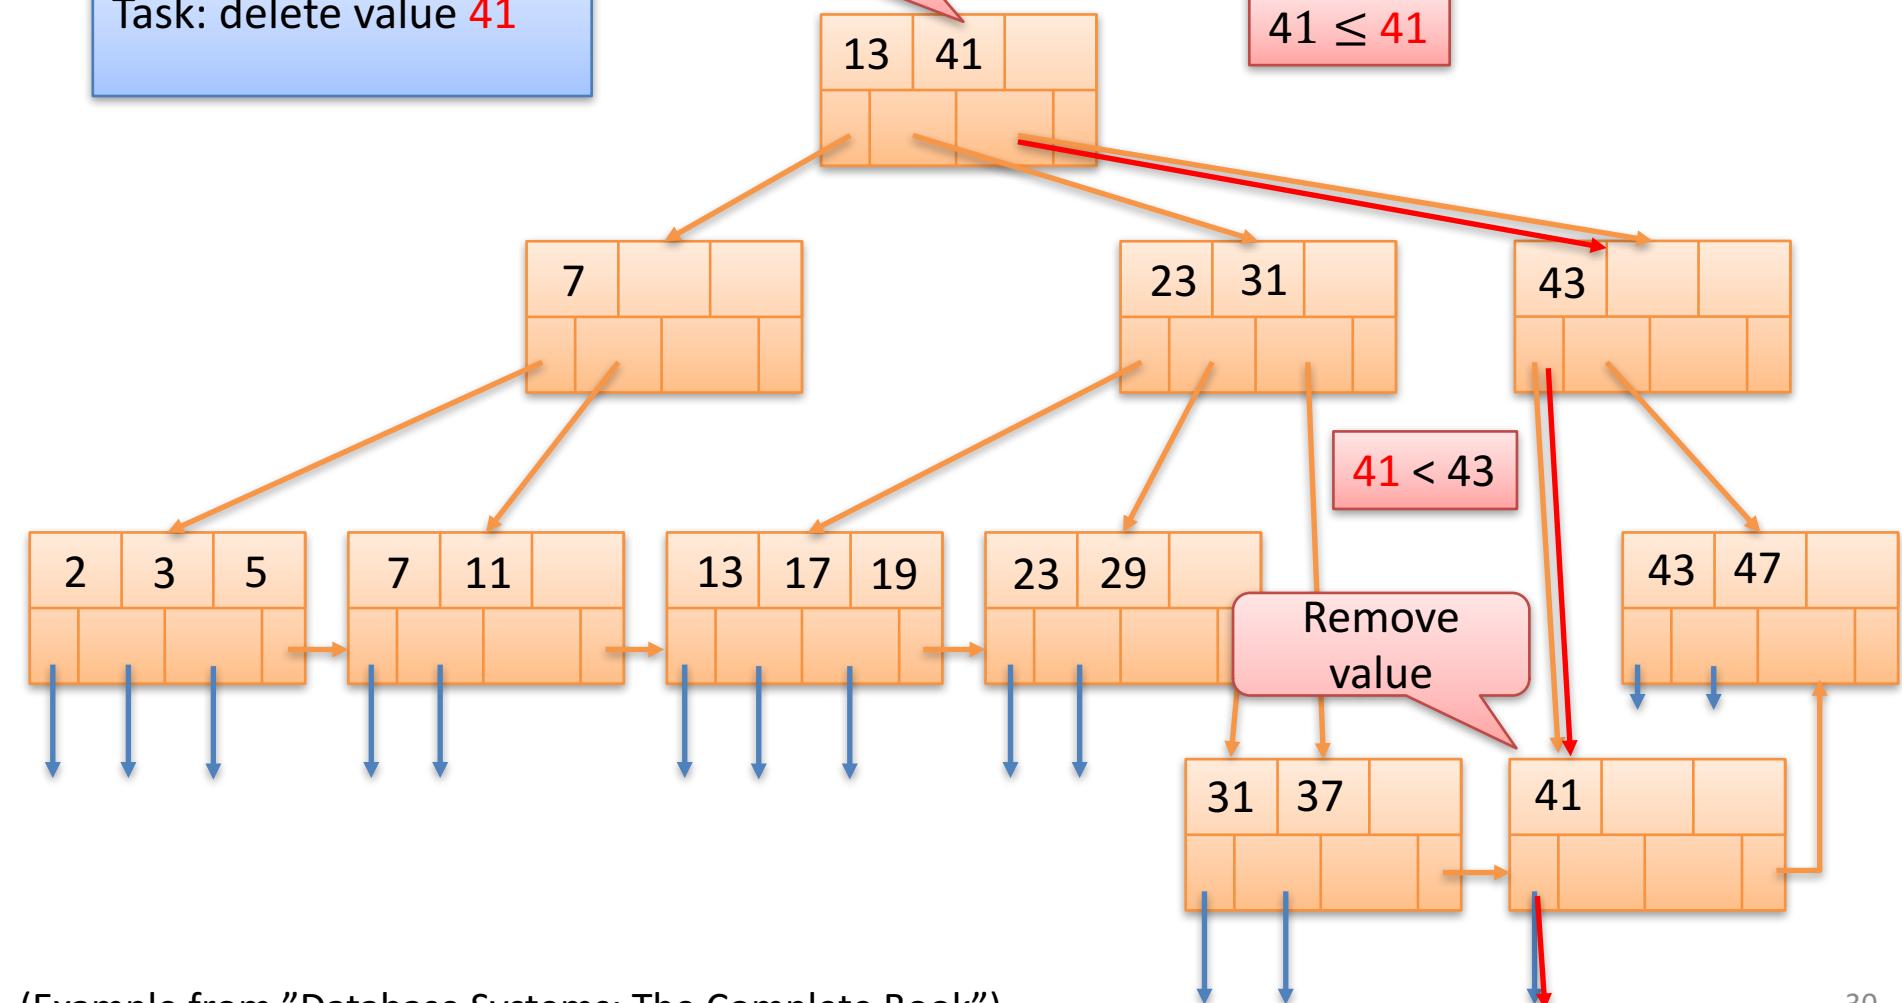
Task: delete value **41**

Find value

$41 \leq 41$

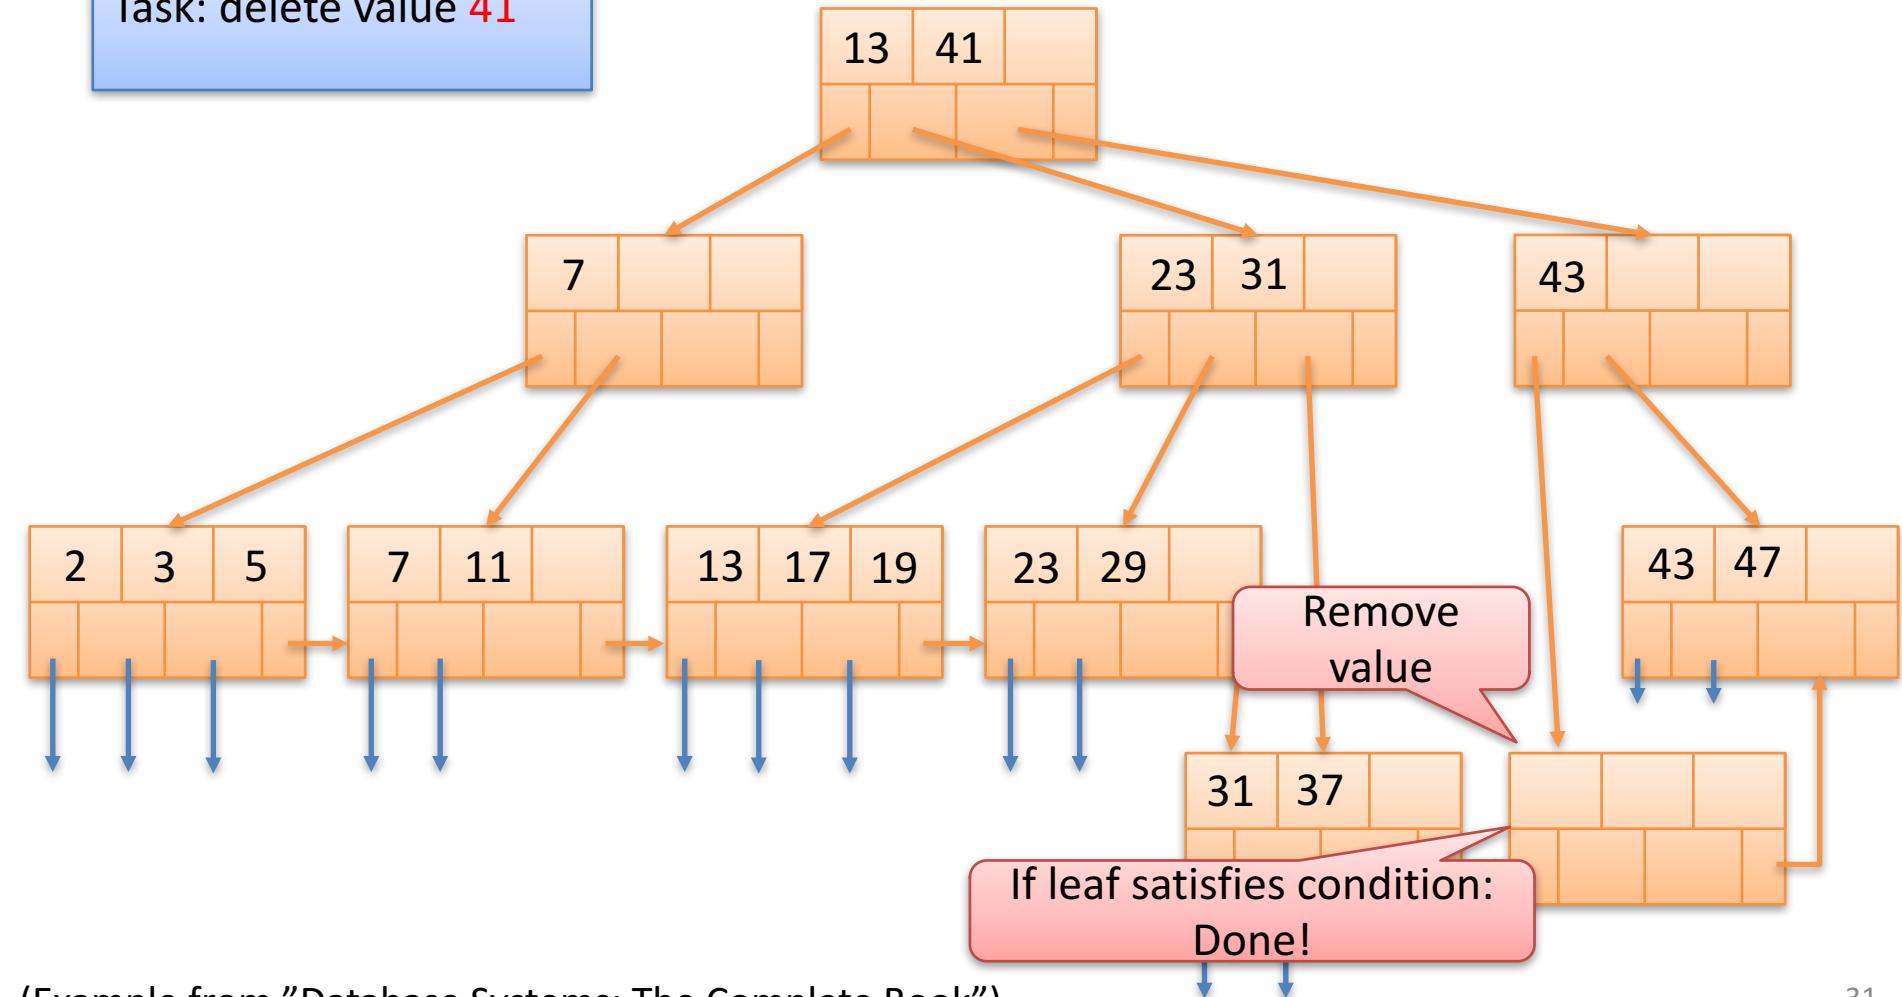
$41 < 43$

Remove value



Deletion part 2

Task: delete value **41**



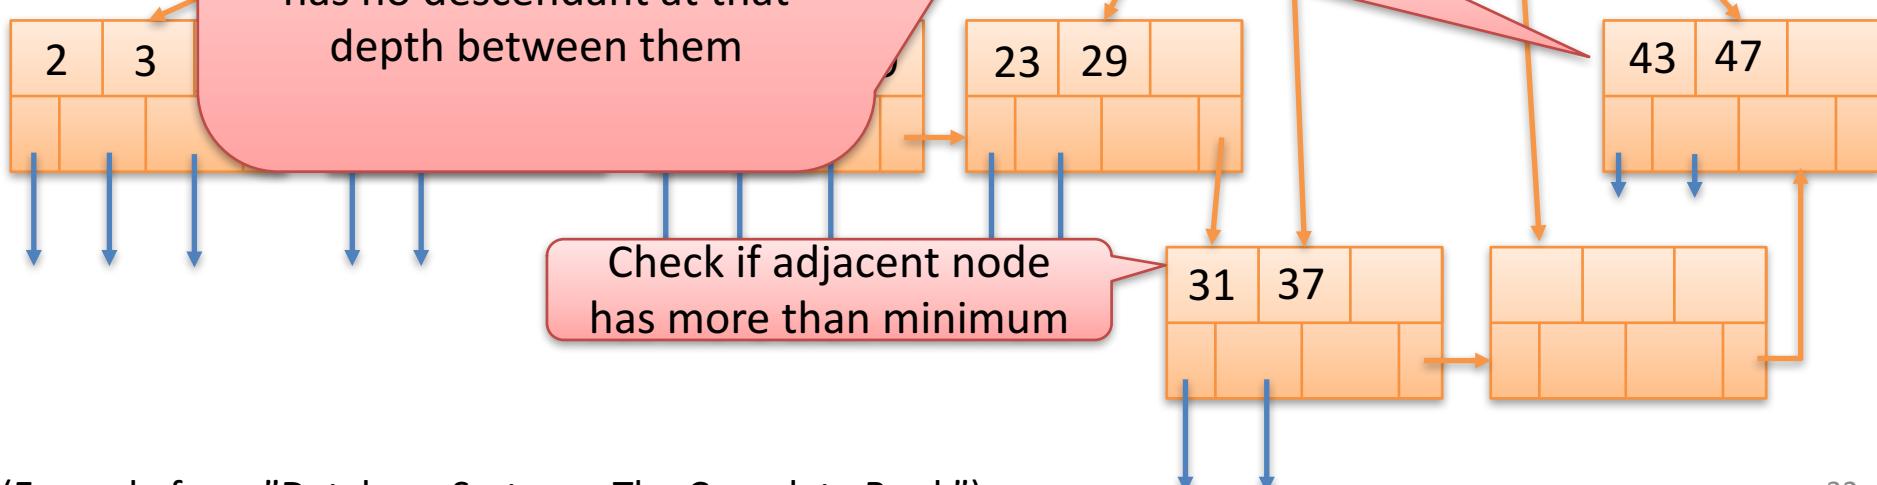
Deletion part 2

Task: delete value **41**

A node is adjacent to another, if they are at the same depth and their common ancestor has no descendant at that depth between them

Check if adjacent node has more than minimum

Check if adjacent node has more than minimum



Deletion part 2

Task: delete value **41**

Clean up!

13 43

7

Clean up!

43

2 3 5

7 11

13 17 19

Steal a pointer

43 47

31

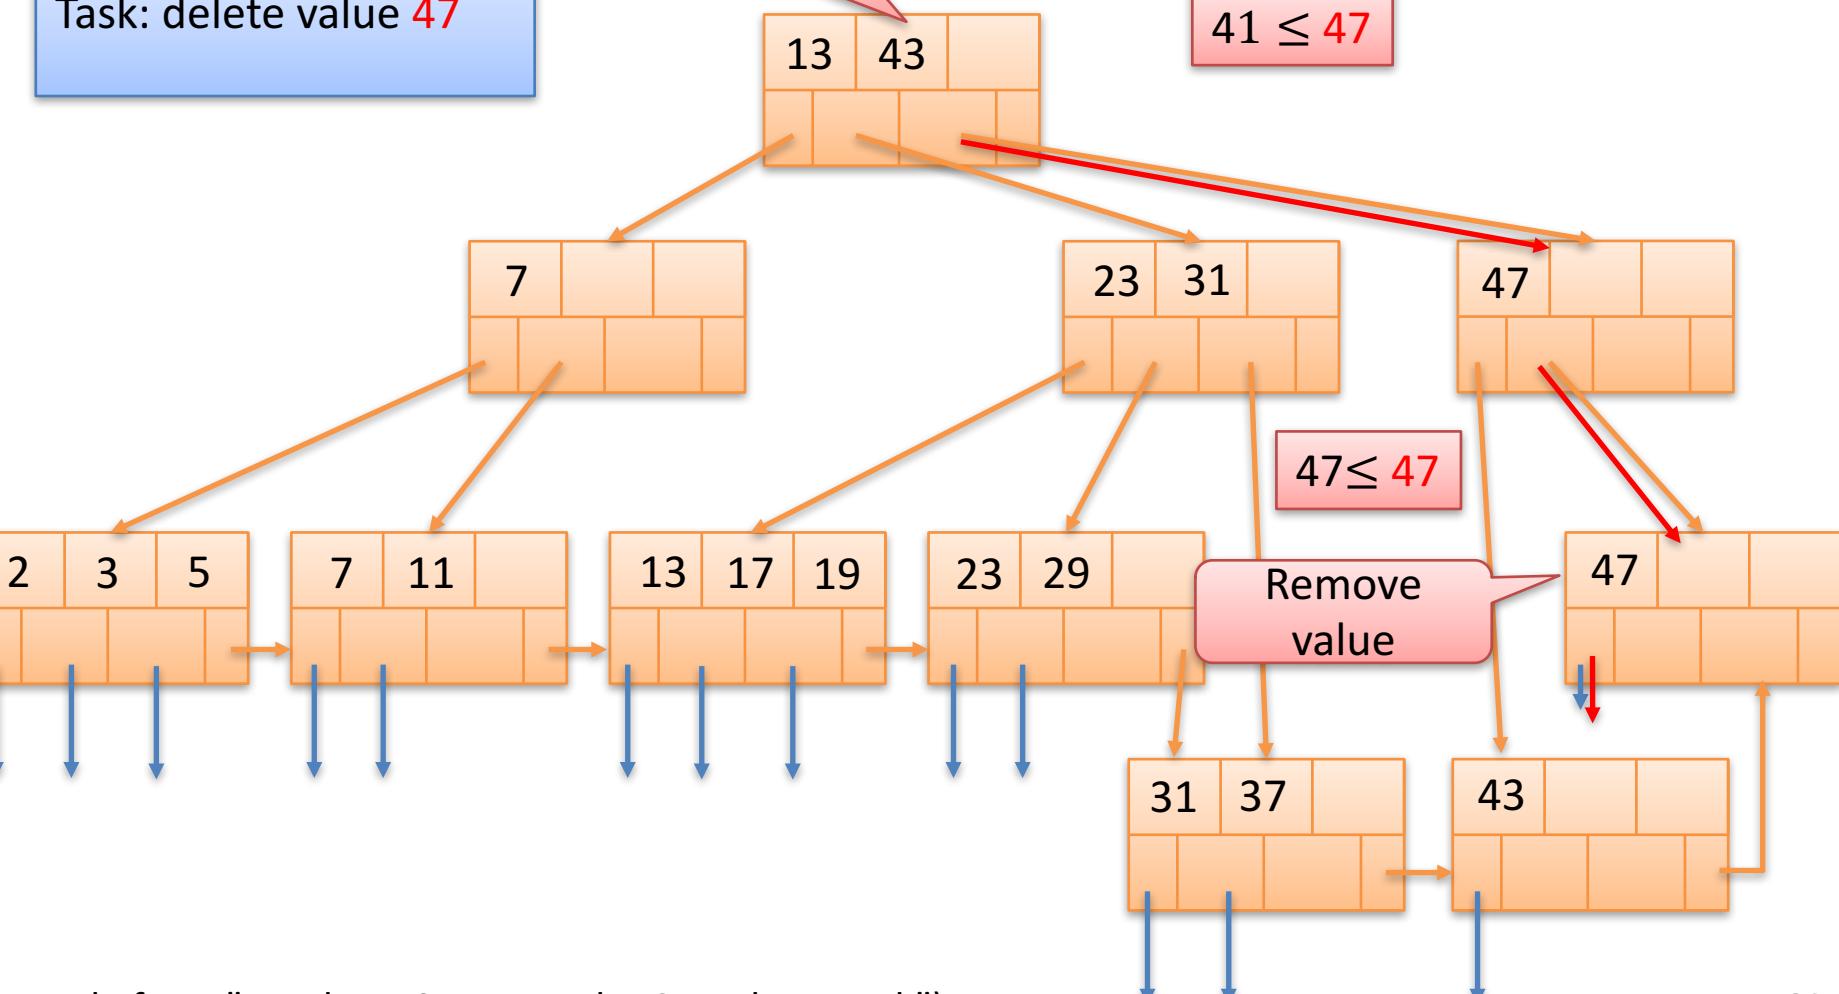
37

Deletion part 3

Task: delete value **47**

Find value

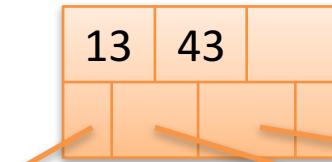
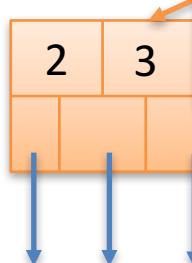
$41 \leq 47$



Deletion part 3

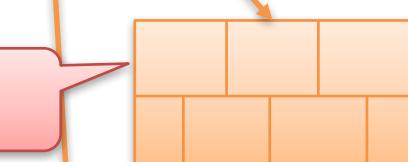
Task: delete value **47**

A node is adjacent to another, if they are at the same depth and their common ancestor has no descendant at that depth between them



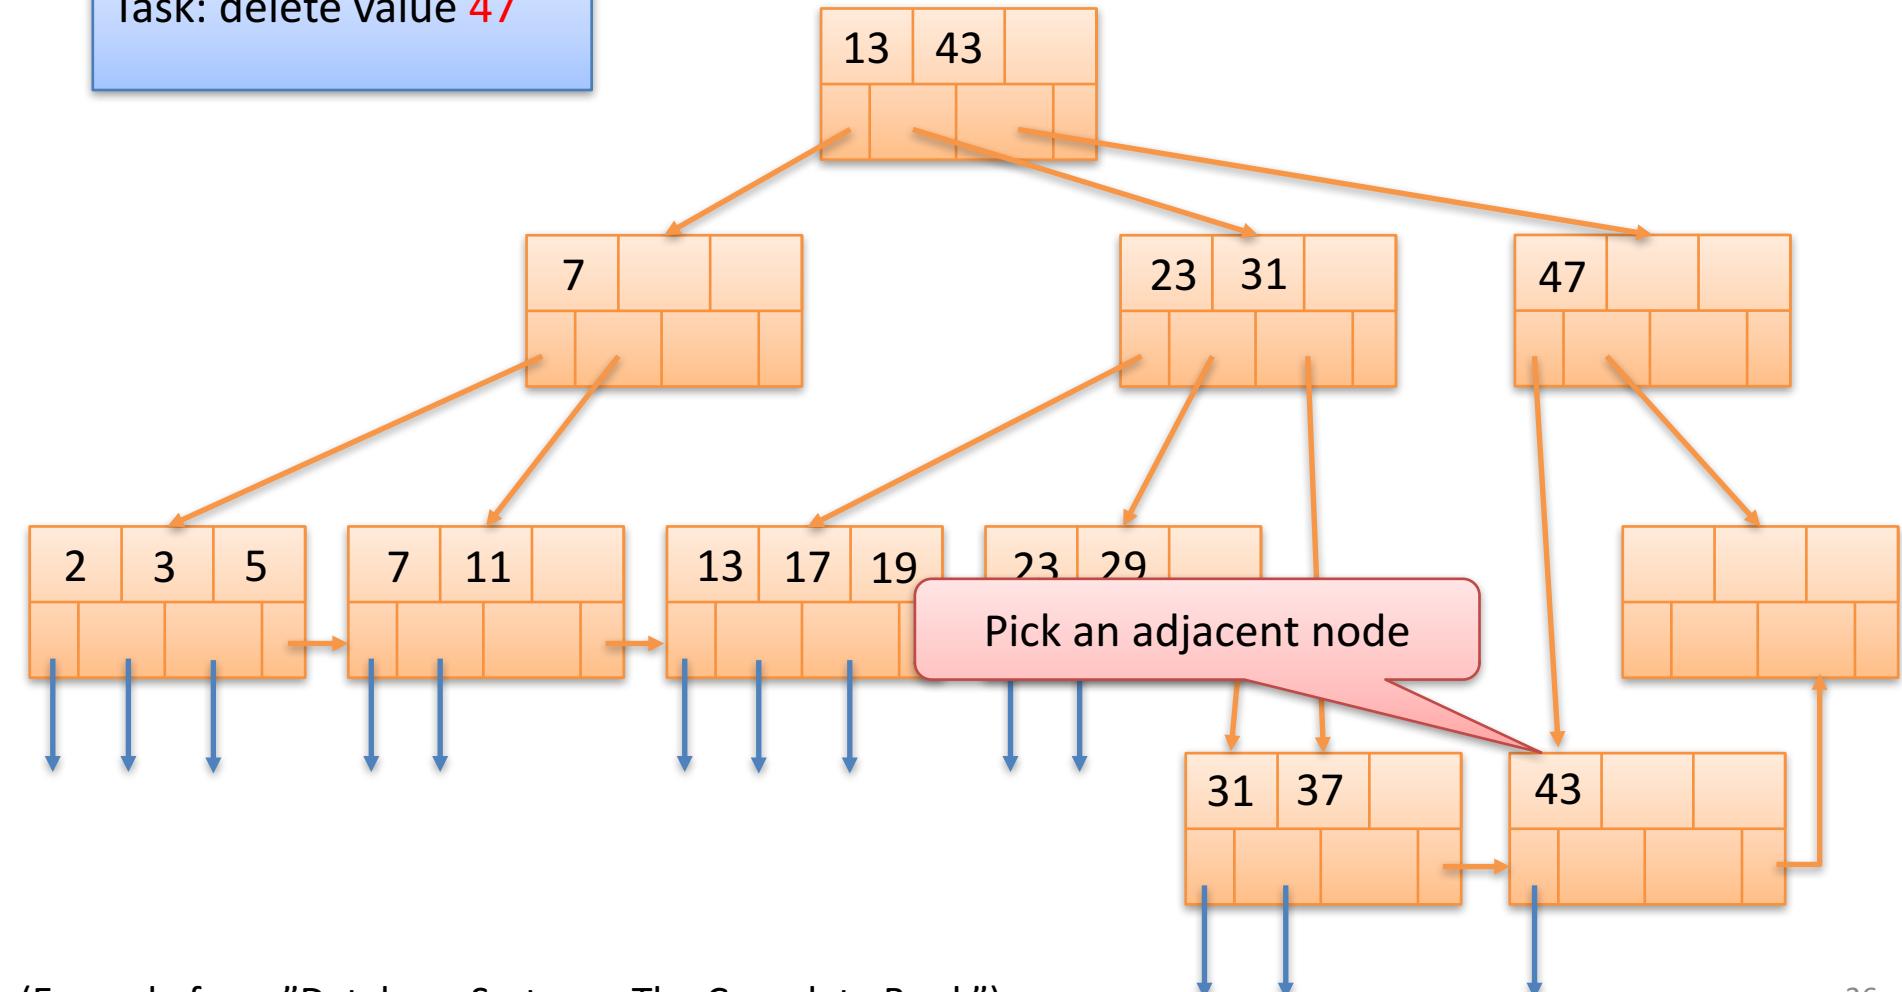
Check if adjacent node
has more than minimum

Remove
value



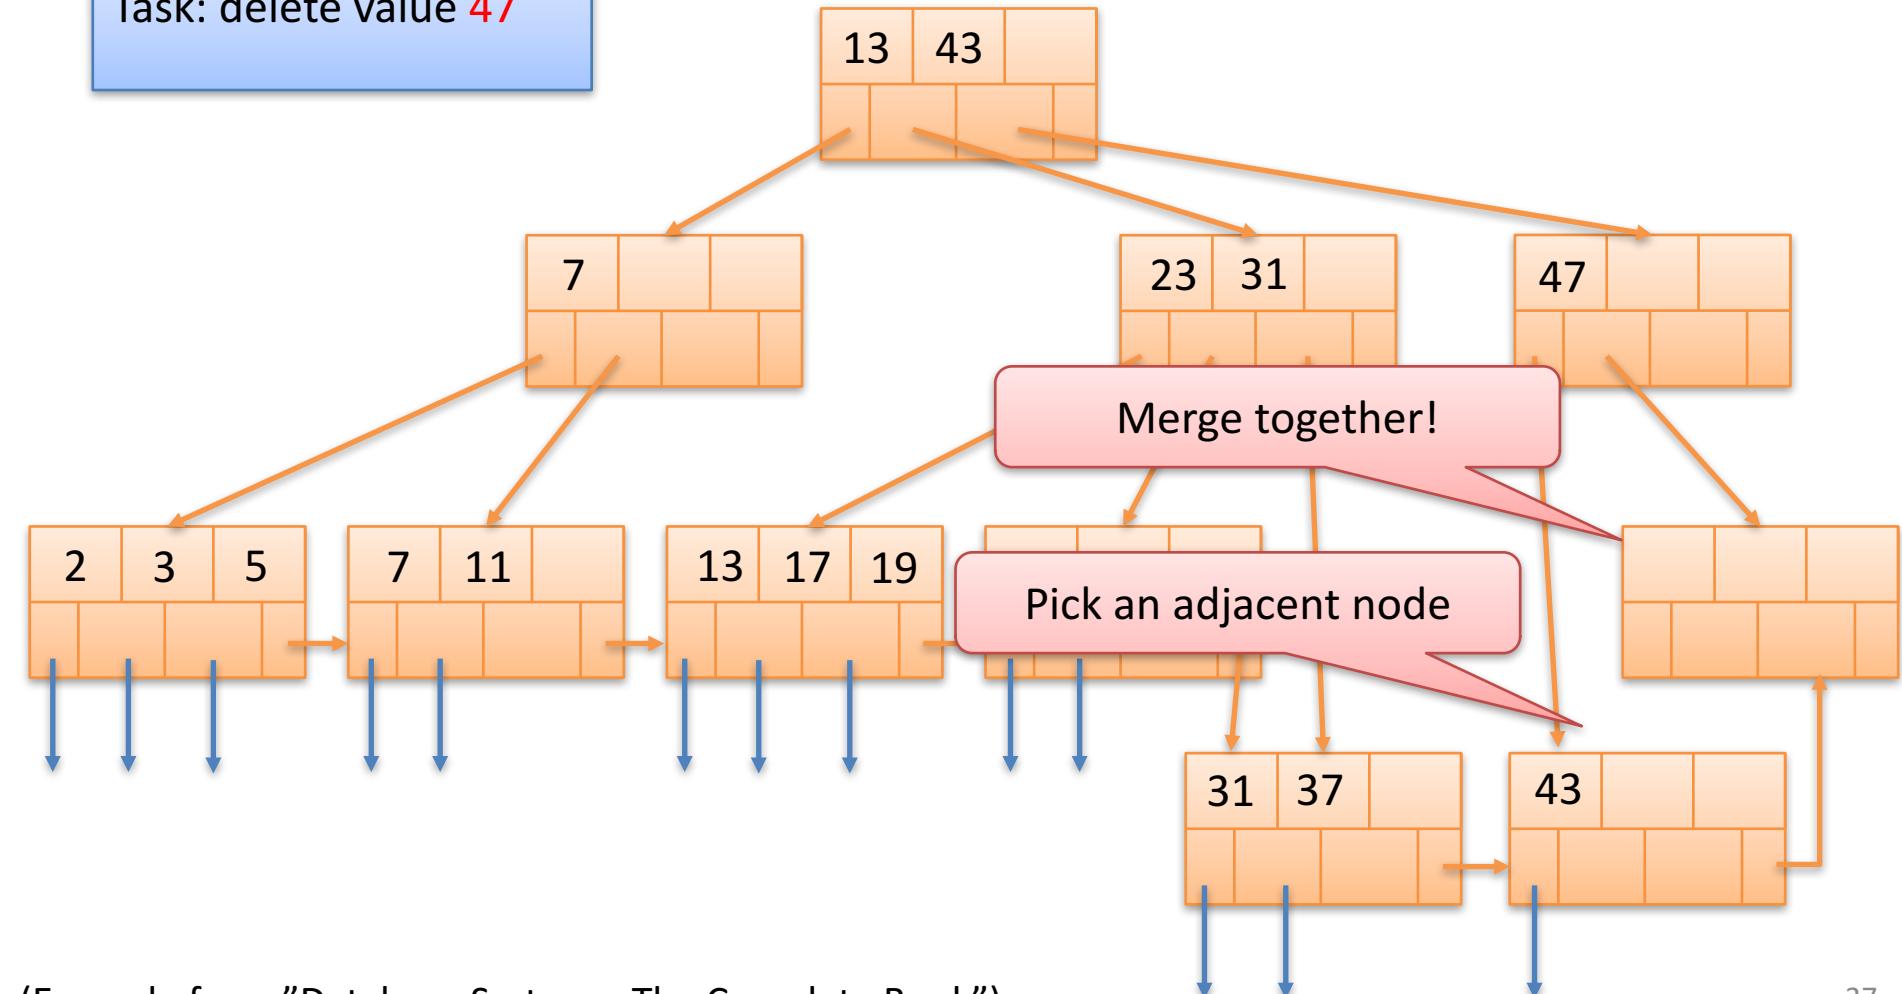
Deletion part 3

Task: delete value **47**



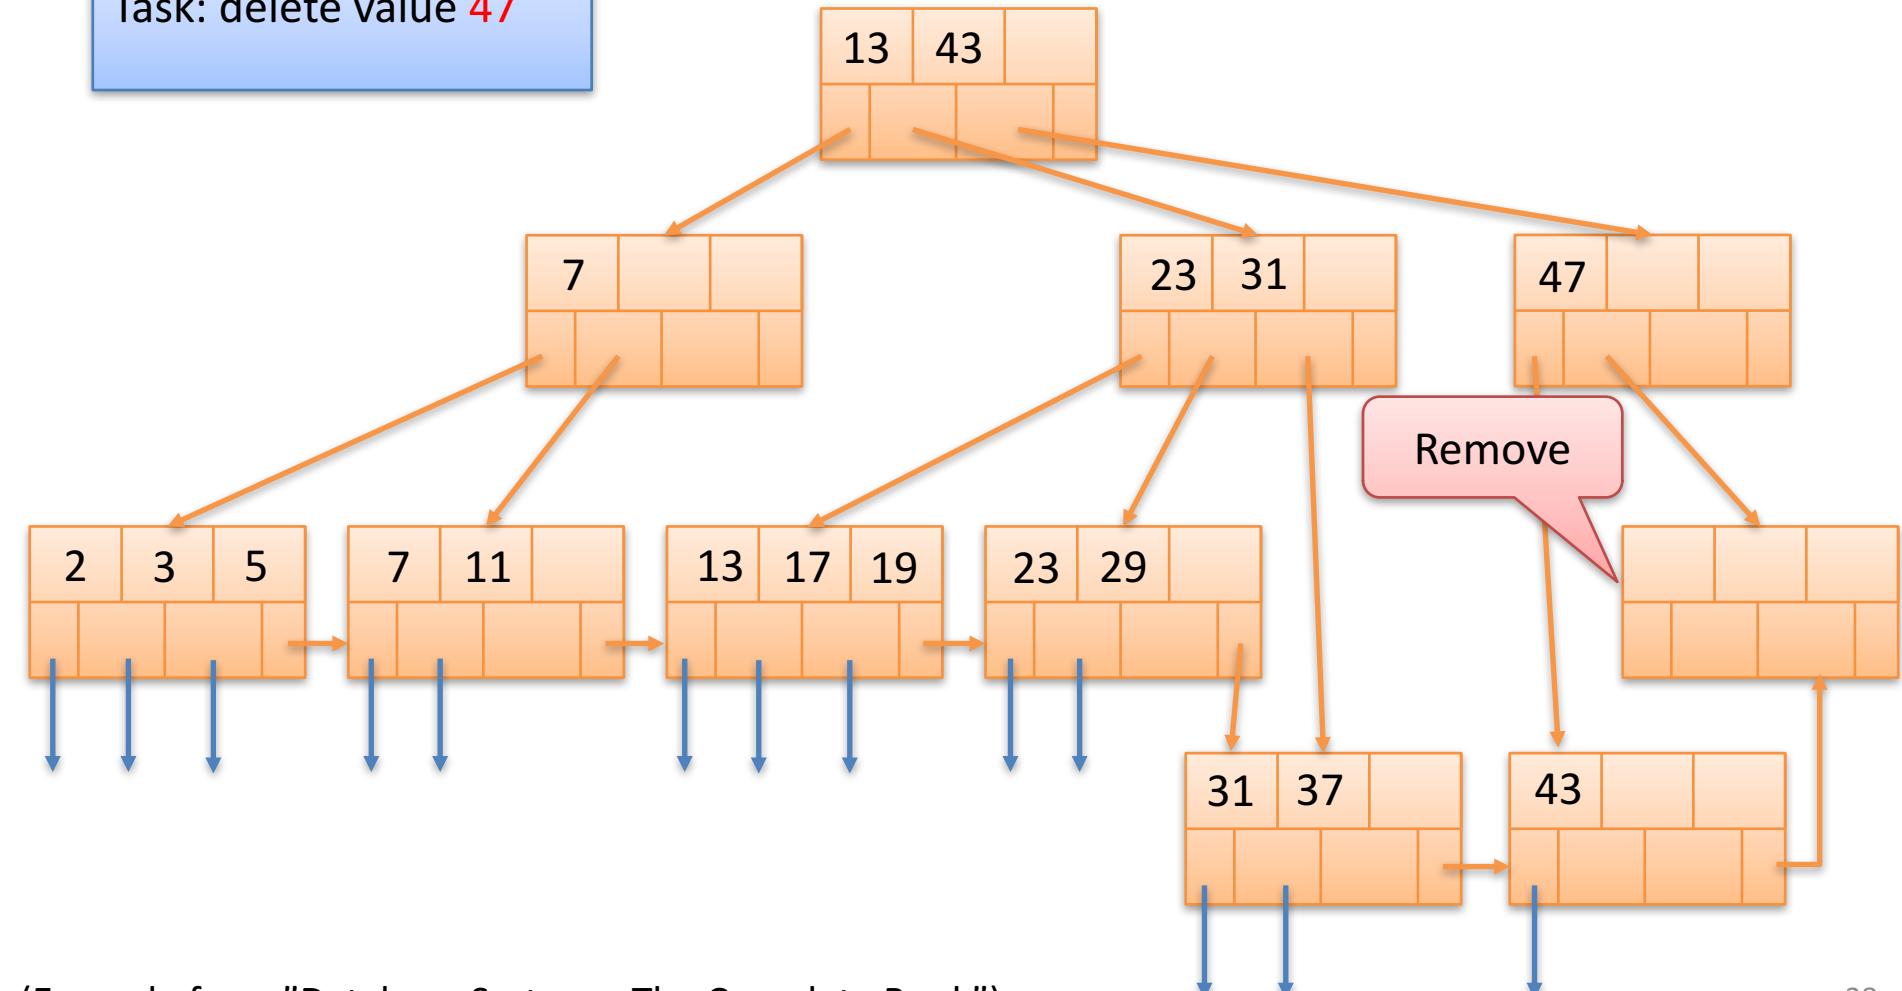
Deletion part 3

Task: delete value **47**



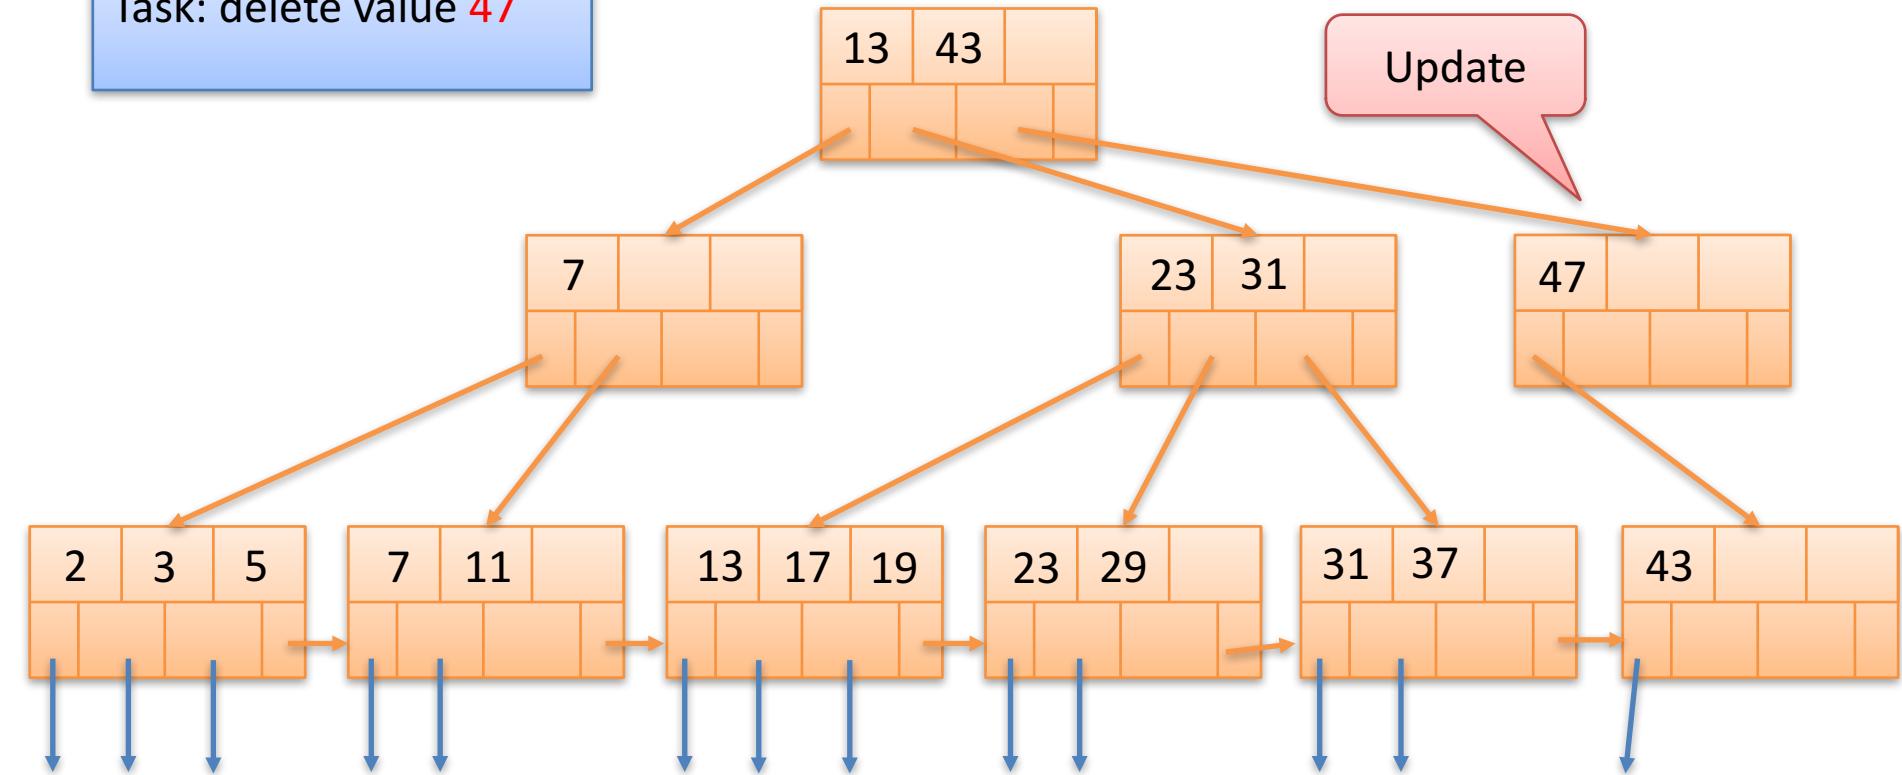
Deletion part 3

Task: delete value **47**



Deletion part 3

Task: delete value **47**

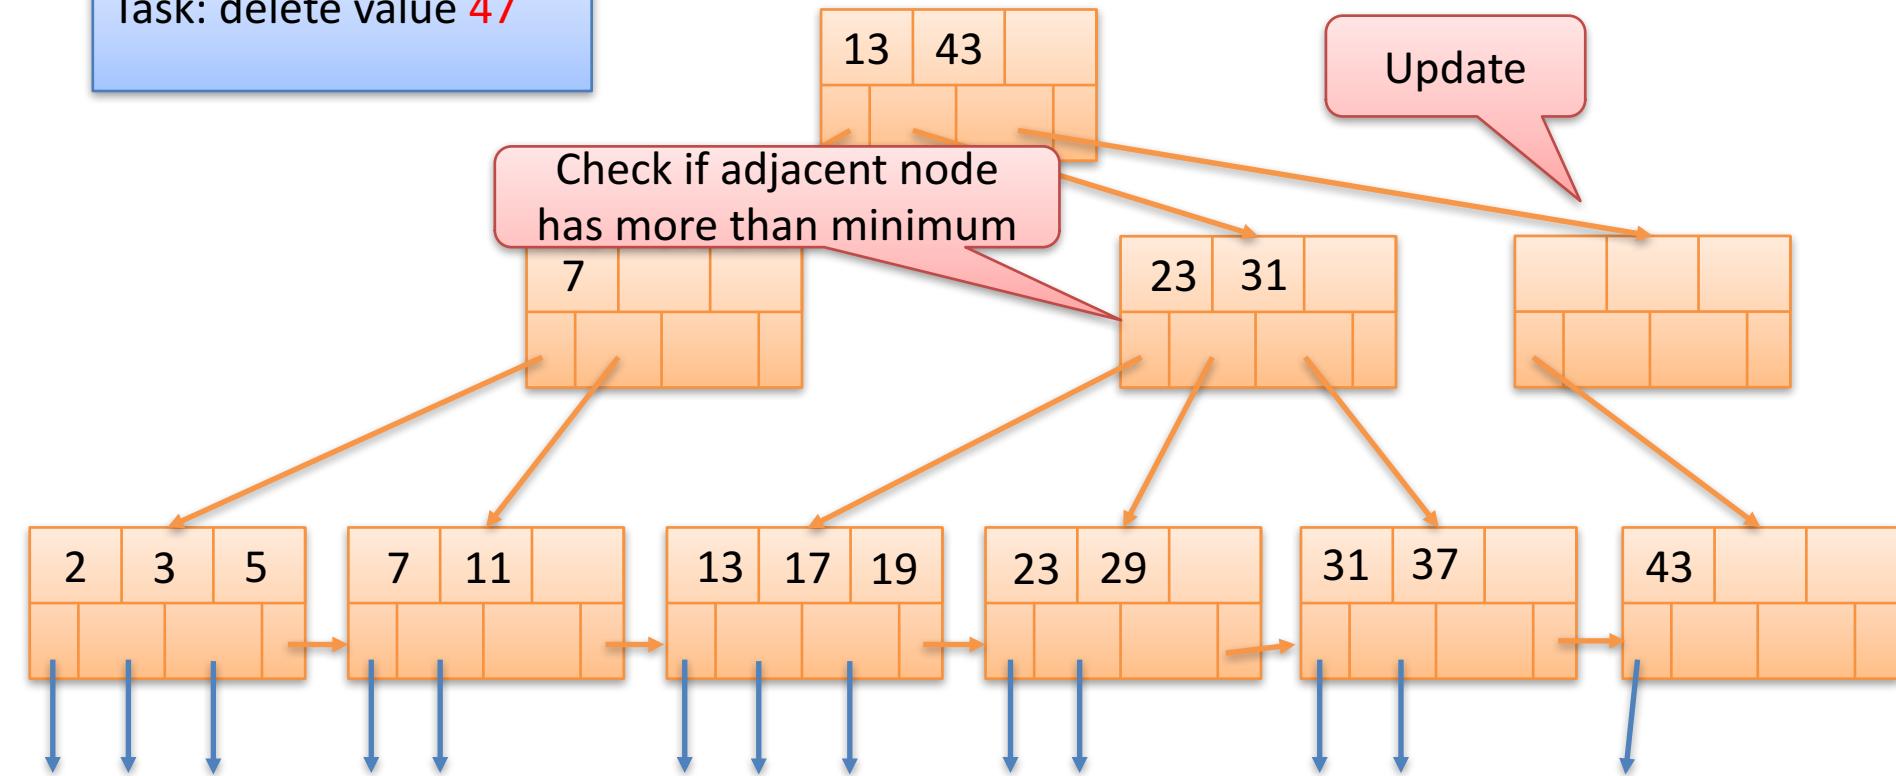


Deletion part 3

Task: delete value **47**

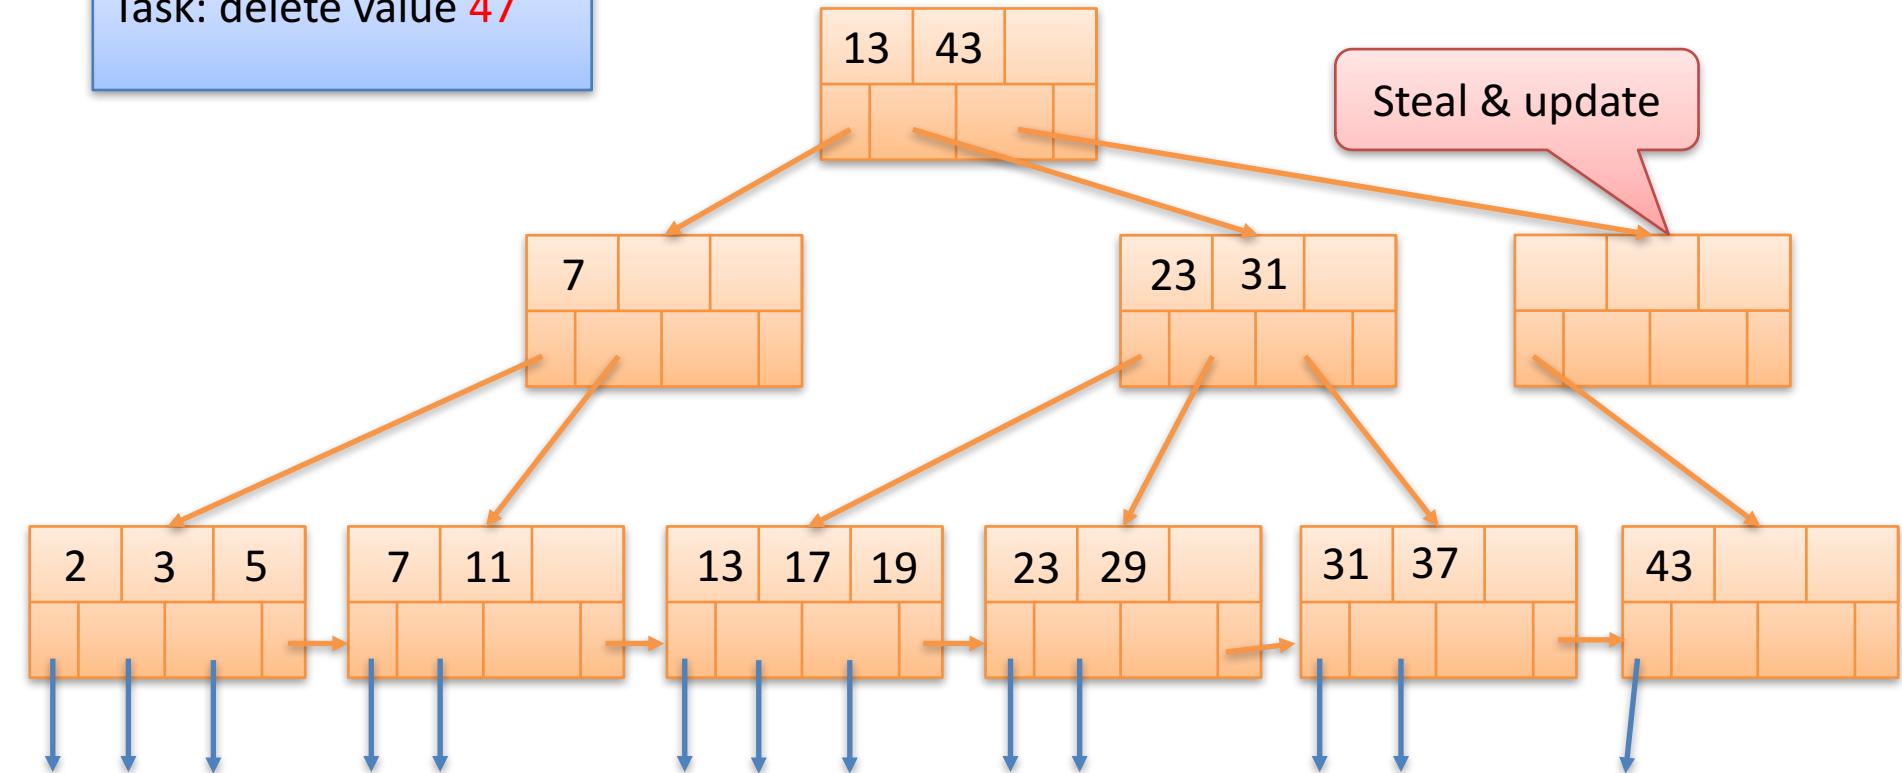
Check if adjacent node
has more than minimum

Update



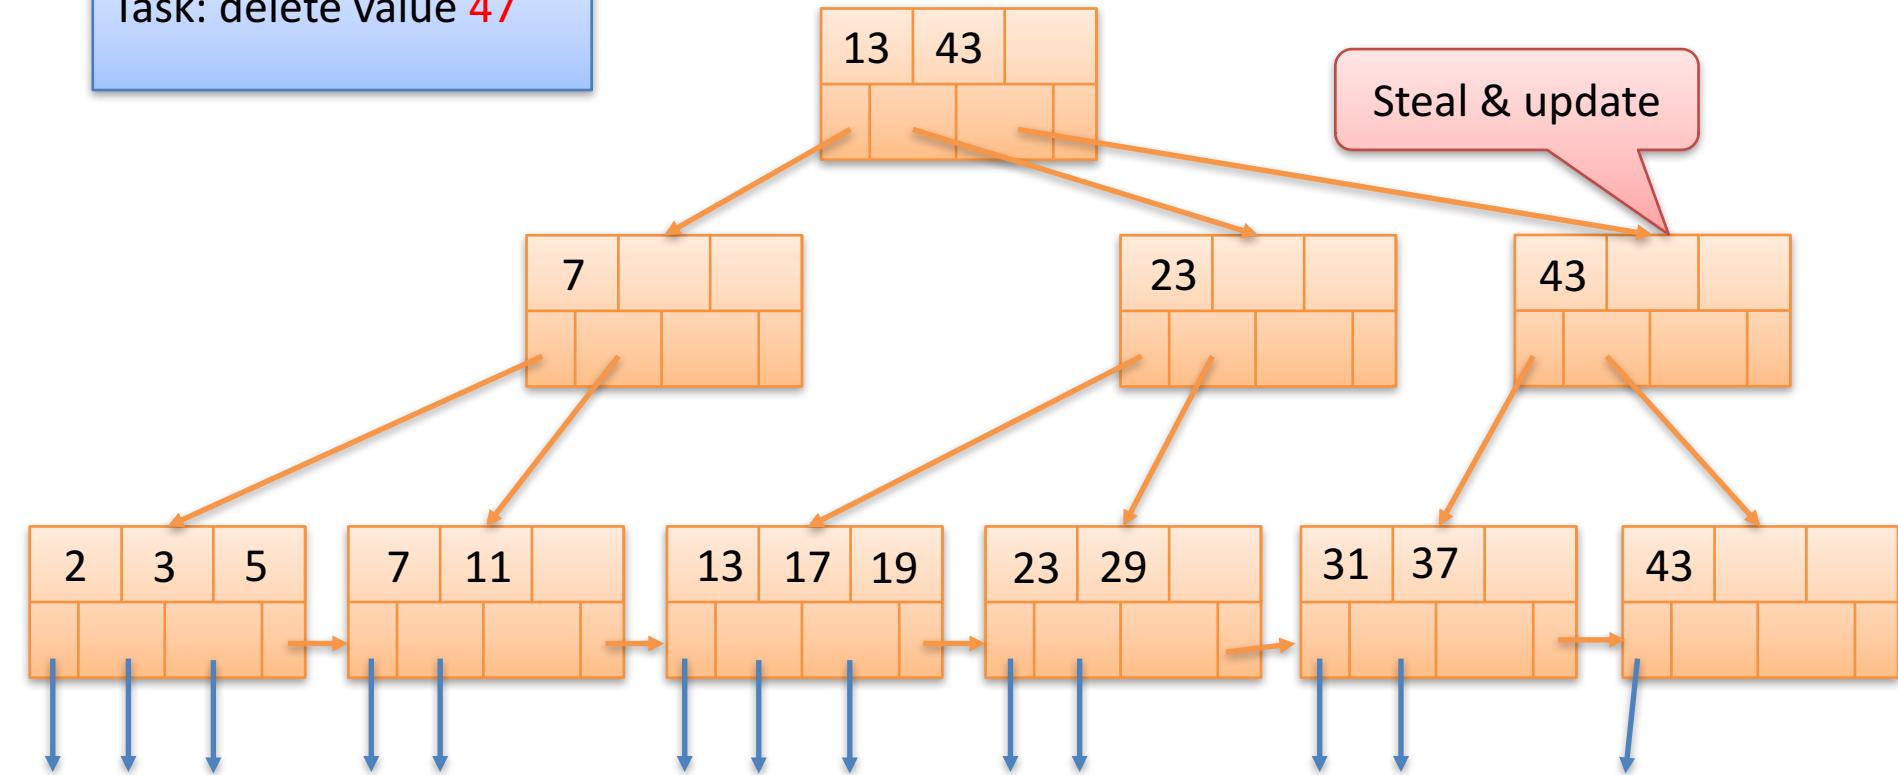
Deletion part 3

Task: delete value **47**



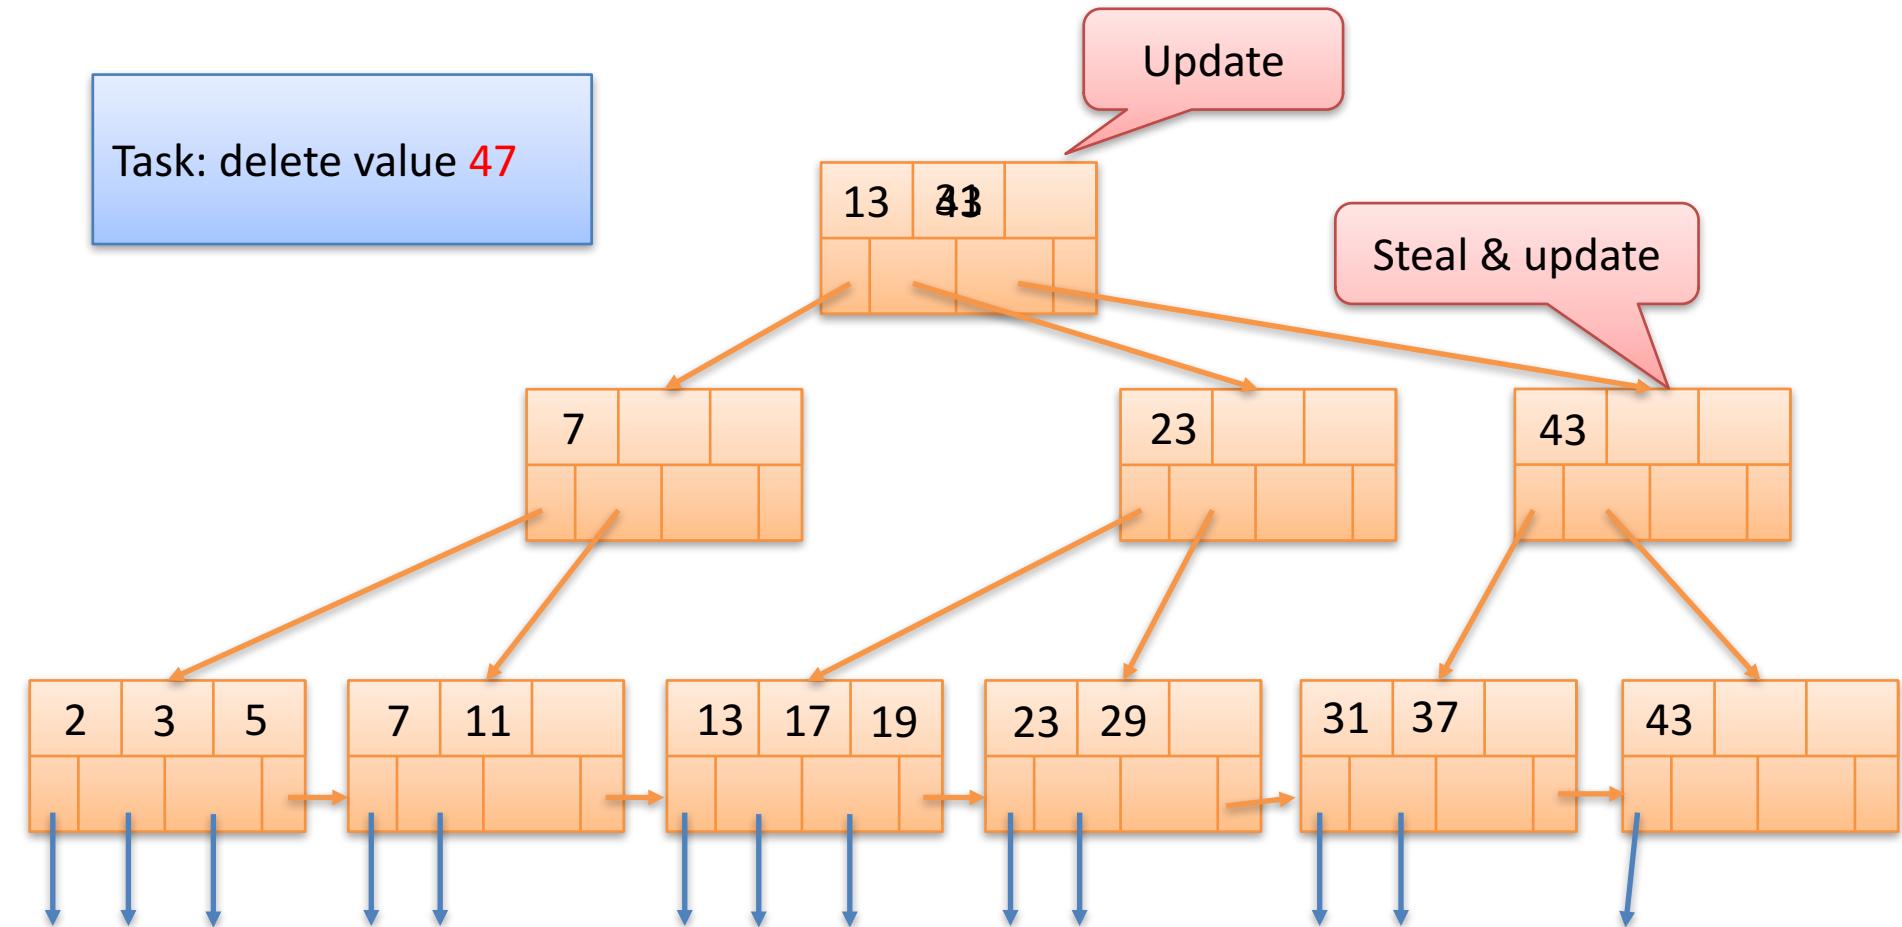
Deletion part 3

Task: delete value **47**



Deletion part 3

Task: delete value **47**



Review: Deletion

- Goal: delete a value/pointer pair
- Procedure:
 - Find the leaf that should contain the value
 - If not there: Done
 - Remove the value/pointer pair
 - Let the current node C
 - Let x be $\begin{cases} 2 & \text{if } C \text{ is root} \\ \left\lceil \frac{n+1}{2} \right\rceil & \text{if } C \text{ is internal node} \\ \left\lceil \frac{n+1}{2} \right\rceil & \text{if } c \text{ is leaf} \end{cases}$
 - If C has above x pointers: Fix ancestors (if necessary) and you are done
 - If C is the root but not a leaf: Remove it (and let the child of the root be the new root)
 - Otherwise, check if an adjacent node has at least $x + 1$ pointers
 - If so: take one, fix ancestors (if necessary) and you are done
 - Otherwise, merge with sibling and go to line 3 with the parent as current node
- The B+ tree **remains balanced!**
- Running time: $O(h \times \log_2 n)$

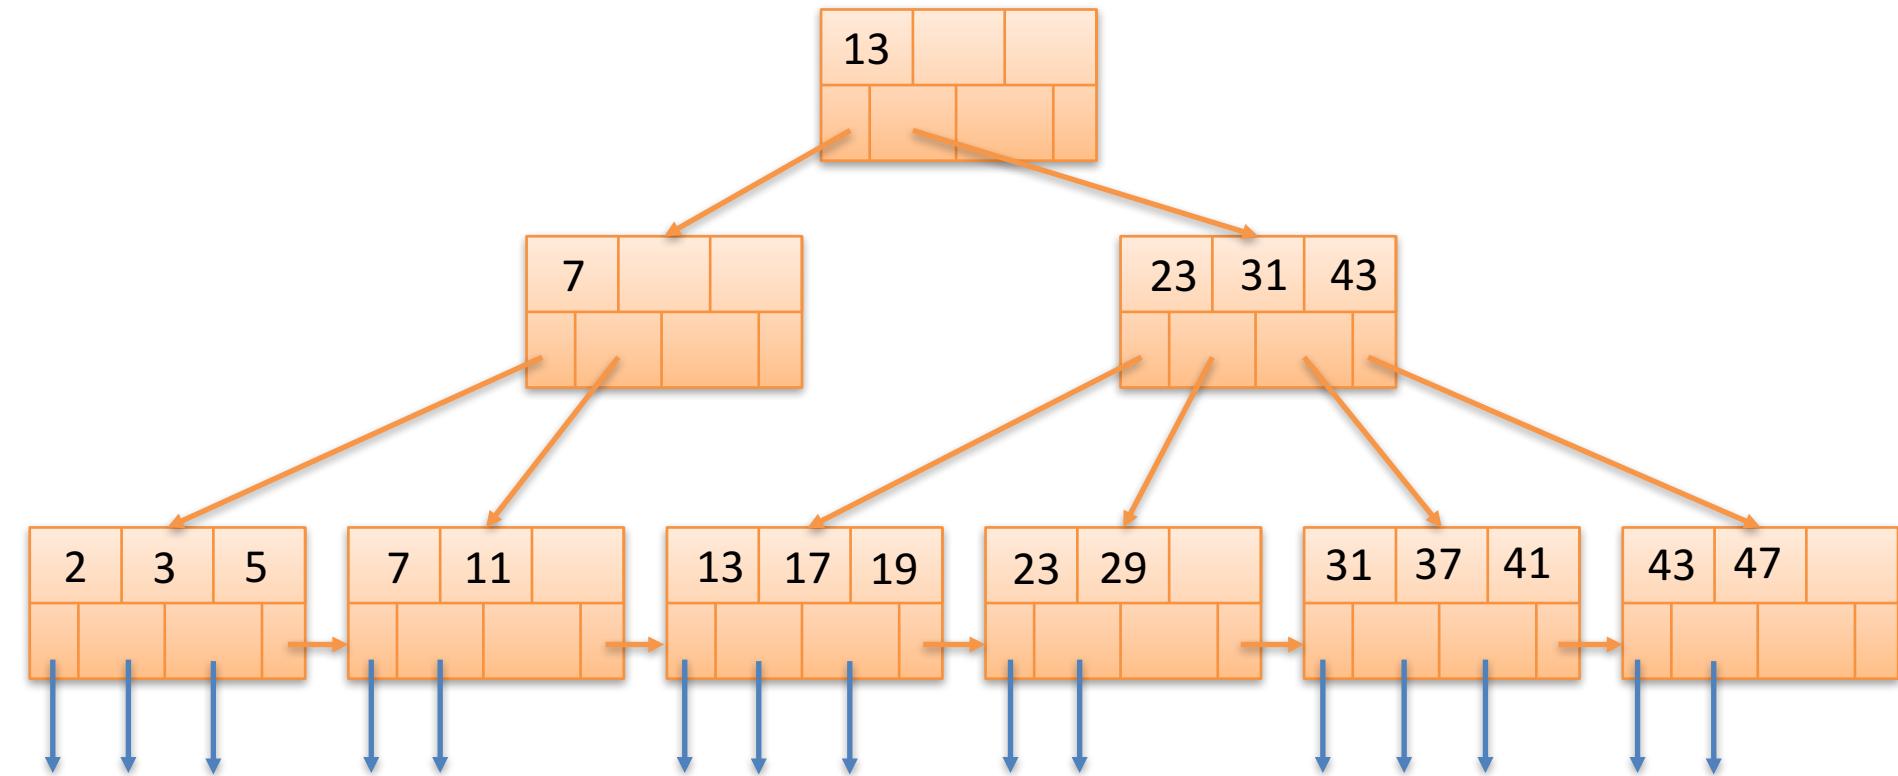
See definition slide 25

Time for a disk operation

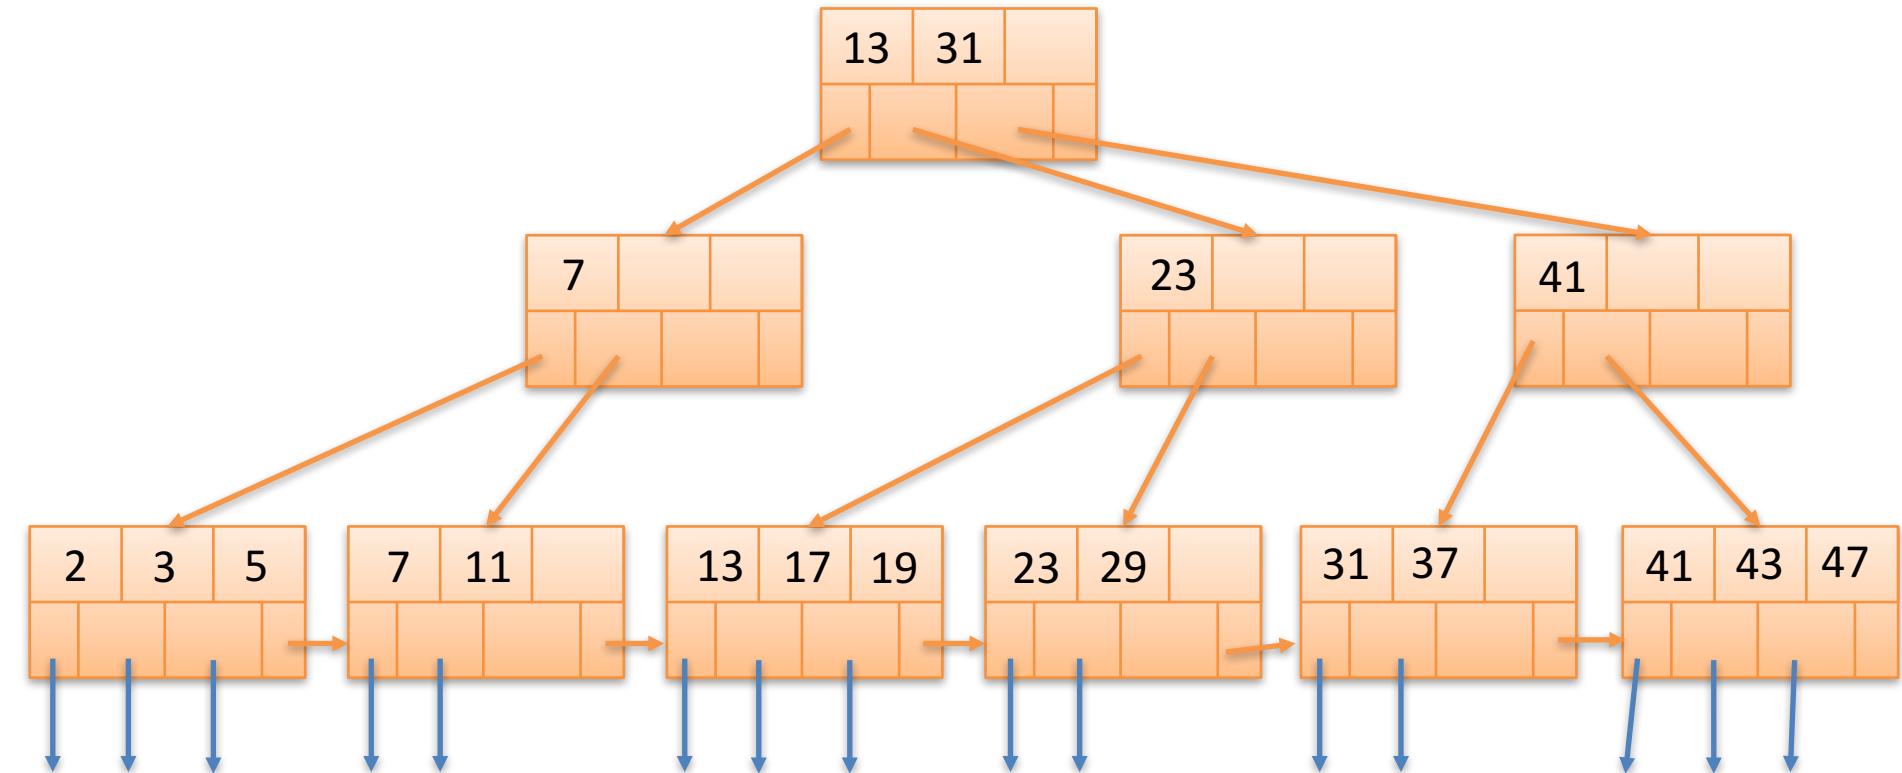
“real” running time $O(h \times D)$

Height of the B+ tree

Before insertion + deletion of 42



After insertion + deletion of 42



Properties of B+ Tree Indexes

- Fast lookups, insertions, deletions in time:

$$n \approx B$$

$$\begin{aligned} O(\text{height of B+ tree} \times \log_2 n) &= O(\log_n N \times \log_2 n) \\ &= O(\log_2 N) \end{aligned}$$

- Remain **balanced**
- **Huge capacity** even with height 3 if blocks large enough

- Block size: 16386 bytes (16 kilobytes)
- Values stored in index: 4 bytes
- Pointers: 8 bytes
- Largest n so that each B+ tree node fits into a block (i.e., $4n + 8(n+1) \leq 16386$) is 1364
- B+ trees with height 3 can store $> n^3 = 2537716544$ values

Properties of B+ Tree Indexes

- Can be implemented **efficiently with respect to number of disk accesses**
 - Number of disk accesses typically $\approx 2 + \text{height of B+ tree}$
- Most of the B+ tree can be kept in memory
 - Upper levels
 - Even with block size of 16384 bytes and $n = 1364$:
 - Level 1 (root) $\approx 16 \text{ KB}$
 - Level 2 (children of root) $\approx (n+1) \times 16 \text{ KB} \approx 21 \text{ MB}$
 - Level 3 $\approx (n+1) \times (n+1) \times 16 \text{ KB} \approx 28 \text{ GB}$

Typically, these are the leaf nodes and can be loaded from disk on demand

Summary

- Indexes allow DBMS to **quickly find and navigate** to desired tuples of a relation
- **Can speed up computation** of selections & joins
- Many **different forms** of indexes
 - B+ trees
 - Hash tables (not covered here)
 - ...
- Trade off:
 - More indexes help answer more queries faster
 - But more indexes make updates slower