


# COMP207

# Database Development

Lecture 19


Digression: The MapReduce Framework

# Relational Databases are Powerful



- Allow us to store, manage, and analyse large datasets
- Some data analysis applications go beyond what such systems can handle...

# Importance of Web Pages



- Graph with all relevant HTML pages on the web
- Task: rank pages by “importance”
  - Basis of Google Search
  - $\approx$  an iterated matrix product (number of rows = number of pages)

**Problem:** millions to billions of web pages

# Finding Patterns in Large Datasets

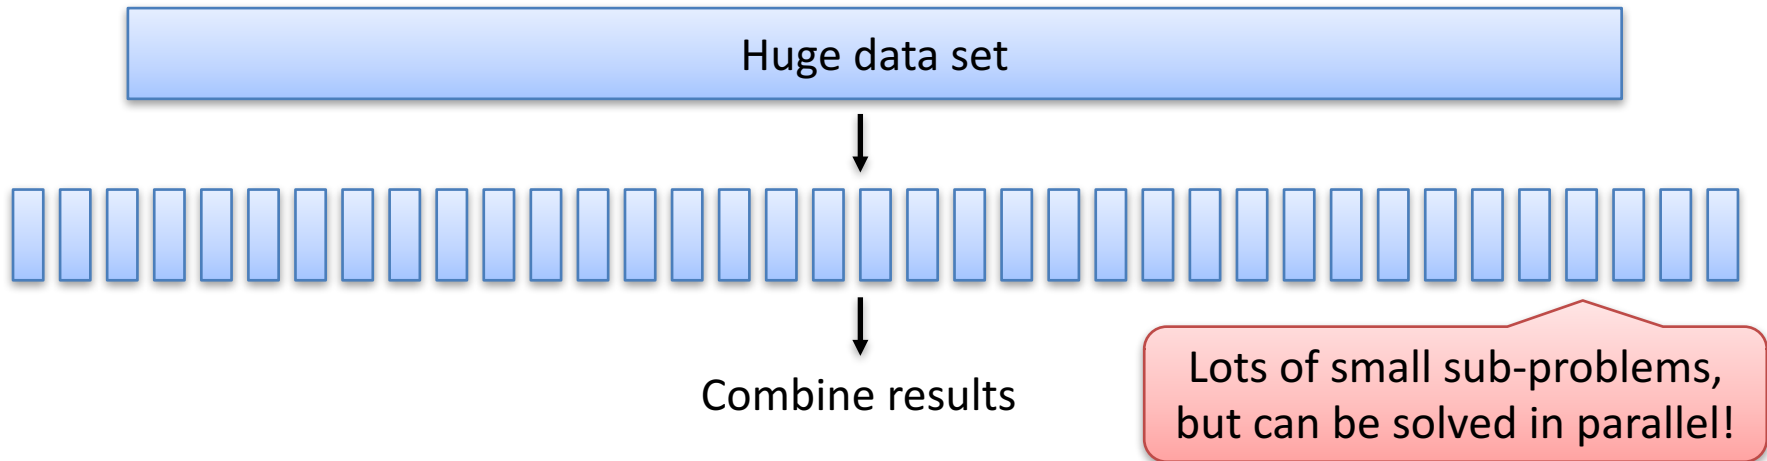
“Big Data”

- Predictive maintenance, e.g.
  - Which parts of a system are likely to fail or cause certain effects?
  - Learn from sensor data, infrastructure data, etc.
- Detect interesting patterns, e.g.
  - Interesting phenomena in space, on earth, under water, in stock markets, etc.
  - Recognise plants with certain diseases
  - Estimate the likelihood of archeological artifacts at a certain site based on archeological data, geological parameters, etc.
- Understand the structure/dynamics of large networks, e.g.
  - Train networks
  - Social networks

In these applications, the data is potentially huge.

# Divide and Conquer

- Often such problems can be solved as follows:
  - Divide data into smaller chunks
  - Perform computations on smaller chunks
  - Combine results



- Idea behind MapReduce!

# MapReduce

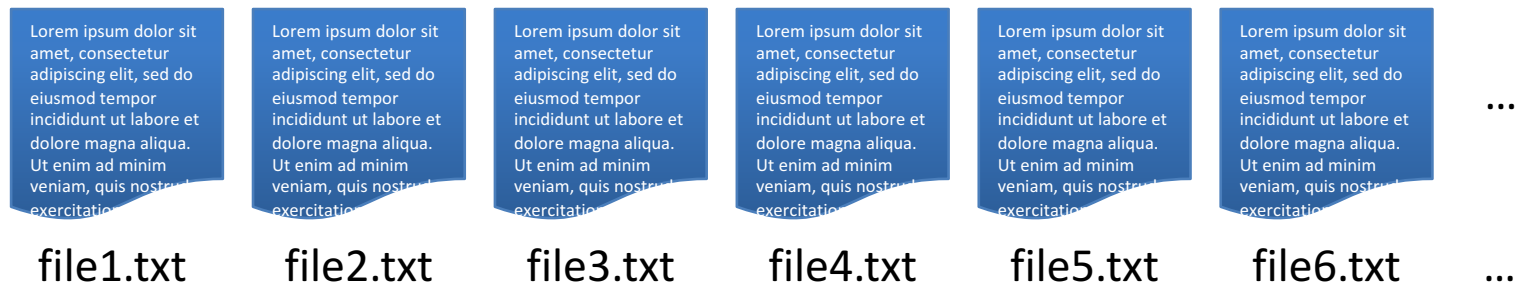
- Programming framework
  - Provides simple means to create programs that can scale to terabytes of data
- Users only implement two methods
  - **Map**: the computation on the smaller chunks of data
  - **Reduce**: how the results are combined to the final result
  - MapReduces manages the rest!
- Implementations:
  - MapReduce: Google's original implementation
  - Apache Hadoop: open-source

# Counting Words in Documents

or: Hello World!

# The Task

- Given: a collection of “not too large” text files

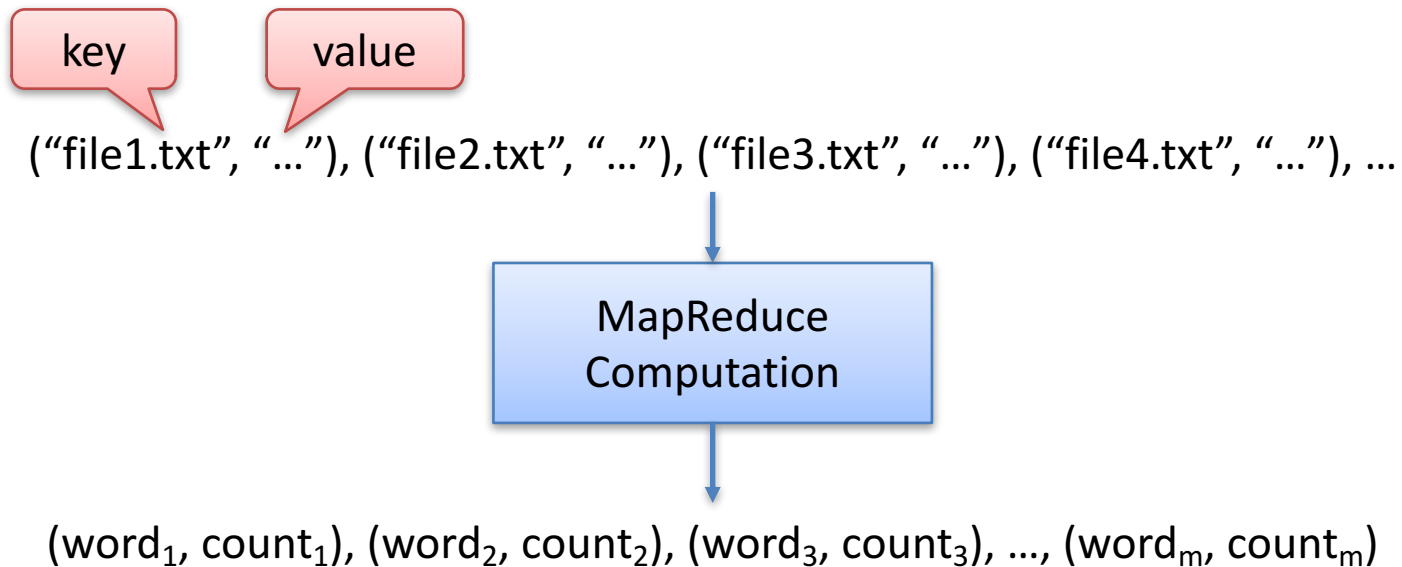


- Task: for each word **w**, count how often **w** occurs in these files



# How To Do This In MapReduce?

- The MapReduce computation:



- To implement this computation, we have to write two functions: **Map & Reduce**


# The Map Function

- The Map function is applied to a single key/value pair and produces a list of zero or more key/value pairs

**Map(String filepath, String contents):**  
for each word **w** in **contents**:  
output pair (w, "1")

- Example:


("file1.txt", "Hello World! Hello UK! Hello Liverpool!")



("Hello", "1"), ("World!", "1"), ("Hello", "1"), ("UK!", "1"), ("Hello", "1"), ("Liverpool!", "1")

# Grouping

- After the Map functions have been applied, we group all values by key



- The list of values for a key is the input to Reduce

# The Reduce Function

- The Reduce function takes a key and a list of values as input and outputs a list of key/value pairs

```
Reduce(String word, Iterator<String> values):
```

```
    int count = 0
```

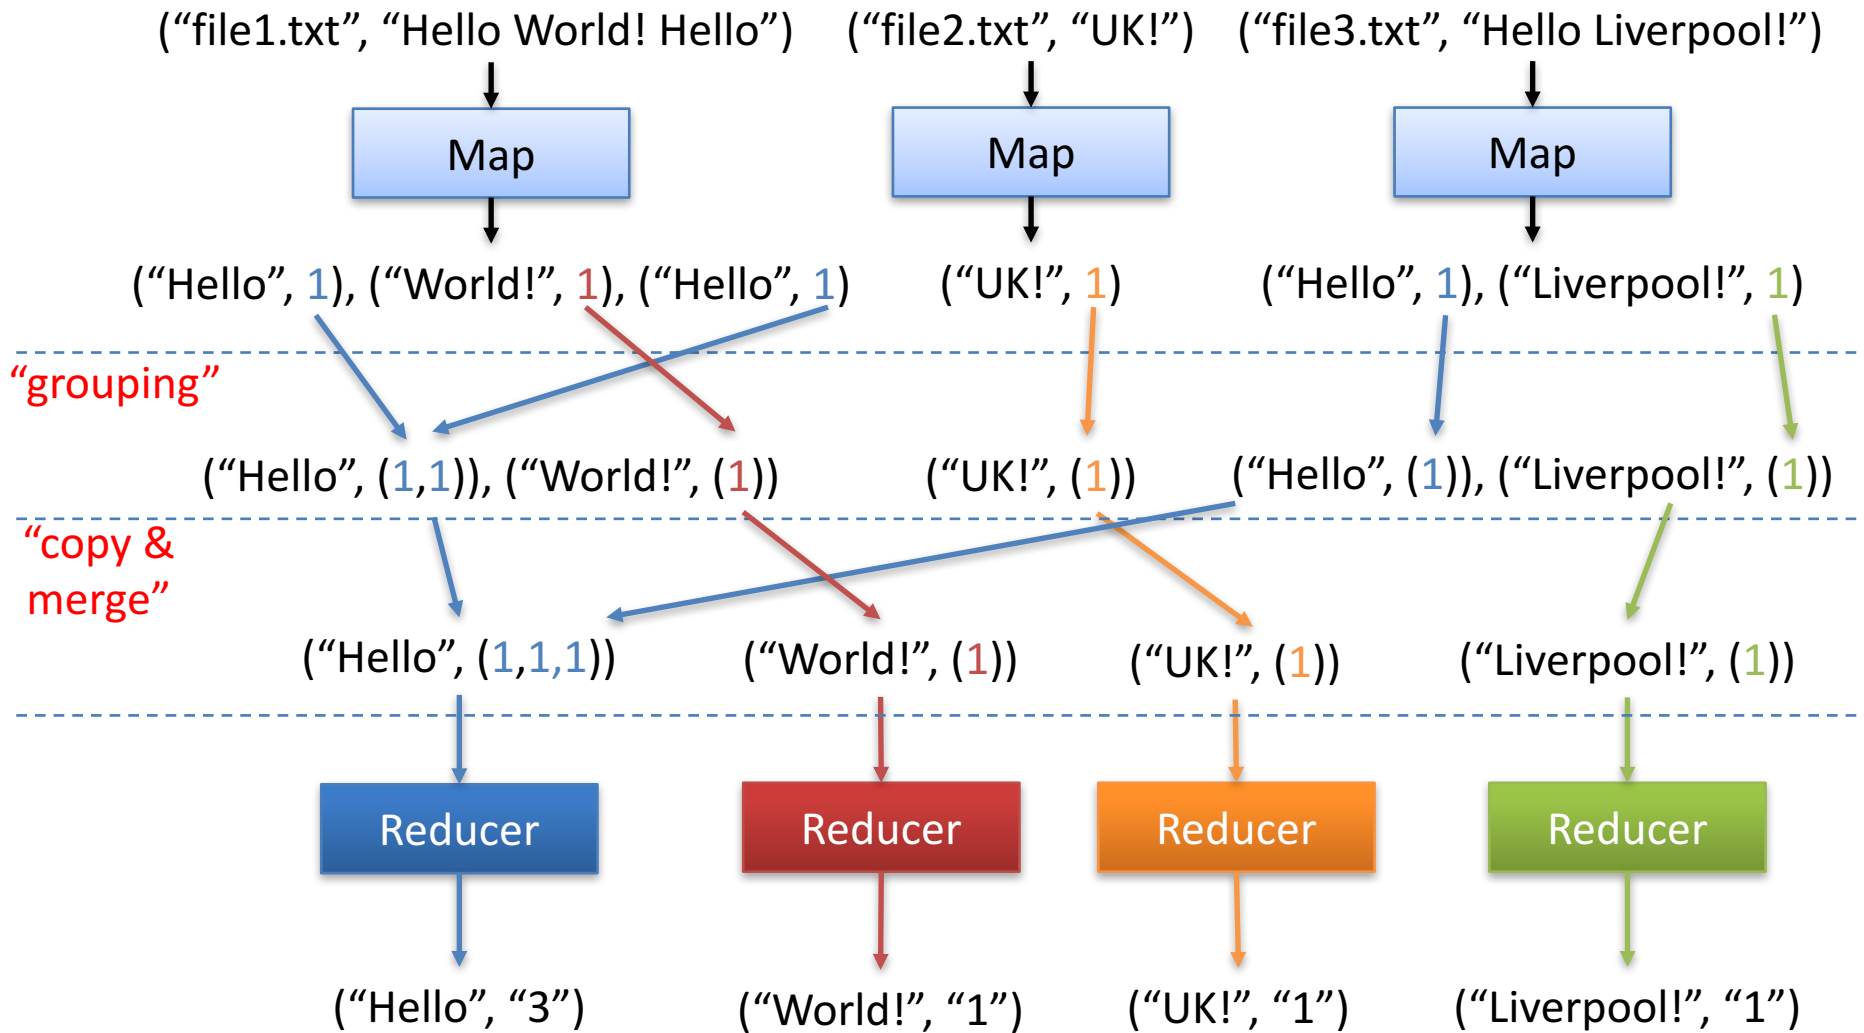
```
    for each v in values:
```

```
        count = count + parseInt(v)
```

```
    output pair (word, toString(count))
```

- Example:
  - Reduce("Hello", ("1","1","1")) outputs ("Hello", "3")

# Putting Things Together



# The Beauty of It...

- We only had to implement these two functions...

**Map(String filepath, String contents):**

for each word **w** in **contents**:

output pair (w, "1")

**Reduce(String word, Iterator<String> values):**

int **count** = 0

for each **v** in **values**:

count = count + parseInt(v)

output pair (**word**, toString(**count**))

- MapReduce takes care of everything else:
  - Parallel execution, including allocation of Map/Reduce to machines
  - Failure

# Matrix multiplication

- Matrix multiplication: Given (r,s)- and (s,t)-matrix  $M$  and  $N$  resp., compute  $P$  s.t.

$$P_{i,k} = \sum_{j=1}^s M_{i,j} N_{j,k}$$

- ( $P$  can be computed in  $O(n^{2.373})$  for  $n = r = s = t$ )
- Easiest (also faster) for MapReduce:  
use 2 MapReduce computations

# First MapReduce

- matrix  $\in \{N, M\}$  and triplet  $\in \mathbb{Z}_+ \times \mathbb{Z}_+ \times \mathbb{R}$

**Map**(String **matrix**, String **triplet**):

let **(i,j,v) = triplet**

if **matrix** = *N* then output pair **(i, (N,j,v))**

else output pair **(j, (M,i,v))**

If  $M_{i,j} = v$  then **(M,"(i,j,v)")** is an input, similar for *N*

**Reduce**(int **no**, Iterator<String> **values**):

for each **(i,j,k)** in **values**:

if **i=M** then:

for each **(i',j',k')** in **values**:

if **i'=N** then:

output pair **((j,j'),  $k \times k'$ )**

Creates the pair **((i,k),  $M_{i,j}N_{j,k}$ )**



# Second MapReduce

```
Map(pair<int,int> pair, double no):  
    output pair (pair,no)
```

```
Reduce(pair<int,int> pair, Iterator<double> nos):  
    double result=0  
    for each no in nos:  
        result=result + no  
    output pair (pair,result)
```

# Relational algebra

- Use MapReduce to implement relational algebra
- Tuples in  $R$  are like  $(R, att_1 = val_1, att_2 = val_2, \dots)$
- Input is  $(\mathbf{t}, \mathbf{t})$  where  $\mathbf{t}$  is some tuple
- Two examples:
  - Selection  $\sigma_c(R)$
  - Natural join  $R \bowtie S$

# Selection $\sigma_c(R)$

**Map**(String **tuple**, String **tupleCopy**):

if **tuple** satisfies **c** then output pair (**tuple**, **tupleCopy**)

**Reduce**(String **tuple**, Iterator<String> **tupleCopies**):

output pair (**tuple**,**tuple**)

- Similar for everything but joins (and intersection and difference): Map does all the work

# Natural join $R \bowtie S$

- Scheme:  $R(A,B)$  and  $S(B,C)$  –  $A,B,C$  are sets of attributes

**Map**(String **tuple**, String **tupleCopy**):

Let **Co** be the attribute and value pairs in **B**  
output pair (**Co**, **tuple**)

**Reduce**(String **Co**, Iterator<String> **tuples**):

for each **r** in **tuples**:

if **r.relation** =  $R$  then:

for each **s** in **tuples**:

if **s.relation** =  $S$  then:


output pair  $(r \bowtie s, r \bowtie s)$

- Similar for other joins

# MapReduce: Closer Look

# Distributed File System

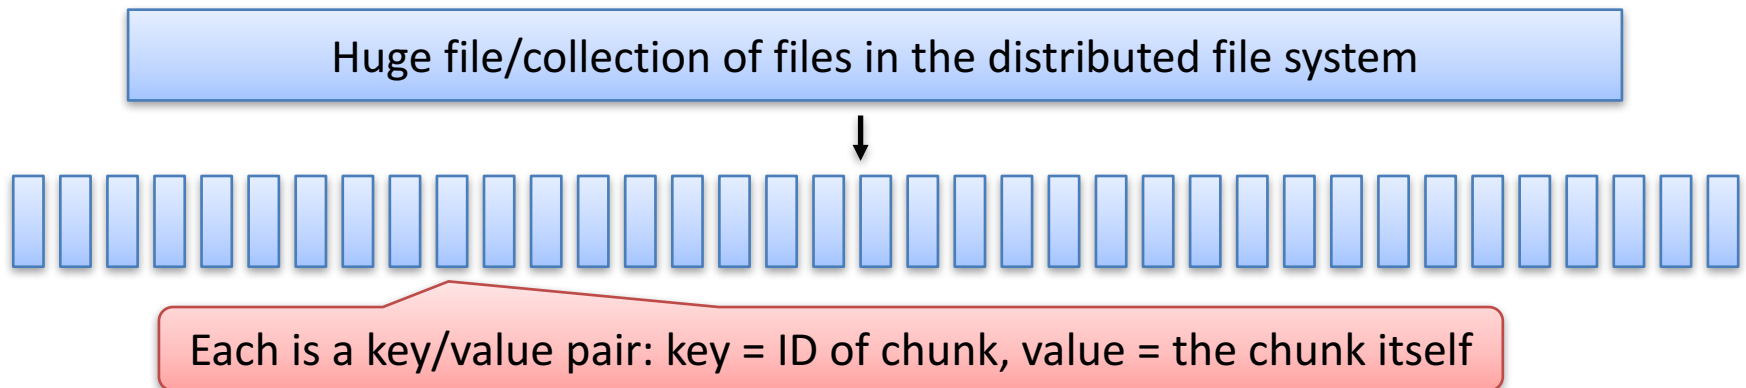
- MapReduce operates on a distributed file system



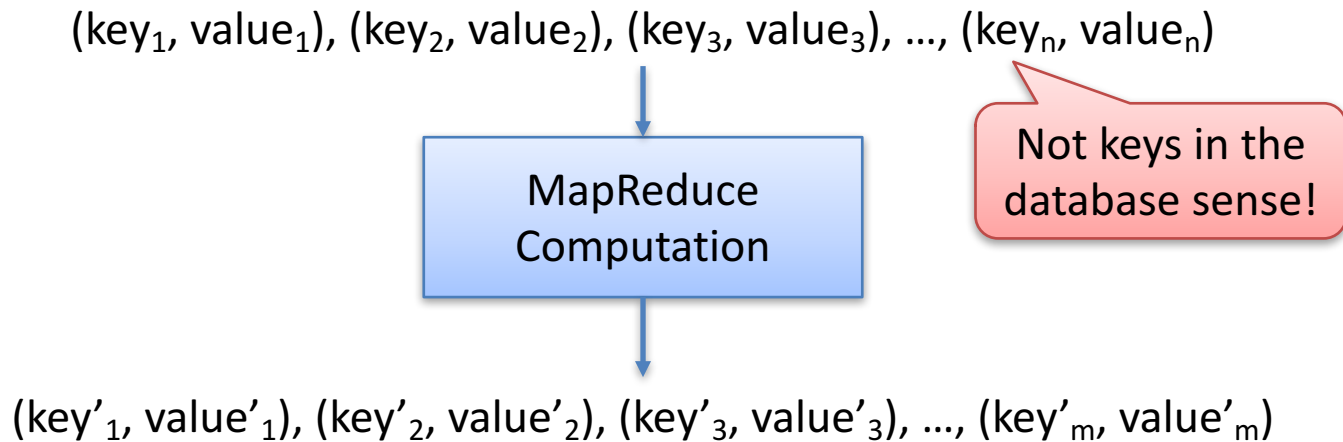
- Similar to distributed DBMS
  - All the data is distributed over the nodes/sites
  - Replication is used for fault-tolerance
- Appears to the user as if it was a regular file system
- Implementations: Google File System, Hadoop File System

# Inputs To MapReduce Programs

- Input files to MapReduce programs can be very large (GBs/TBs in size)
  - Collections of web pages, links in a social network, Twitter feeds, stock market data, ...
  - Satellite images, data from scientific experiments, ...
- MapReduce splits these into chunks (~16-64 MB) and provides an ID for each chunk (e.g., the filename)



# MapReduce Computations



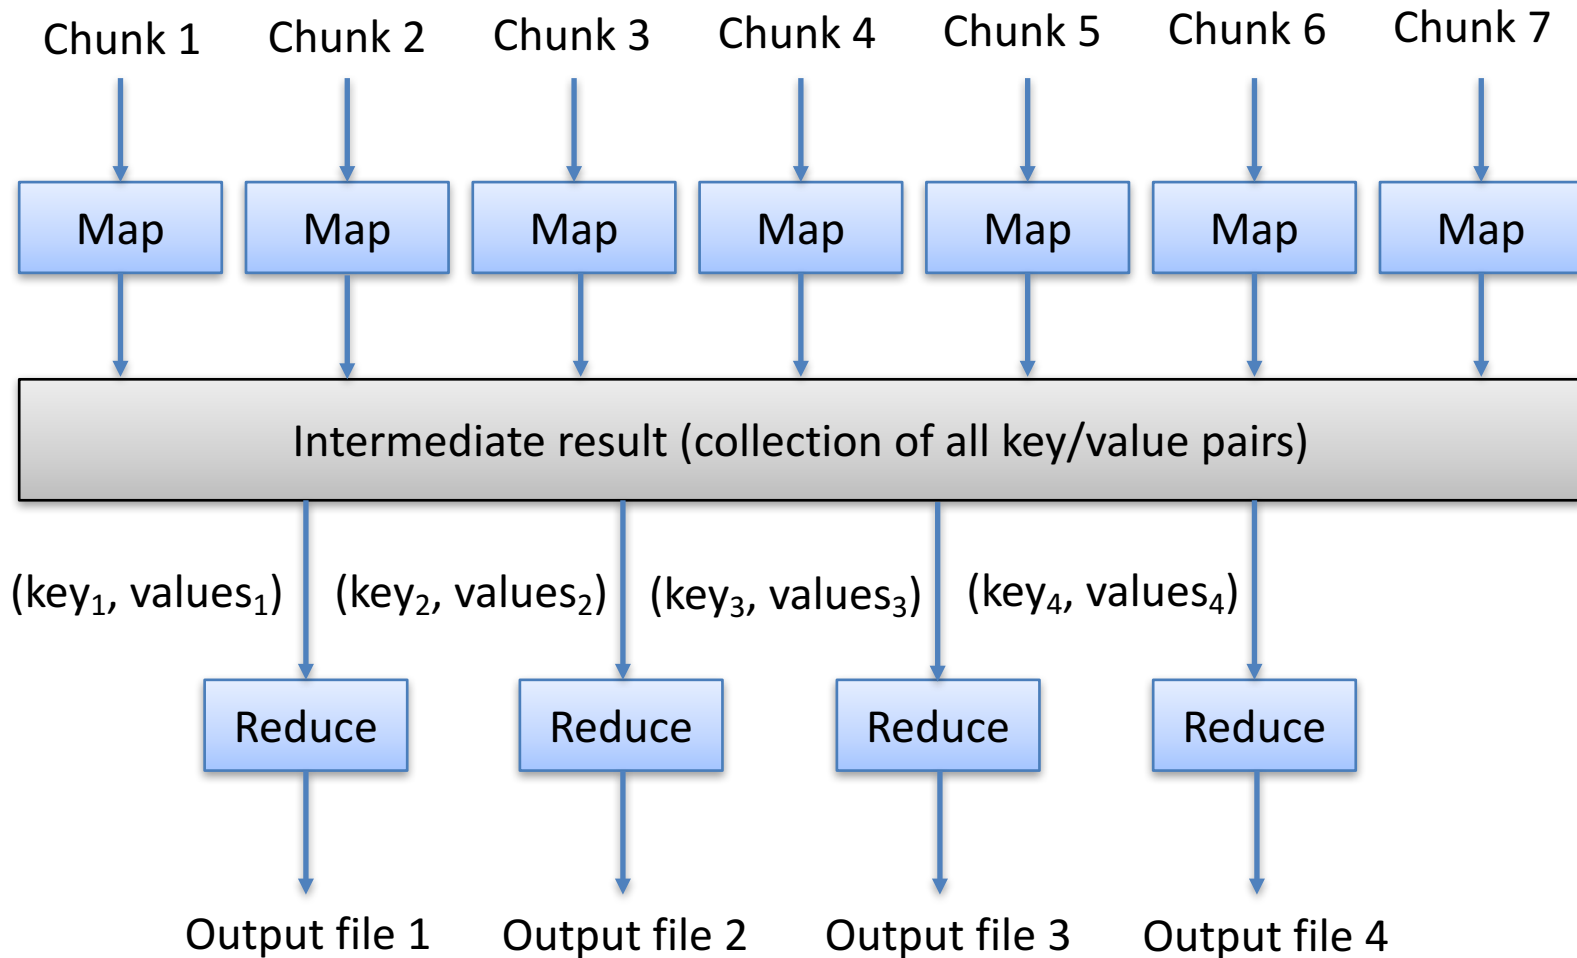
- Keys/values can be arbitrary objects
  - Keys are typically small (e.g., integers, strings, etc.)
  - Values might be larger (e.g., entire text documents, images, etc.)
- Each MapReduce computation is expressed in terms of two functions: **Map & Reduce**



# Implementation

- **Map**(String key, String value):
  - Returns a list of key/value pairs
- **Reduce**(String key, Iterator<String> values):
  - Returns a list of key/value pairs
- In practice also some additional code to set:
  - Locations of input and output files
  - Tuning parameters (e.g., number of machines, memory per Map/Reduce task, etc.)

# Execution



# Expressiveness

- MapReduce can be used to compute a number of interesting functions on large datasets
  - Inverted index: for each word, return the list of all documents that contain it
  - Operators of relational algebra
  - Matrix multiplication → PageRank
- MapReduce is not a universal parallelism framework
  - Not every problem that is parallelisable can be expressed and solved nicely in MapReduce
- Bottle-neck: communication & disk access

# Summary

- MapReduce is a framework for solving problems on large volumes of data
  - Simple implementation: Map & Reduce
  - MapReduce takes care of the rest
- Many interesting problems can be solved using MapReduce
- Big Data tools are based on MapReduce

*Thanks for listening!*

*... this will not be covered in the exam.*