

COMP207

Database Development

Lecture 25

Beyond Relational Data: NoSQL Databases (Part 2)

NoSQL Databases

- NoSQL databases:
 - **Distributed**
 - Focus on **availability, high performance & fault-tolerance**
 - Give up some characteristic features of relational DBMS such as full ACID transactions



Recall the CAP Theorem

Key-Value Stores

- Collection of table, tables = key-value pairs

Key	Value
User 1	{ "name": "John Smith", "items": { "1": { "name": "product 1", "quantity": 3 }, "2": { "name": "product 2", "quantity": 1 }, ... } }
...	...

Essentially an index!

- Simple access mechanism:
 - **find(k)**: returns value for key **k**
 - **write(k, v)**: inserts value **v** under key **k**
- Fast due to index on key, no further indexes & no transactions

Example Voldemort Session

- Terminal:


```
> bin/voldemort-shell.sh test tcp://localhost:6666
Established connection to test via [tcp://localhost:6666]
> put "hello" "world"
> get "hello"
version(1:2) ts:1511089049077: "world"
> delete "hello"
> get "hello"
null
```

- Java:

```
client = ... // connect to server
client.put("hello", "world");
Versioned<String> val = client.get("hello");
System.out.println("Value for 'hello': " + val);
client.delete("hello");
```

DynamoDB & Voldemort Techniques

- Techniques:
 - Distributed hash table
 - Replication
 - Versioning and conflict-resolution using vector clocks
- Access:
 - Reads return local data
 - Writes update local data first, which will then be propagated
- Questions:
 - Why not assign a node to each number?
 - Why reuse nodes?



Document Stores

Document Stores

- Databases that store collections of “documents”
 - Document = semistructured data associated with an object ID

Collection “restaurants”

ObjectID	Value
1028434	<pre>{ "name": "Morris Park Bake Shop", "address": { "coord": [-73.856077, 40.848447], "street": "Morris Park Ave", "zipcode": "10462" }, "grades": [{ "date": 1393804800000, "grade": "A", "score": 2 }, { "date": 1299715200000, "grade": "B", "score": 14 }] }</pre>
...	<pre>{ "name": "Wendy'S", "address": { "coord": [-73.961704, 40.662942], ... } ... }</pre>
...	...

often assigned automatically

Documents represented in JSON

Typical Use Case

- Restaurant information web service

Rules out key-value stores


- Filter collection of restaurants by location, ratings, etc.
- Data associated with restaurants might vary
- User requests typically require all data associated with a restaurant, but do not need to join data from different restaurants

Collection "restaurants"

ObjectID	Value
1028434	<pre>{ "name": "Morris Park Bake Shop", "address": { "coord": [-73.856077, 40.848447], "street": "Morris Park Ave", "zipcode": "10462" }, "grades": [{ "date": 1393804800000, "grade": "A", "score": 2 }, ...</pre>
...	<pre>{ "name": "Wendy'S", ...</pre>

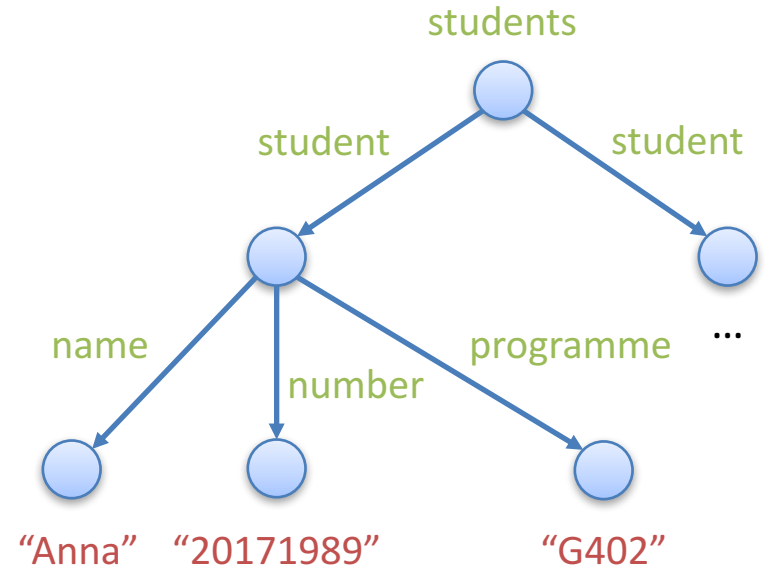
JSON vs XML

```
<?xml version="1.0" standalone="yes">
<students>
  <student>
    <name>Anna</name>
    <number>20171989</number>
    <programme>G402</programme>
  </student>
  <student>
    <name>John</name>
    <number>20174378</number>
    <programme>G702</programme>
  </student>
  ...
</students>
```



JSON vs XML

```
{"students": [  
  {"name": "Anna",  
   "number": 20171989,  
   "programme": "G402"},  
  {"name": "John",  
   "number": 20174378,  
   "programme": "G702"},  
  ...  
]}
```



Case Study: MongoDB

<https://www.mongodb.com>

- Stores documents in variant of JSON (previous slide)
 - Creating/managing collections

```
db.createCollection( "students" )
```

- Insert/update/delete documents

```
db.students.insert(  
  {name: "Anna", studentID: 20171989, year: 2})
```

- Finding documents

```
db.students.find( {year: 2} )  
db.students.find( {year: { $lt: 3 } } ).sort( {year: 1} )
```

More flexible than
key-value stores

- Aggregation, ...

Indexes in MongoDB

- Also: indexes on one or more fields of documents (in contrast to key-value DBs)

Collection "students"

ObjectID	Value
1028434	{ "name": "Anna", "studentID": "1234", "year": 2, "programme": "G402" }
...	...

- `db.students.createIndex({ name: 1 })`
- `db.students.createIndex({ programme: 1, year: -1 })`

MongoDB Techniques

- Sharding (horizontal fragmentation)
 - Collections are split into **horizontal fragments**
 - Based on **shard key**: indexed field, exists in all documents

ObjectID	Value
1028434	{ "name": "Anna", "studentID": "1234", "year": 2 }
...	...

- Replication
 - Horizontal fragments of collections are replicated
 - Master-slave approach:
 - Primary copy (master) and secondary copies (replicas)
 - Updates: on master, then delegated to replicas
 - Reads: delegated to master (default) or to any replica

Limited Transaction Support

- ACID for transactions affecting single documents, and the multi-document version is being developed
- Concurrency control:
 - Multi-granularity locks at global/database/collection level, optimistic scheduling at lower levels
 - If you need two-phase commit, must implement it yourself (or use a distributed relational DBMS), or wait until they implement it
 - Limited isolation:
 - Updates can be “locked” so that other operations see their effect only after completion
 - Doesn’t work with sharding yet
 - See <https://docs.mongodb.com/manual/faq/concurrency>

Other Types of NoSQL Databases

Column Stores

- Google BigTable, Apache Hbase, ...

Diagram illustrating a Column Store structure with annotations:

- Table name:** student
- Column family:** contact
- Column qualifier:** address
- Cell:** Anna
- Row:** 123

	name		contact			
	first	last	address	city	postcode	email
123	Anna
456	Ben
789	Chloe
...

- Table names/column families fixed, column qualifiers vary
- Columns referenced as <column family>:<column qualifier>

Case Study: Hbase

<https://hbase.apache.org/>

- Uses Hadoop Distributed File System
 - also mentioned in MapReduce lecture
 - Creating table

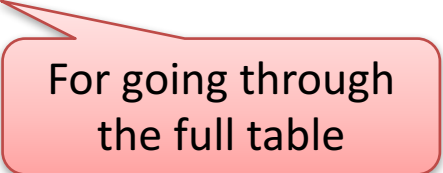
```
create 'STUDENT', 'Name', 'Grade', 'Programme'
```

- Inserting rows

```
put 'STUDENT', 'row1', 'Name:Fname', 'Anna'  
put 'STUDENT', 'row1', 'Grade:COMP207', '100'
```

- Finding documents

```
get 'STUDENT', 'row1'  
scan 'STUDENT'
```



For going through
the full table

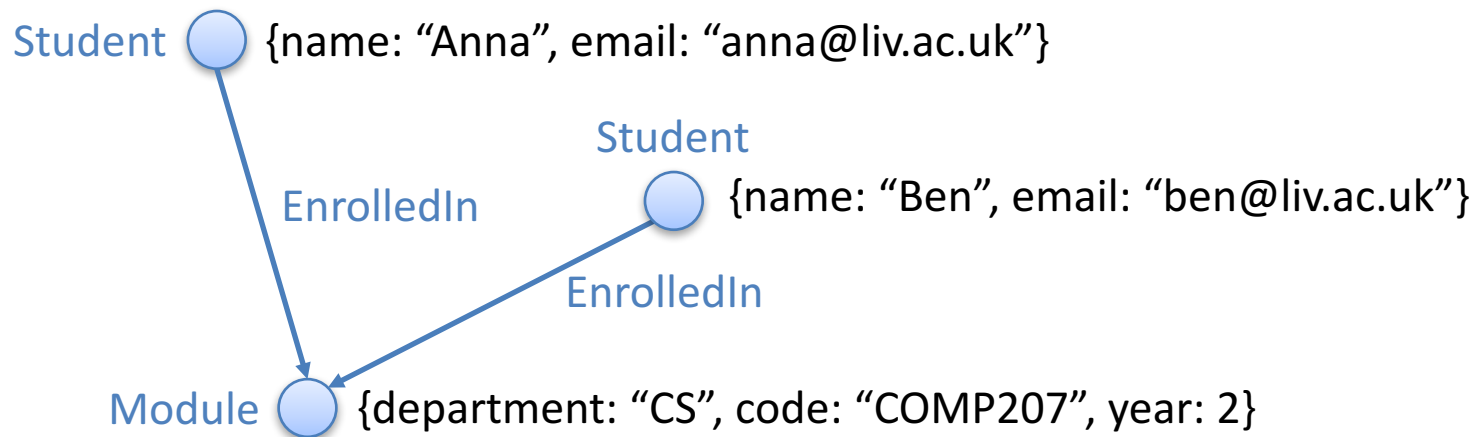
Hbase Techniques

- Uses two levels of fragmentation:
 - Top level: rows are divided into regions (i.e., ranges of rows)
 - Bottom level: regions store different column families in different nodes (i.e., computers)
- No transaction support
- Each item has a time stamp and one can access past versions of the database (if setup)

```
> get 'STUDENT', 'row1'
COLUMN          CELL
name:first      timestamp=1511139129384, value=Anna
name:last       timestamp=1511139140333, value=...
```

Graph Databases

- Store data as a graph:



- Many systems: e.g., Neo4j
 - Data is accessed using SQL-like path query language
 - Indexes, ...

Summary

- NoSQL databases are the newest members to the data management family
 - Simpler/different data models and query language
 - Weaken ACID properties
- Main use cases:
 - Applications where speed, availability, scalability is crucial, and ACID properties can be weakened
 - Applications where the different way of accessing data (cell-wise as in a column store, as a graph, ...) is crucial
- Not a replacement for relational databases