

COMP207  
Database Development  
Tutorial 1c  
Precedence Graphs

Precedence graphs

# Precedence Graph

time	T1	T2	T3
1			
2			
3			
4			

1. Scheduler produces schedule Sa:

Sa: r1(X);w2(X);w1(X);w3(X);




2. Set up a table

# Precedence Graph

time	T1	T2	T3
1			
2			
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);

3. Node for each transaction






# Precedence Graph

time	T1	T2	T3
1	r(X)		
2			
3			
4			

Read in the first operation of the schedule. No action can be taken yet – there is nothing to compare the operation to

Sa: r1(X);w2(X);w1(X);w3(X);




4. Process Sa:



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);






Read in the next operation of Sa

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);






Begin to test for serialisability by looking for conflicting operations

Q1: Same item?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);






Q1: Same item? **YES** - both use X



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);






Q1: Same item? **YES**

Q2: Different Transactions?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);






Q1: Same item? **YES**

Q2: Different Transactions? **YES** (T1 and T2)

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? YES




Q2: Different Transactions? YES

Q3: Is one of the operations a write?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**




Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**

Q2: Different Transactions? **YES**




Q3: Is one of the operations a write? **YES**

**Three YES** means these are conflicting operations

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Same item? **YES**

Different Transactions? **YES**


Is one of the operations a write? **YES**

These are conflicting operations:  
Produce an edge (arrow) from the  
earlier operation to the later one

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3			
4			

Sa: r1(X);w2(X);w1(X);w3(X);




Same item? **YES**

Different Transactions? **YES**

Is one of the operations a write? **YES**



These are conflicting operations  
Create an edge on the graph



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);





- Q1: Same item?
- Q2: Different Transactions?
- Q3: Is one of the operations a write?



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**


Q2: Different Transactions?

Q3: Is one of the operations a write?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**


Q2: Different Transactions? **YES**

Q3: Is one of the operations a write?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**



**Three YES** means these are conflicting operations



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**

Q2: Different Transactions? **YES**


Q3: Is one of the operations a write? **YES**

These are conflicting operations  
Create an edge

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			

Sa: r1(X);w2(X);w1(X);w3(X);




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**



These are conflicting operations  
Create an edge



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item?



Q2: Different Transactions?

Q3: Is one of the operations a write?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**



Q2: Different Transactions?

Q3: Is one of the operations a write?

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**

Q2: Different Transactions? **YES**



Q3: Is one of the operations a write?



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

Sa: r1(X);w2(X);w1(X);w3(X);



Q1: Same item? **YES**


Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

Sa: r1(X);w2(X);w1(X);w3(X);




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

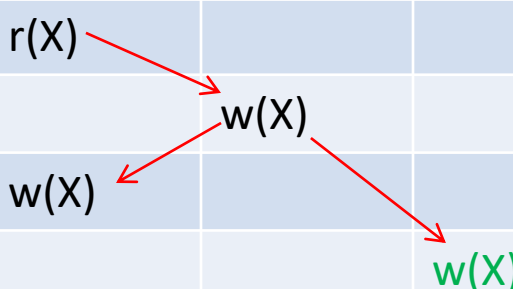
Q3: Is one of the operations a write? **YES**

**Three YES** means these are conflicting operations




# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)



Sa: r1(X);w2(X);w1(X);w3(X);




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

These are conflicting operations  
Create an edge



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)


Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

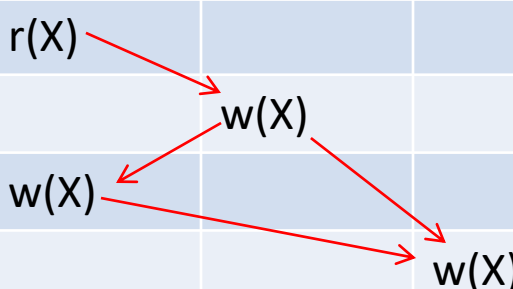
These are conflicting operations  
Create an edge

Sa: r1(X);w2(X);w1(X);w3(X);



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

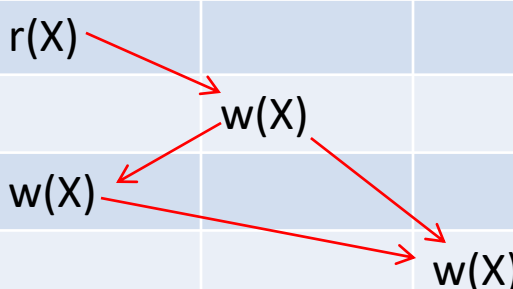
**Three YES** means these are conflicting operations  
Create an edge

Sa: r1(X);w2(X);w1(X);w3(X);



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)



Q1: Same item? **YES**


Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

**Three YES** means these are conflicting operations

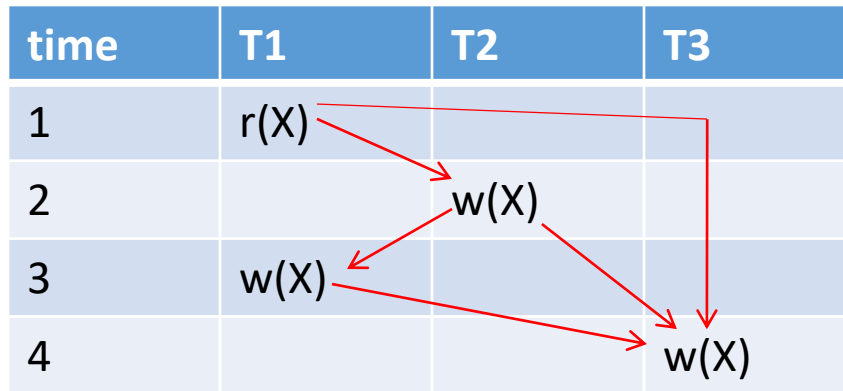
Create an edge

Sa: r1(X);w2(X);w1(X);w3(X);



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)




Q1: Same item? **YES**

Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

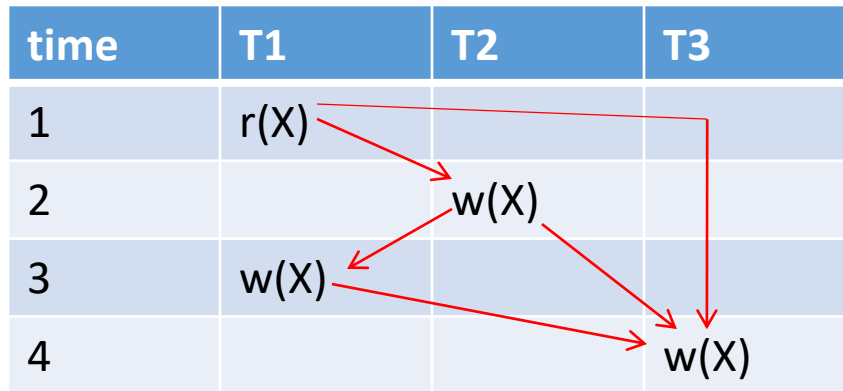
**Three YES** means these are conflicting operations  
Create an edge

Sa: r1(X);w2(X);w1(X);w3(X);



# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)



Q1: Same item? **YES**


Q2: Different Transactions? **YES**

Q3: Is one of the operations a write? **YES**

**Three YES** means these are conflicting operations

Create an edge

Sa: r1(X);w2(X);w1(X);w3(X);






# Precedence Graph

time	T1	T2	T3
1	r(X)		
2		w(X)	
3	w(X)		
4			w(X)

If needed, duplicate edges can be reduced to one only for clarity as in T1 to T3.

We cannot reduce the edges between T1 and T2 as they are different edges – note the direction

Sa: r1(X);w2(X);w1(X);w3(X);



# Precedence Graph Analysis

- Schedule  $S_a$  produced by the scheduler is obviously NOT a serial schedule as different transactions have been interleaved with each other – hence it is a concurrent schedule
- This means there may be conflict between operations in the schedule that could give ‘wrong’ results
- However, there are combinations of conflicting operations that will still give the ‘correct’ result, i.e. they will work just as if they had been executed serially (so  $S_a$  is ‘Conflict Serialisable’)

# Precedence Graph Analysis

- To see if  $S_a$  is conflict serialisable, produce a Precedence Graph and analyse the output
- For  $S_a$  there is a cycle in the graph
  - $T1 \rightarrow T2$  and  $T2 \rightarrow T1$
  - So the schedule is not conflict serialisable
  - It may give the 'wrong' result
- If NO cycle is present
  - The graph is 'acyclic'
    - The schedule is conflict serialisable (to a serial schedule)

# Precedence Graph Analysis

- If a schedule is conflict serialisable, we are saying it is equivalent to executing the transactions in it as if they were done in a serial manner
- But what serial schedule is this equivalent to?
  - If we have a schedule with three transactions in it (T1, T2 and T3) then is the conflict equivalent to running T1-T2-T3 or T2-T3-T1 or T1-T3-T2 etc?

# View-Equivalence - Complexity

Algorithms are available to test for a schedule being view-serializable or not.

Algorithmic Complexity is characterised by ‘Computational Complexity’ – the total amount of elementary computation required by the algorithm as a function of the size of the data set to which the algorithm applies.

Many algorithms have Polynomial Complexity – complexity expressed as a polynomial function of data set size – ‘tractable problem’

# View-Equivalence - Complexity

- The problem of view-serialisability is NP-hard
  - **Non-deterministic Polynomial-time** hard
    - Polynomial time algorithms are reasonable to compute
    - But... The time to run the algorithm grows too fast to compute exact solutions in all cases
  - Unlikely to find an efficient polynomial time algorithm
  - NP-Complete problems have no known solution algorithm with a polynomial complexity
- Linear Complexity is a special case of Polynomial Complexity. NP-Complete problems are a class of problems that are 'intractable problems'

# View Equivalence

- Less stringent definition of schedule equivalence
- Schedules consisting of the same operations are view equivalent when:
  - 1. In  $S_a$ , if  $T_1$  reads initial value of item then in  $S_b$ ,  $T_1$  must also read the initial value of the same item
  - 2. In  $S_a$ , if  $T_1$  reads an item written by  $T_2$ , then in  $S_b$ ,  $T_1$  must also read the same item written by  $T_2$
  - 3. In  $S_a$ , if  $T_1$  performs the final write on an item, then in  $S_b$ ,  $T_1$  must also perform the final write on the same item

# View Equivalence

- A schedule is view serialisable when it is view equivalent to a serial schedule
- Every conflict serialisable schedule is view serialisable
- A view serialisable schedule may not be conflict serialisable
  - These break the constrained write rule
    - Can't write an item unless it's been read first
      - "Blind Write"
    - There is an unconstrained write assumption that does allow blind writes



# View Equivalence

- Creates all possible combinations of the transactions
- Given:  $S_x = r_1(X); r_2(Y); r_2(Y); w_2(X); w_3(Y); r_1(X);$ 
  - $S_a = T_1 T_2 T_3$
  - $S_b = T_1 T_3 T_2$
  - $S_c = T_2 T_1 T_3$
  - $S_d = T_2 T_3 T_1$
  - $S_e = T_3 T_1 T_2$
  - $S_f = T_3 T_2 T_1$

# View Equivalence

- Given:  $S_x = r_1(x); r_2(Y); r_2(Y); w_2(X); w_3(Y); r_1(X);$
- A schedule is view serialisable if:
  - 1. A transaction that reads the initial data in a schedule is the same transaction that reads the initial data in another schedule
    - T1 reads data before T2 – so T1 comes before T2
      - Remove  $S_c$ ,  $S_d$  and  $S_f$

# View Equivalence

- Given:  $r_1(x); r_2(Y); r_2(Y); w_2(X); w_3(Y); r_1(X);$
- A schedule is view serialisable if:
  - 2. A transaction that reads an item after another transaction writes it in a schedule is the same transaction that reads an item after another transaction writes it in another schedule
    - T1 reads data after T2 writes – so T2 comes before T1
      - Remove  $S_a, S_b, S_e$
    - But ... In test1, we show T1 must come before T2
    - This produces a cycle so the schedule  $S_x$  is NOT view serialisable

# View Equivalence

- Given:  $S_y = w_3(Z); r_2(X); w_2(Y); r_1(Z); w_3(Y); w_1(Y);$ 
  - Note this is a 'blind write' example
- A schedule is view serialisable if:
  - 3. A transaction that writes the final value for an item in a schedule is the same transaction that writes the final value for that item in another schedule
    - T1 writes the final value – so T1 comes after T2 and T3
      - If it doesn't, then T2 or T3 will overwrite the value of Y
    - Remove  $S_a, S_b, S_c$  and  $S_e$
    - $S_d = T_2, T_3, T_1$  or  $S_f = T_3, T_2, T_1$  are **view equivalent** schedules